



# Kafka workshop 2025



Somkiat Puisungnoen

Search

Somkiat | Home

Update Info 1 View Activity Log 10+ ...

Timeline About Friends 3,138 Photos More

When did you work at Opendream? X

... 22 Pending Items

Post Photo/Video Live Video Life Event

What's on your mind?

Public Post

Intro

Software Craftsmanship

Software Practitioner at สยามชานาญกิจ พ.ศ. 2556

Agile Practitioner and Technical at SPRINT3r

Somkiat Puisungnoen 15 mins · Bangkok · ...

Java and Bigdata



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

Page

Messages

Notifications 3

Insights

Publishing Tools

Settings

Help ▾



somkiat.cc

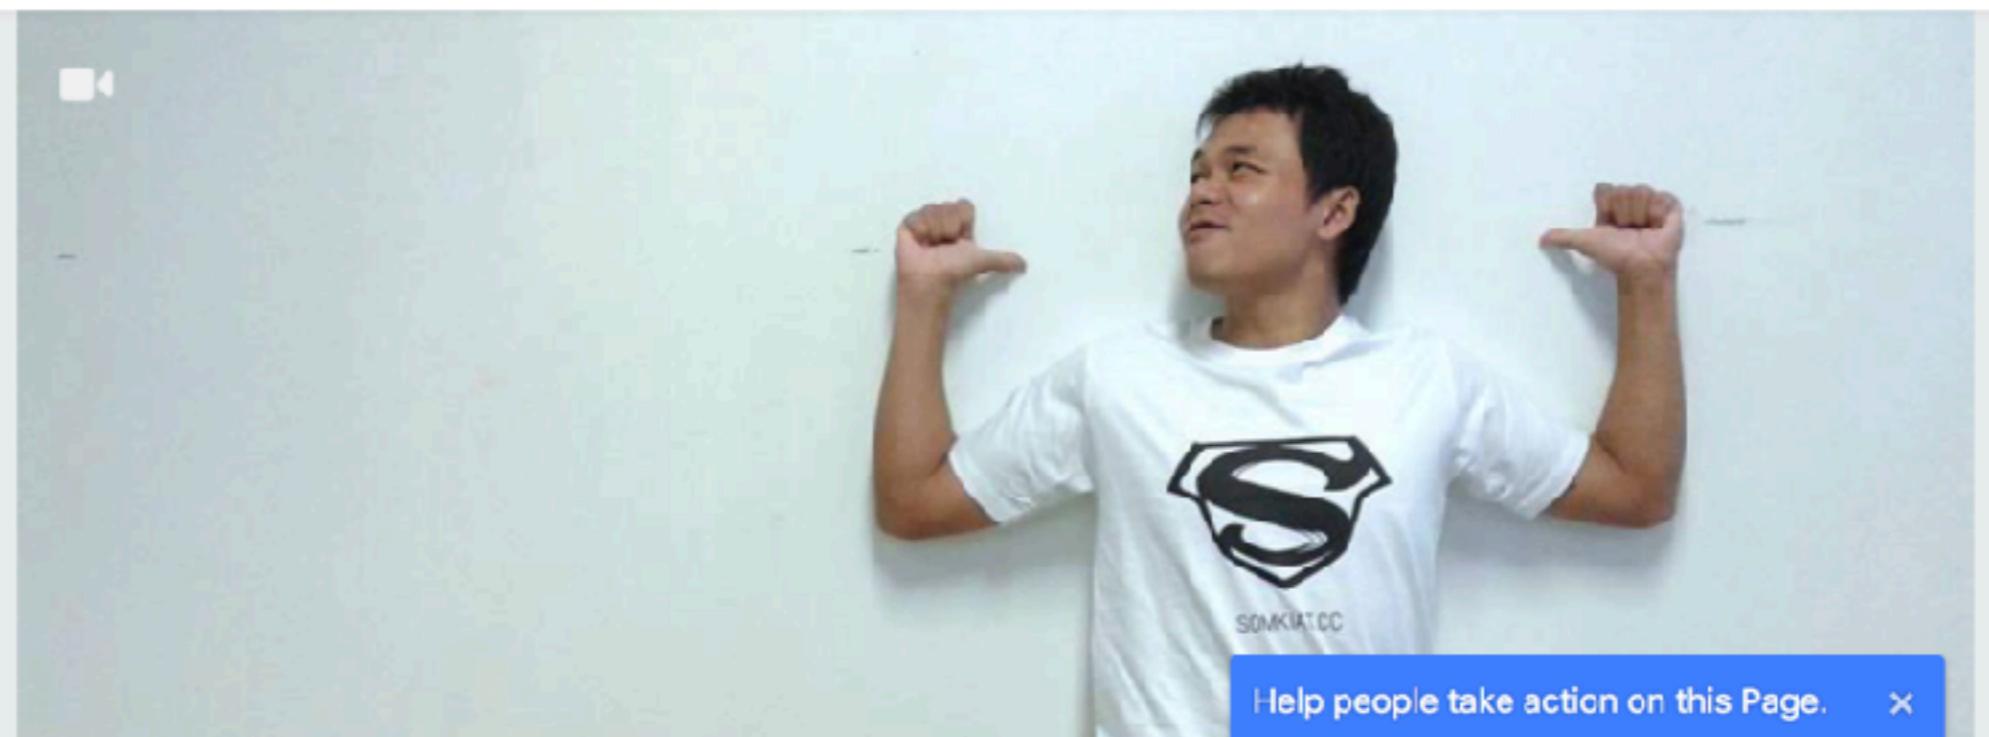
@somkiat.cc

Home

Posts

Videos

Photos



Liked

Following

Share

...

Help people take action on this Page. 

+ Add a Button



Kafka

3

**[https://github.com/up1/  
course-kafka-2025](https://github.com/up1/course-kafka-2025)**



# Topics (1)

- Why Apache Kafka ?
- What is Apache Kafka ?
- Architecture
- Topologies and Tools
- Use cases



# Topics (2)

- Working with Kafka cluster
- Broker, Producer, Consumer
- Monitoring
- Performance testing



# Topics (3)

- Event-driven with Kafka
- Develop with C#
- Testing
- Monitoring system



# Learning paths

## Part 1

Why, What ?

Starting Kafka

Architecture

Install and config

## Part 2

Configuration

Producers

Consumers

Monitoring

## Part 3

Event-driven

Develop with C#



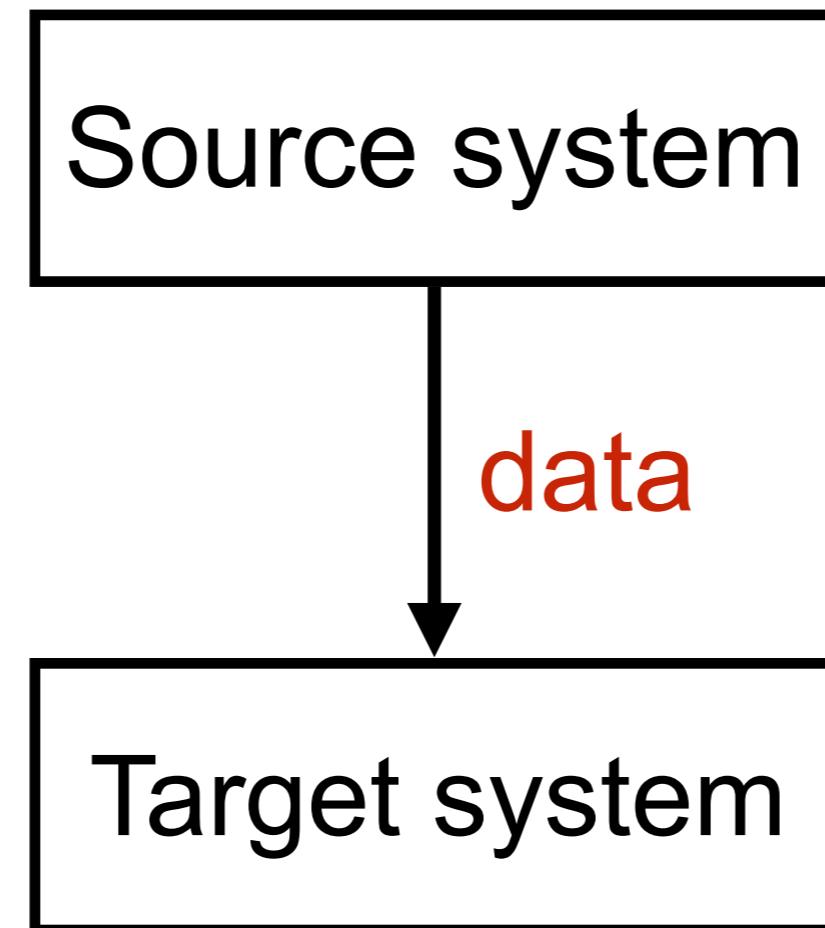


<https://kafka.apache.org/>



# Why ?

Starting with simple



# Why ?

More source systems

More target systems

More protocols

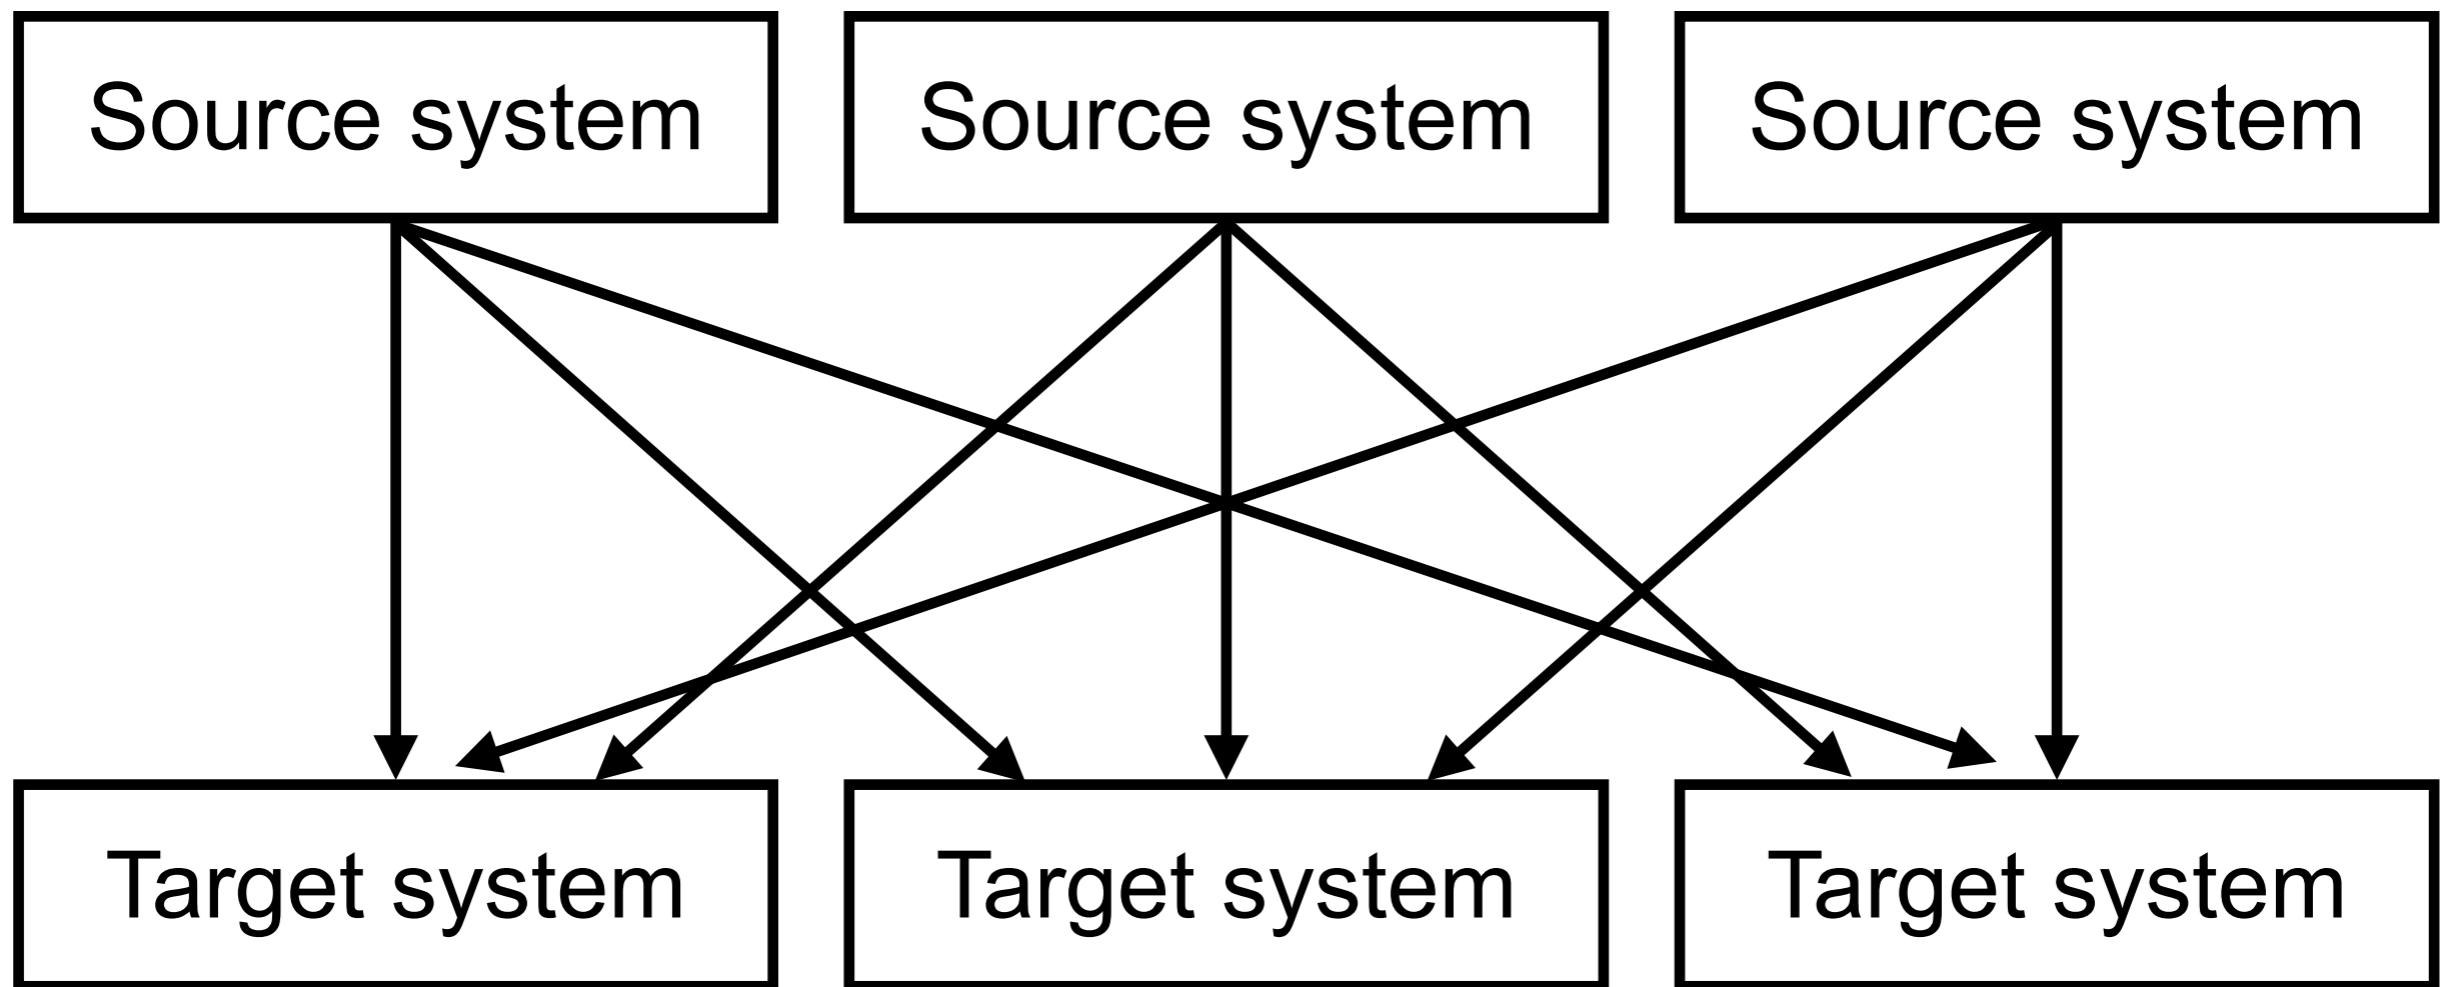
More data formats and schemas

More loads/traffics



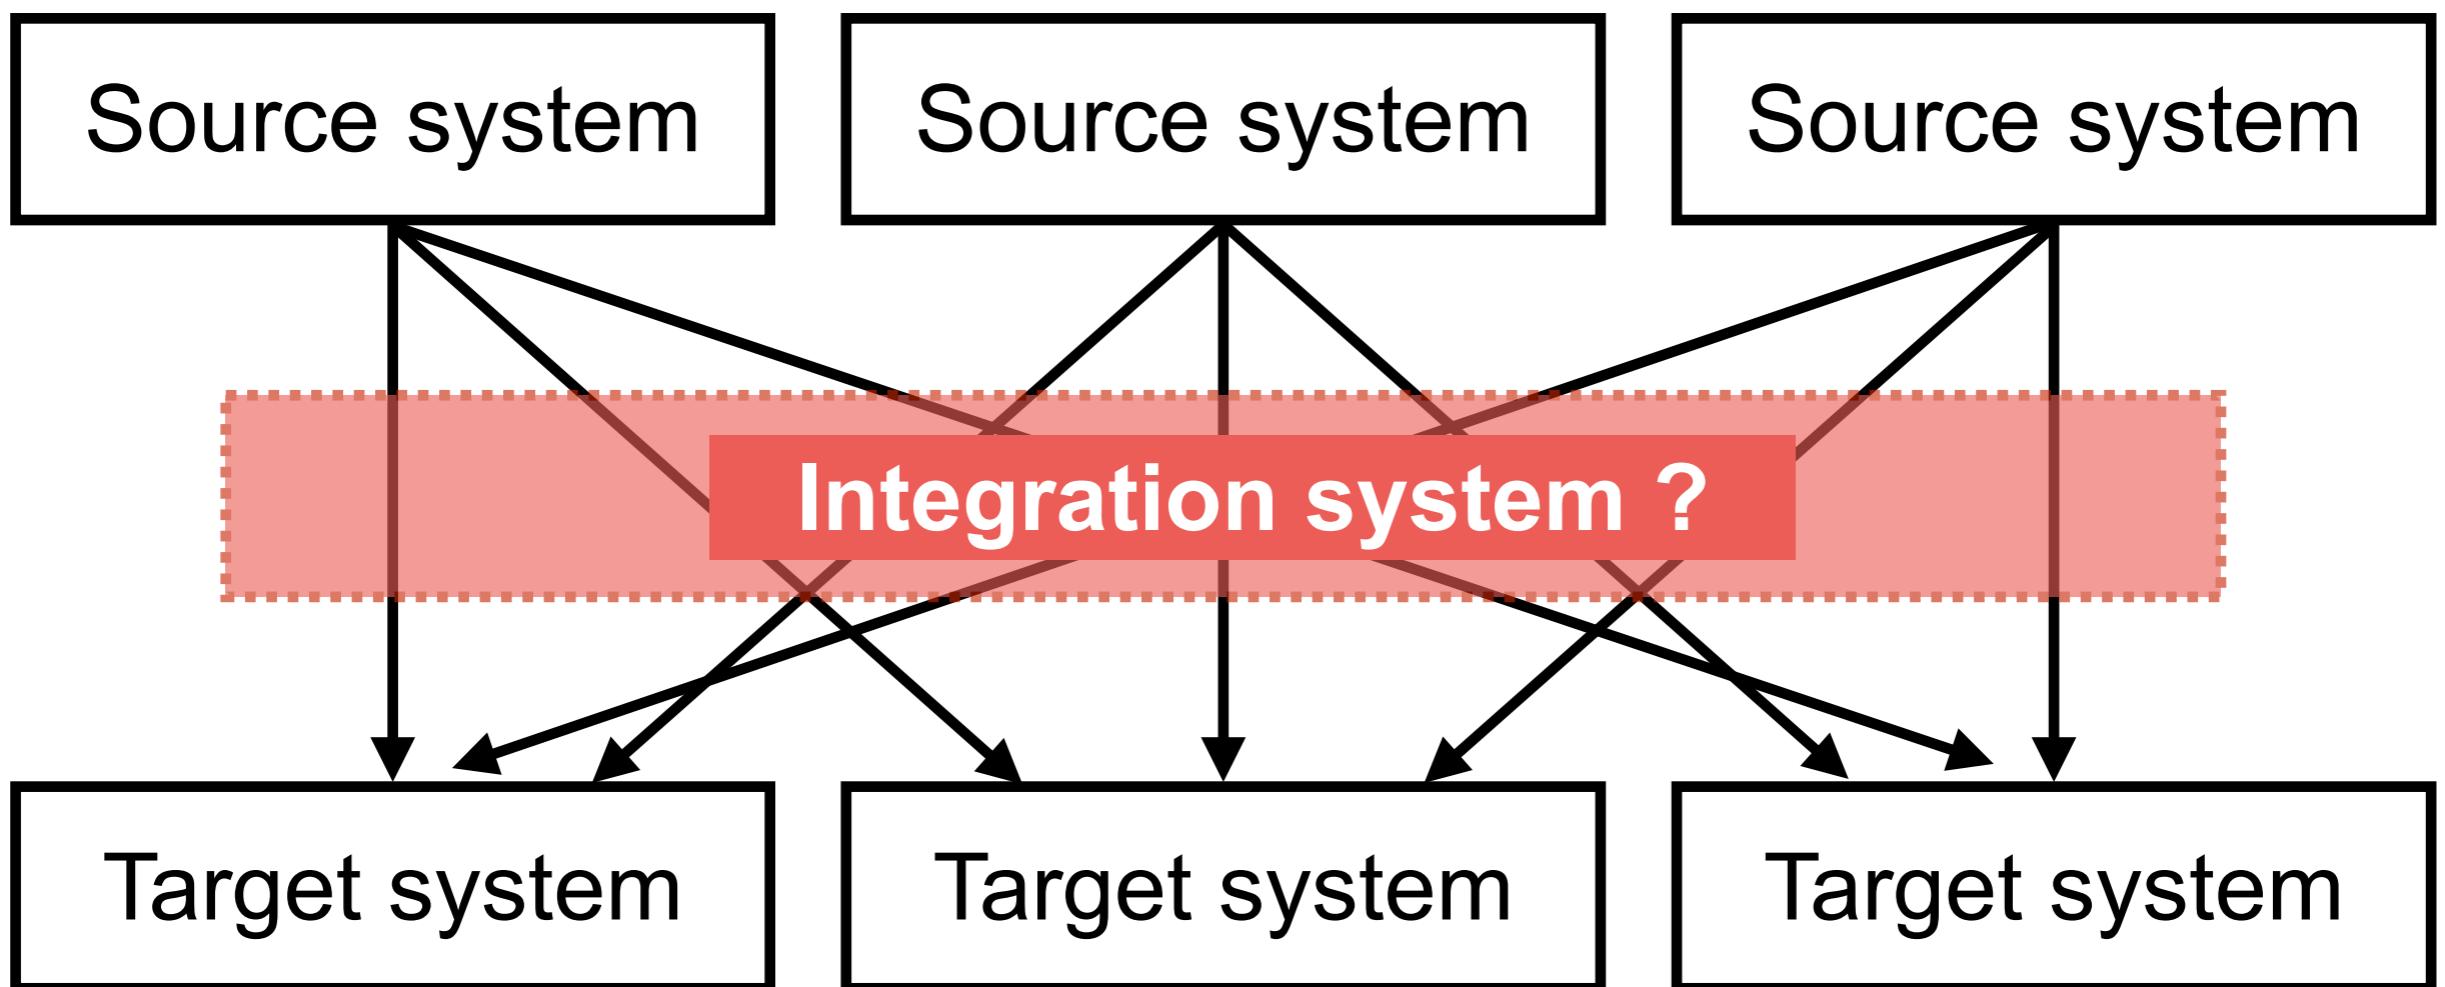
# Why ?

Complex system and coupling !!



# Why ?

Need integration system



# Integration ?

Protocols

Data format

Data schema and evolution

TCP  
HTTP  
REST  
JDBC

CSV  
JSON  
Avro  
Binary



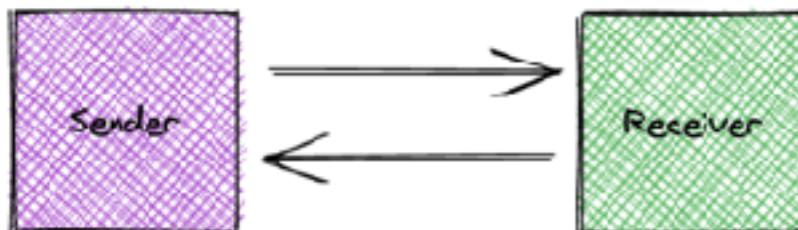
# Communication ?

## Sync vs Async Communication

Request and wait for a response  
Fire and forget with messages/events

@boyney123

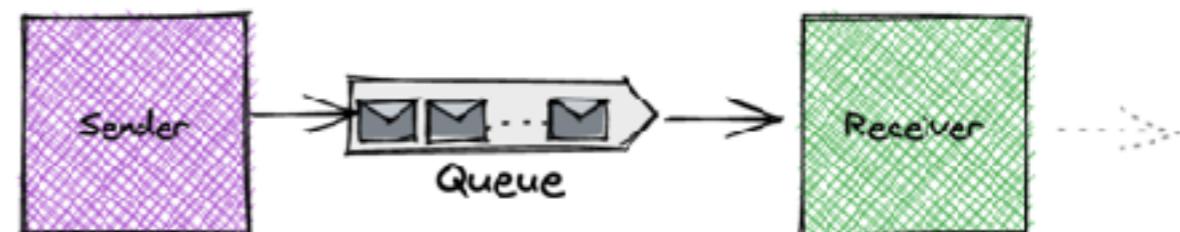
### Synchronous



Request-response model

Send a request and wait for some response  
Example: API requests

### Asynchronous



Fire and forget model

Send message/event and forget  
Example: Event Driven Architecture

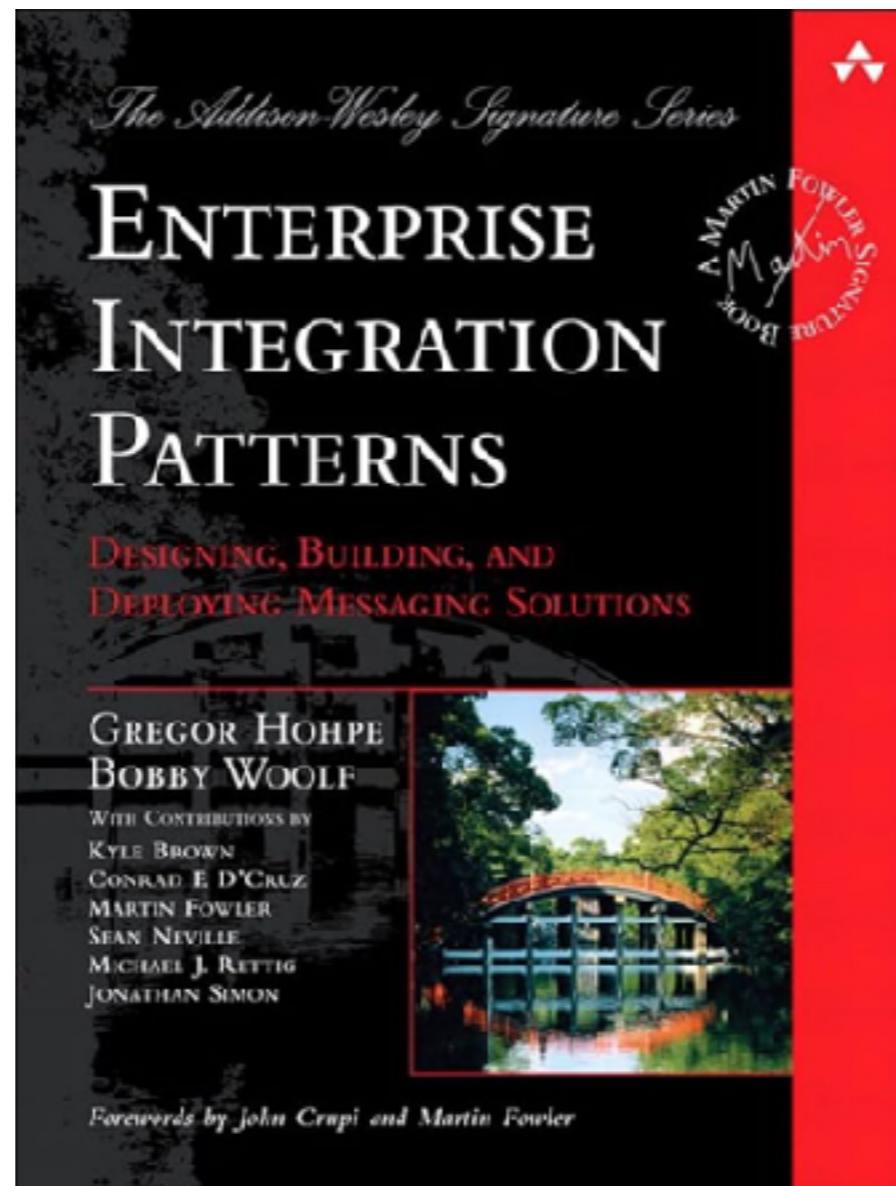
<https://eda-visuals.boyney.io/visuals/sync-vs-async>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Enterprise Integration Patterns



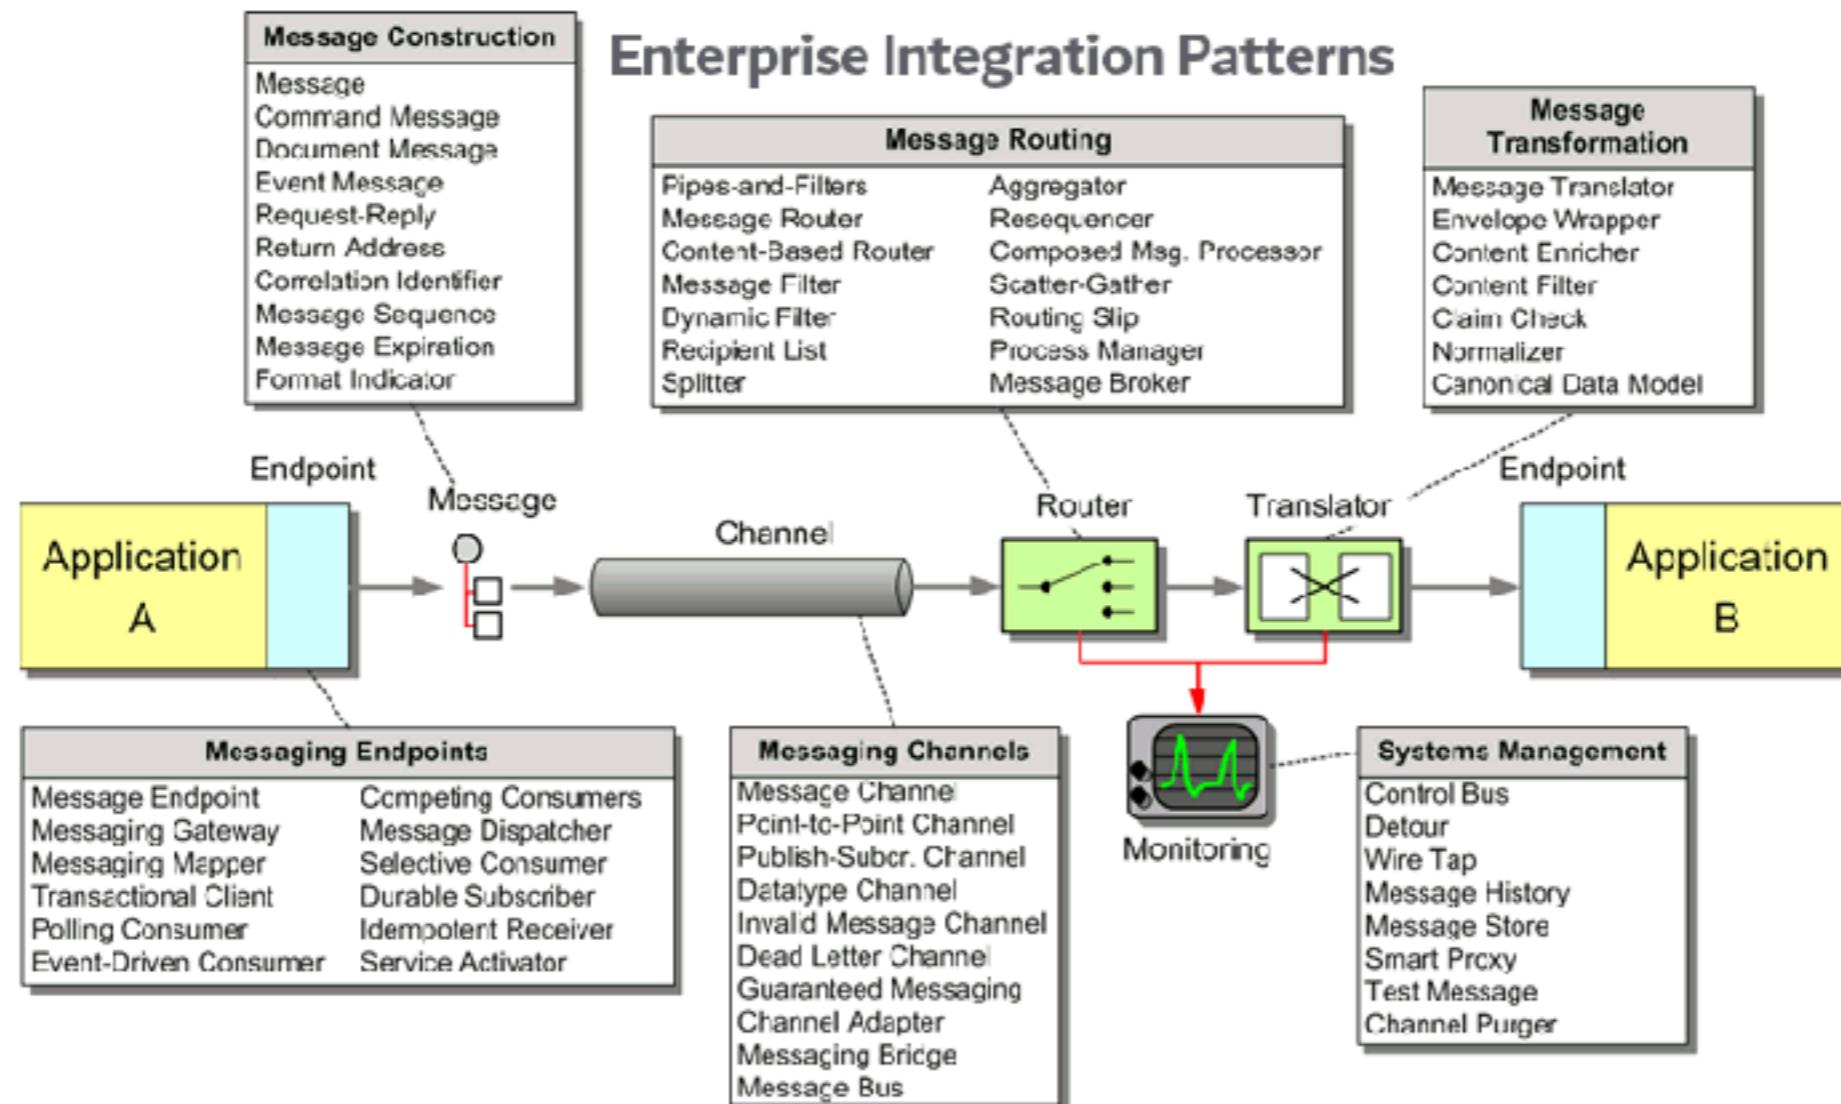
<https://www.enterpriseintegrationpatterns.com/>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Enterprise Integration Patterns



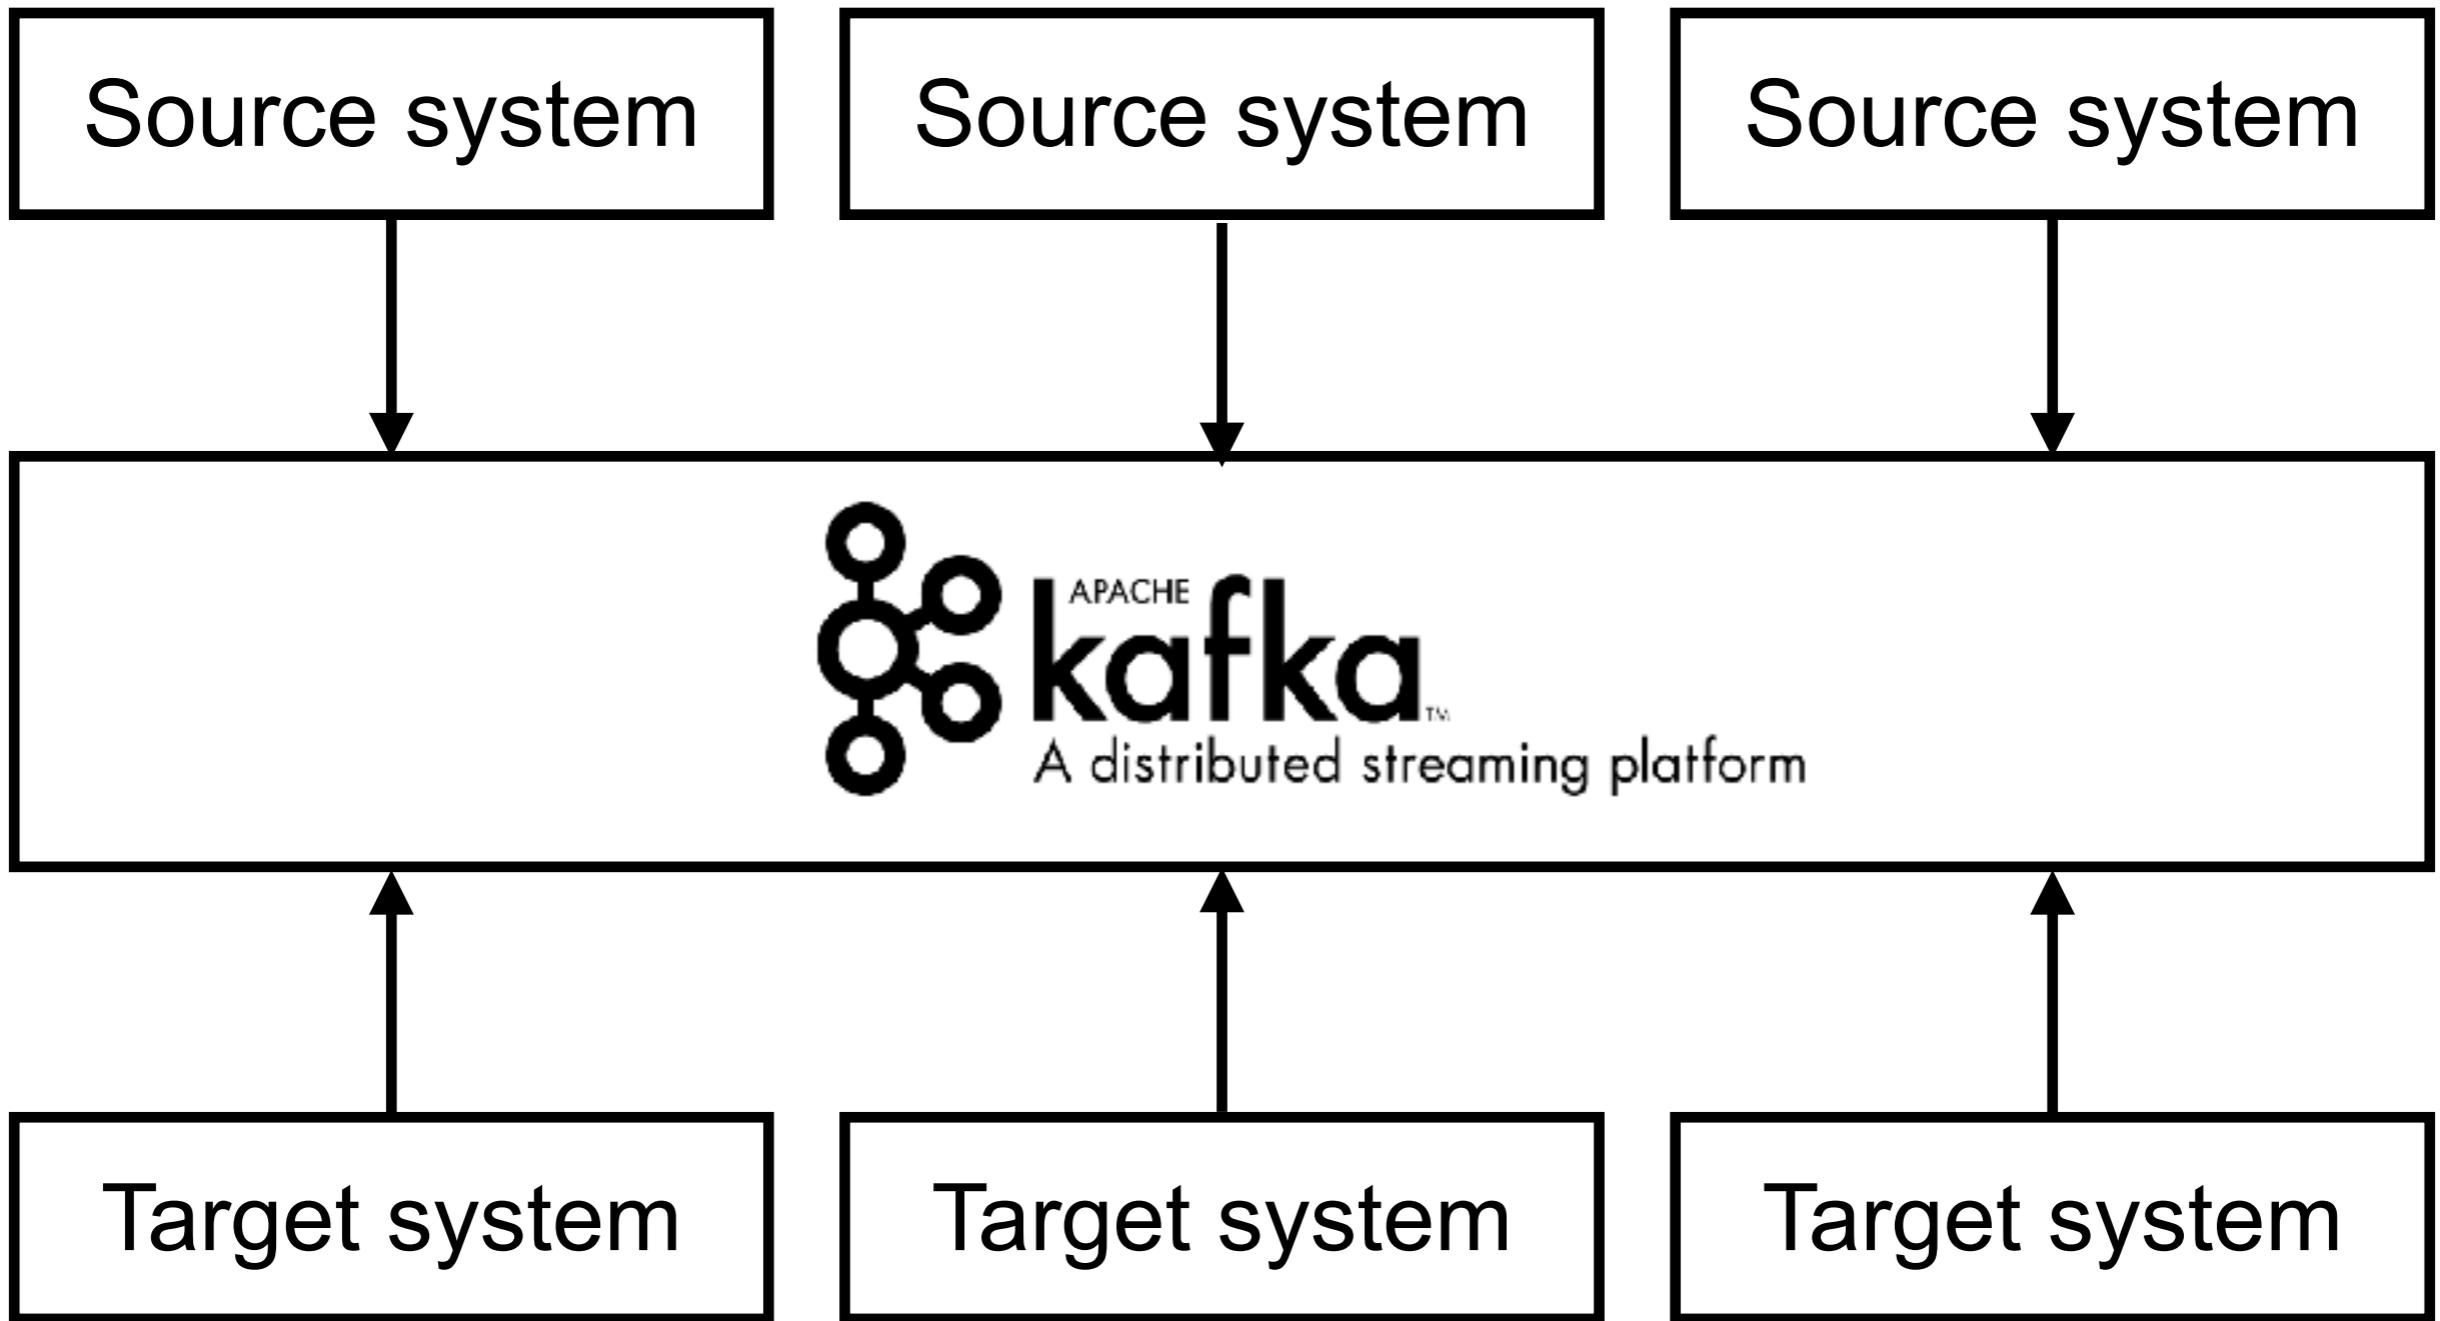
<https://www.enterpriseintegrationpatterns.com/>



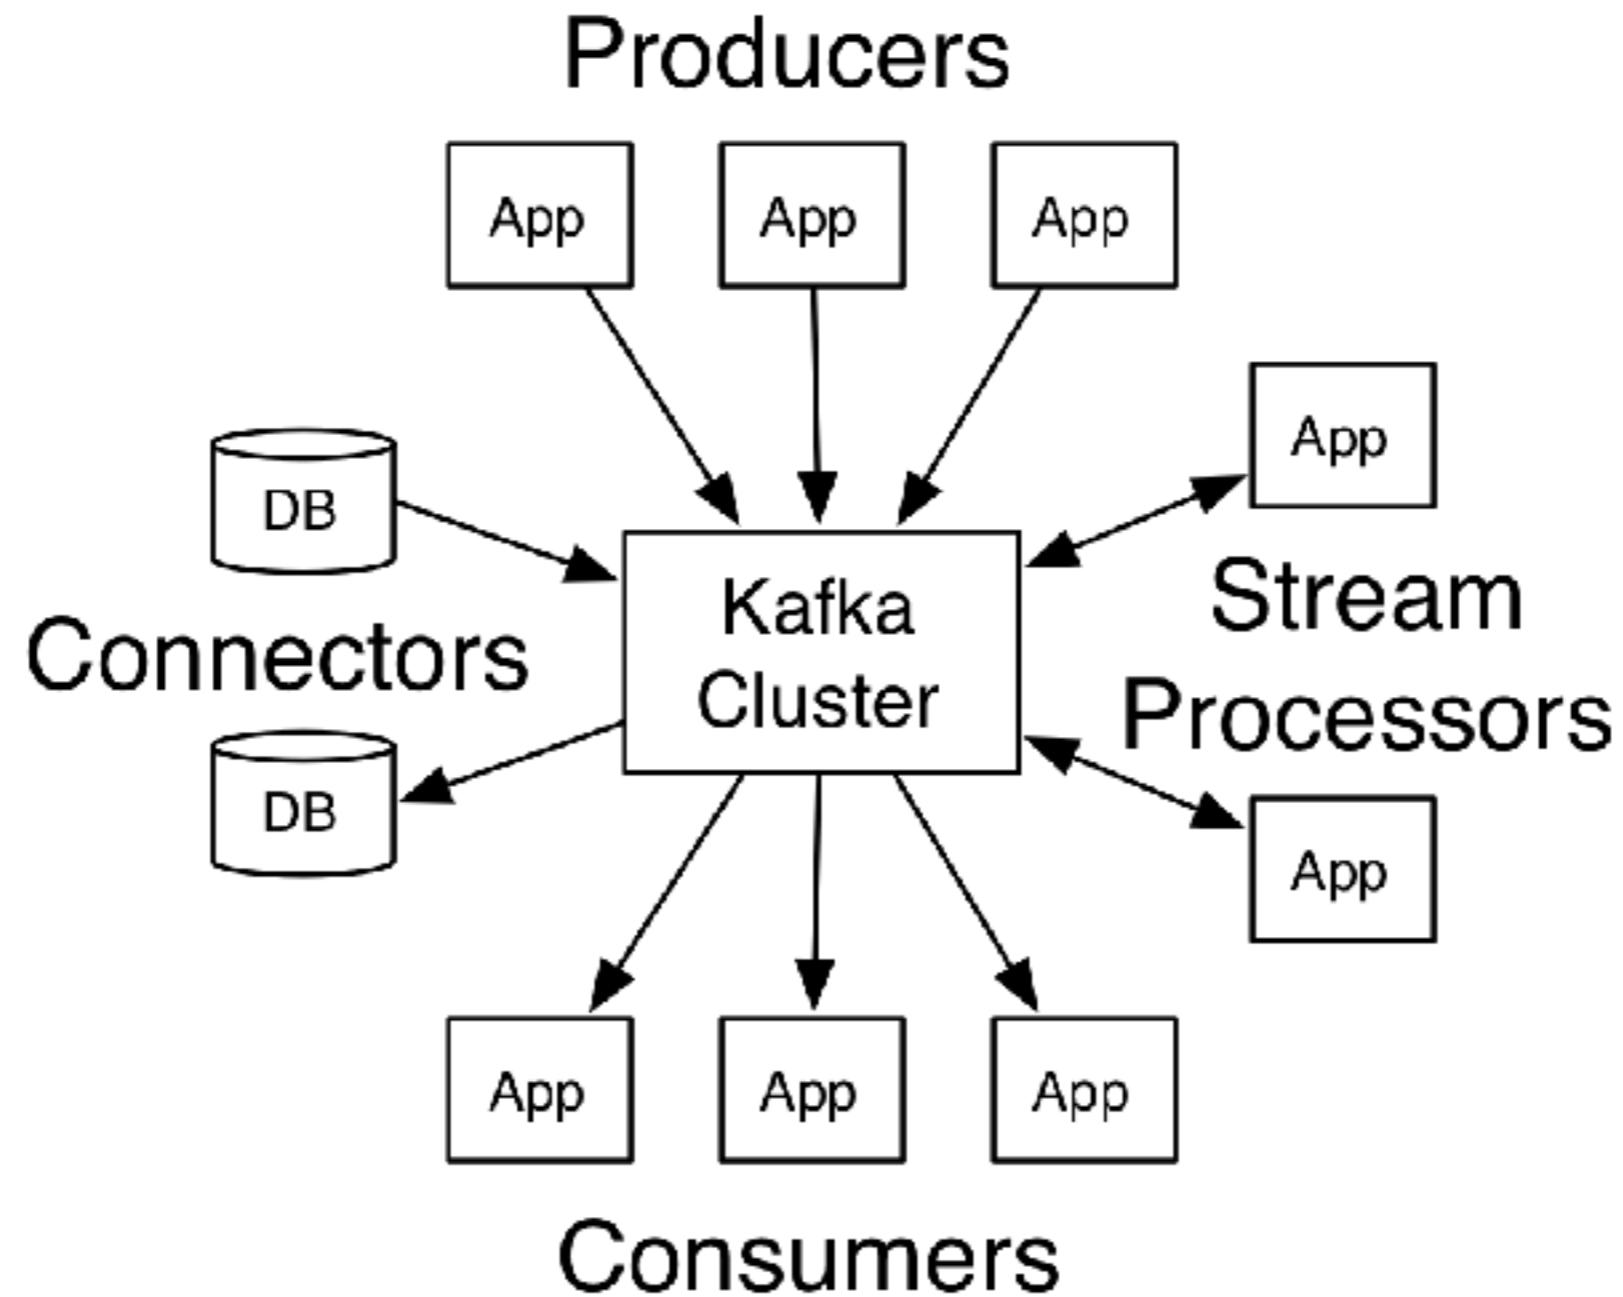
Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Apache Kafka



# Apache Kafka



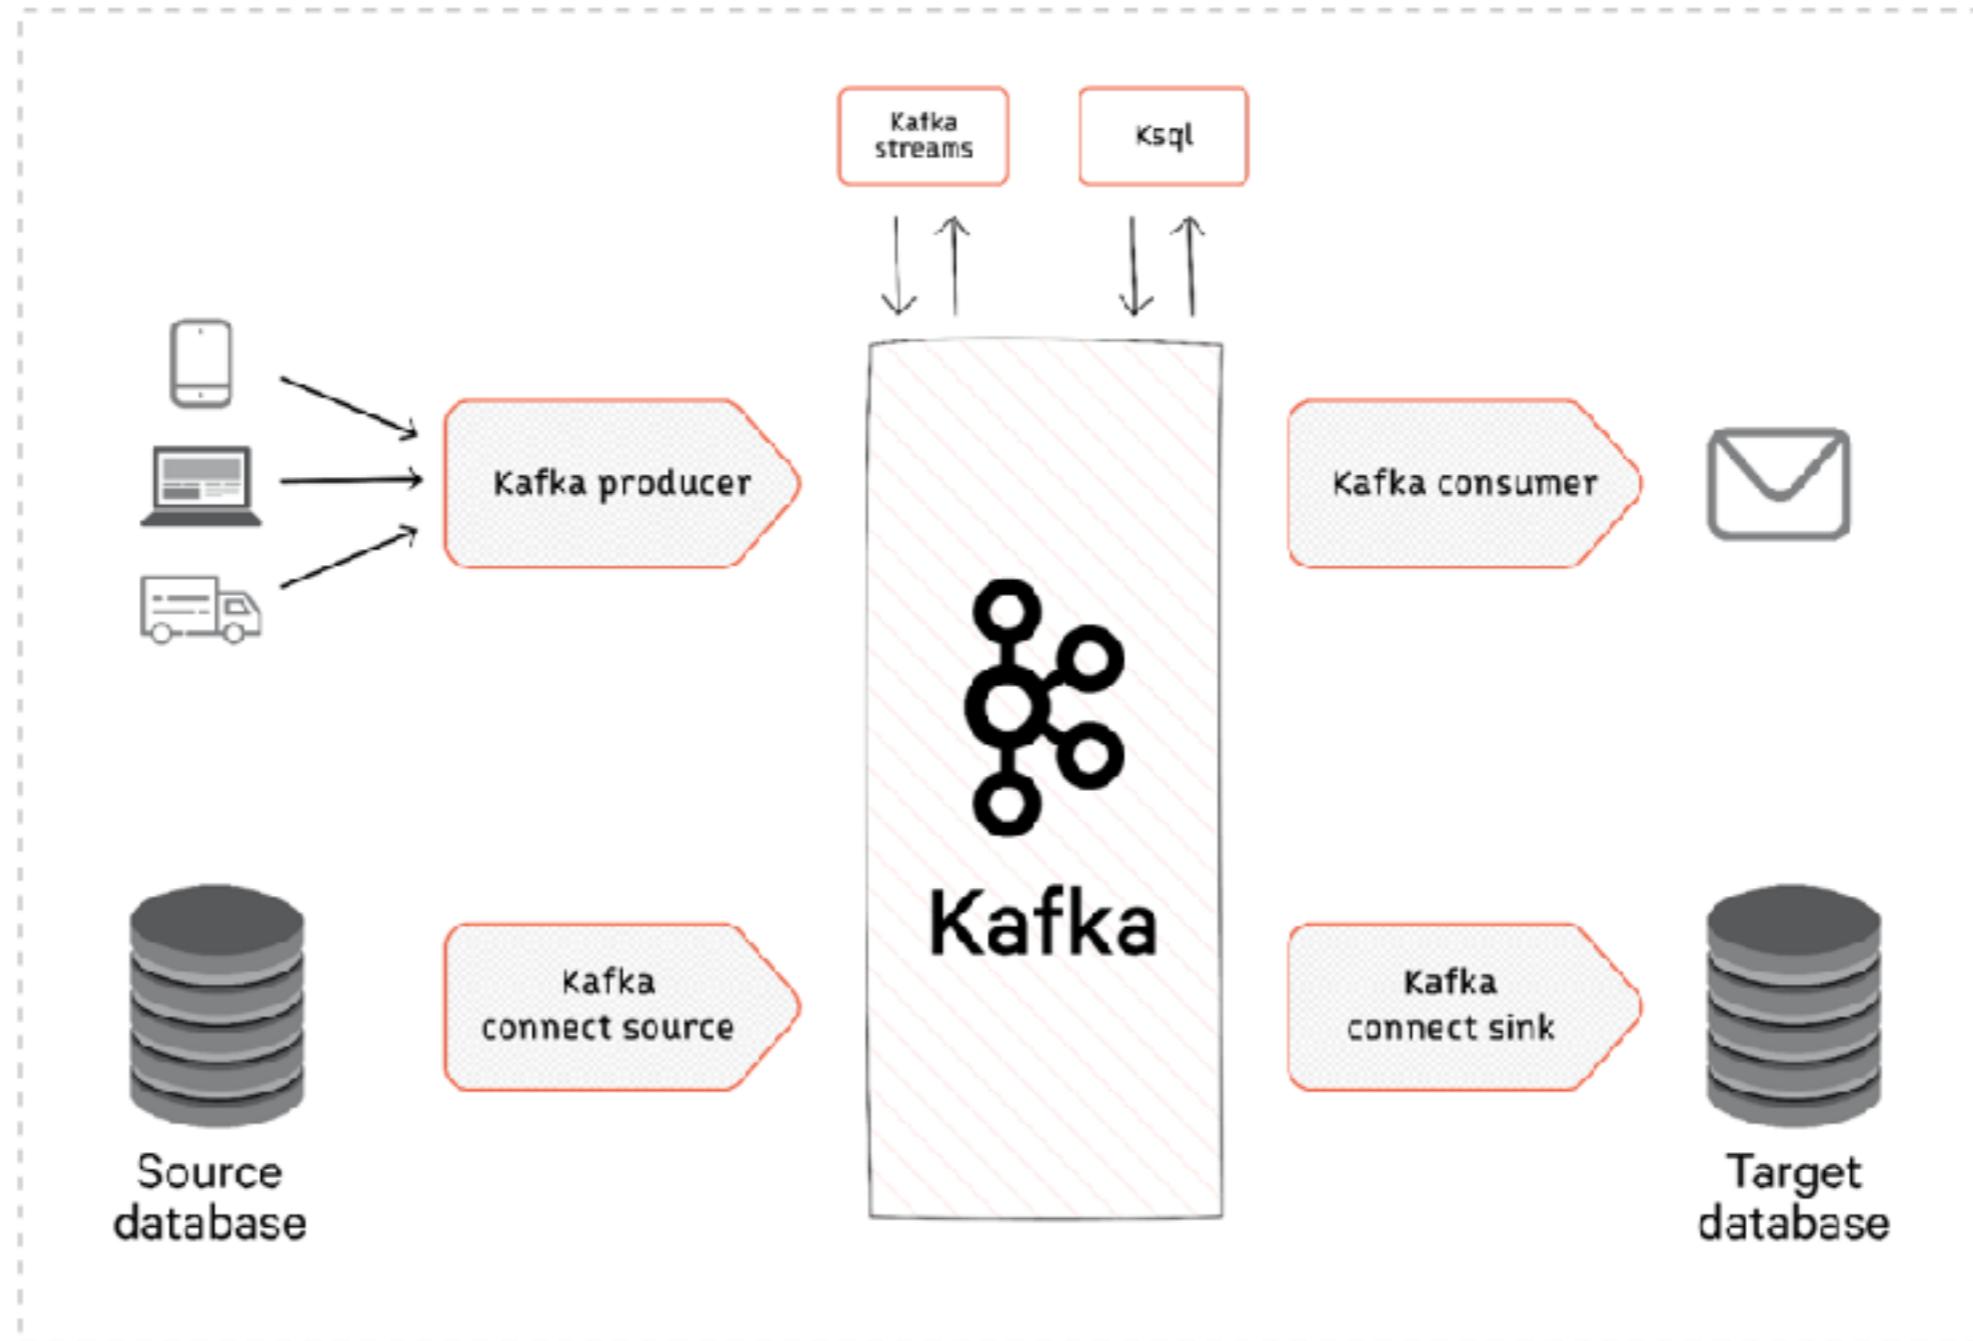
<https://kafka.apache.org/intro.html>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Apache Kafka



<https://www.redpanda.com/guides/kafka-alternatives-kafka-api>



# What Apache Kafka ?

Created by LinkedIn

Open source project (maintain by Confluent)

Distributed system

Resilient architecture, **fault tolerant**

Horizontal scalability

High performance/ low latency

High throughput

Low latency



# Goals ?

Data Integration

Decoupled system

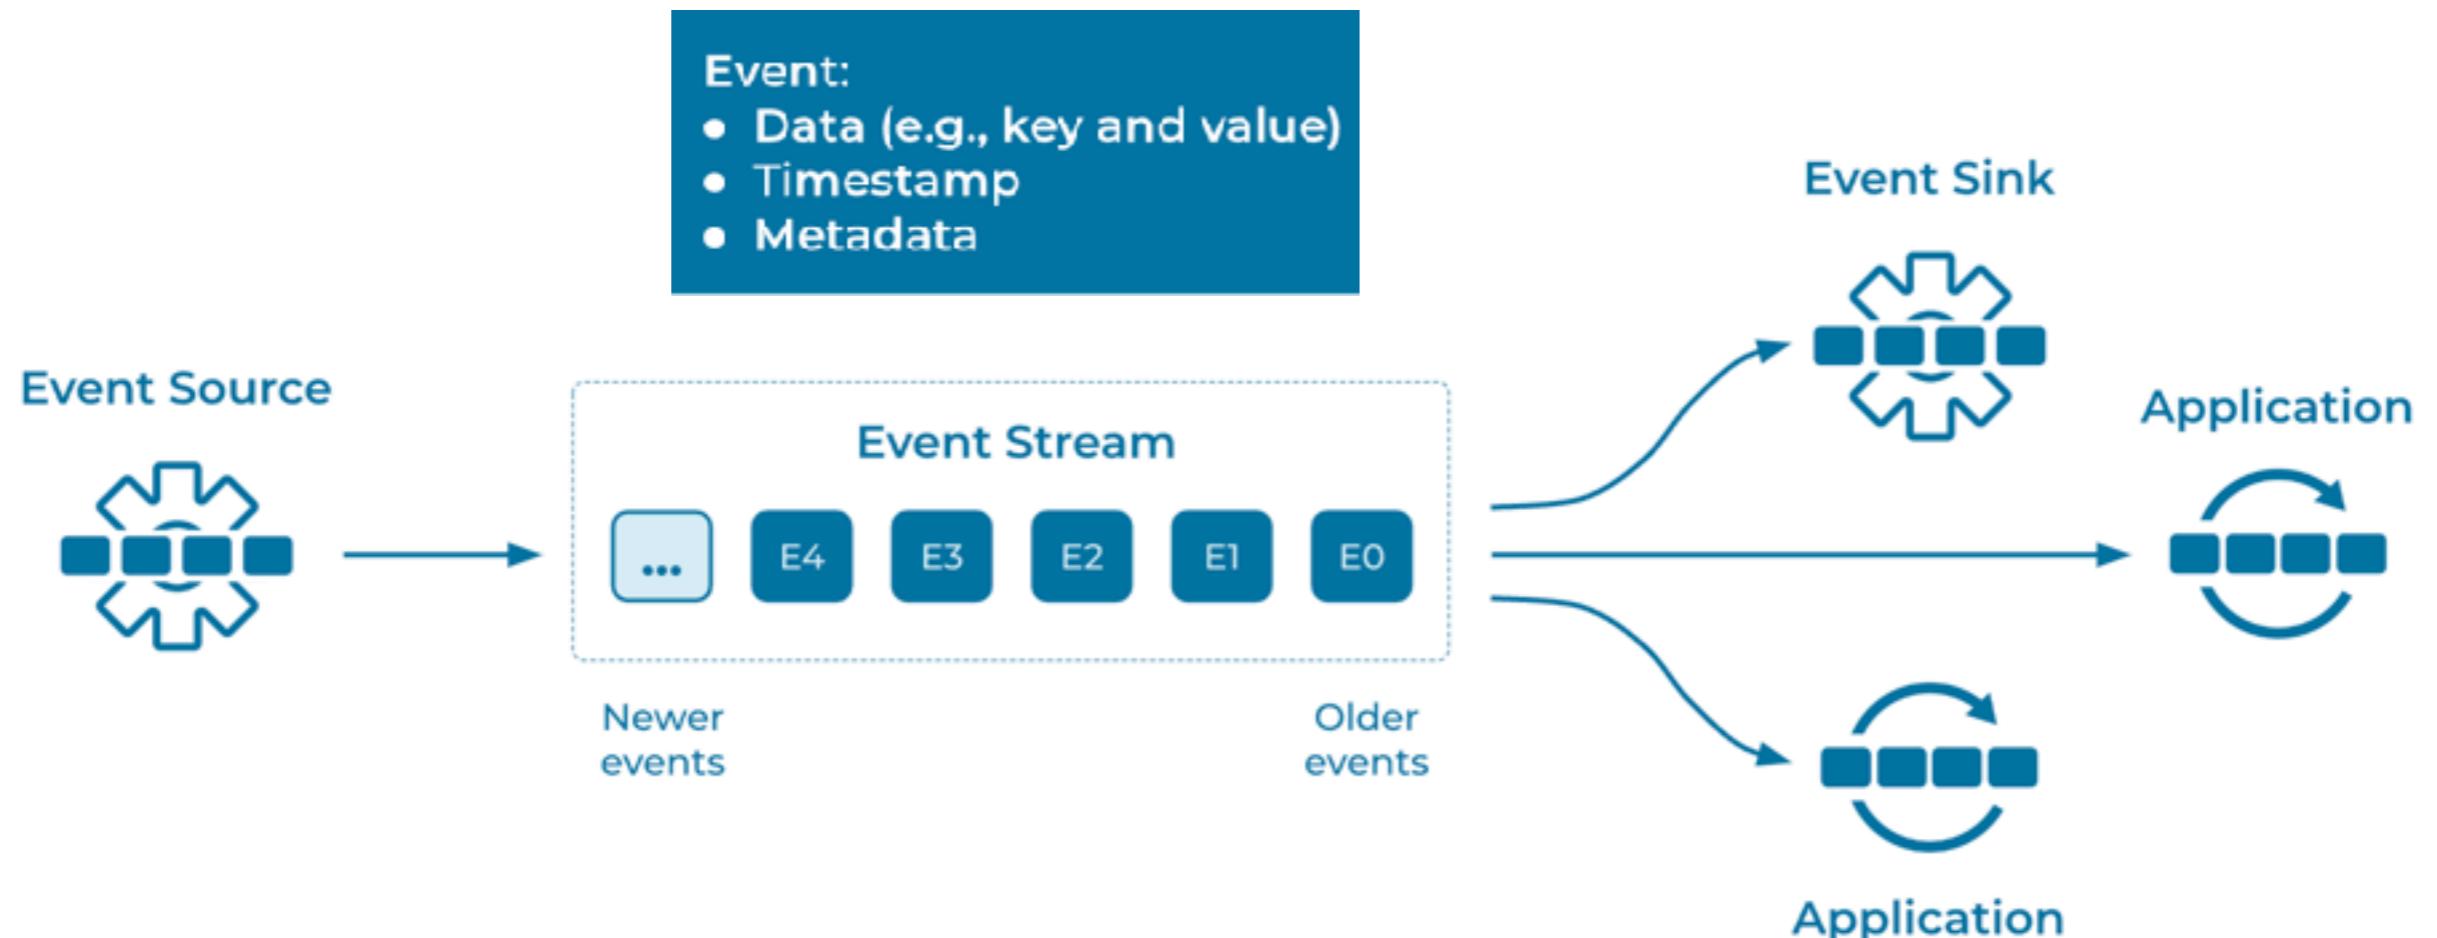
Open-source  
**distributed event streaming platform**



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Event Streaming



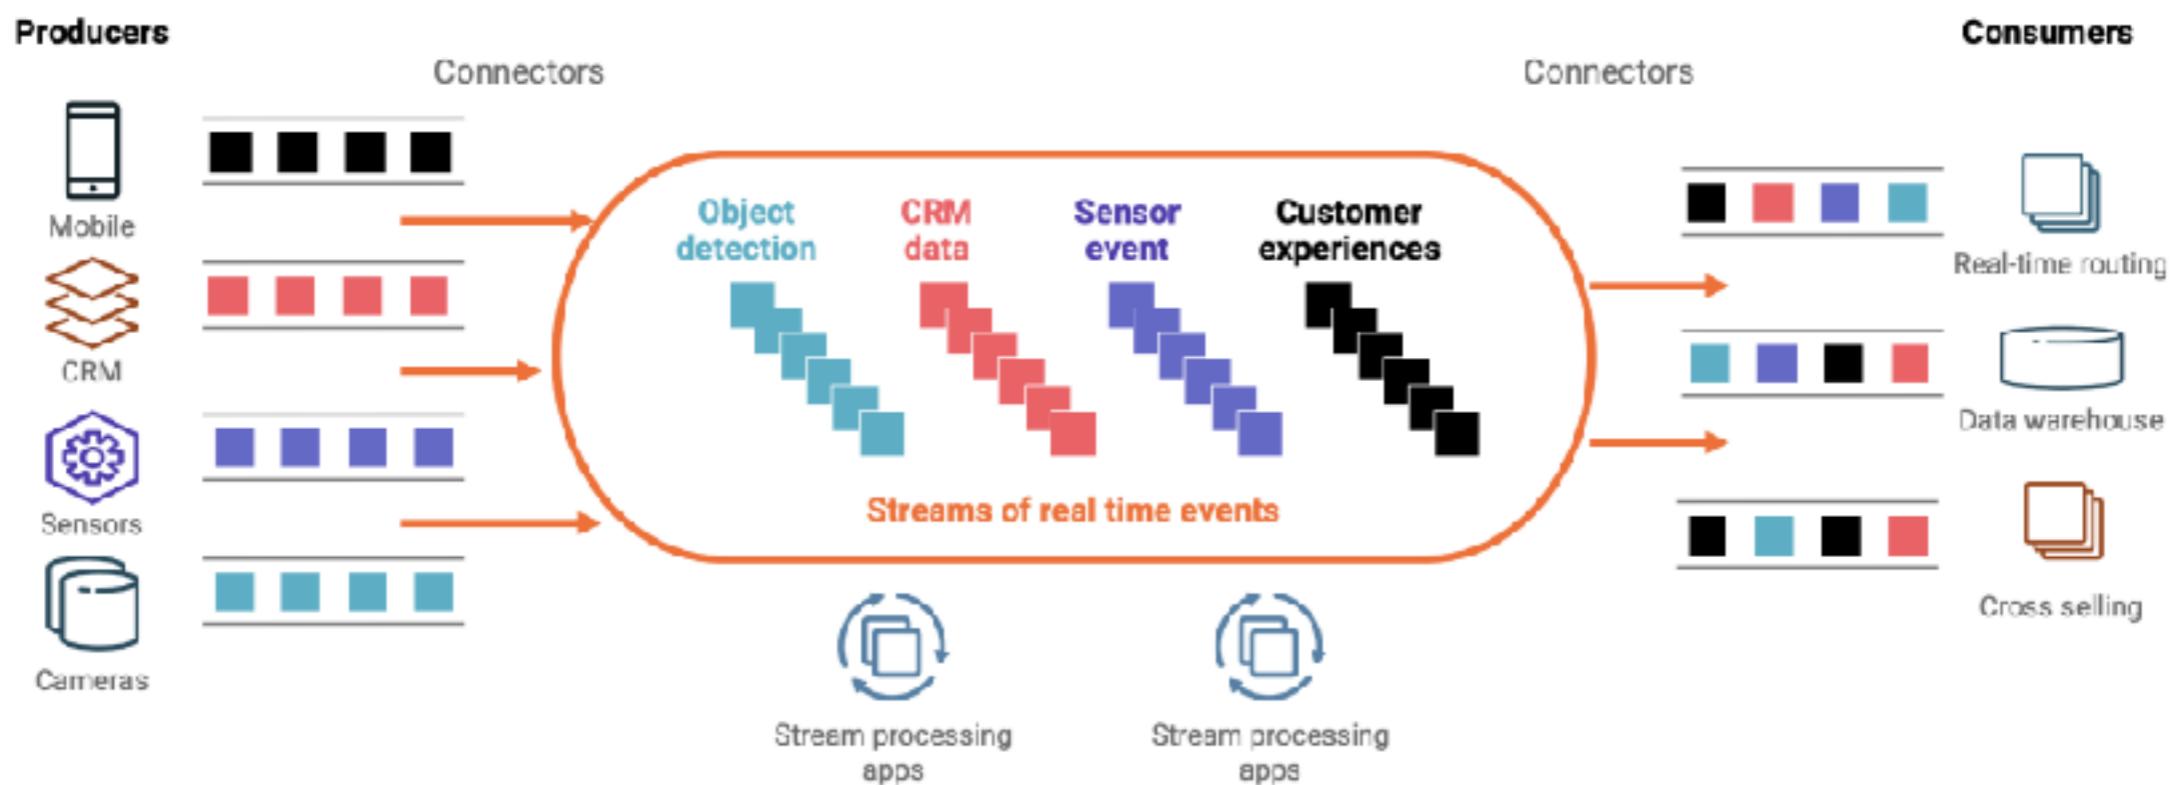
<https://developer.confluent.io/patterns/event-stream/event-stream/>



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Event Streaming



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Use cases

Messaging system

Activity tracking

Gather metric from different locations

Gather application logs

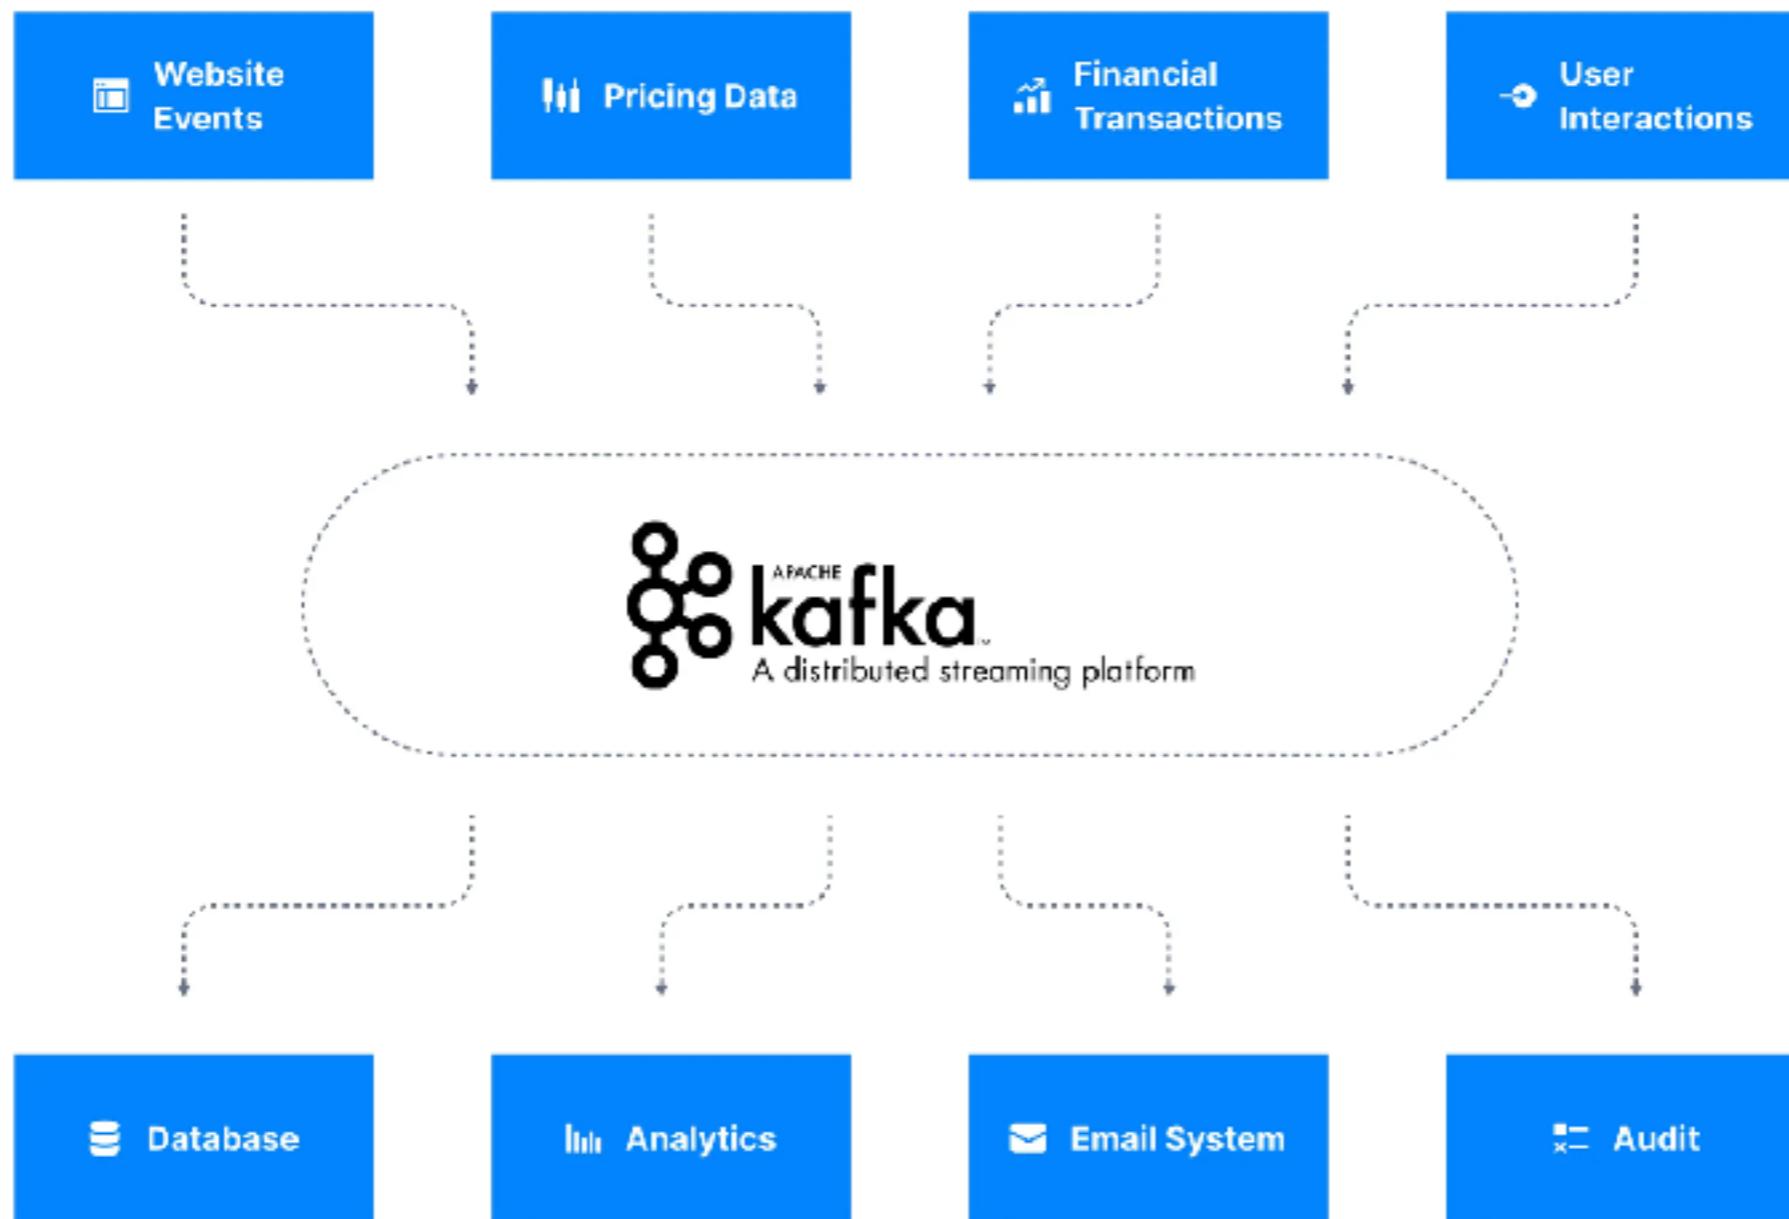
Stream processing

Decoupling of system dependencies

Integration with others system/technologies



# Use cases



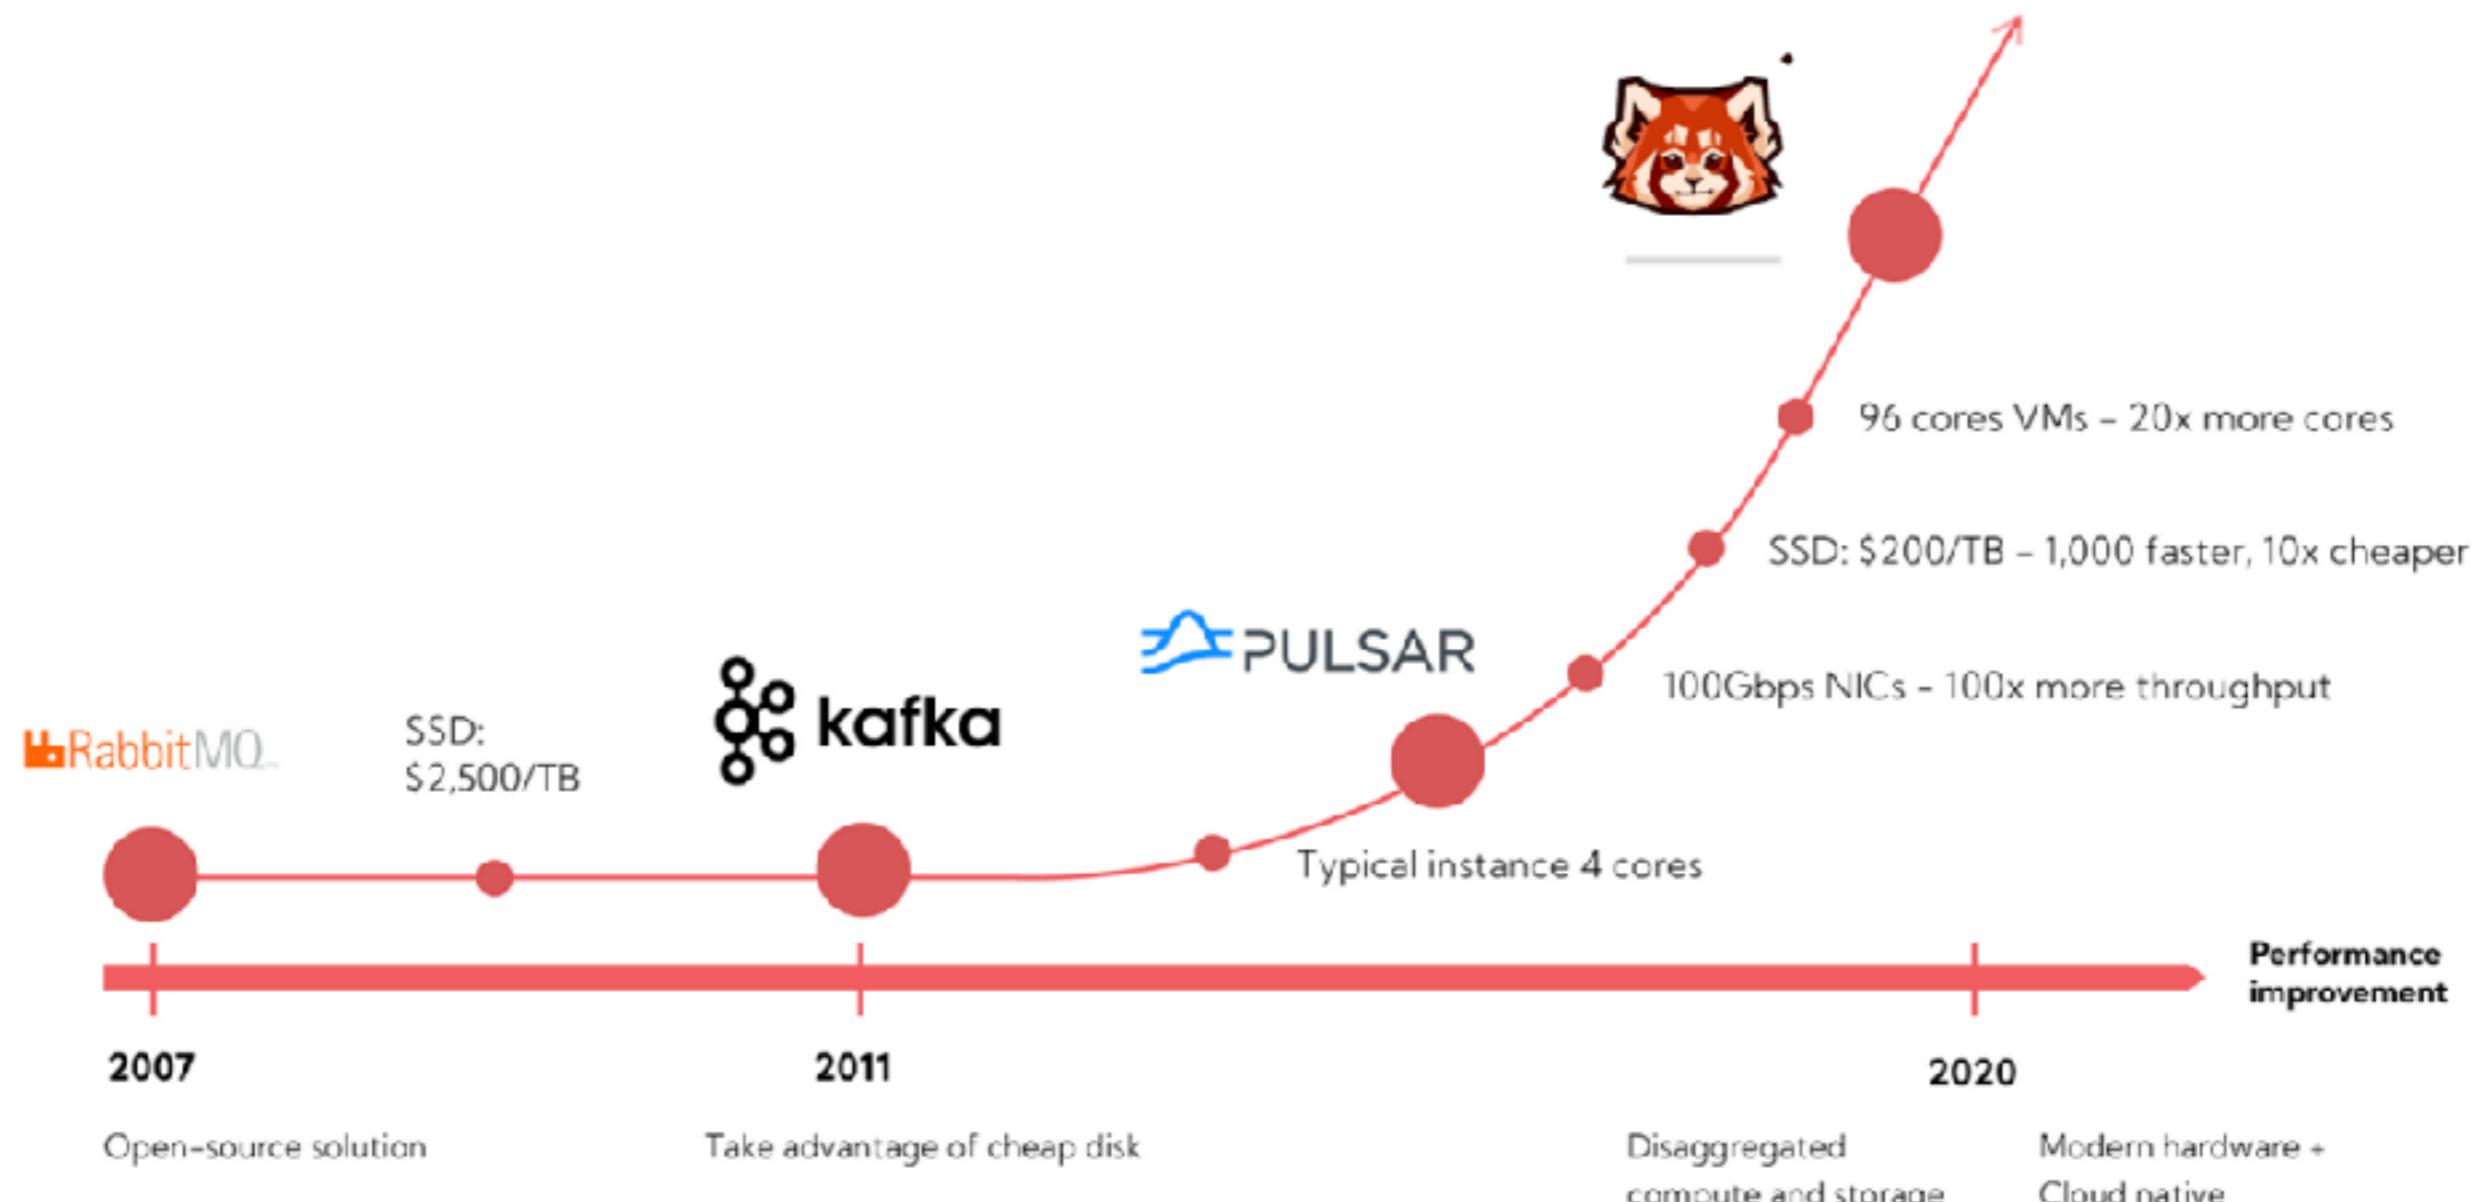
<https://learn.conduktor.io/kafka/what-is-apache-kafka/>



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Alternative for Kafka



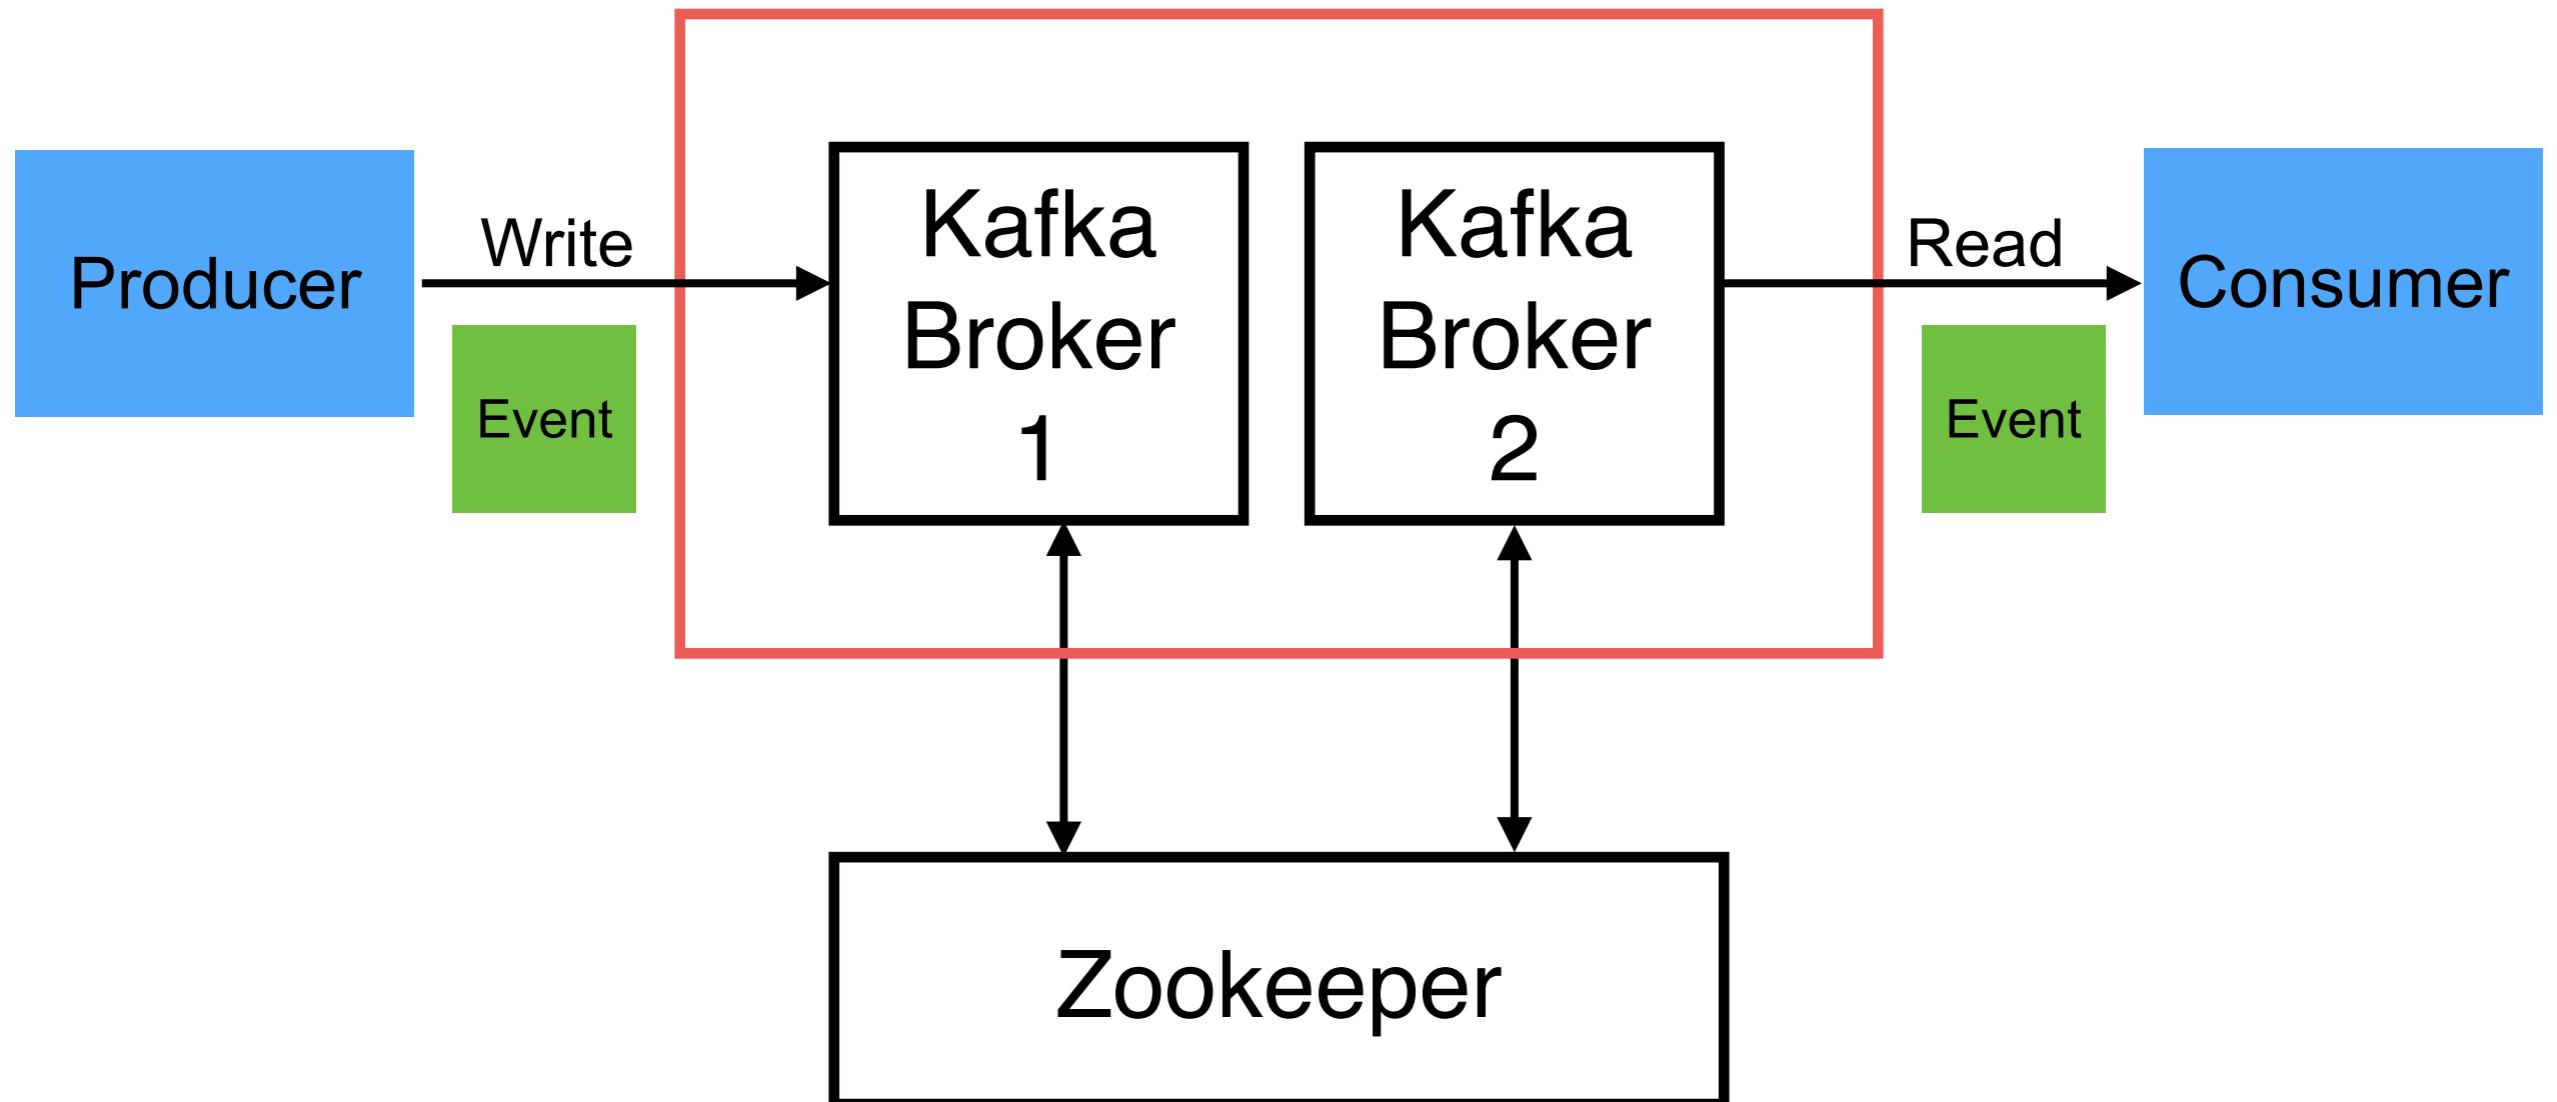
<https://www.redpanda.com/guides/kafka-alternatives-kafka-api>



# Architecture of Kafka

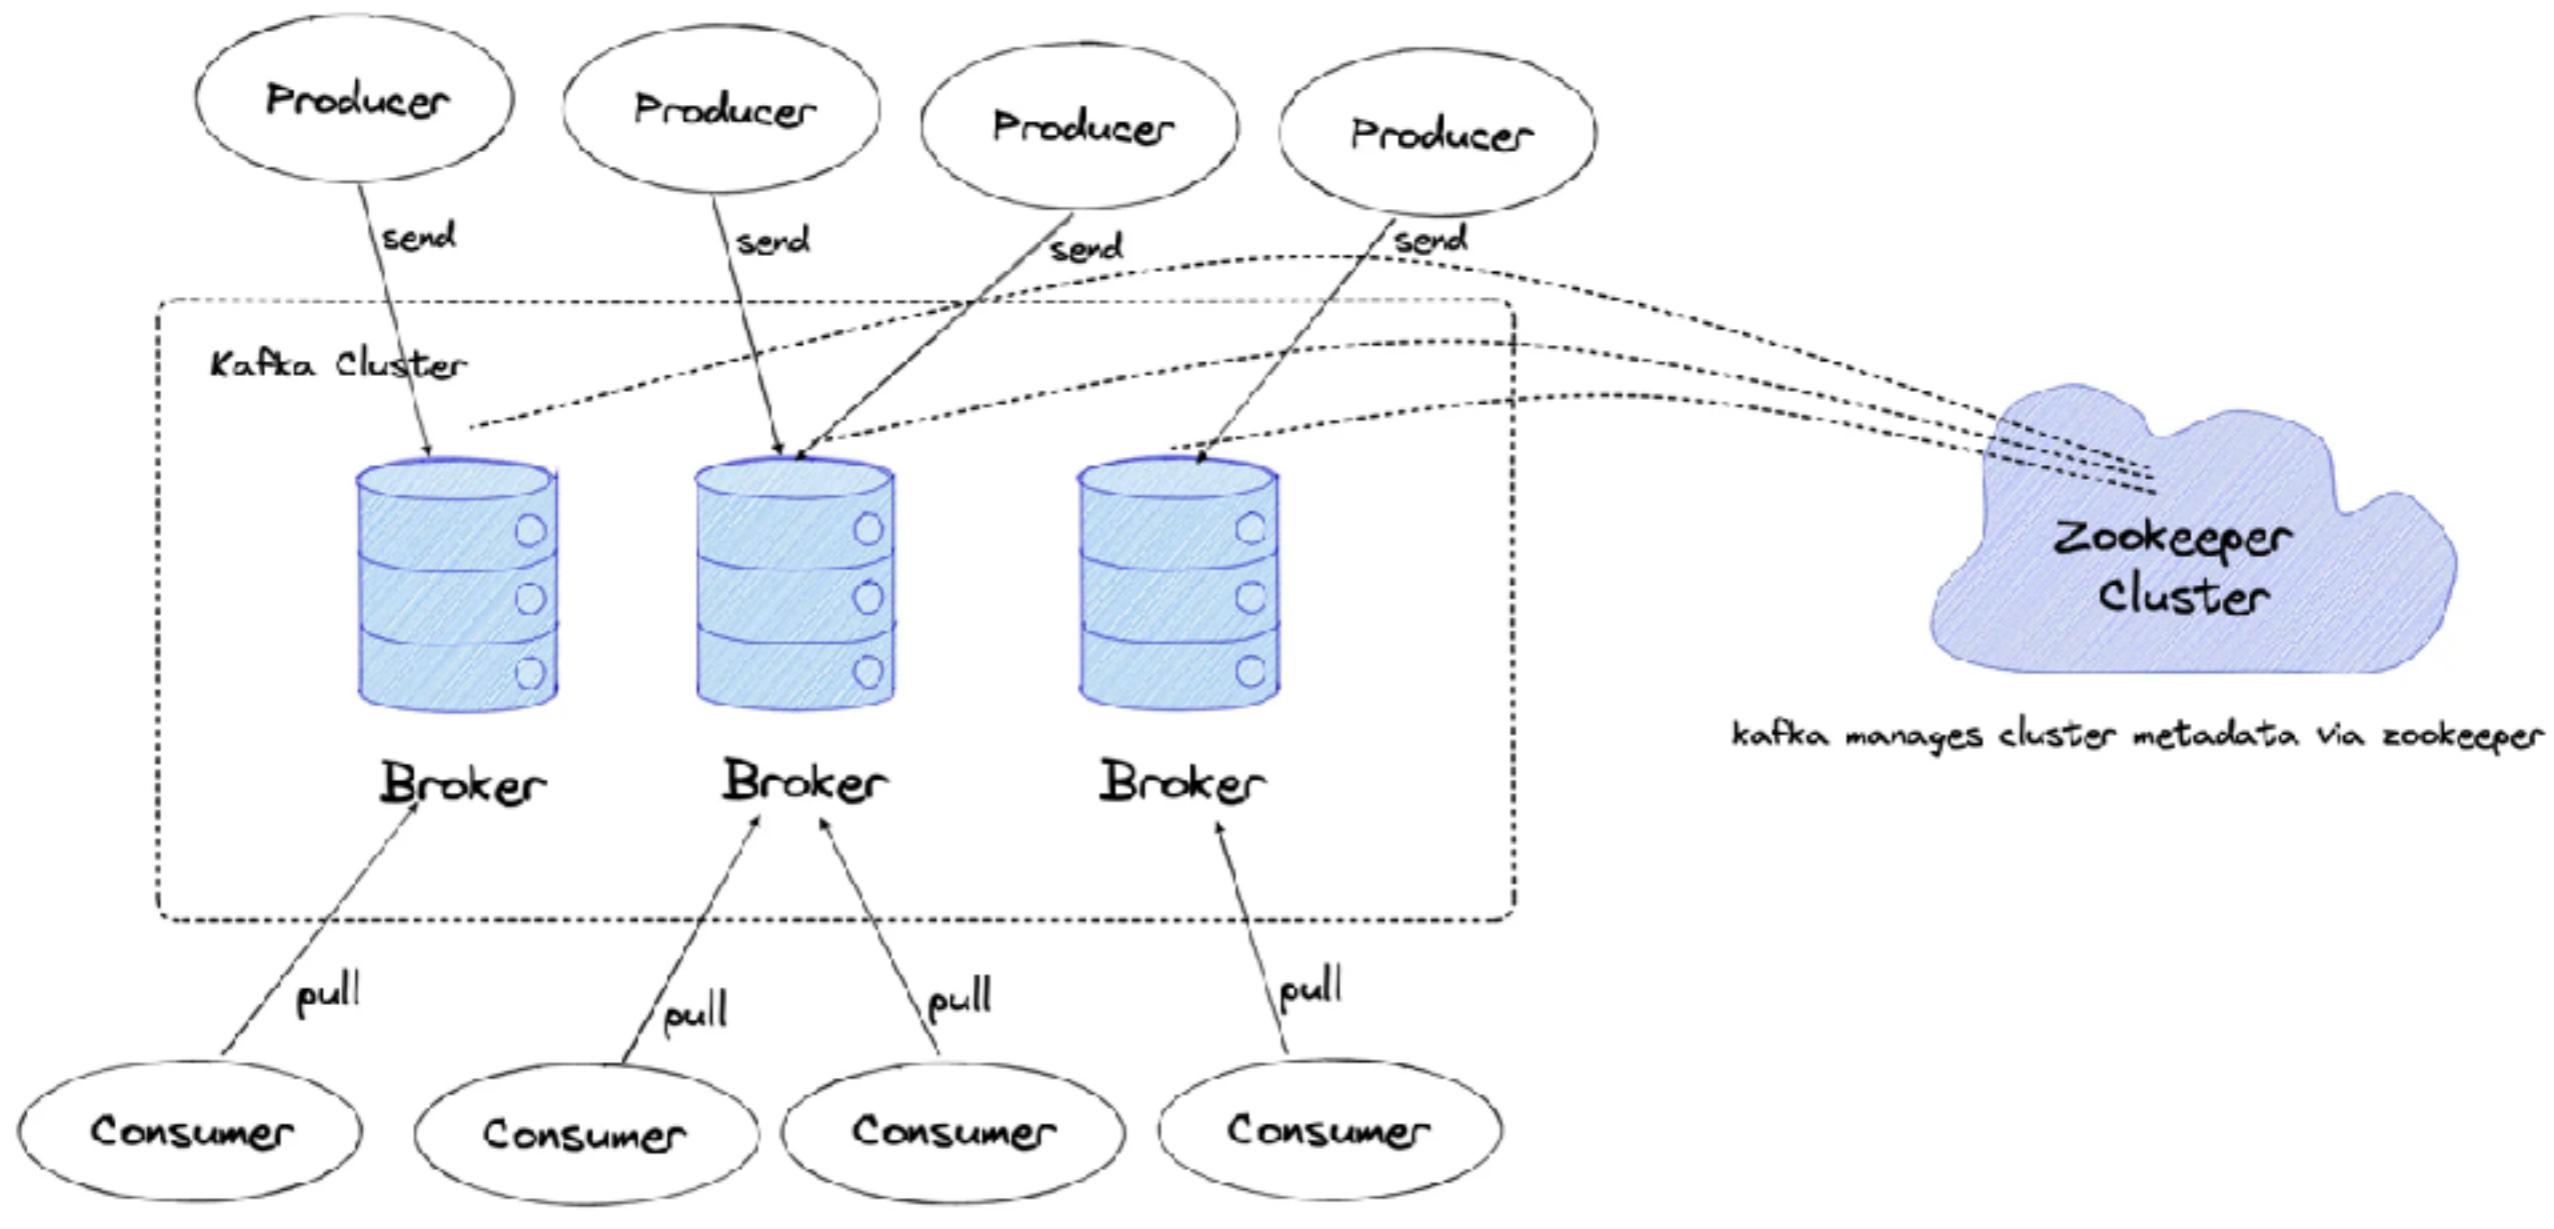


# Architecture (<4)



# Architecture (<4)

The producer sends the message to the broker



<https://api7.ai/blog/why-kafka-needs-an-api-gateway>



Kafka

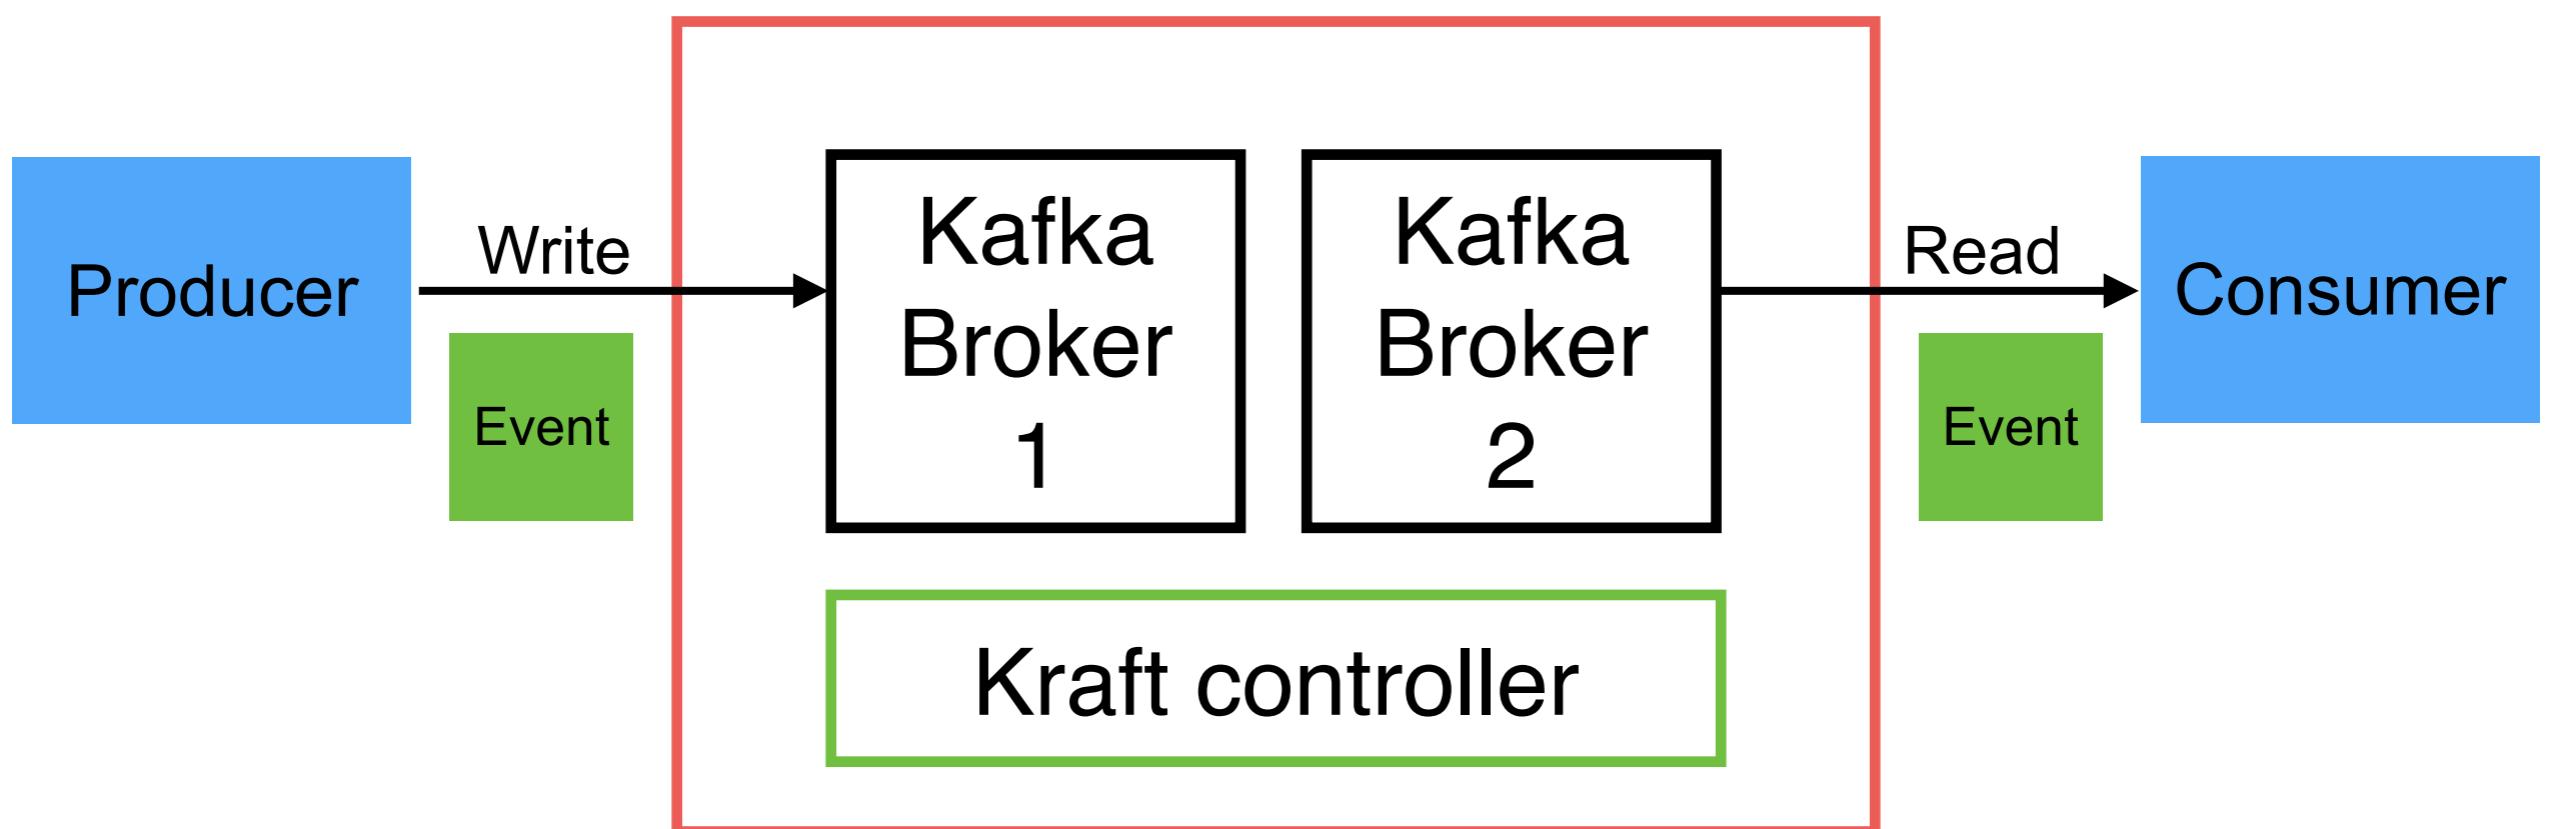
© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Architecture (4+)

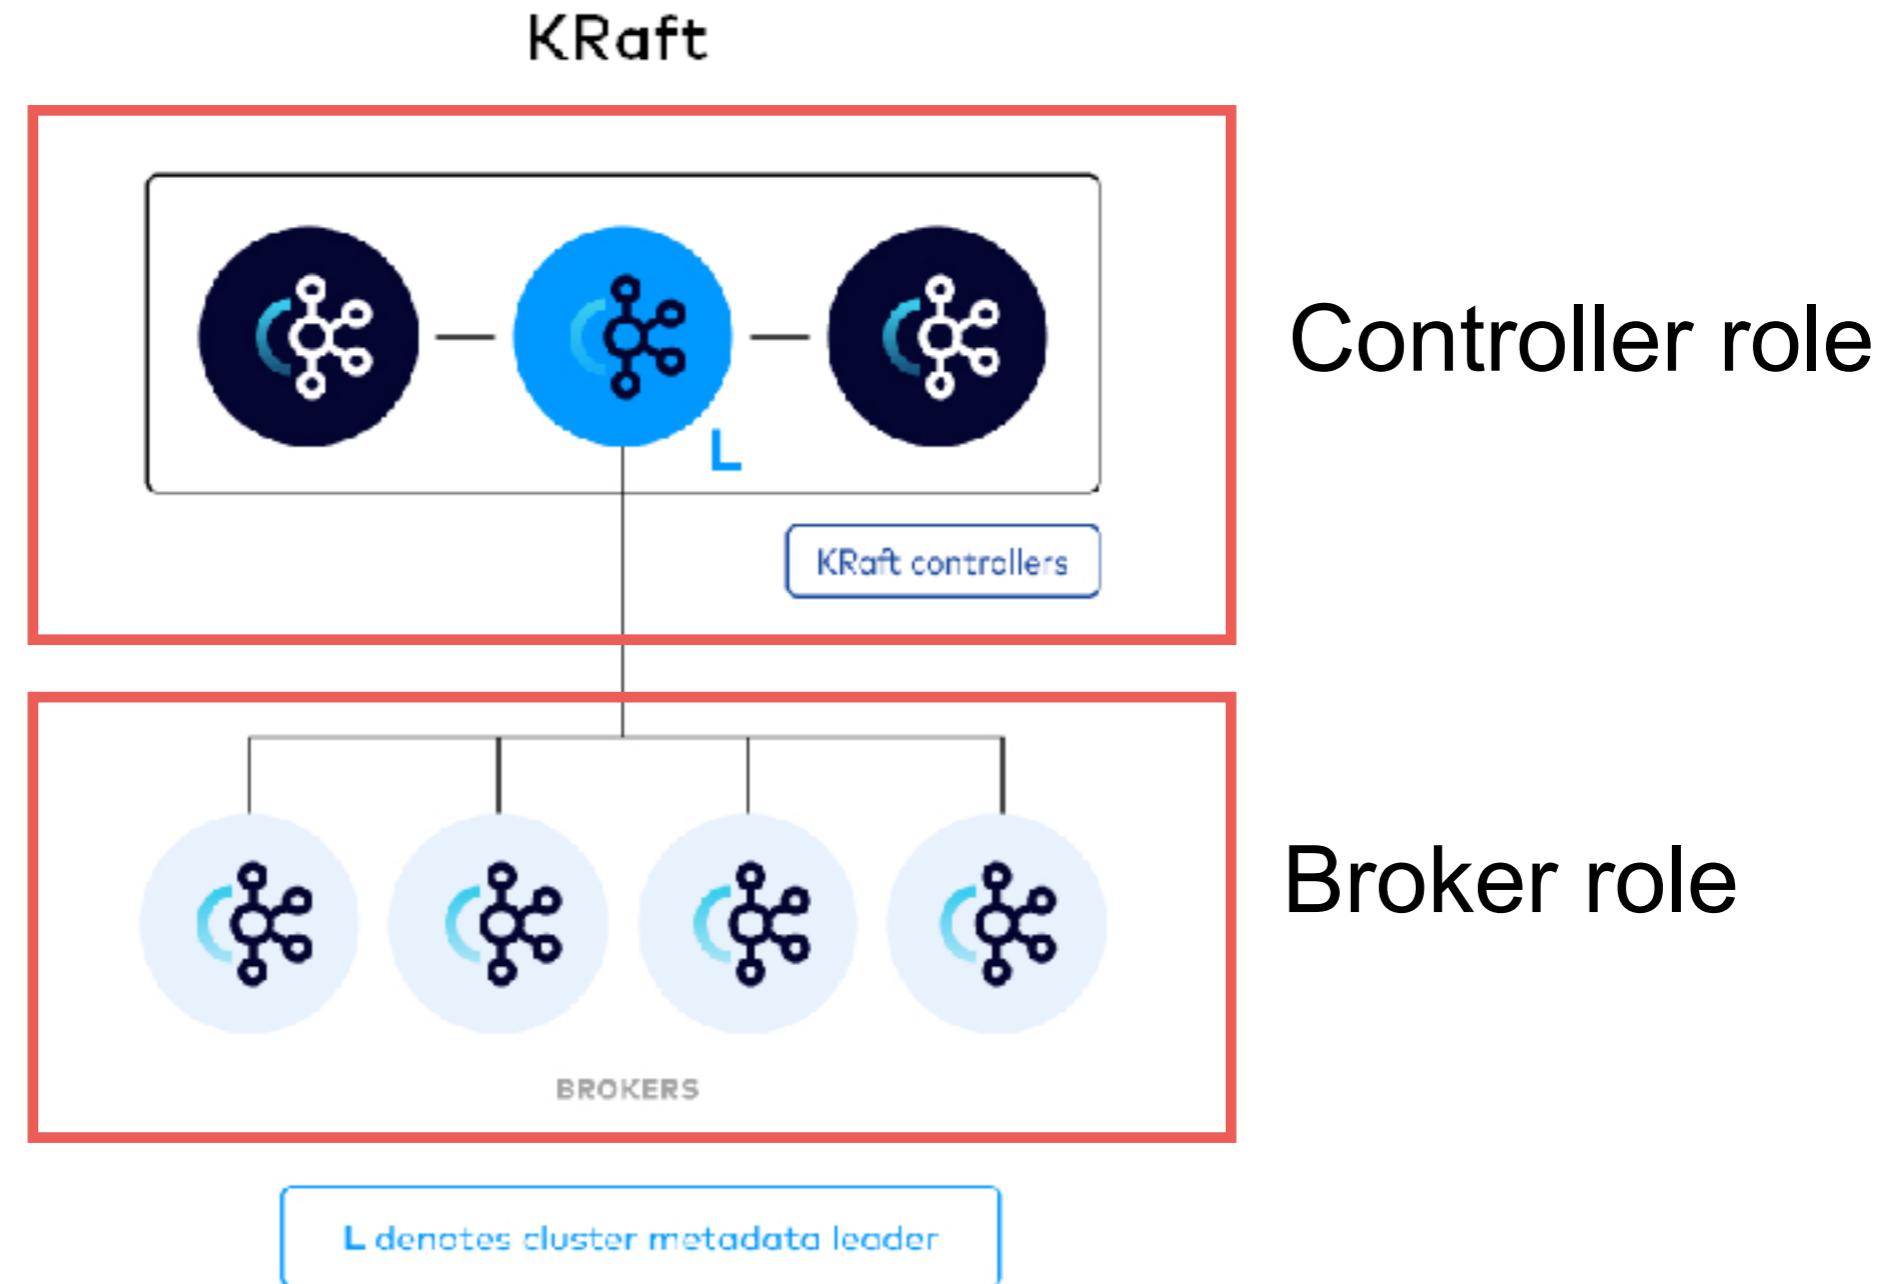
KRaft (Kafka Raft Metadata Mode)

Remove Zookeeper

Simplify deployment, management and scaling



# Kraft architecture



<https://docs.confluent.io/platform/current/kafka-metadata/kraft.html>

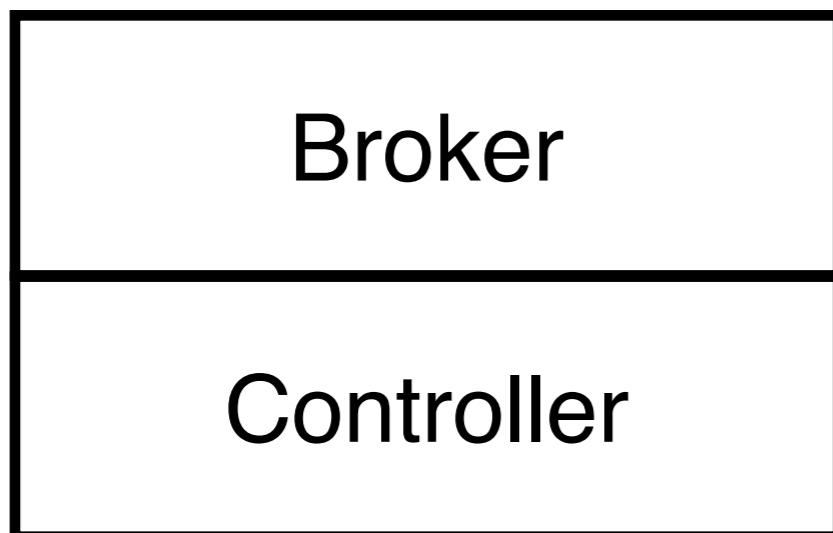


Kafka

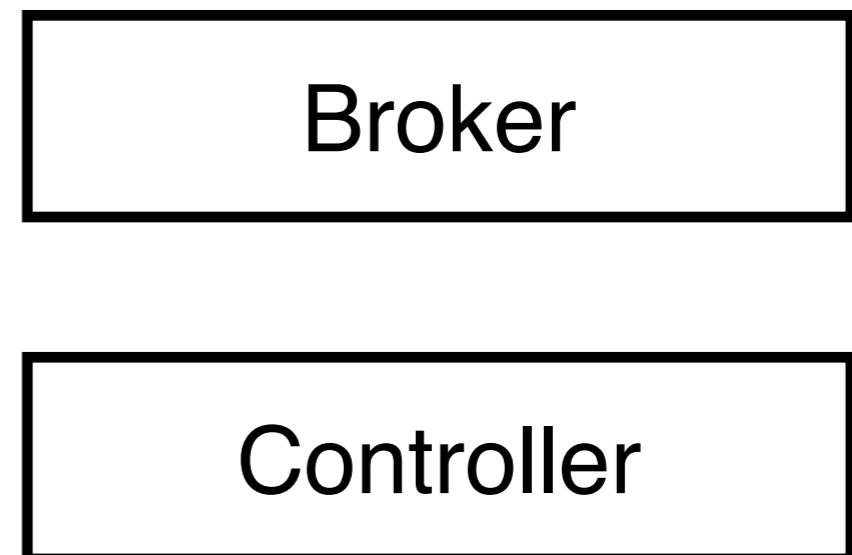
© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Kraft architecture

## Combined mode



## Isolated mode

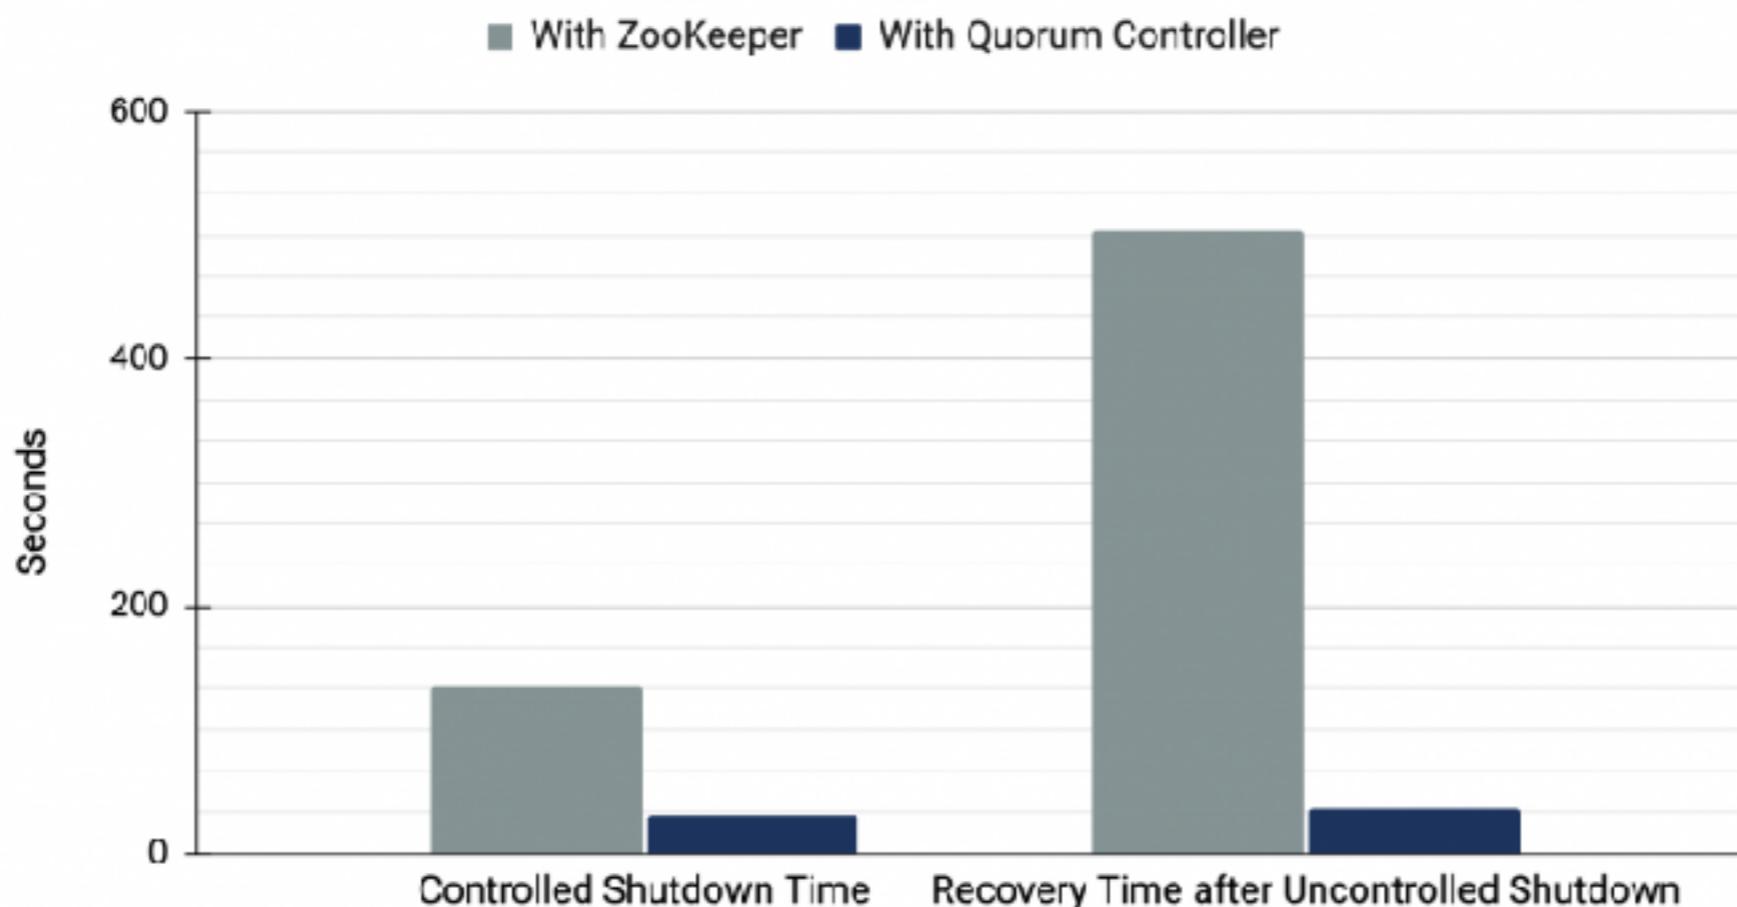


<https://docs.confluent.io/platform/current/kafka-metadata/kraft.html>



# Performance !!

Timed Shutdown Operations In Apache Kafka with 2 Million Partitions  
*Faster is better*



<https://docs.confluent.io/platform/current/kafka-metadata/kraft.html>

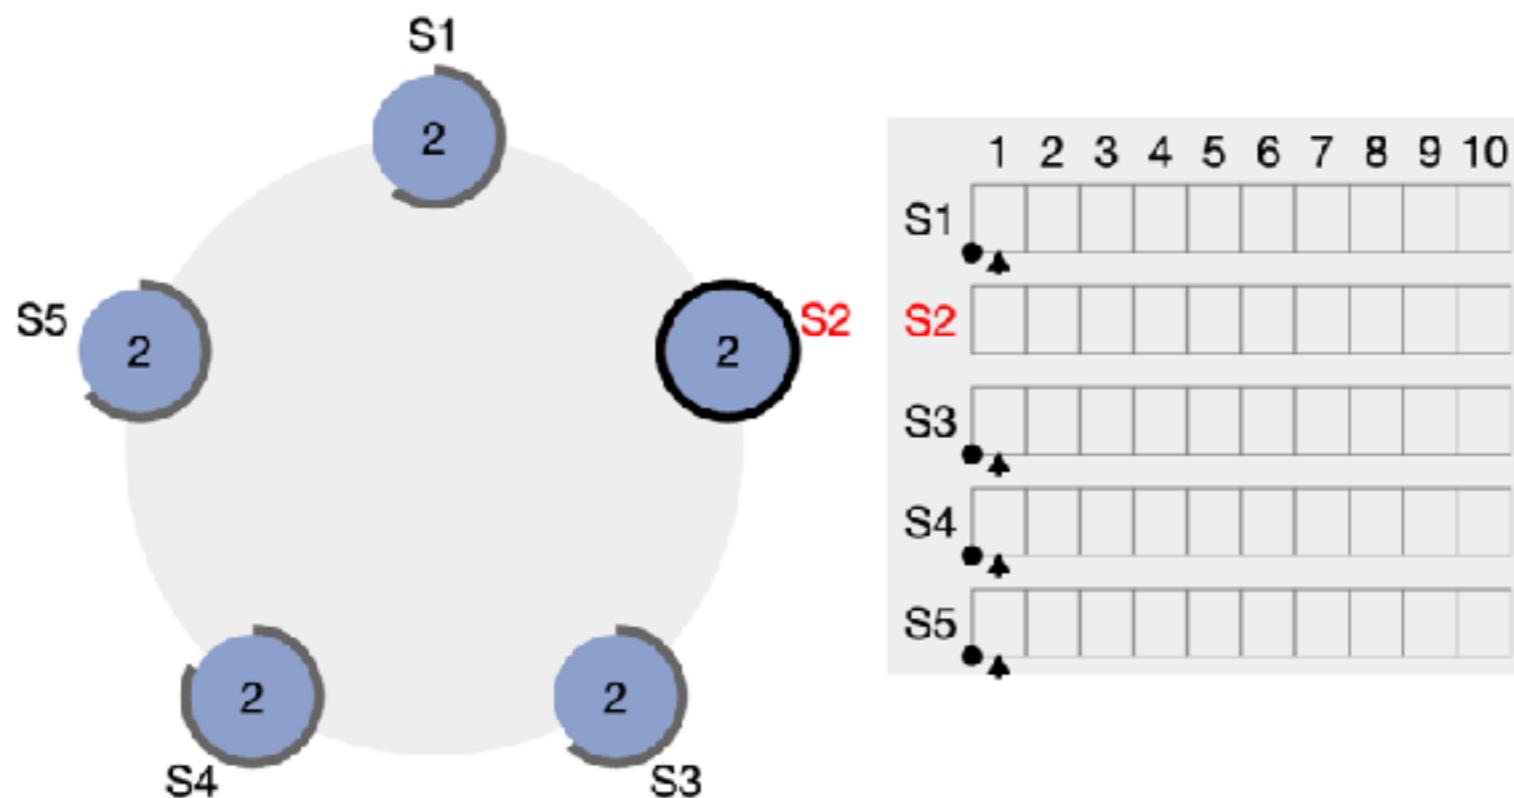


Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Raft Consensus algorithm

Manage state machine replication in distributed system  
Maintain consistent state of cluster



<https://raft.github.io/>



# Key point of Raft

Minimum server = 3  
Odd number of servers



# Kafka Fundamentals



# Kafka Fundamentals

Event/record/message

Producers

Broker

Topic

Partition and offset

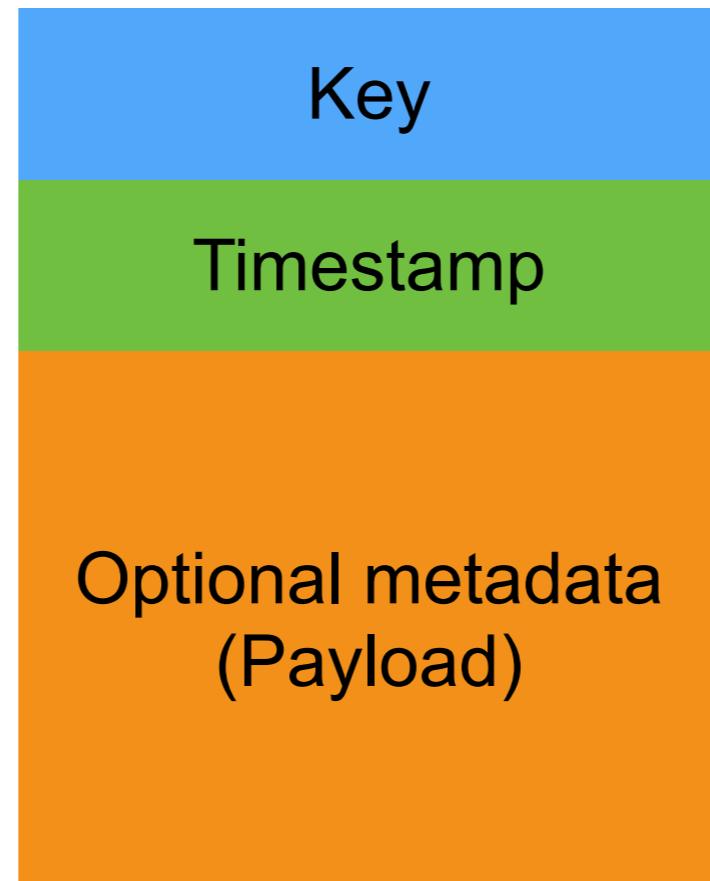
Consumer

Consumer groups

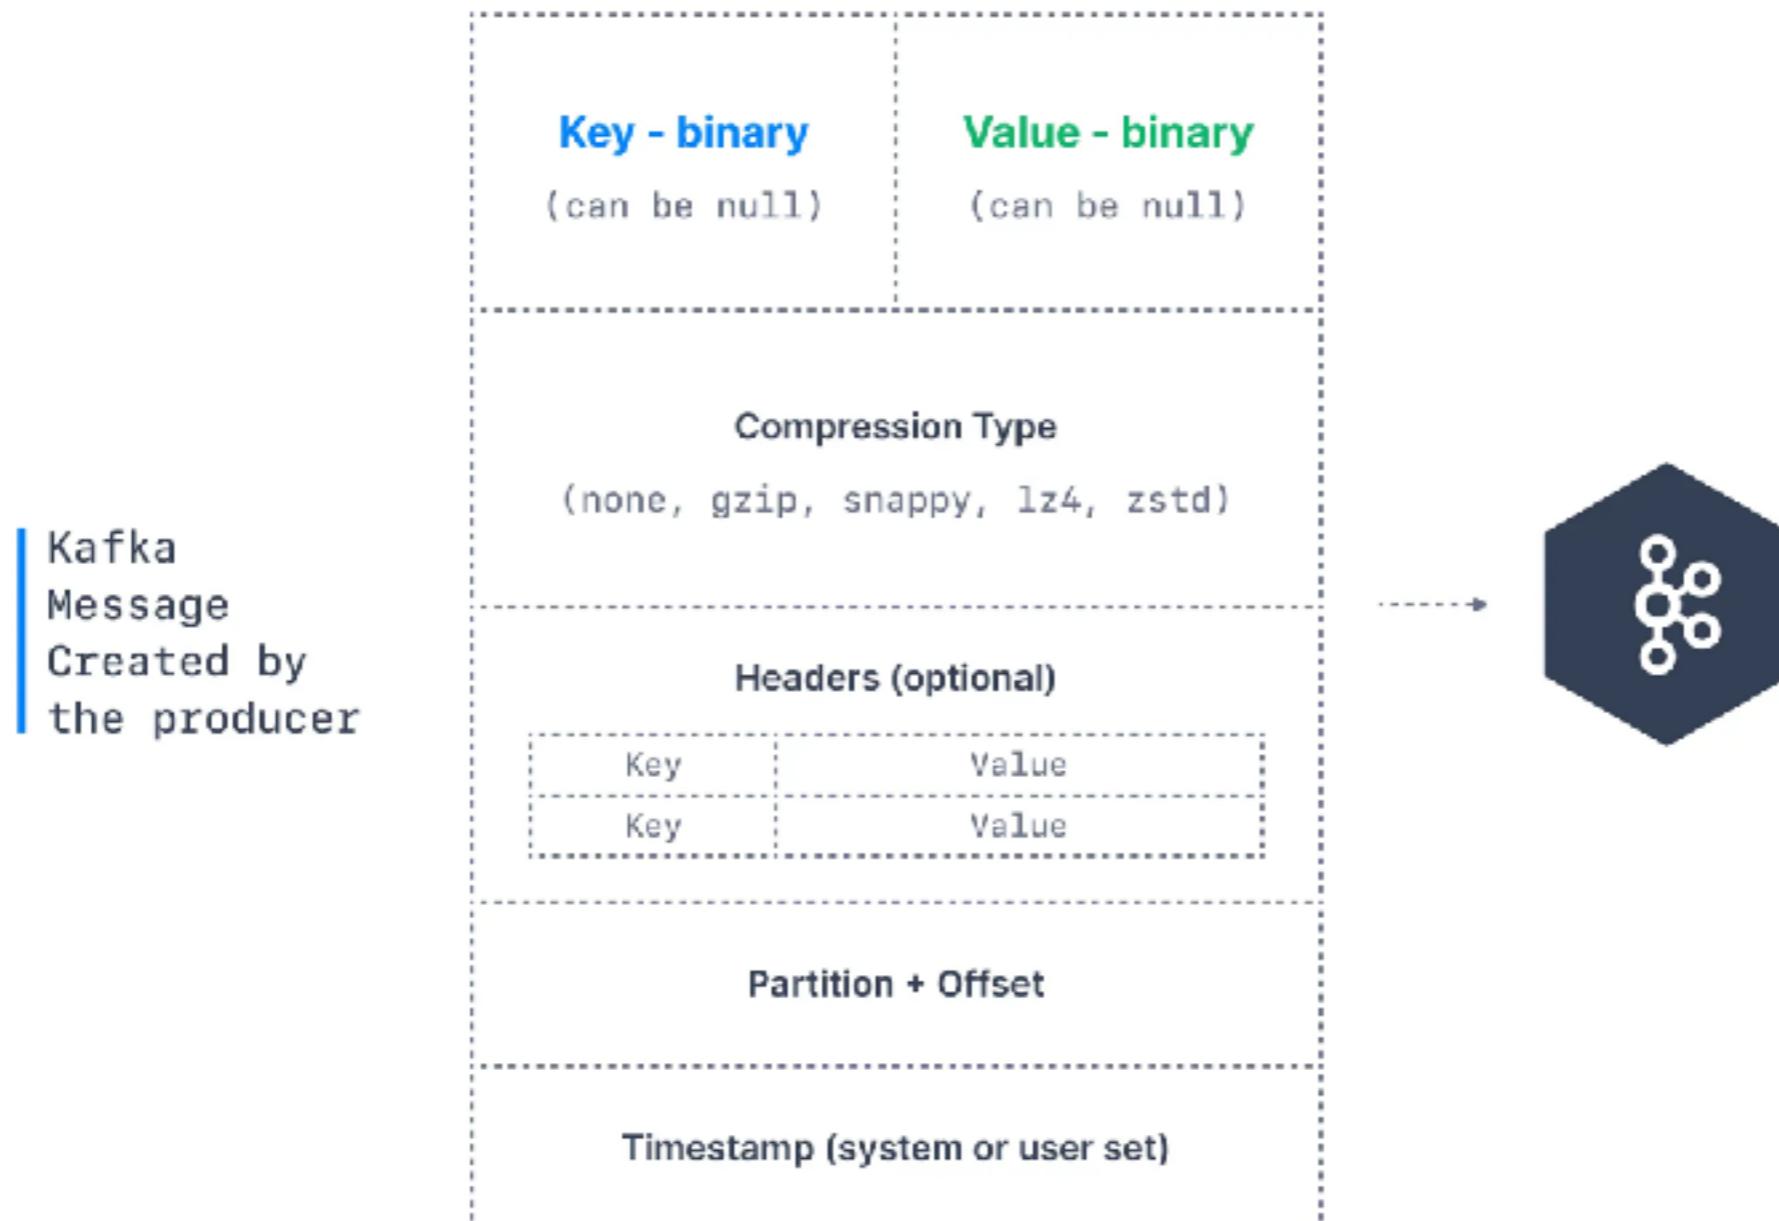
Controller



# Event



# Event Anatomy



# Topics, partitions and offsets

## Topics

Stream of data

Similar to a table in database

You can have many topics as you want

A topic is identified by **name**

Topic



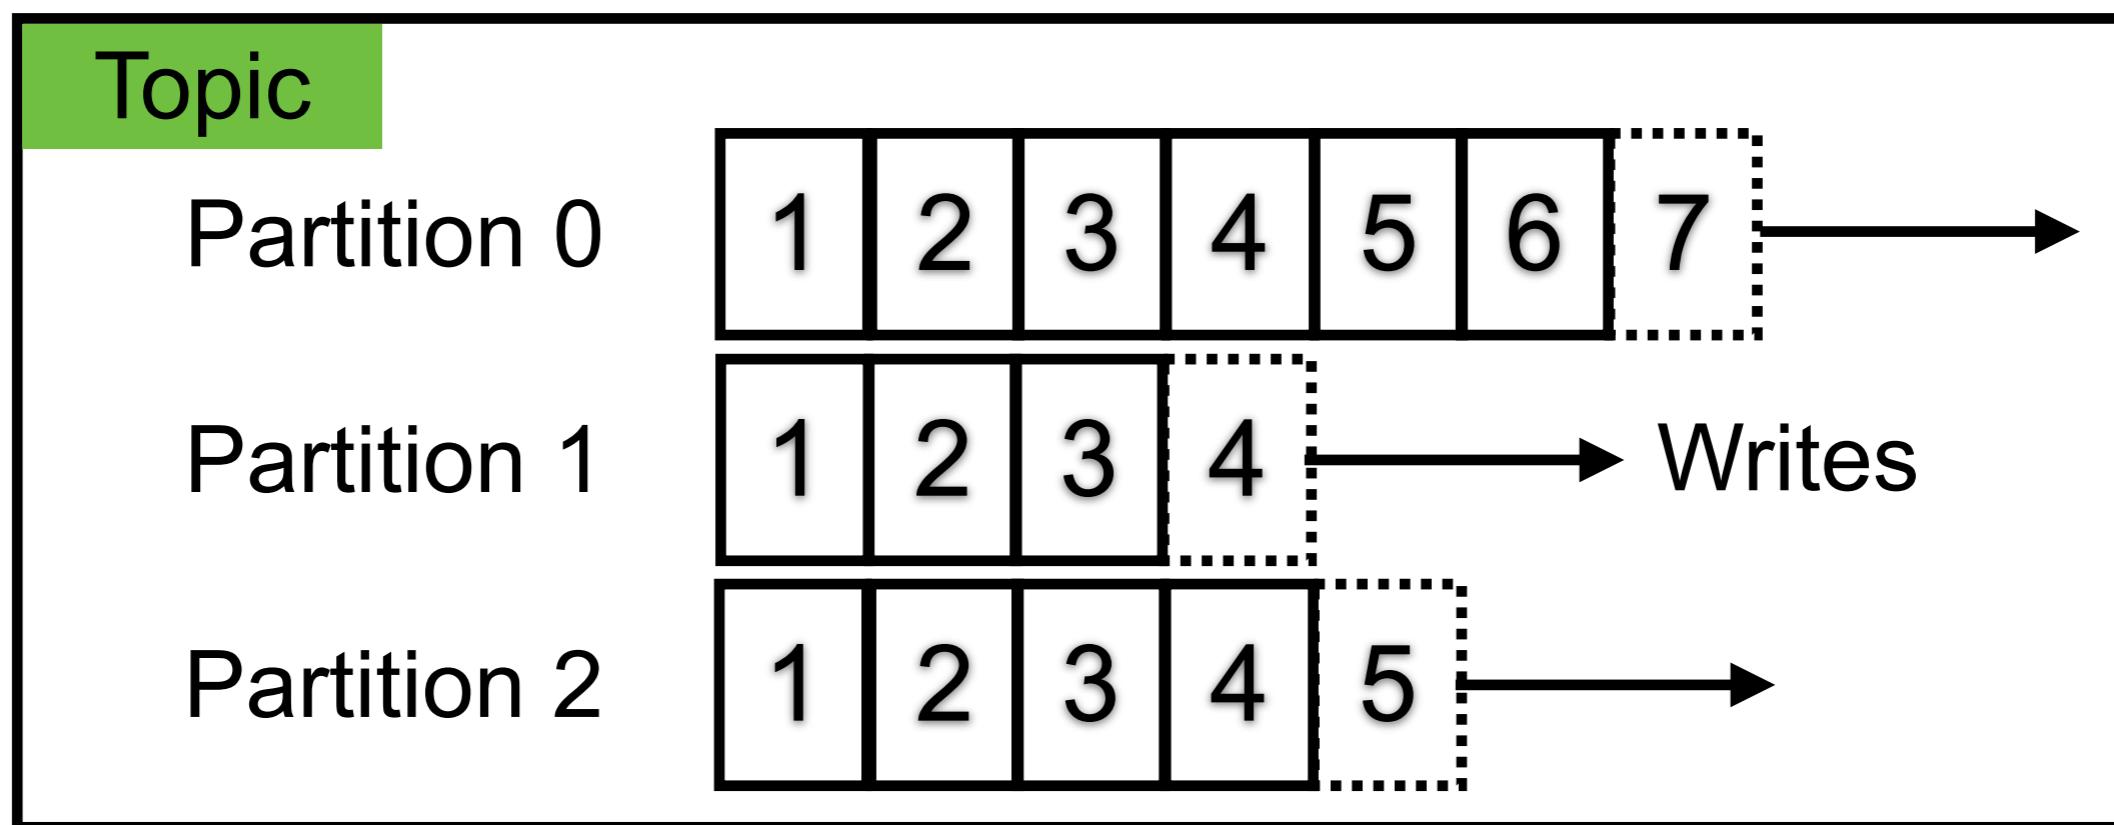
# Topics, partitions and offsets

## Partitions

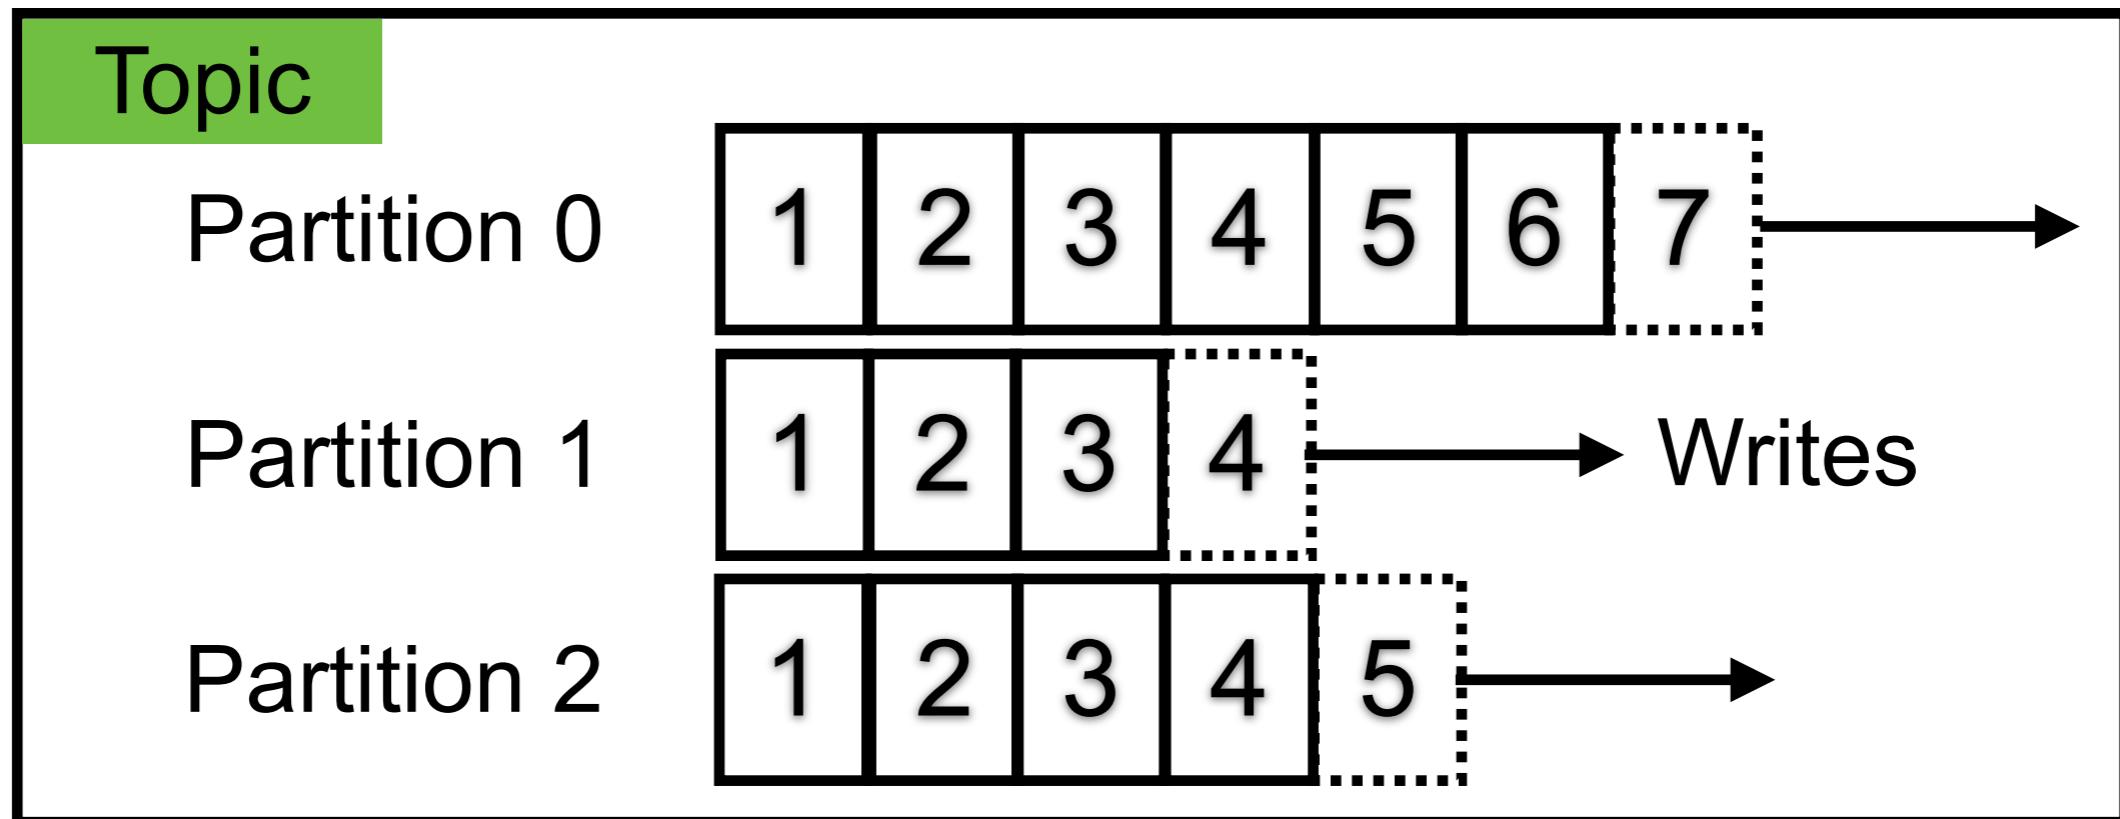
Topics are split in partitions

Each partition is **ordered**

Each message in partition get incremental id (**offset**)



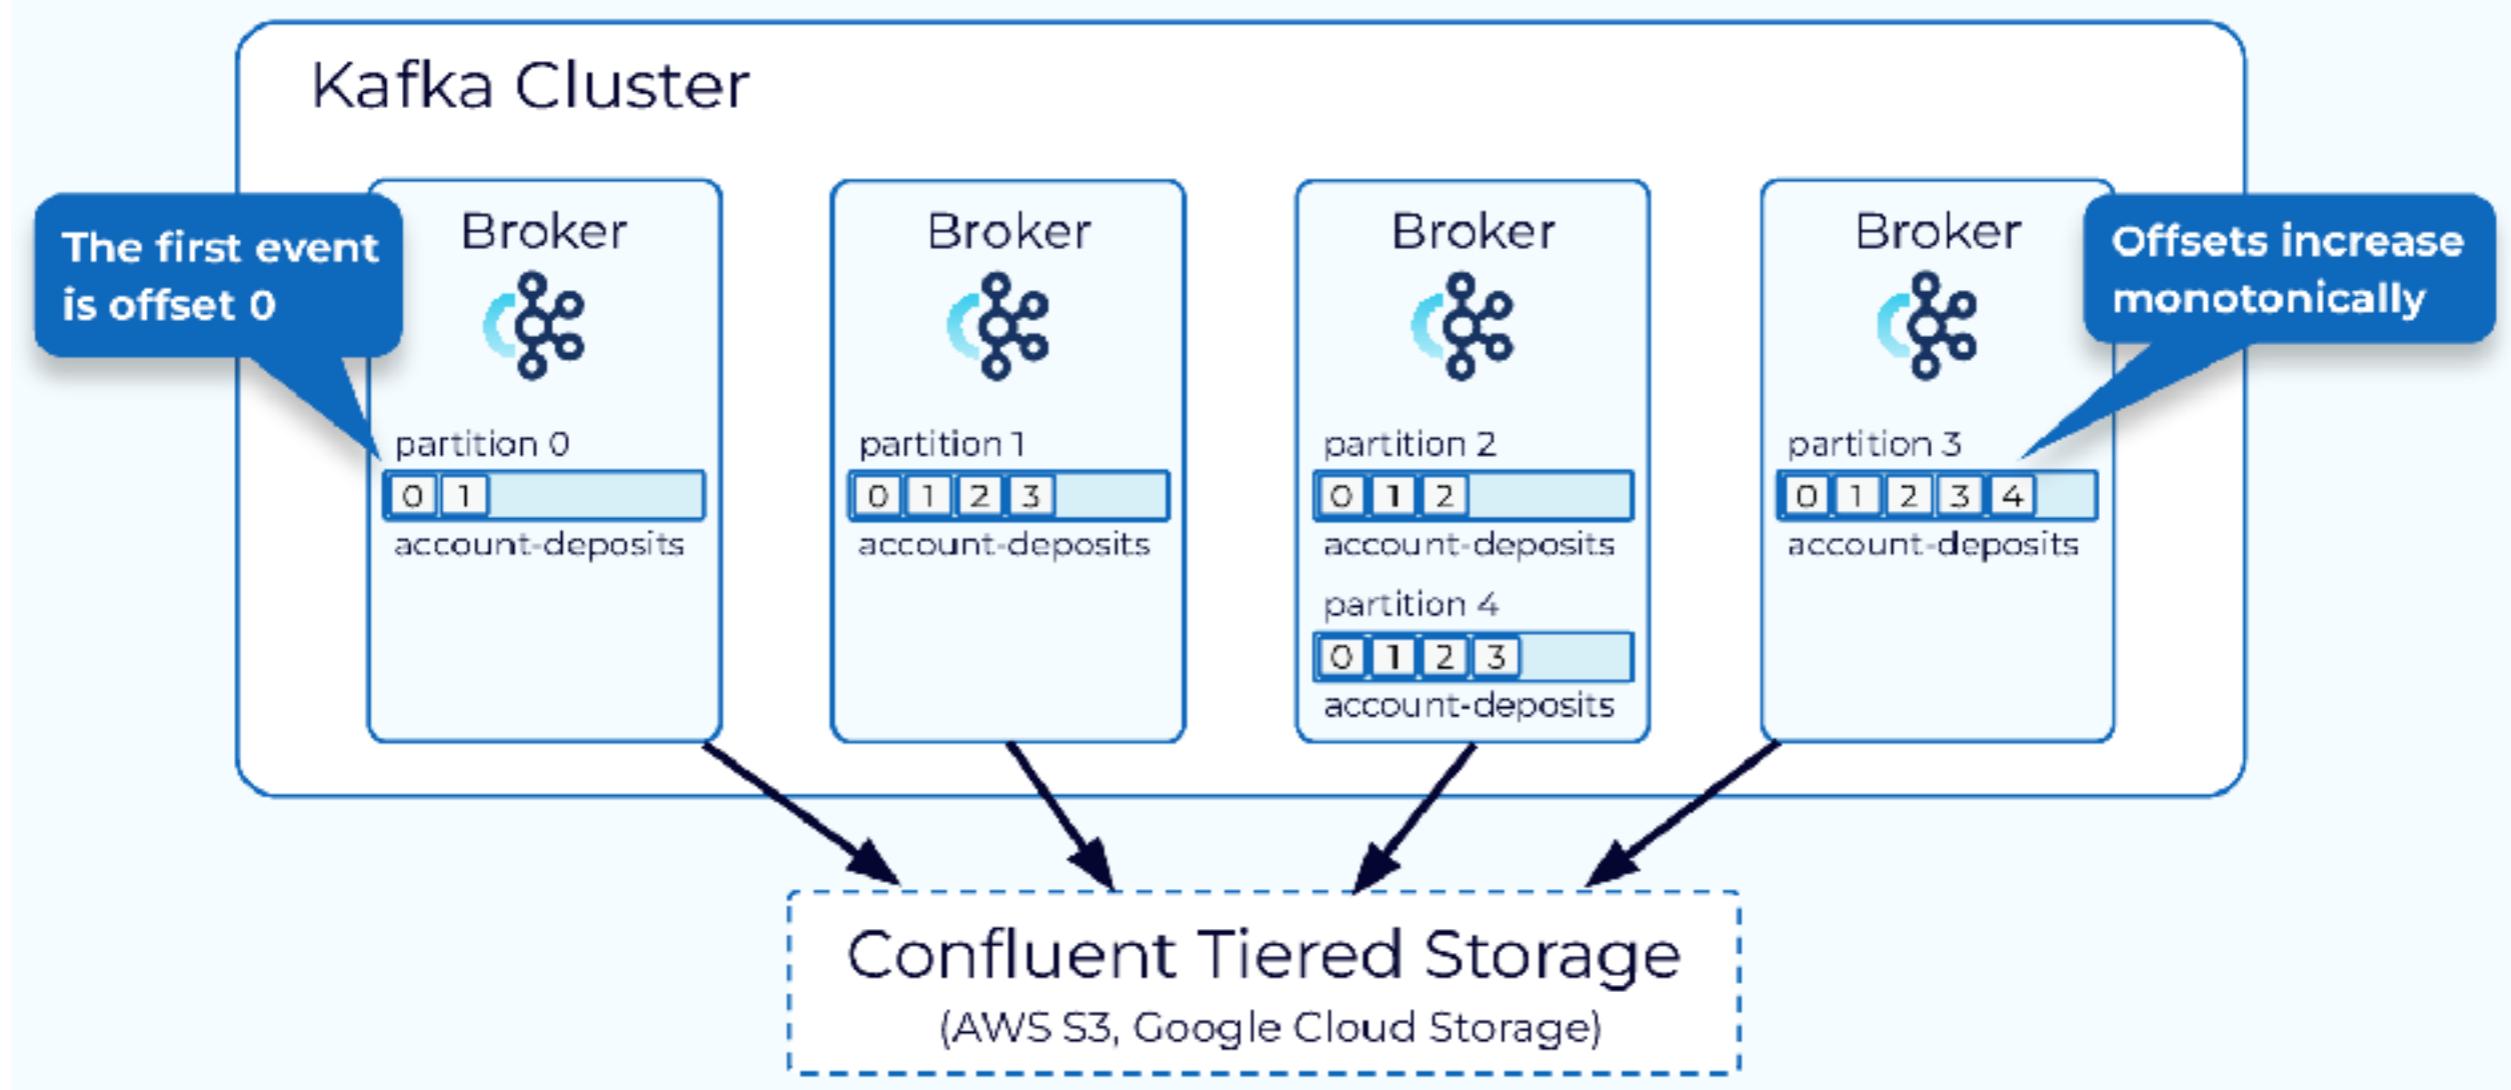
# Topics, partitions and offsets



Order is guaranteed only within partition  
Data is keep in limited time (**Default = 1 week**)  
Data is **immutable** (can't be changed)  
Data is assigned **randomly** to a partition



# Topics, partitions and offsets



<https://developer.confluent.io/courses/architecture/get-started/>



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# More partitions ?

Higher is your data throughput

More open files

Longer downtime

More RAM is consumed by clients

4000  
per broker

200,000  
per cluster

50 brokers  
per cluster

<https://api7.ai/blog/why-kafka-needs-an-api-gateway>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Brokers and topics

## Brokers

Kafka cluster is composed of multiple brokers (servers)

Each broker is identified by ID (integer)

Each broker contains certain topic partitions

After connecting any broker (**bootstrap broker**),  
you will connected to the entire cluster

Broker 1

Broker 2

Broker 3



# Kafka broker discovery

Every Kafka broker is called “bootstrap server”

You only need to connect to one broker, and you will connected to the entire cluster

Broker 1  
(bootstrap)

Broker 2  
(bootstrap)

**Kafka cluster**

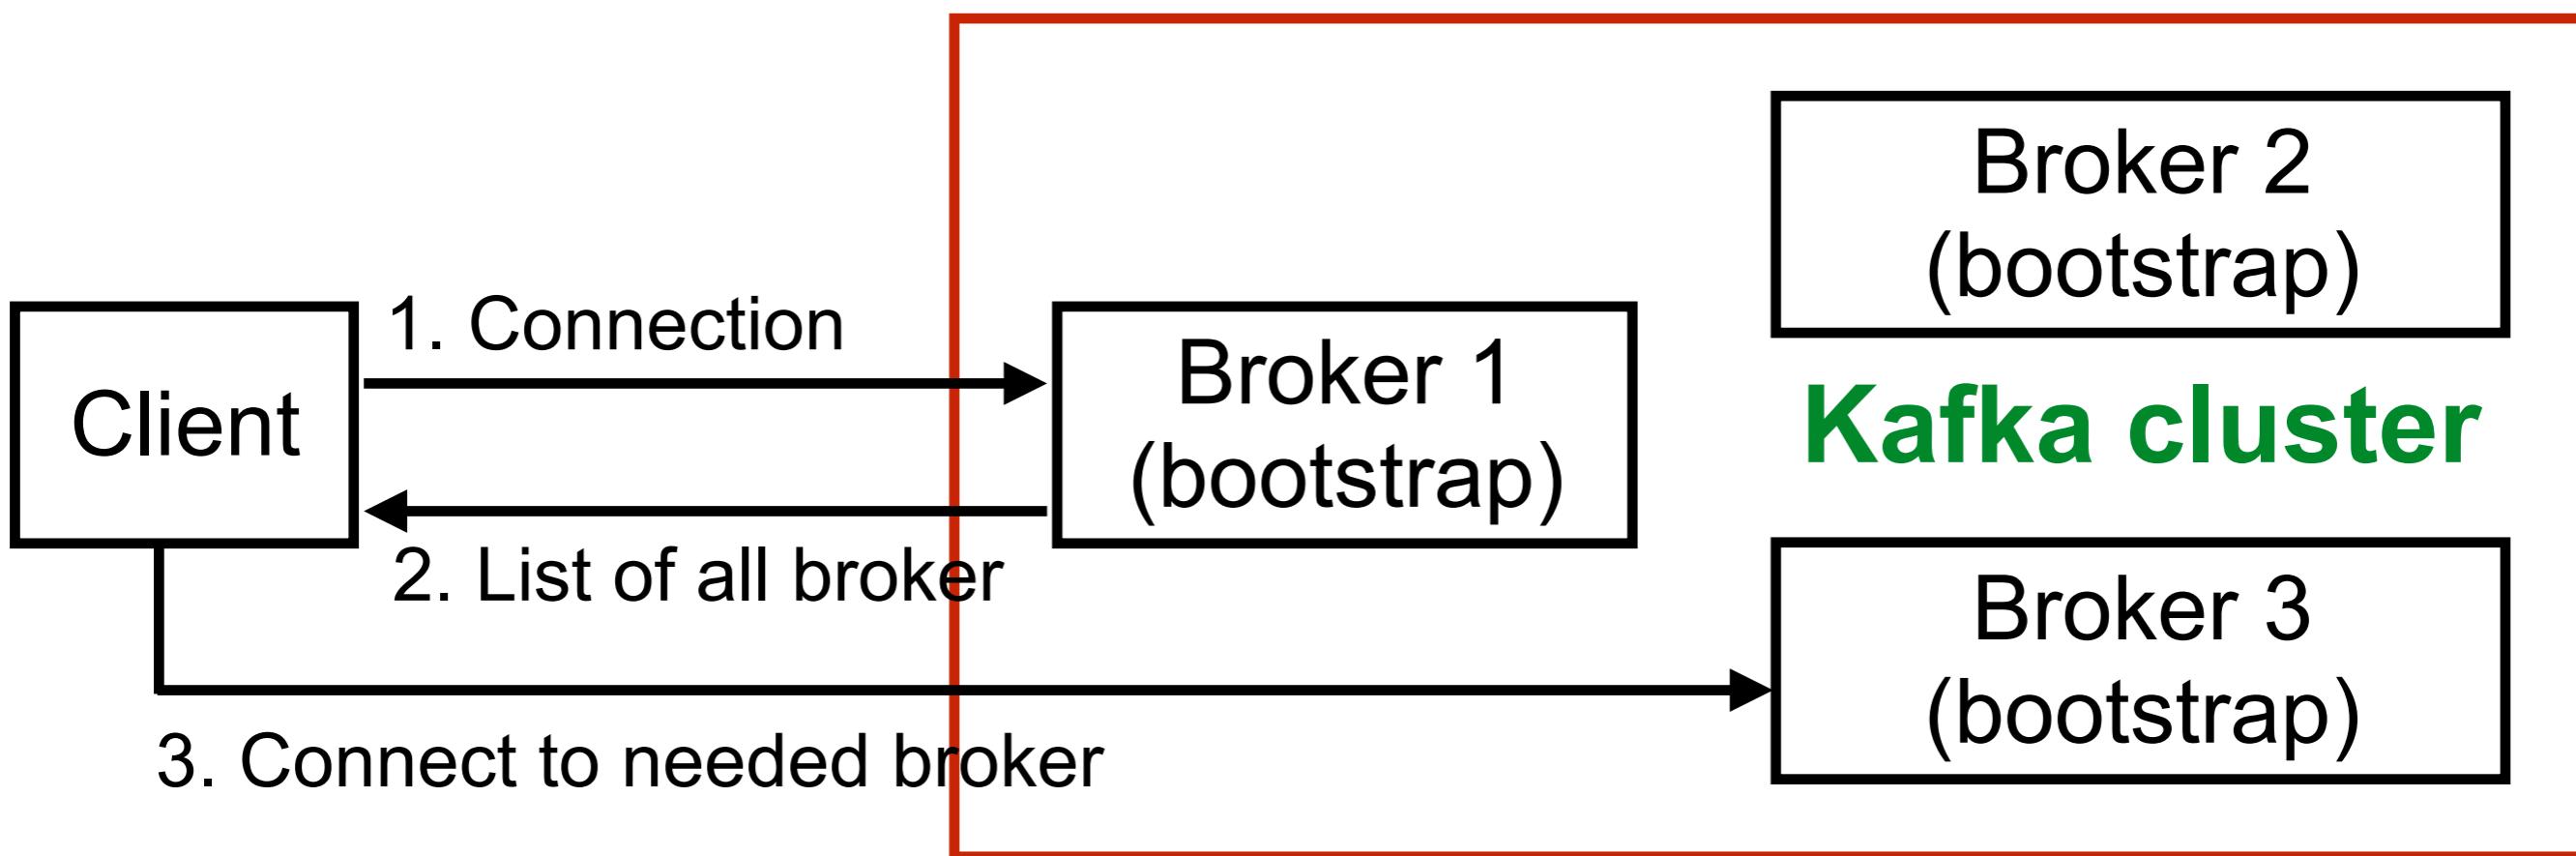
Broker 3  
(bootstrap)

Broker 4  
(bootstrap)



# Kafka broker discovery

Each broker knows about all brokers, topics and partitions (**metadata**)



# Apache Zookeeper

Kafka can not work without Zookeeper !!



APACHE  
**ZooKeeper**<sup>TM</sup>

<https://zookeeper.apache.org/>



# Apache Zookeeper

Zookeeper **manages** brokers

Zookeeper help in performing **leading election** for partitions

Zookeeper sends **notifications** to Kafka in case of changes (new topic, broker die)

Zookeeper by design operates with a odd number of servers (1, 3, 5, 7)



# Data in Zookeeper

Kafka's data operations	Format
Broker metadata	ID, hostname
Topic metadata	Topic name, partition count, replica count
Partition assignment	Leader partition
Consumer group metadata	Consumer group name
Cluster metadata	Active broker, list of topics, partitions
Leader election	Select leader of partition

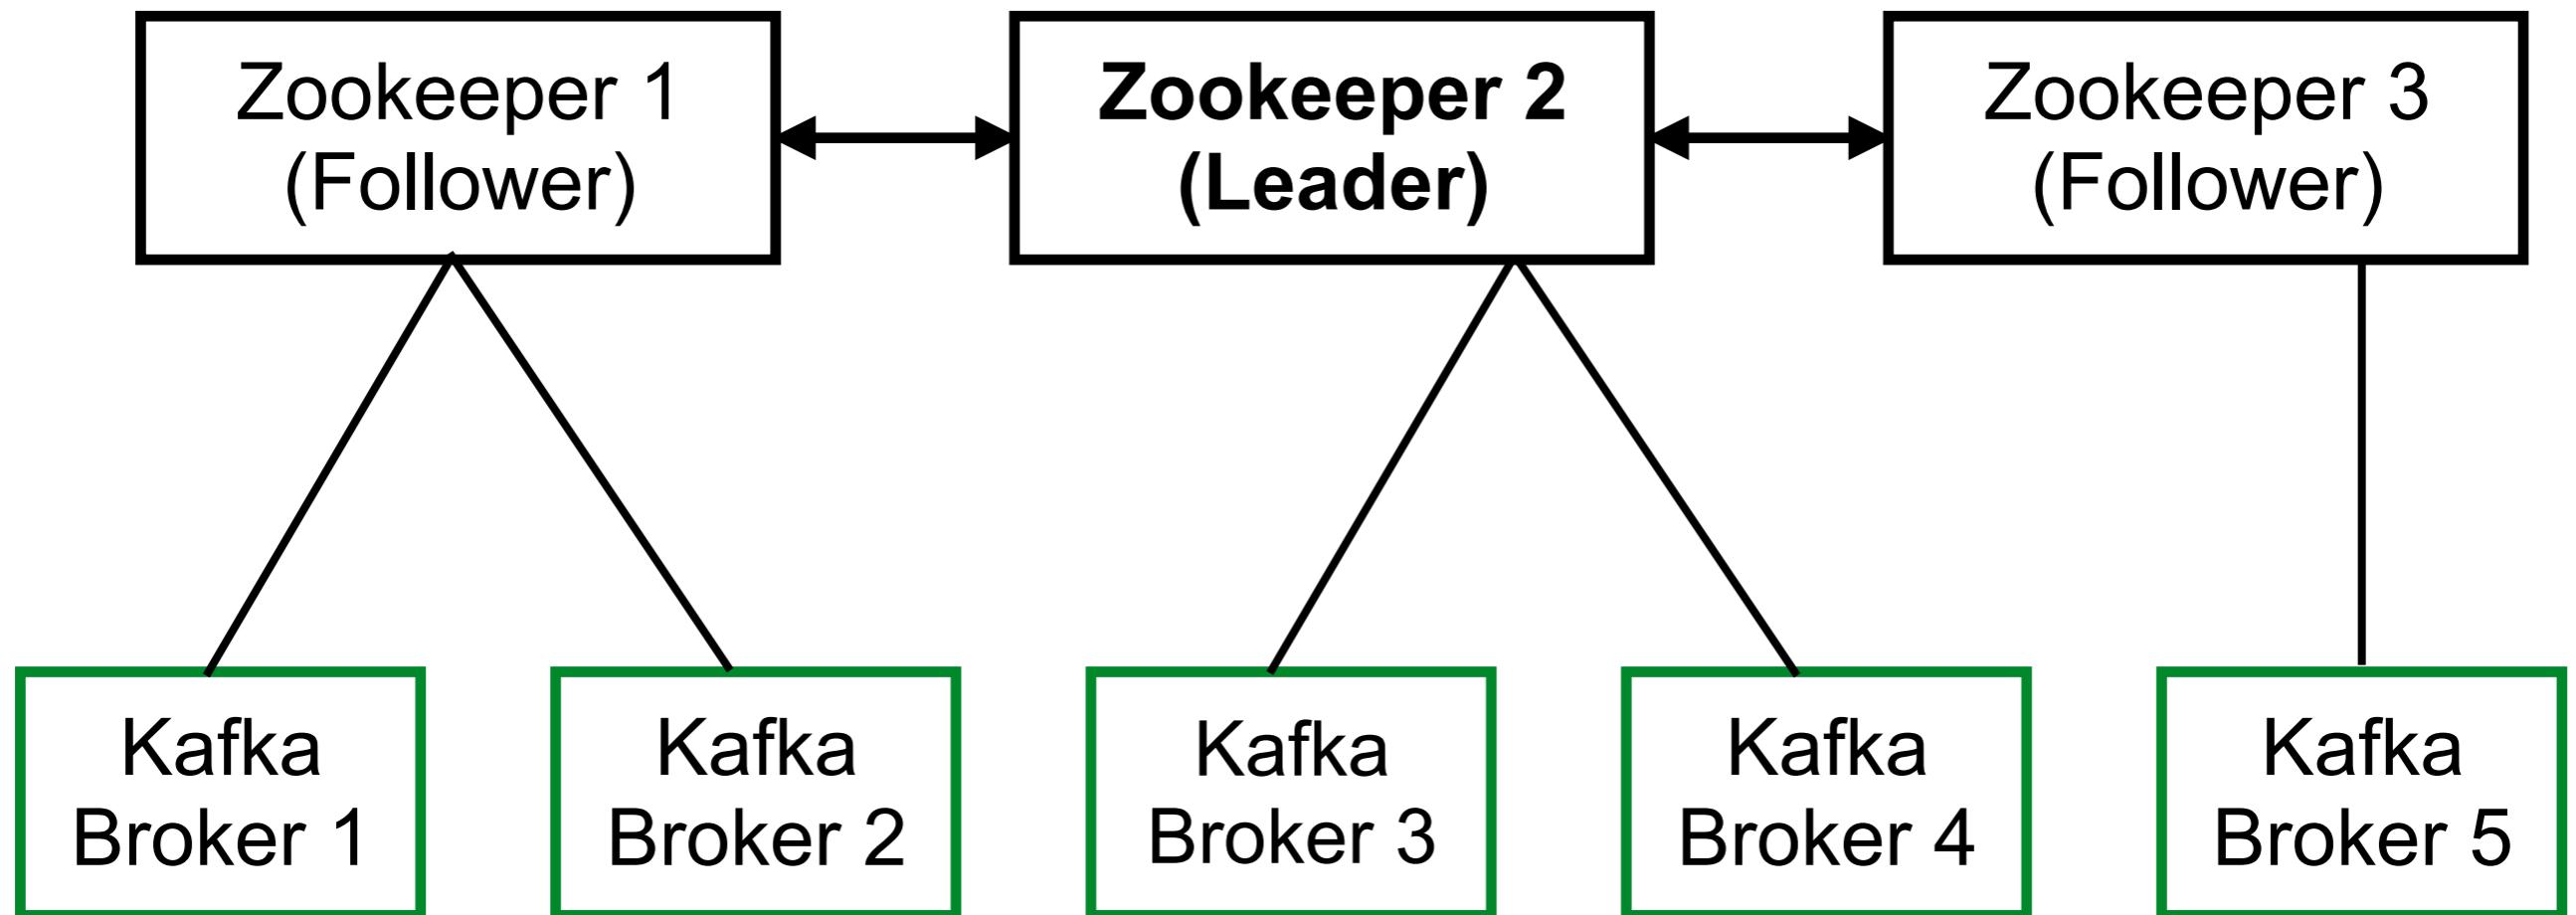


# Zookeeper architecture

Leader for write  
Follows for read



# Zookeeper architecture



# Brokers and topics

## Brokers

Good number to start id 3 brokers

Broker 1

Broker 2

Broker 3



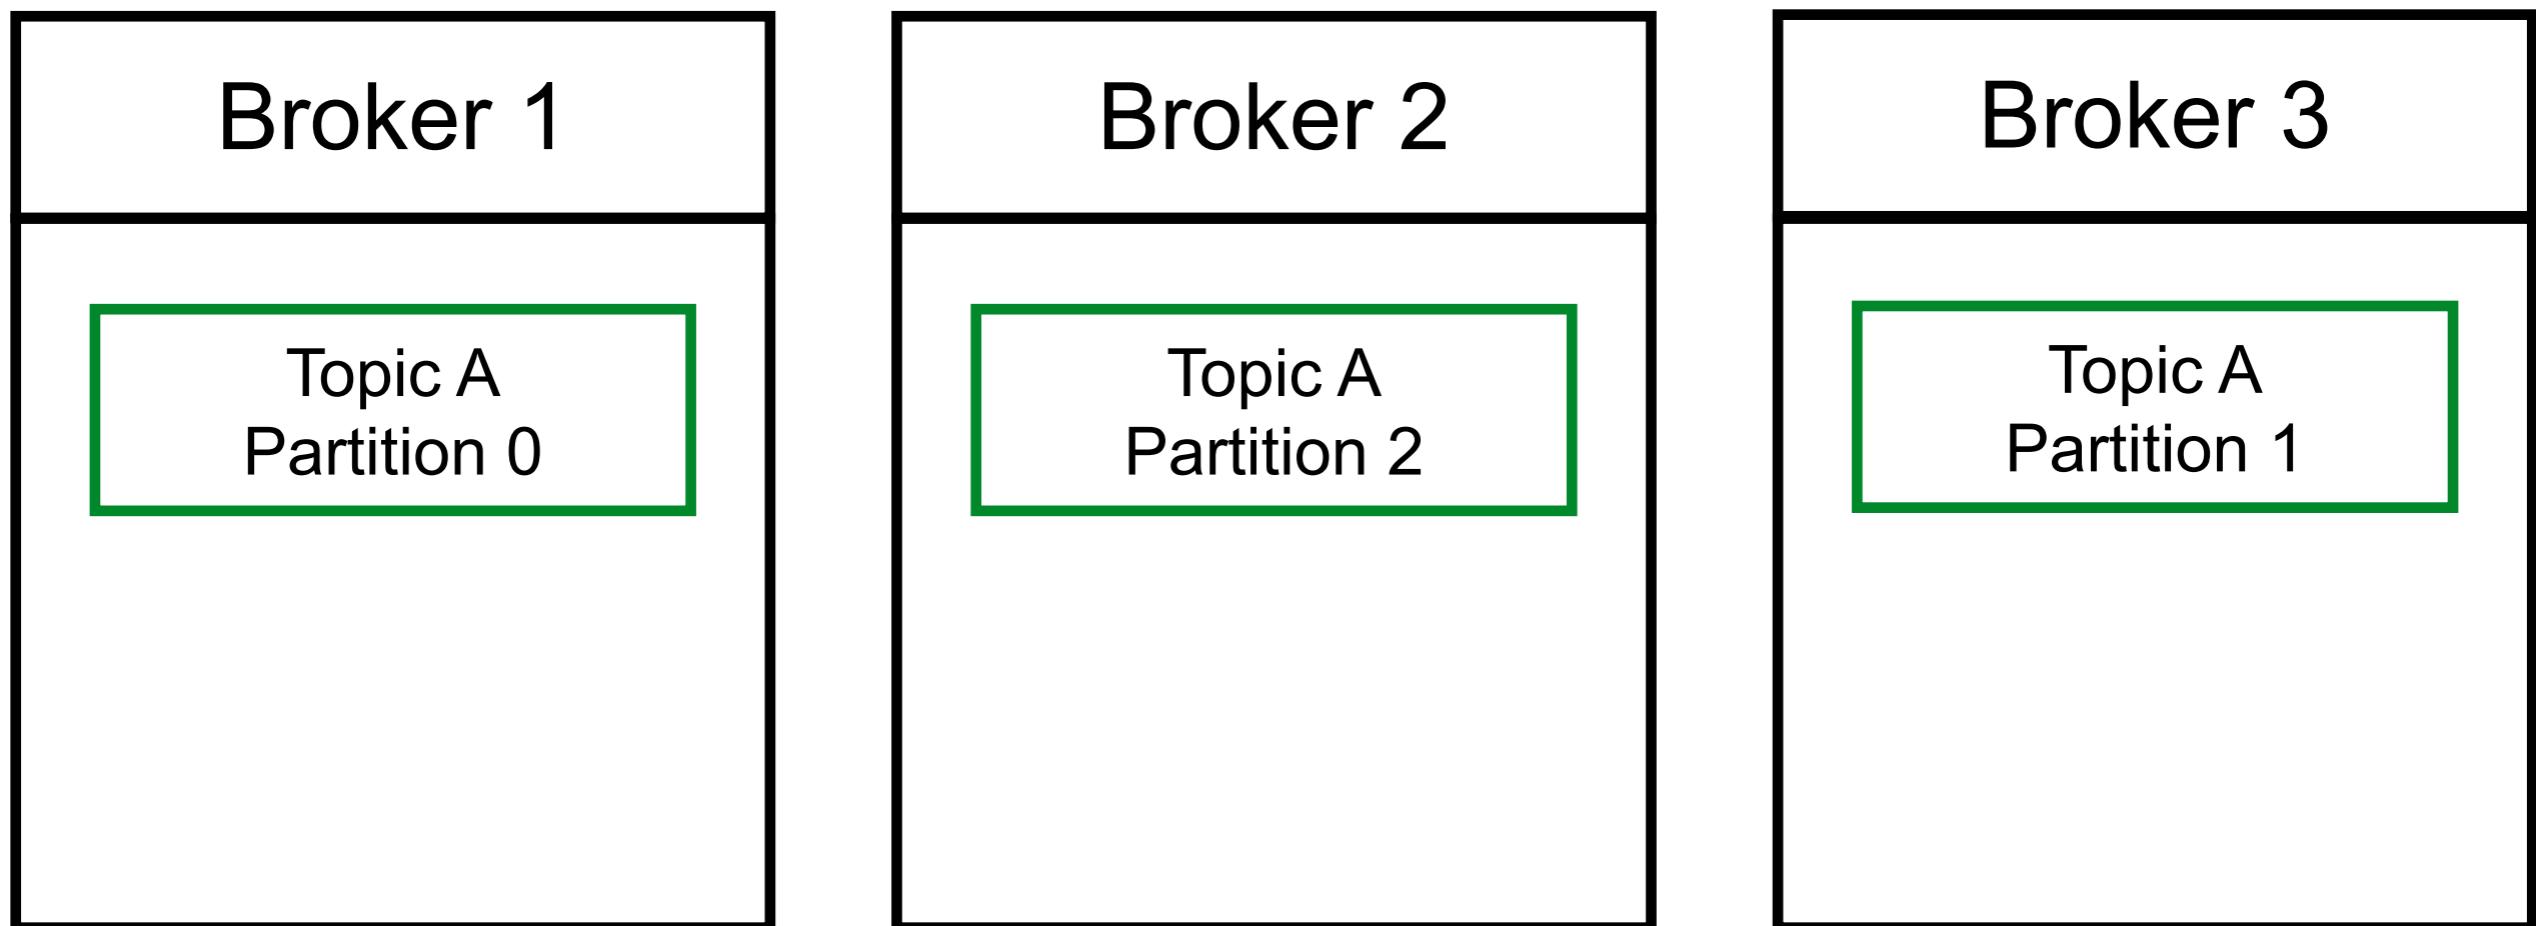
Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Brokers and topics

## Example

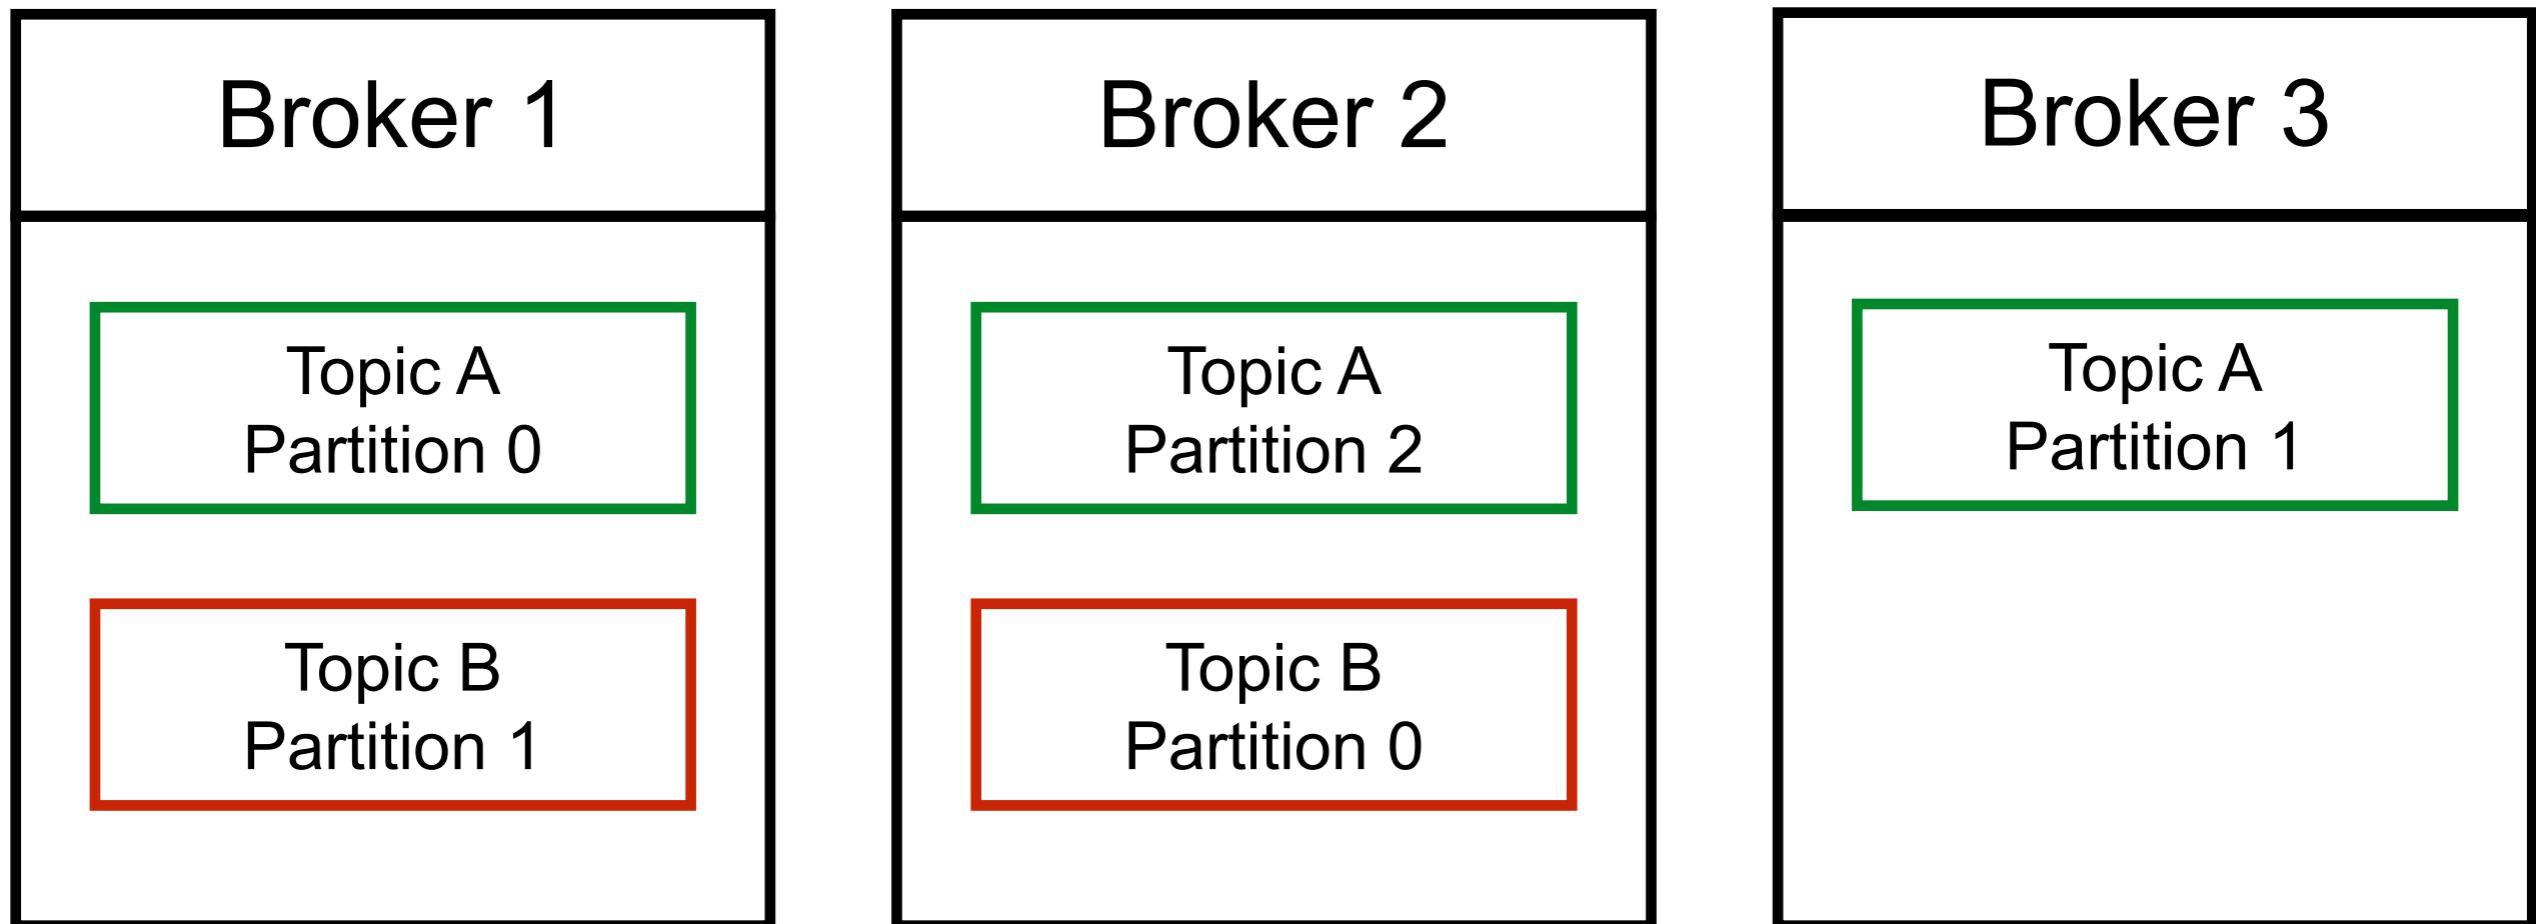
Topic A with 3 partitions



# Brokers and topics

## Example

Topic B with 2 partitions



# Topic with replication factor

Topic should have a replica factor  $> 1$  (2-3)

**replica factor < no. of brokers**

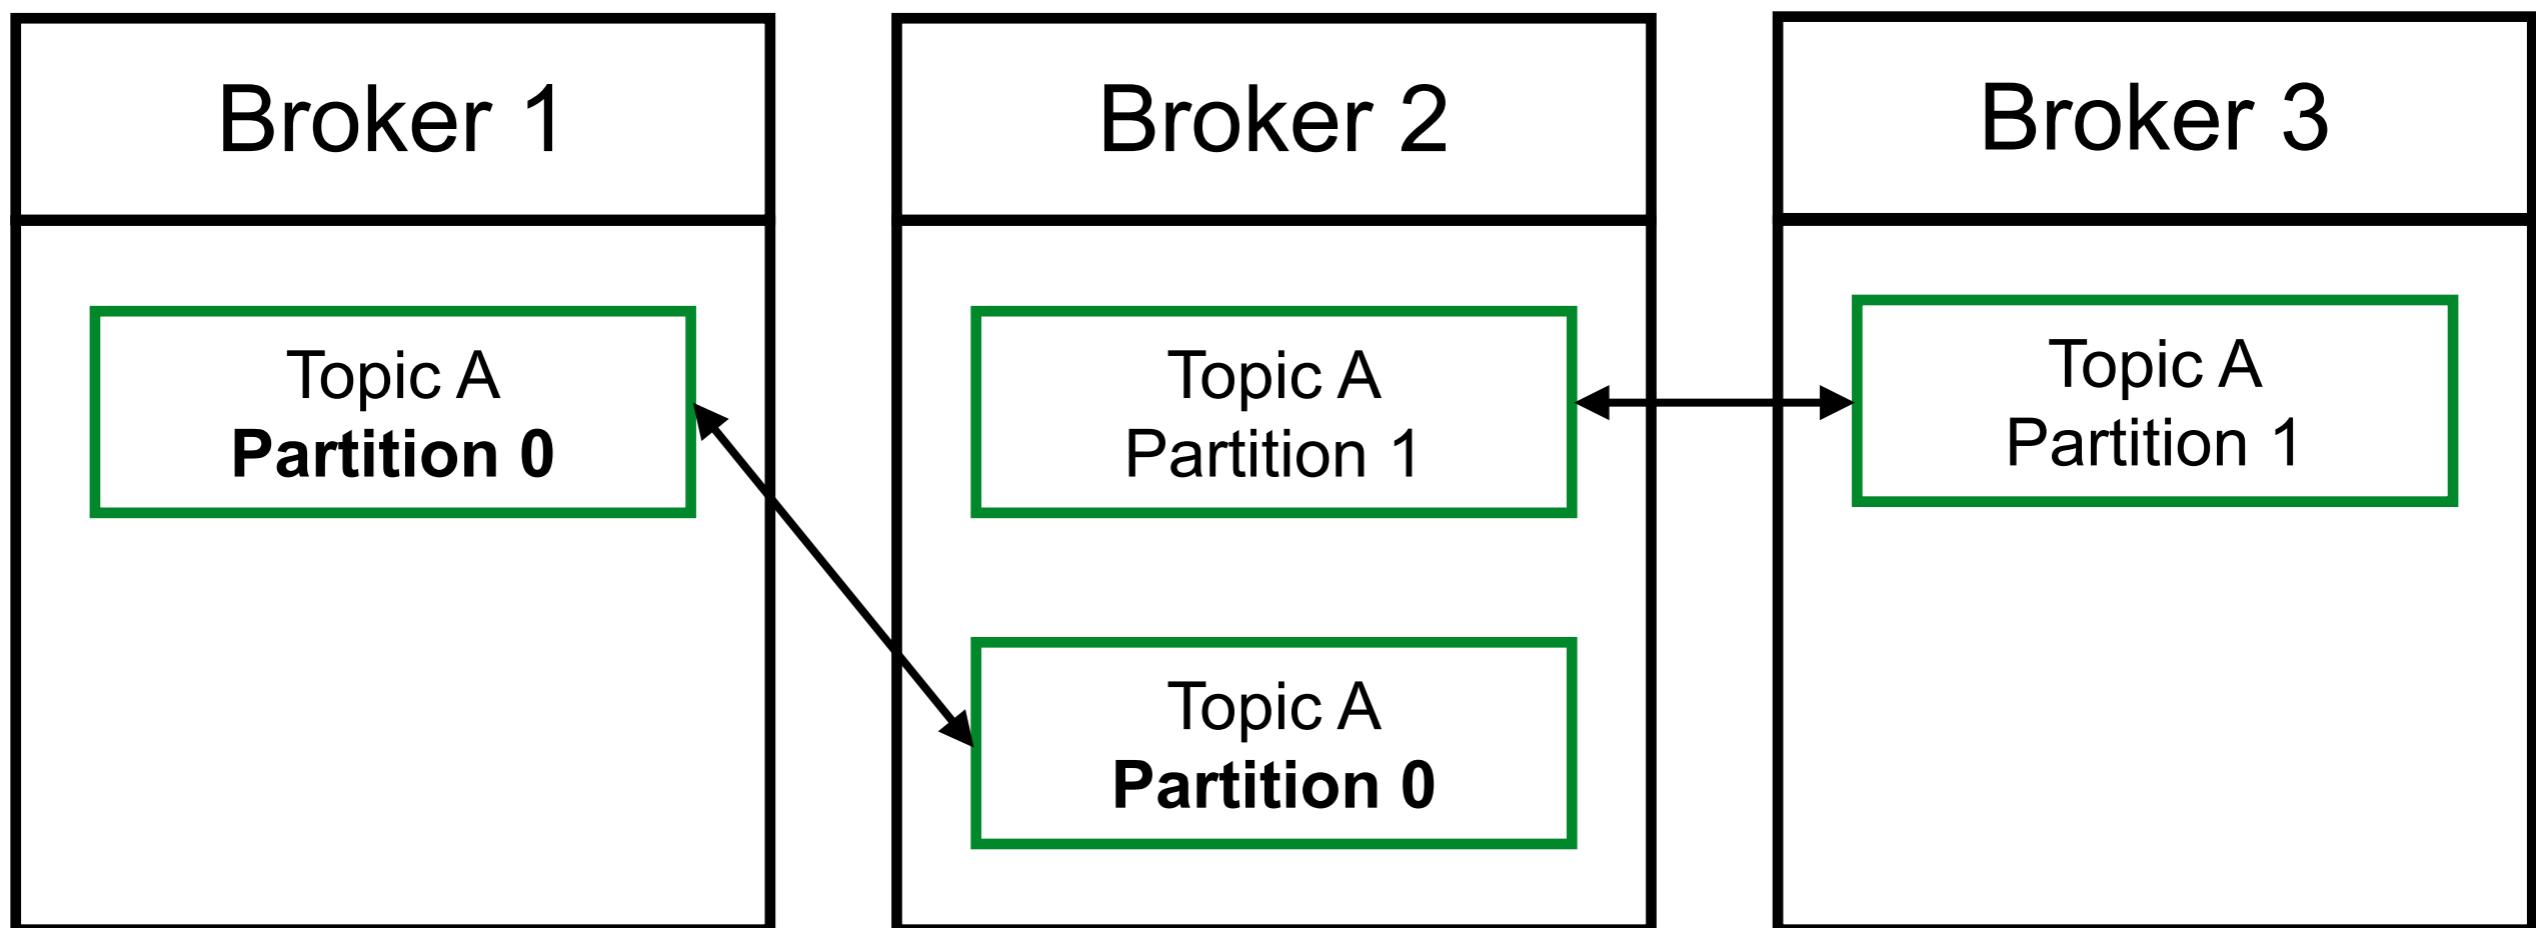
This way if a broker down, another broker can serve the data



# Topic with replication factor

## Example

Topic A with 2 partitions and replication factor = 2



# **Leader for a partition**

At any time only **one Broker** can be leader for partition

**Only leader partition** can receive and serve data for a partition

Other brokers will synchronize the data

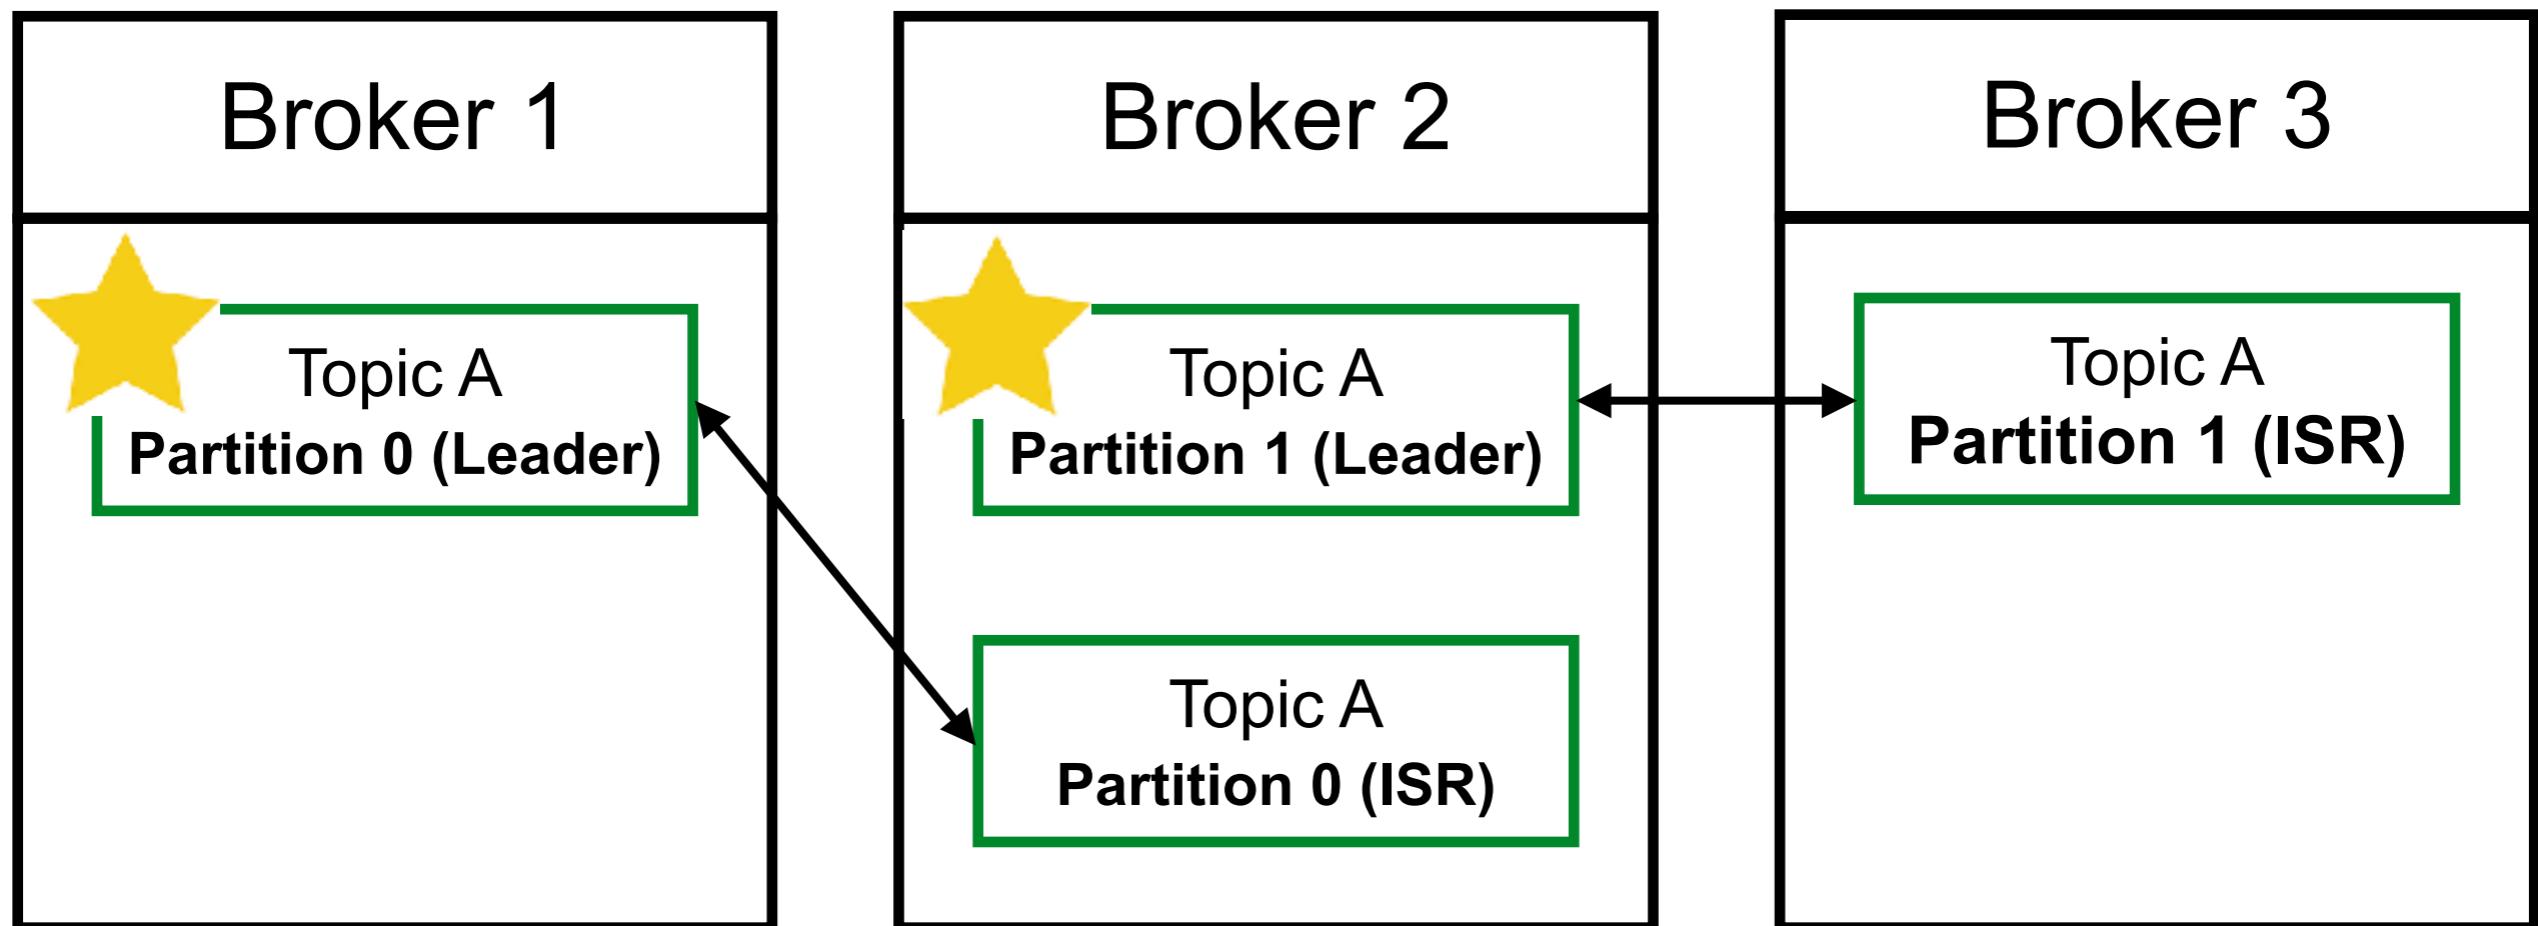
Other partition called **in-sync replica (ISR)**



# Leader for a partition

## Example

Topic A with 2 partitions and replication factor = 2

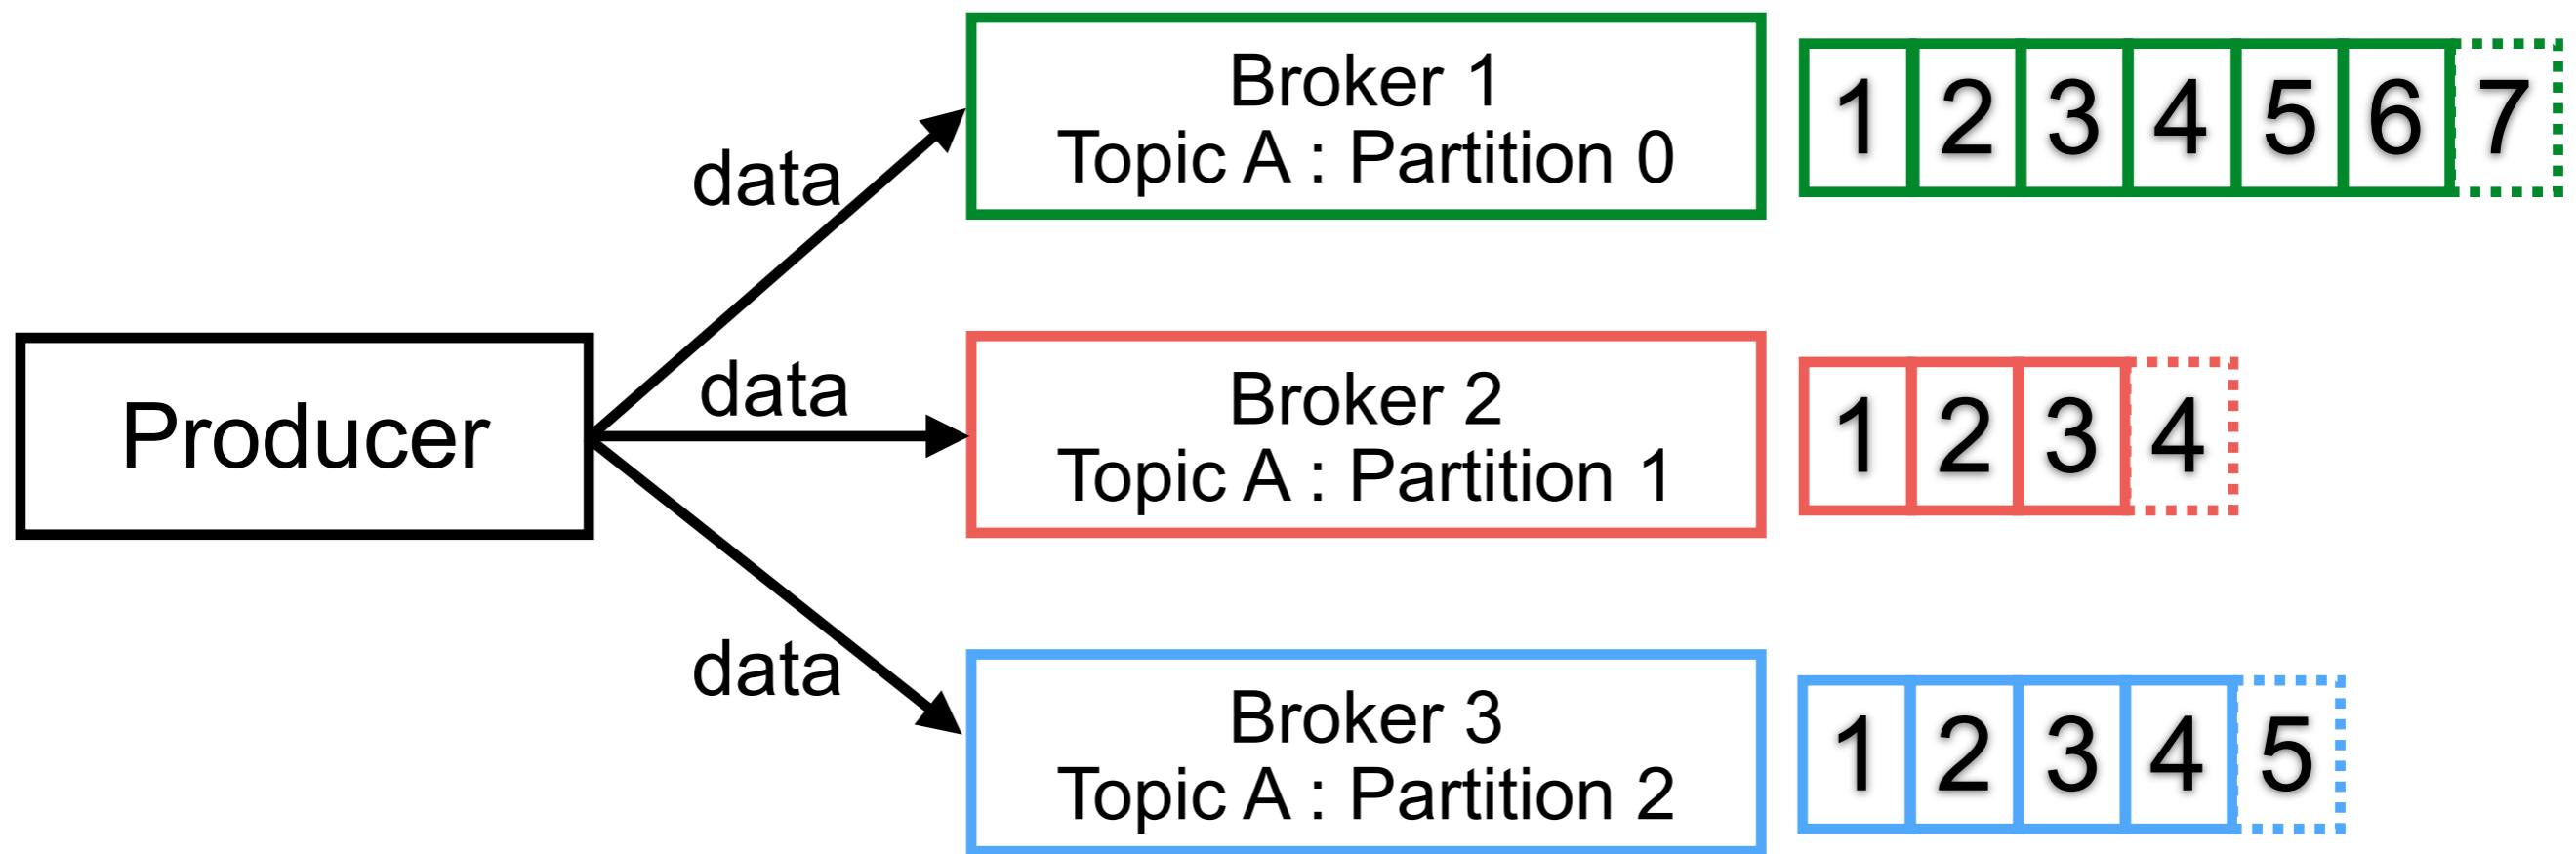


# Producers

Producers write data to topics  
Automatically know to broker and partition to write  
When broker failures, producers will automatically recover



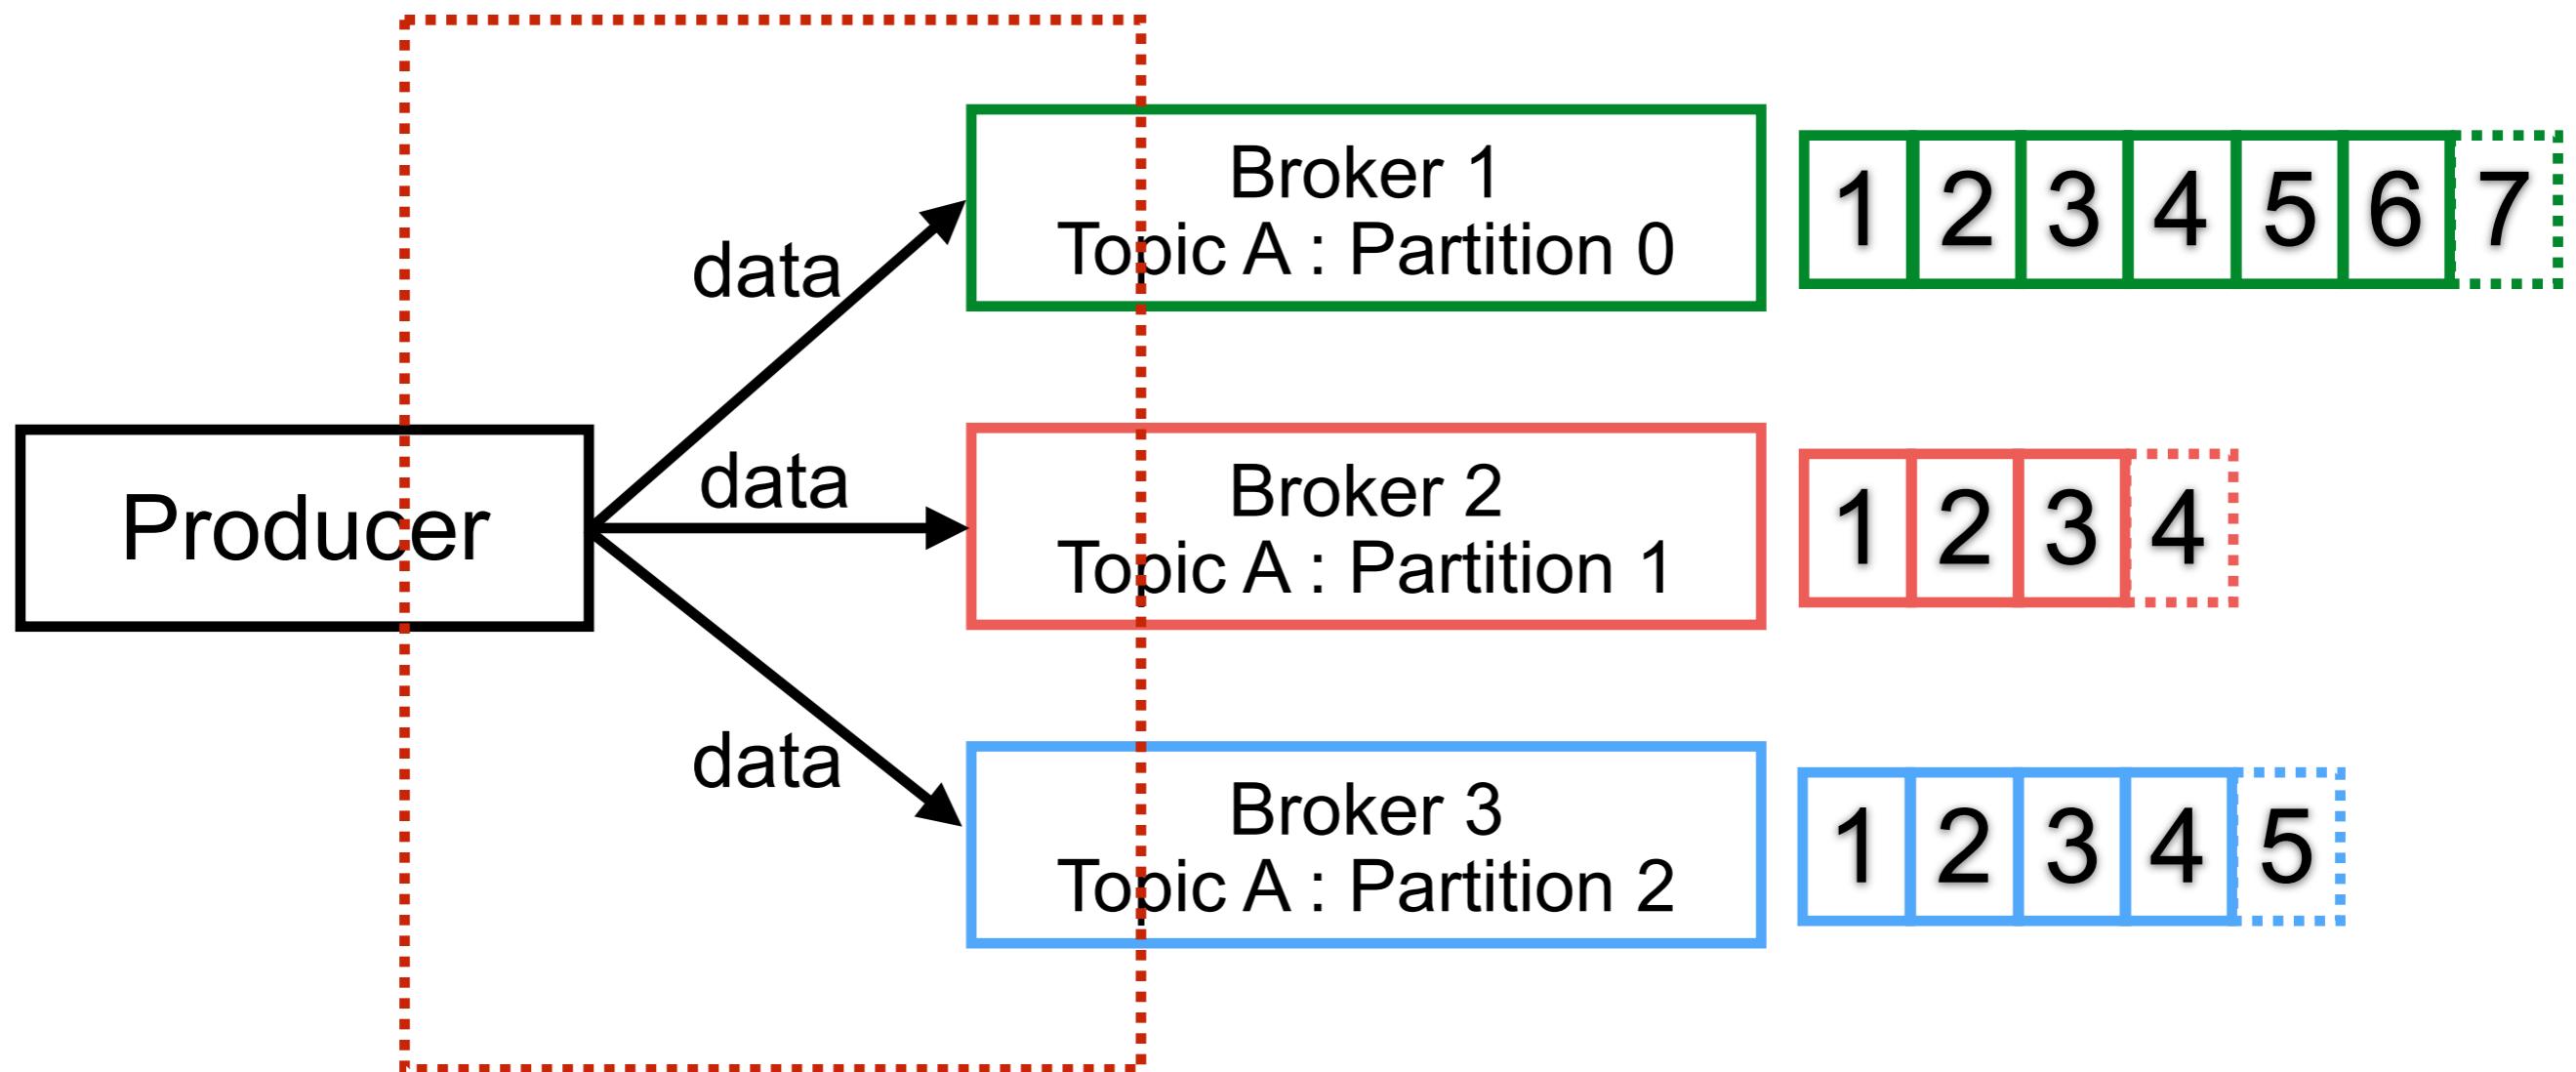
# Producers



\*\*\* *The load is balanced to many brokers (no. of partitions)* \*\*\*



# Producers issue to write data !!



# Producers with acknowledgment

## acks=0

Producer not wait for acknowledgment

Possible data loss

## acks=1

Producer will wait for **leader** acknowledgment

Limited data loss

## acks=all

Producer will wait for **leader + ISR** acknowledgment

No data loss

<https://docs.confluent.io/platform/current/clients/producer.html>

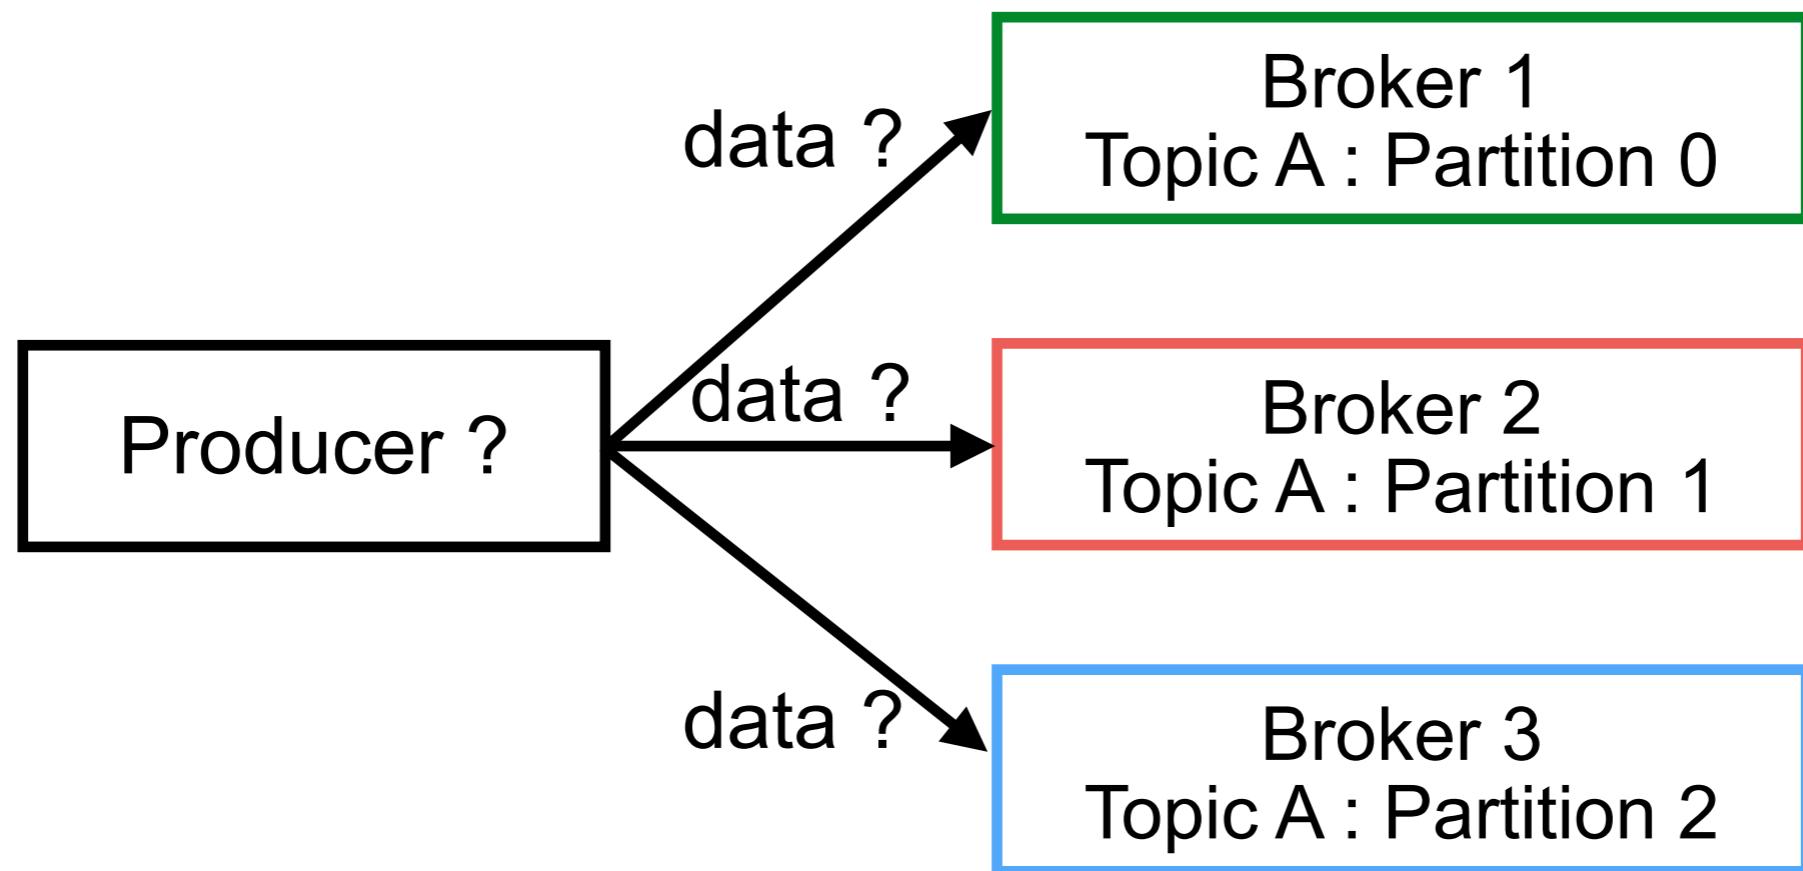


Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Producers with message keys

Producers can choose to sent a **key** with data  
**Key = null**, data is sent **Sticky partitioner**



# How to assign to partition ?

```
/**  
 * Compute the partition for the given record.  
 *  
 * @param topic The topic name  
 * @param numPartitions The number of partitions of the given {@code topic}  
 * @param key The key to partition on (or null if no key)  
 * @param keyBytes serialized key to partition on (or null if no key)  
 * @param value The value to partition on or null  
 * @param valueBytes serialized value to partition on or null  
 * @param cluster The current cluster metadata  
 */  
  
public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[] valueBytes, Cluster cluster,  
                     int numPartitions) {  
    if (keyBytes == null) {  
        return stickyPartitionCache.partition(topic, cluster);  
    }  
    return BuiltInPartitioner.partitionForKey(keyBytes, numPartitions);  
}  
  
public void close() {}
```

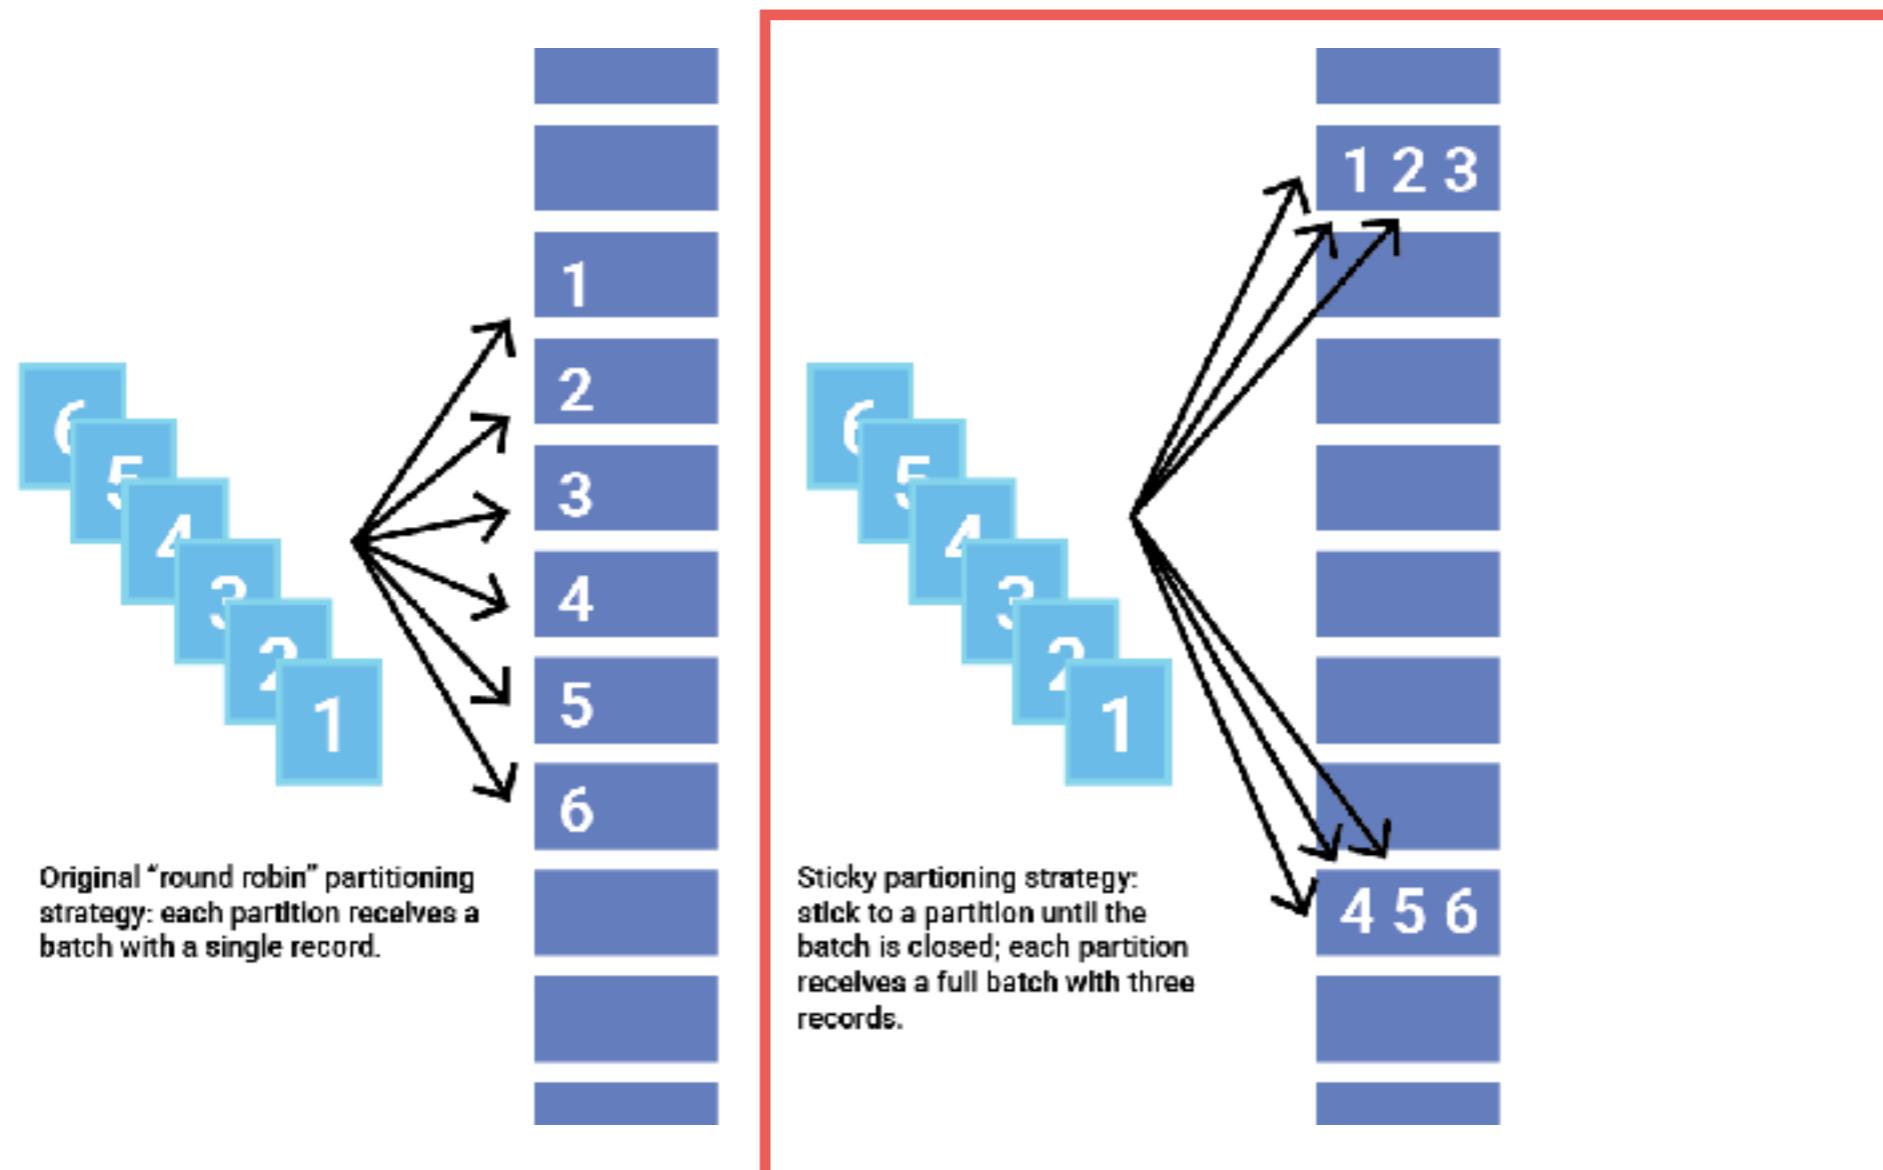
<https://github.com/apache/kafka/blob/trunk/clients/src/main/java/org/apache/kafka/clients/producer/internals/DefaultPartitioner.java>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# How to assign to partition ?



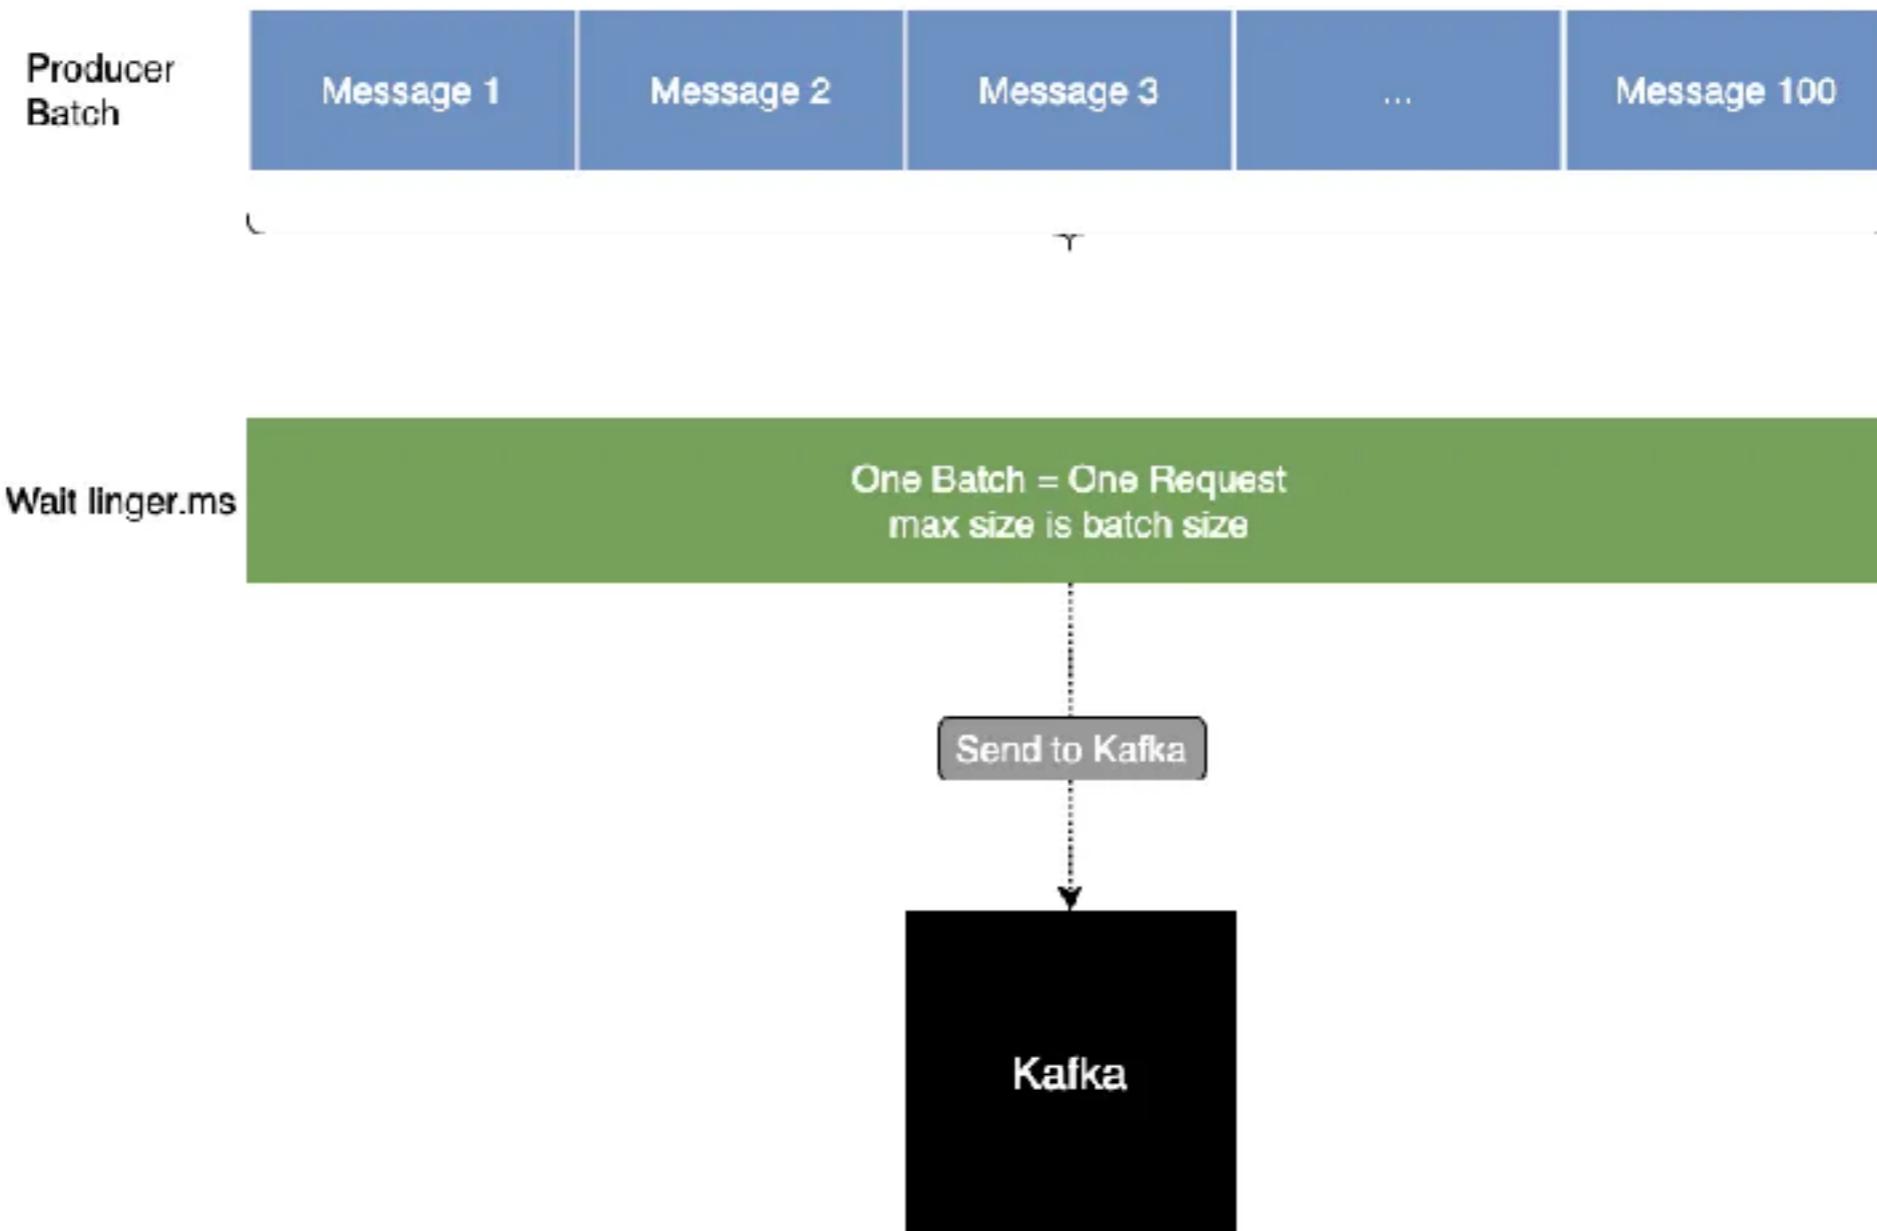
batch.size and linger.ms



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Batch size and linger.ms



<https://docs.confluent.io/platform/current/installation/configuration/producer-configs.html>

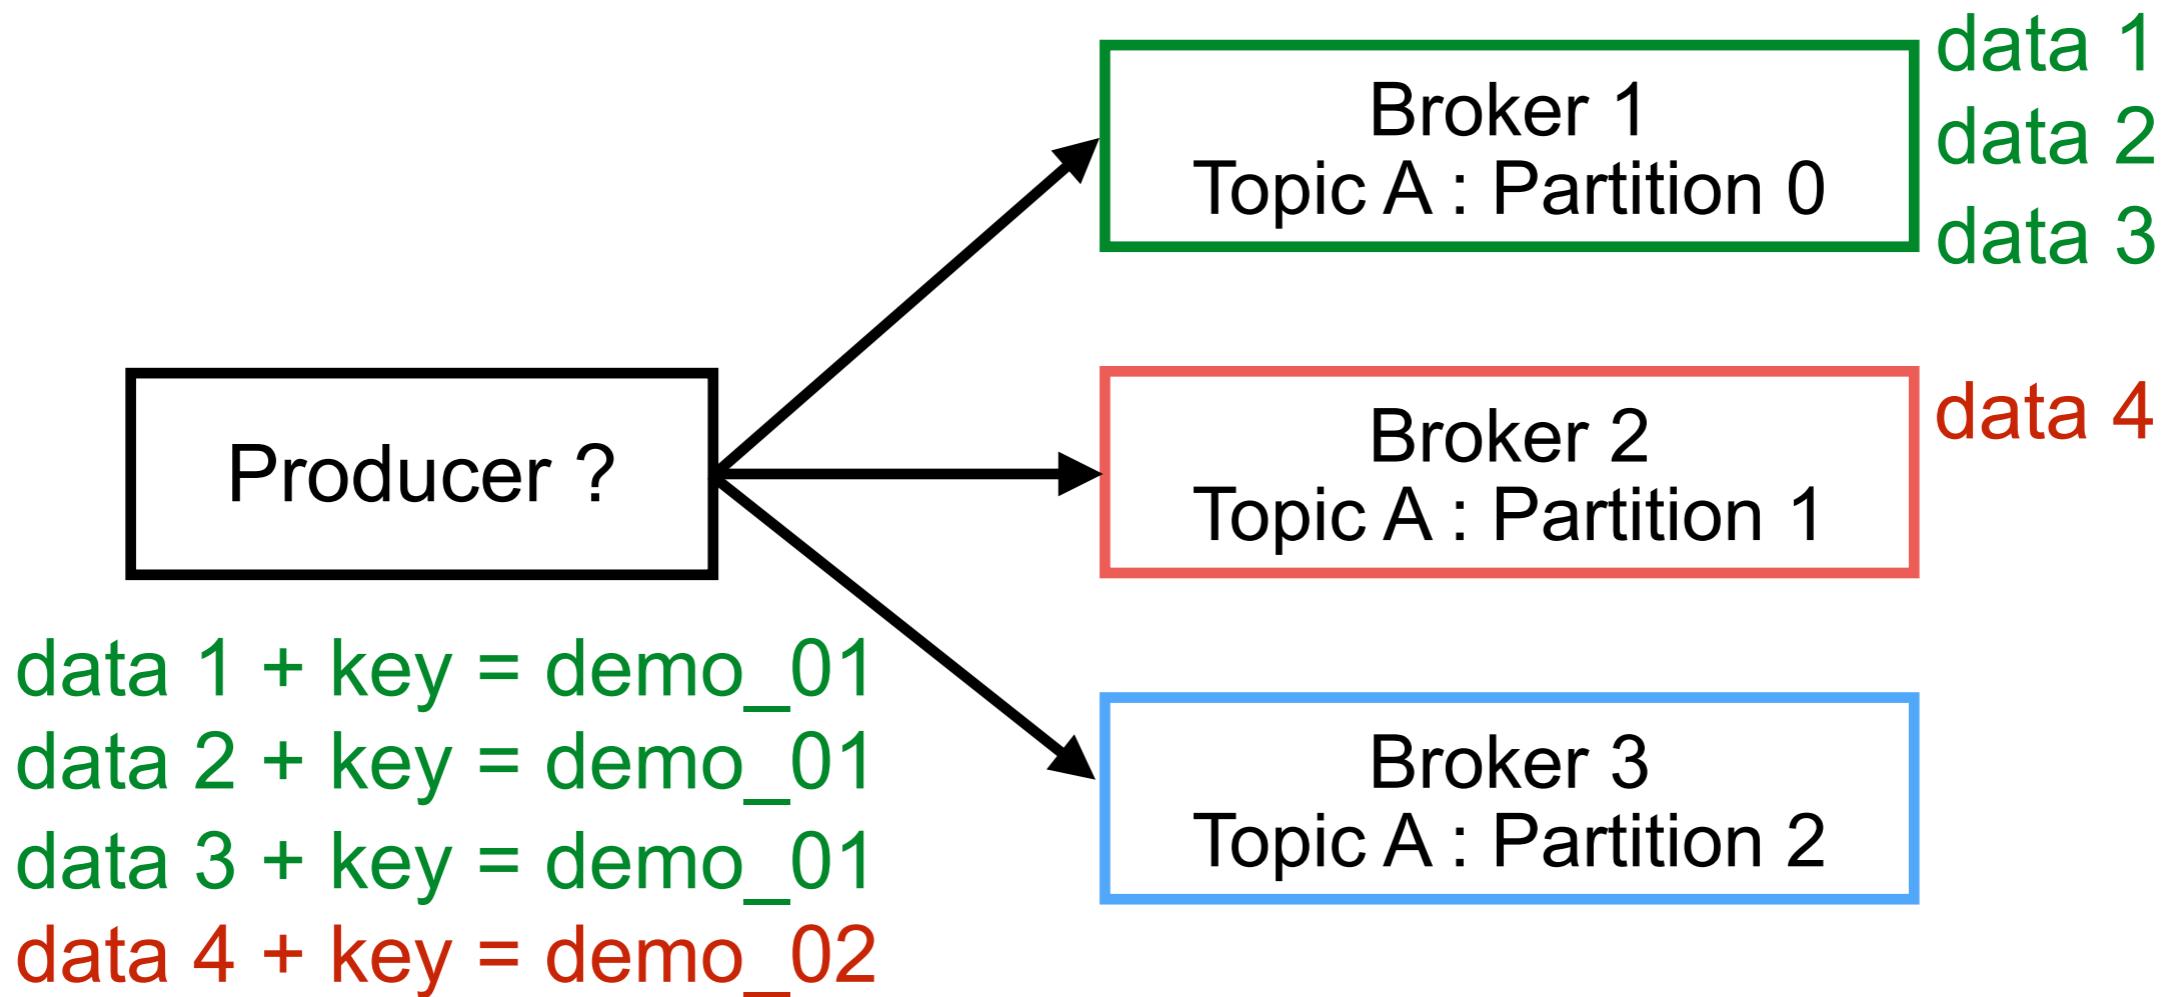


Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Why use message keys ?

You need message ordering



# How to assign to partition ?

```
326     /*
327      * Default hashing function to choose a partition from the serialized key bytes
328      */
329     public static int partitionForKey(final byte[] serializedKey, final int numPartitions) {
330         return Utils.toPositive(Utils.murmur2(serializedKey)) % numPartitions;
331     }
332 }
```

## Murmur2 algorithm (hashing function)

<https://github.com/apache/kafka/blob/trunk/clients/src/main/java/org/apache/kafka/clients/producer/internals/BuiltInPartitioner.java>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

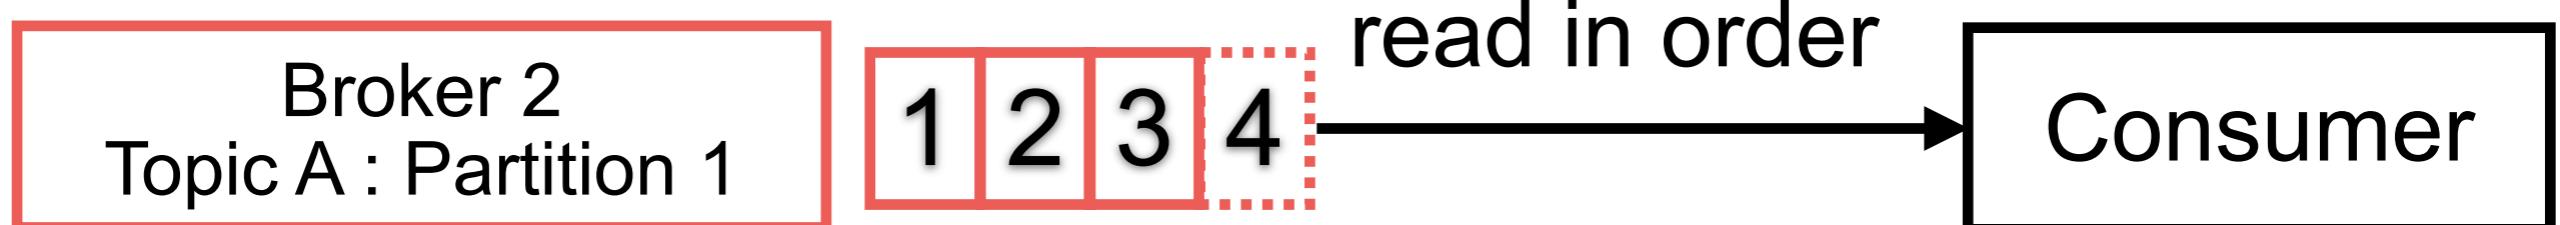
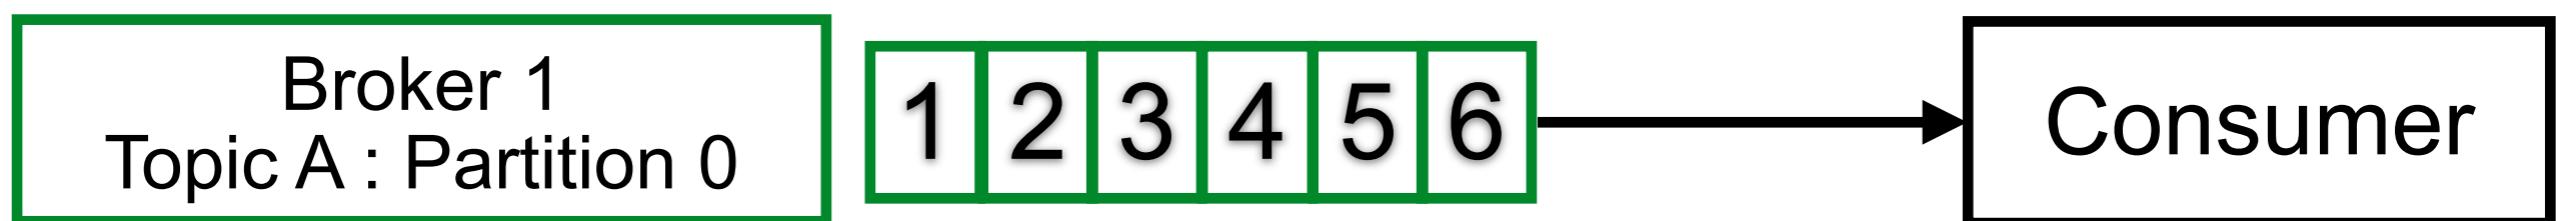
# Consumers

Consumers read data from topic

Consumers know which broker to read from

Data will read in order **within each partition**

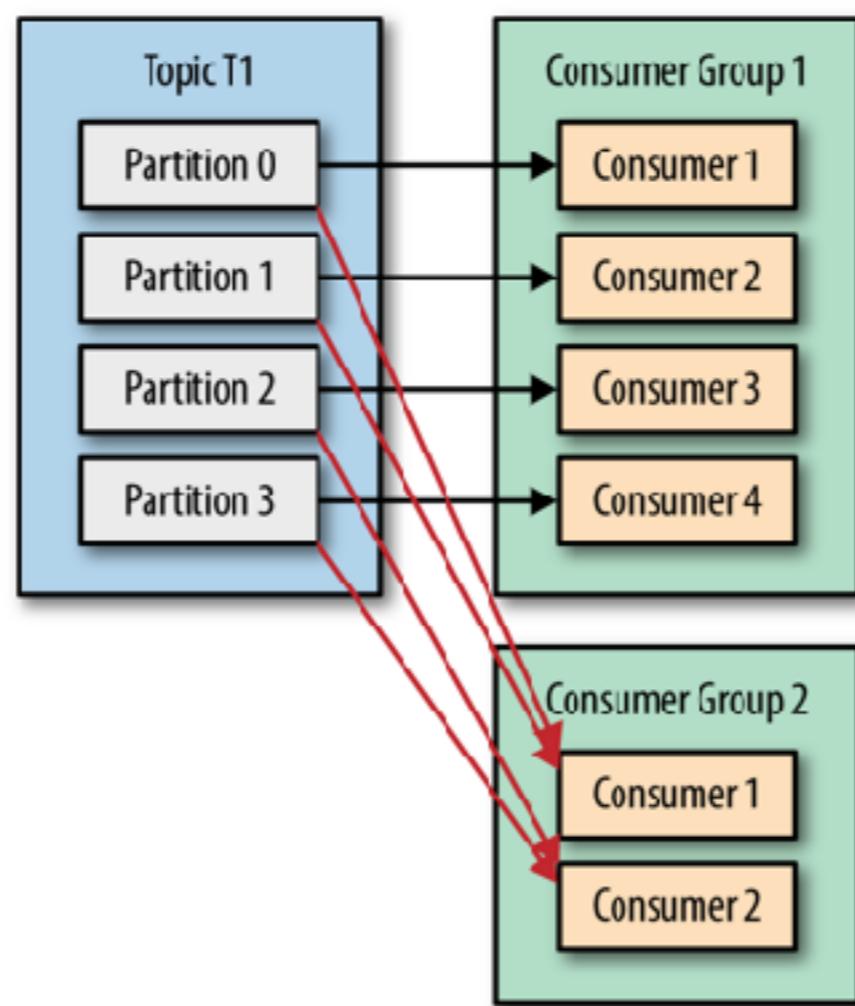
When broker **failures**, consumer know how to recover



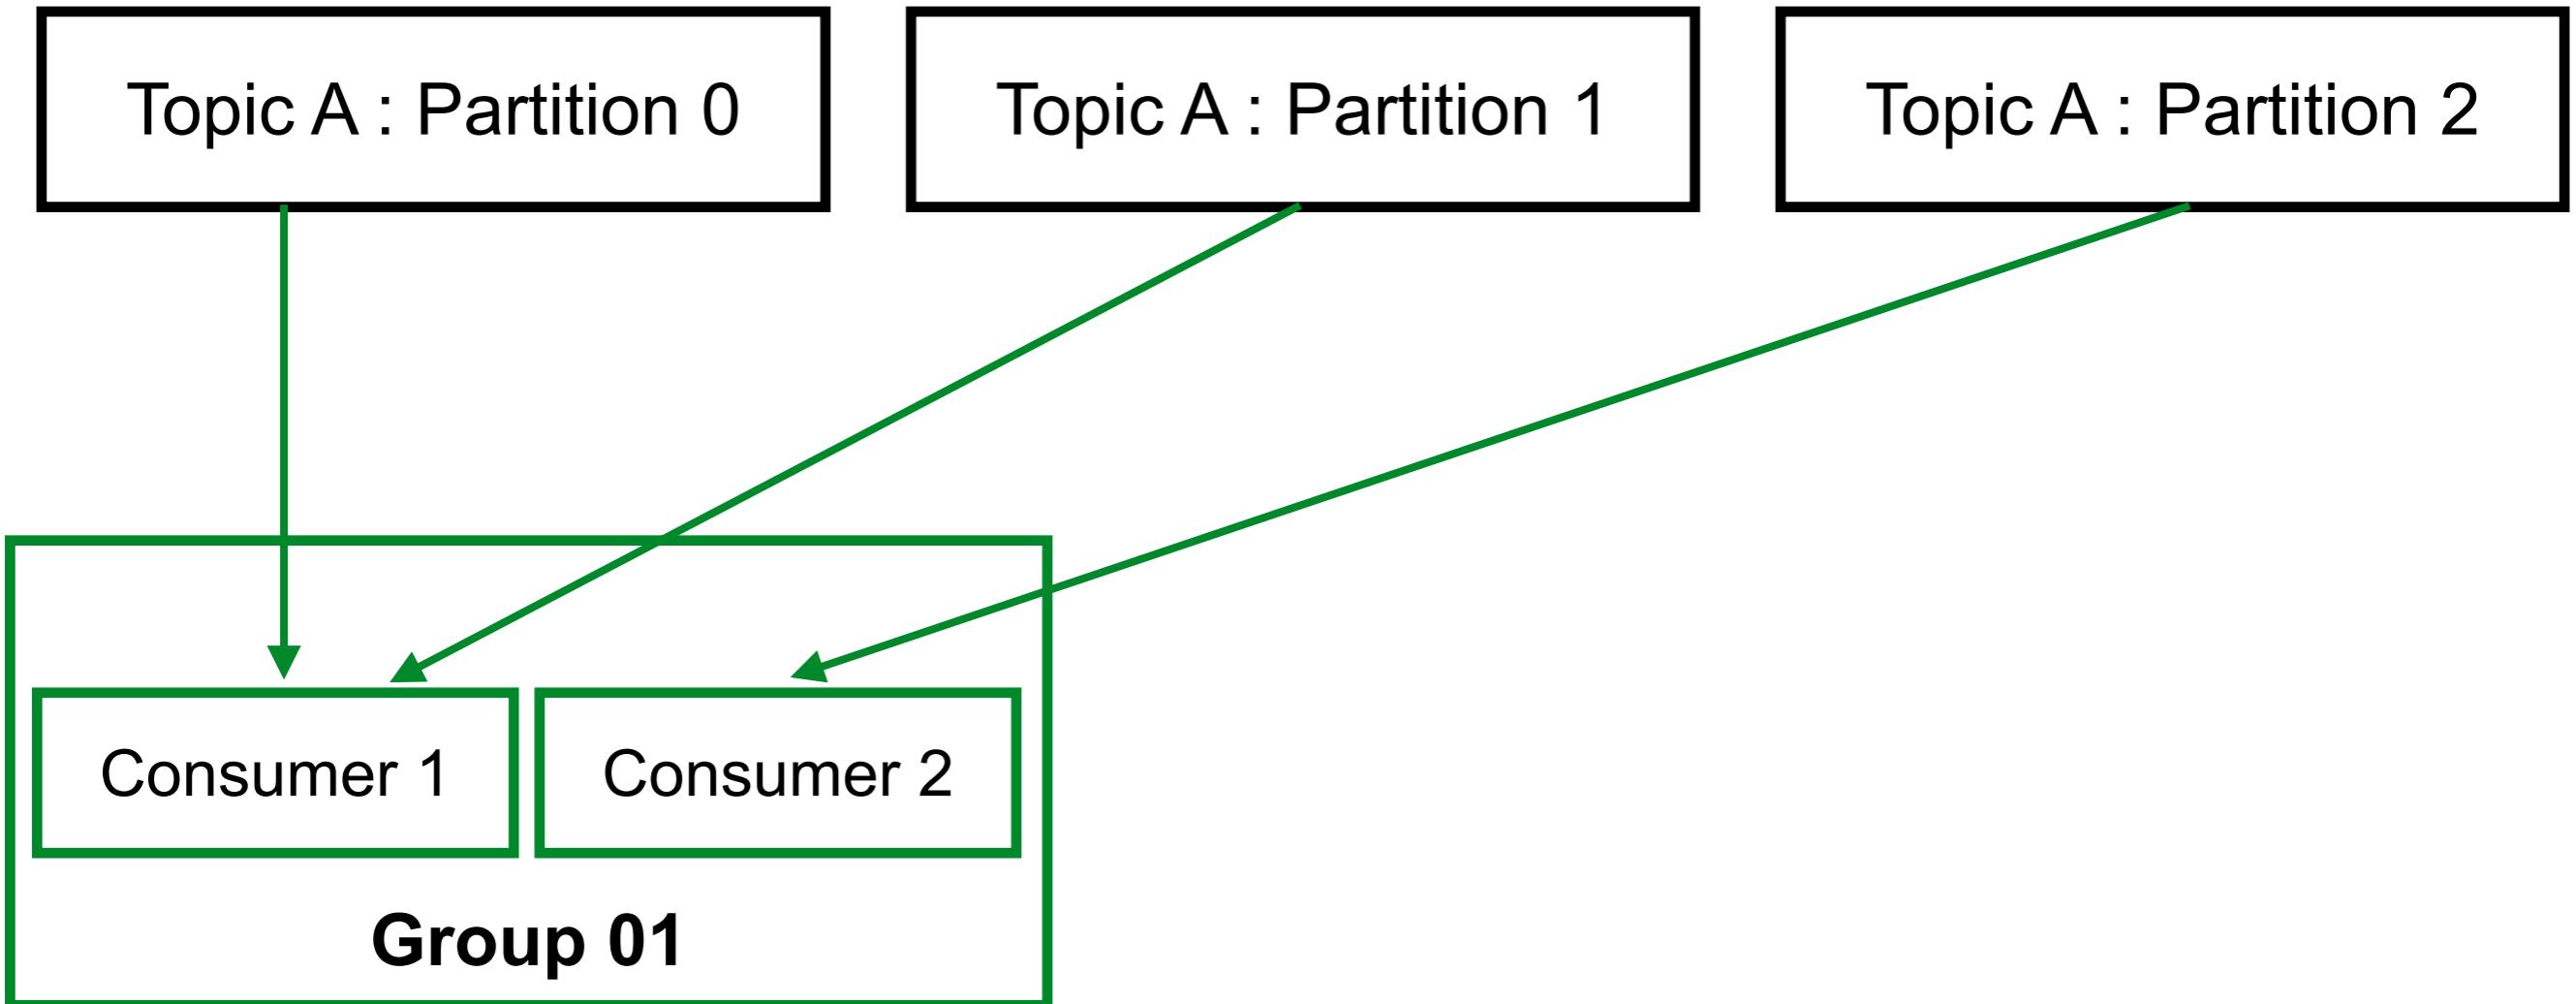
# Consumer groups

Consumers read data in consumer groups

Each consumer in a group read from exclusive partitions



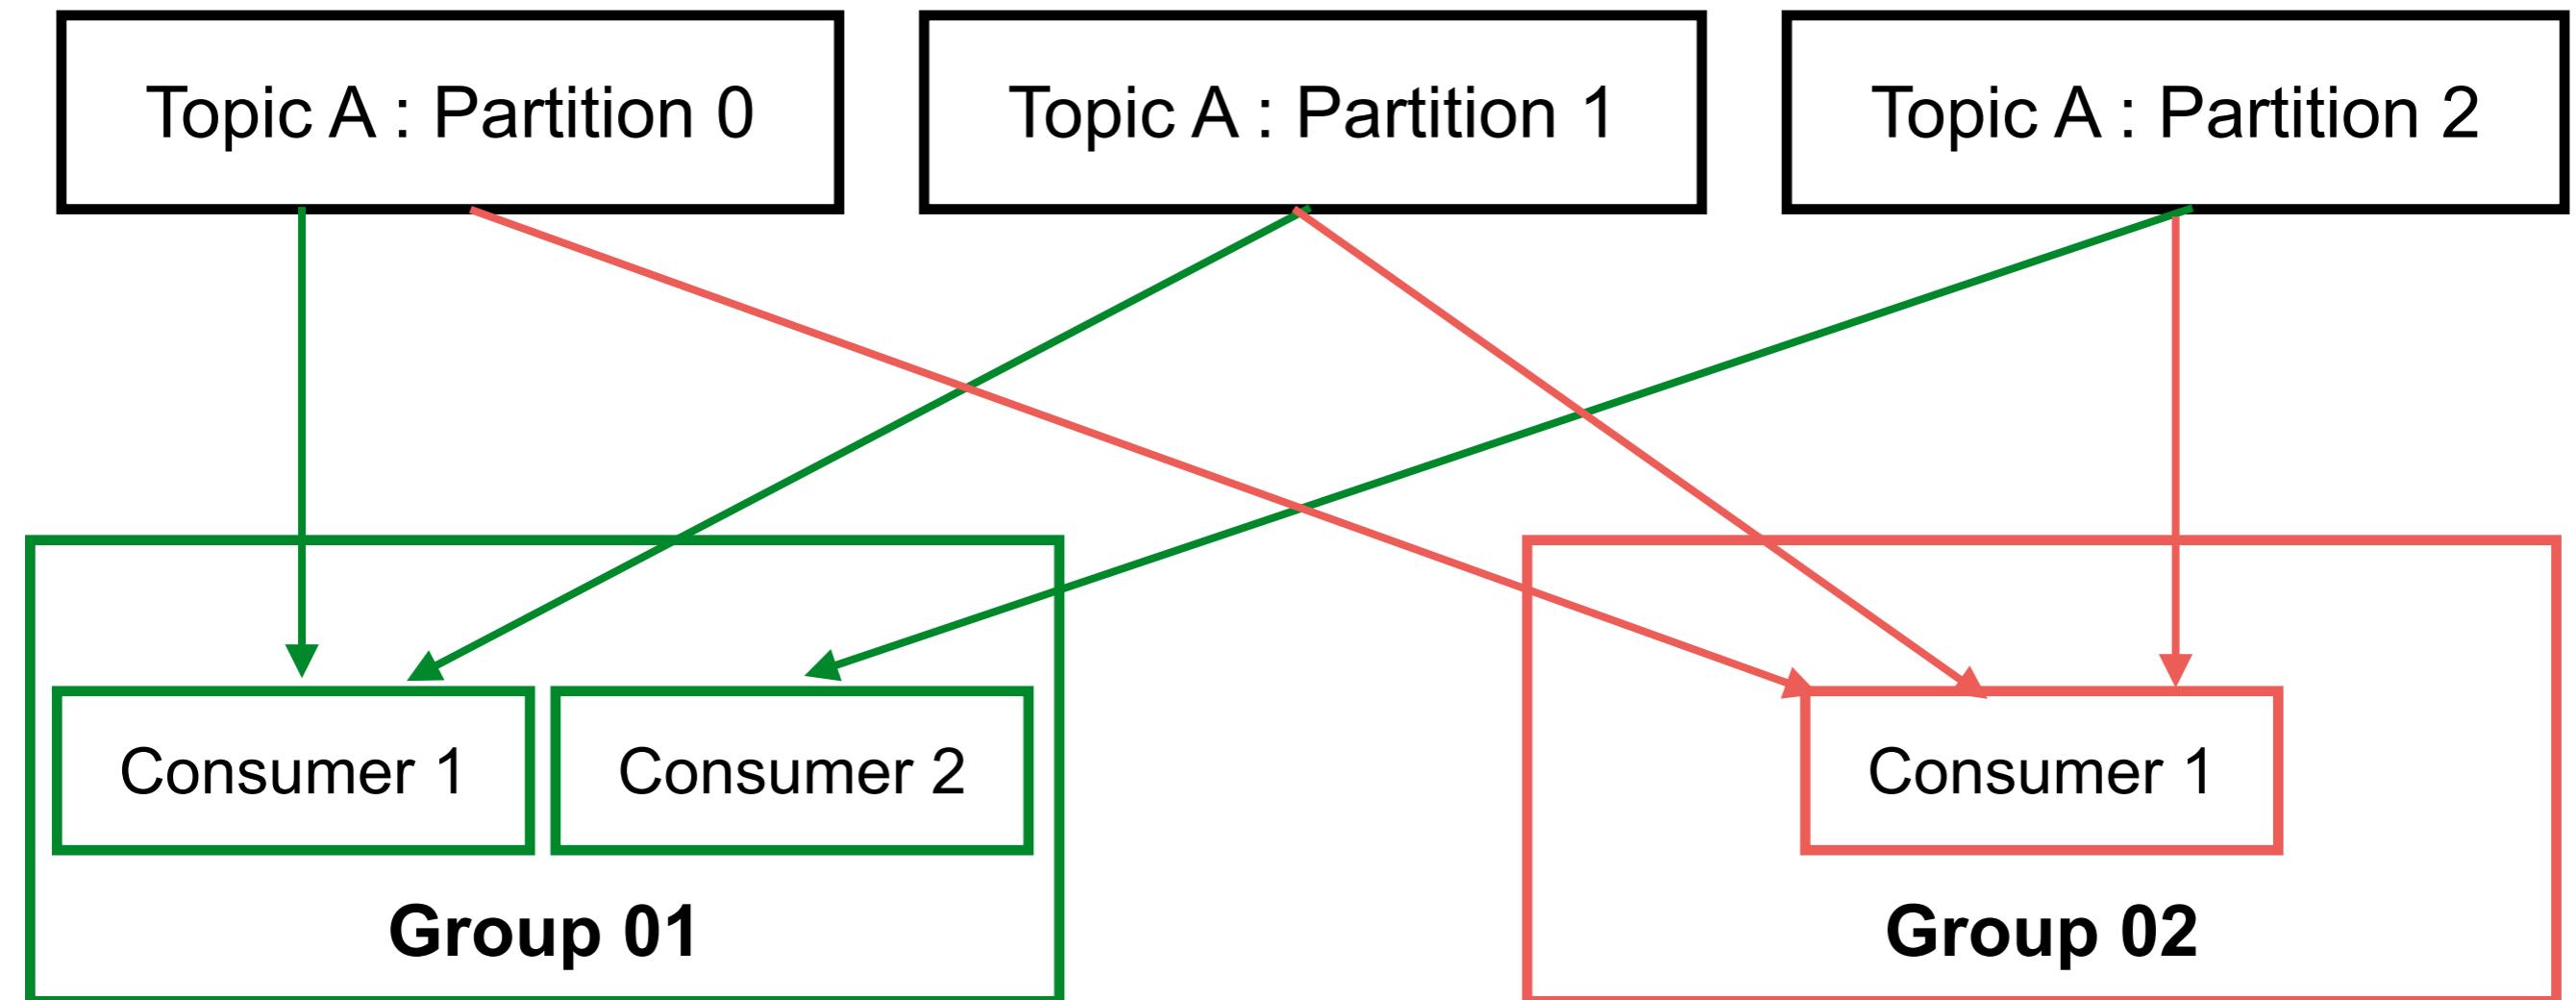
# Consumer groups



Consumer will automatically use a GroupCoordinator  
ConsumerCoordinator to assign a consumer to a partition.



# Consumer groups



# Partition assignment strategies

Range (default)

Round Robin

Sticky

Customize (AbstractPartitionAssignor)

<https://docs.confluent.io/platform/current/installation/configuration/consumer-configs.html#partition-assignment-strategy>

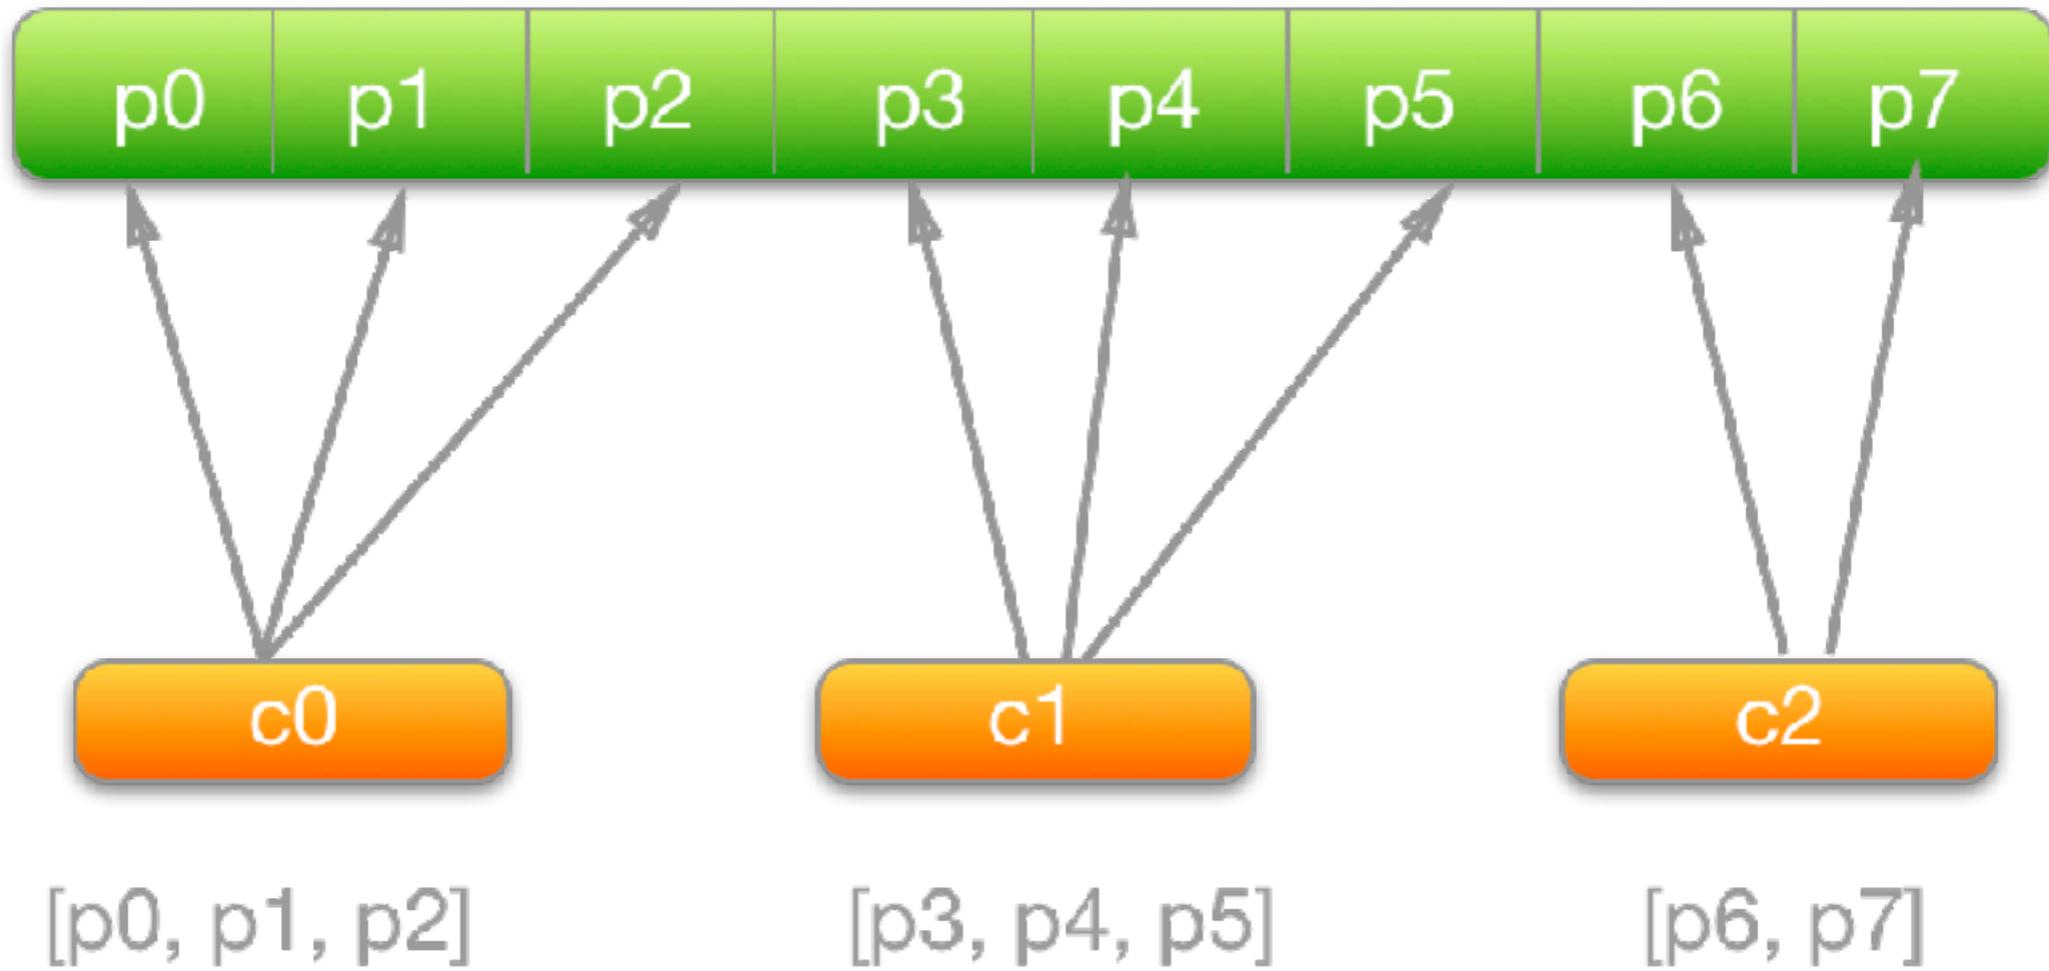


Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

75

# Range assignor



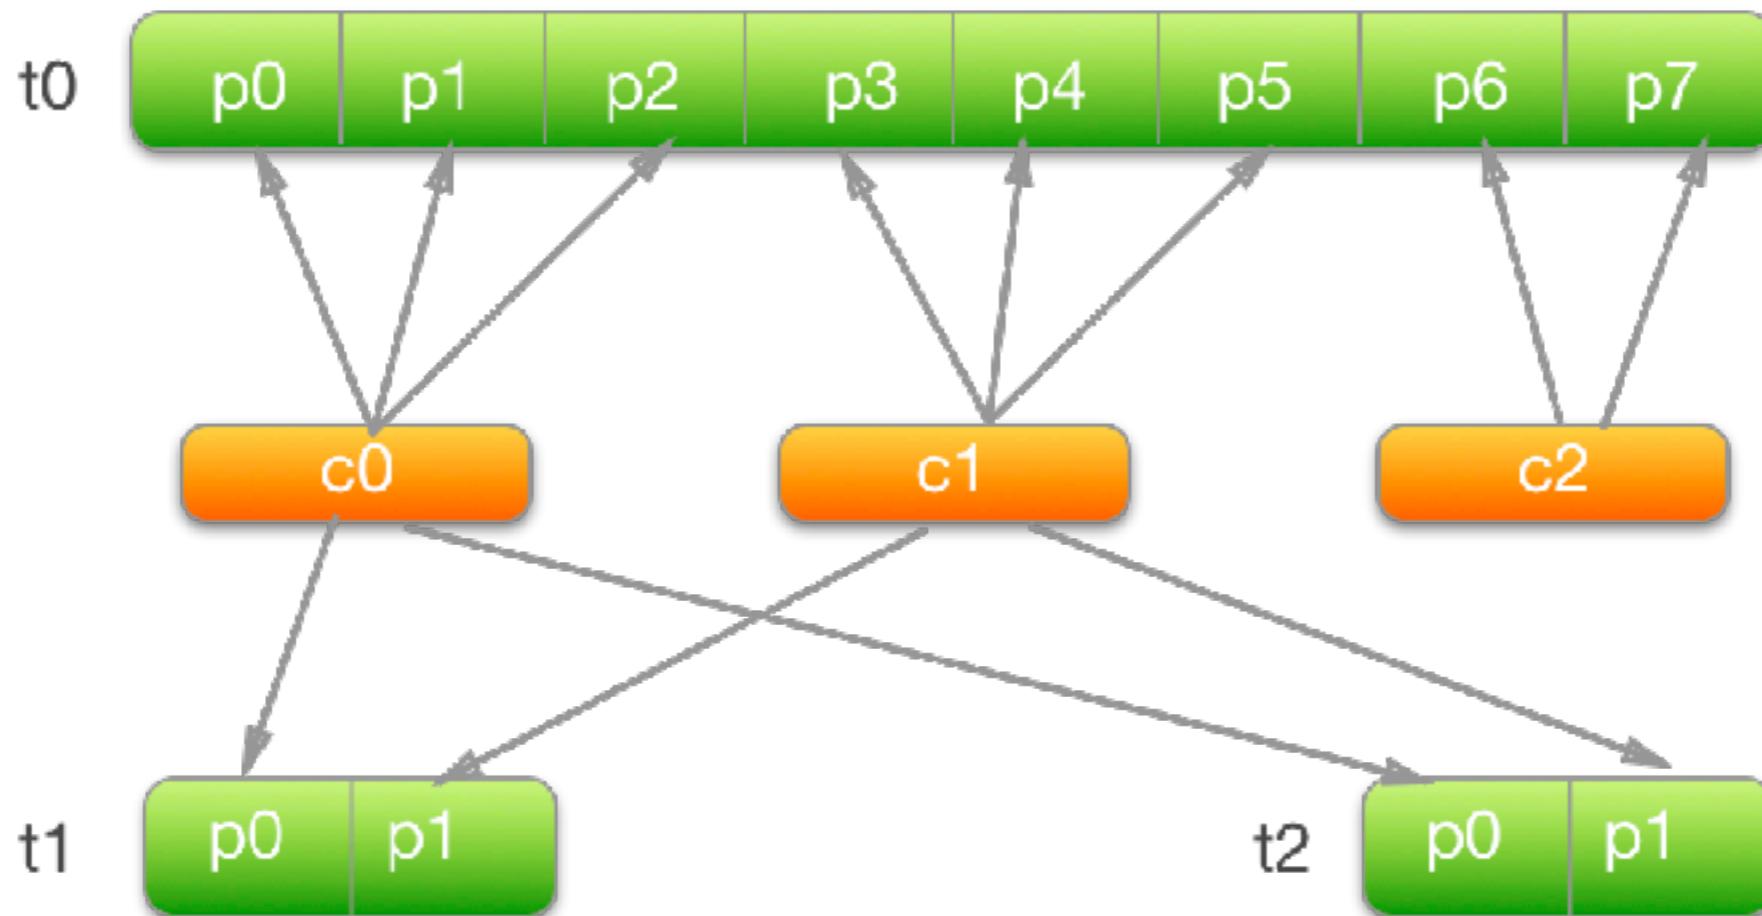
[org.apache.kafka.clients.consumer.RangeAssignor](#)



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Range assignor



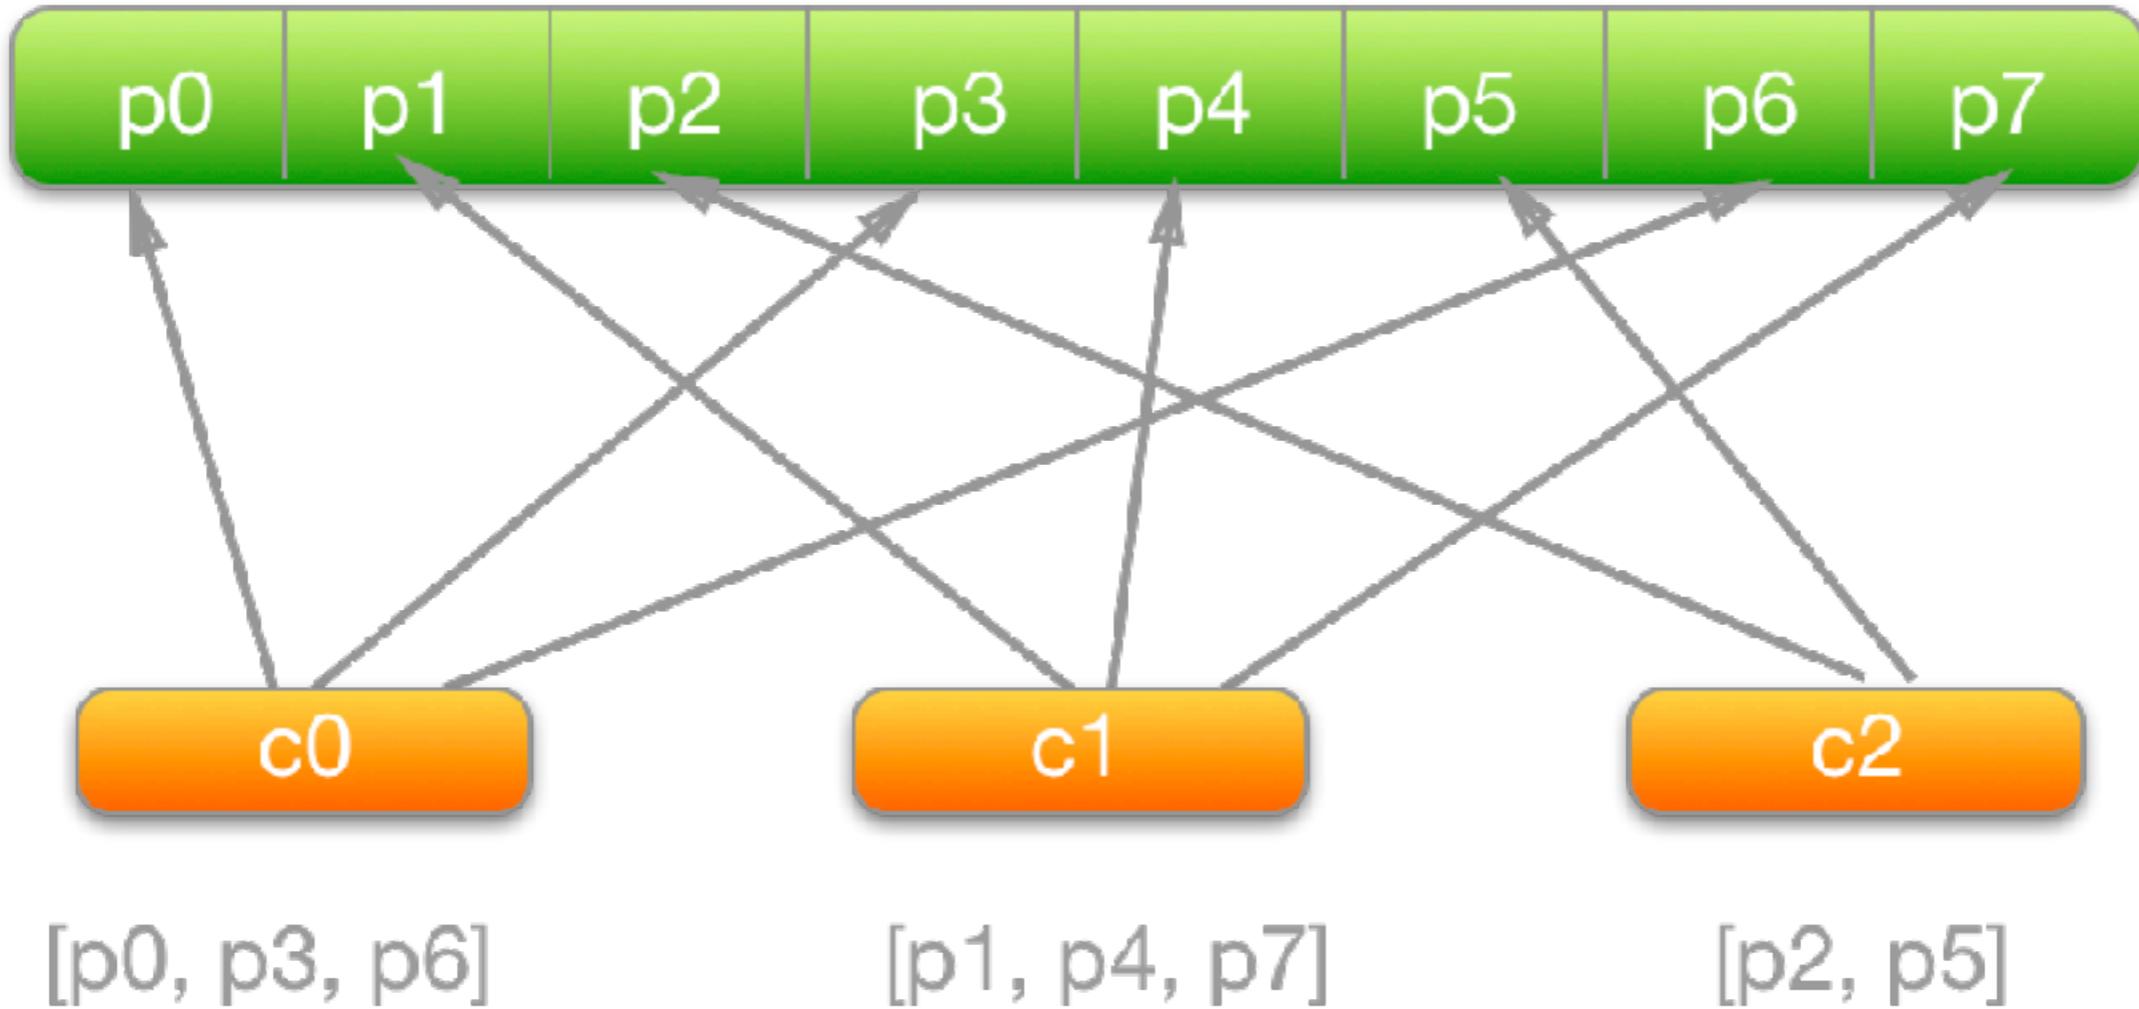
c0 -> [t0p0, t0p1, t0p2, t1p0, t2p0]

c1 -> [t0p3, t0p4, t0p5, t1p1, t2p1]

c2 -> [t0p6, t0p7]



# Round Robin assignor



*org.apache.kafka.clients.consumer.RoundRobinAssignor*

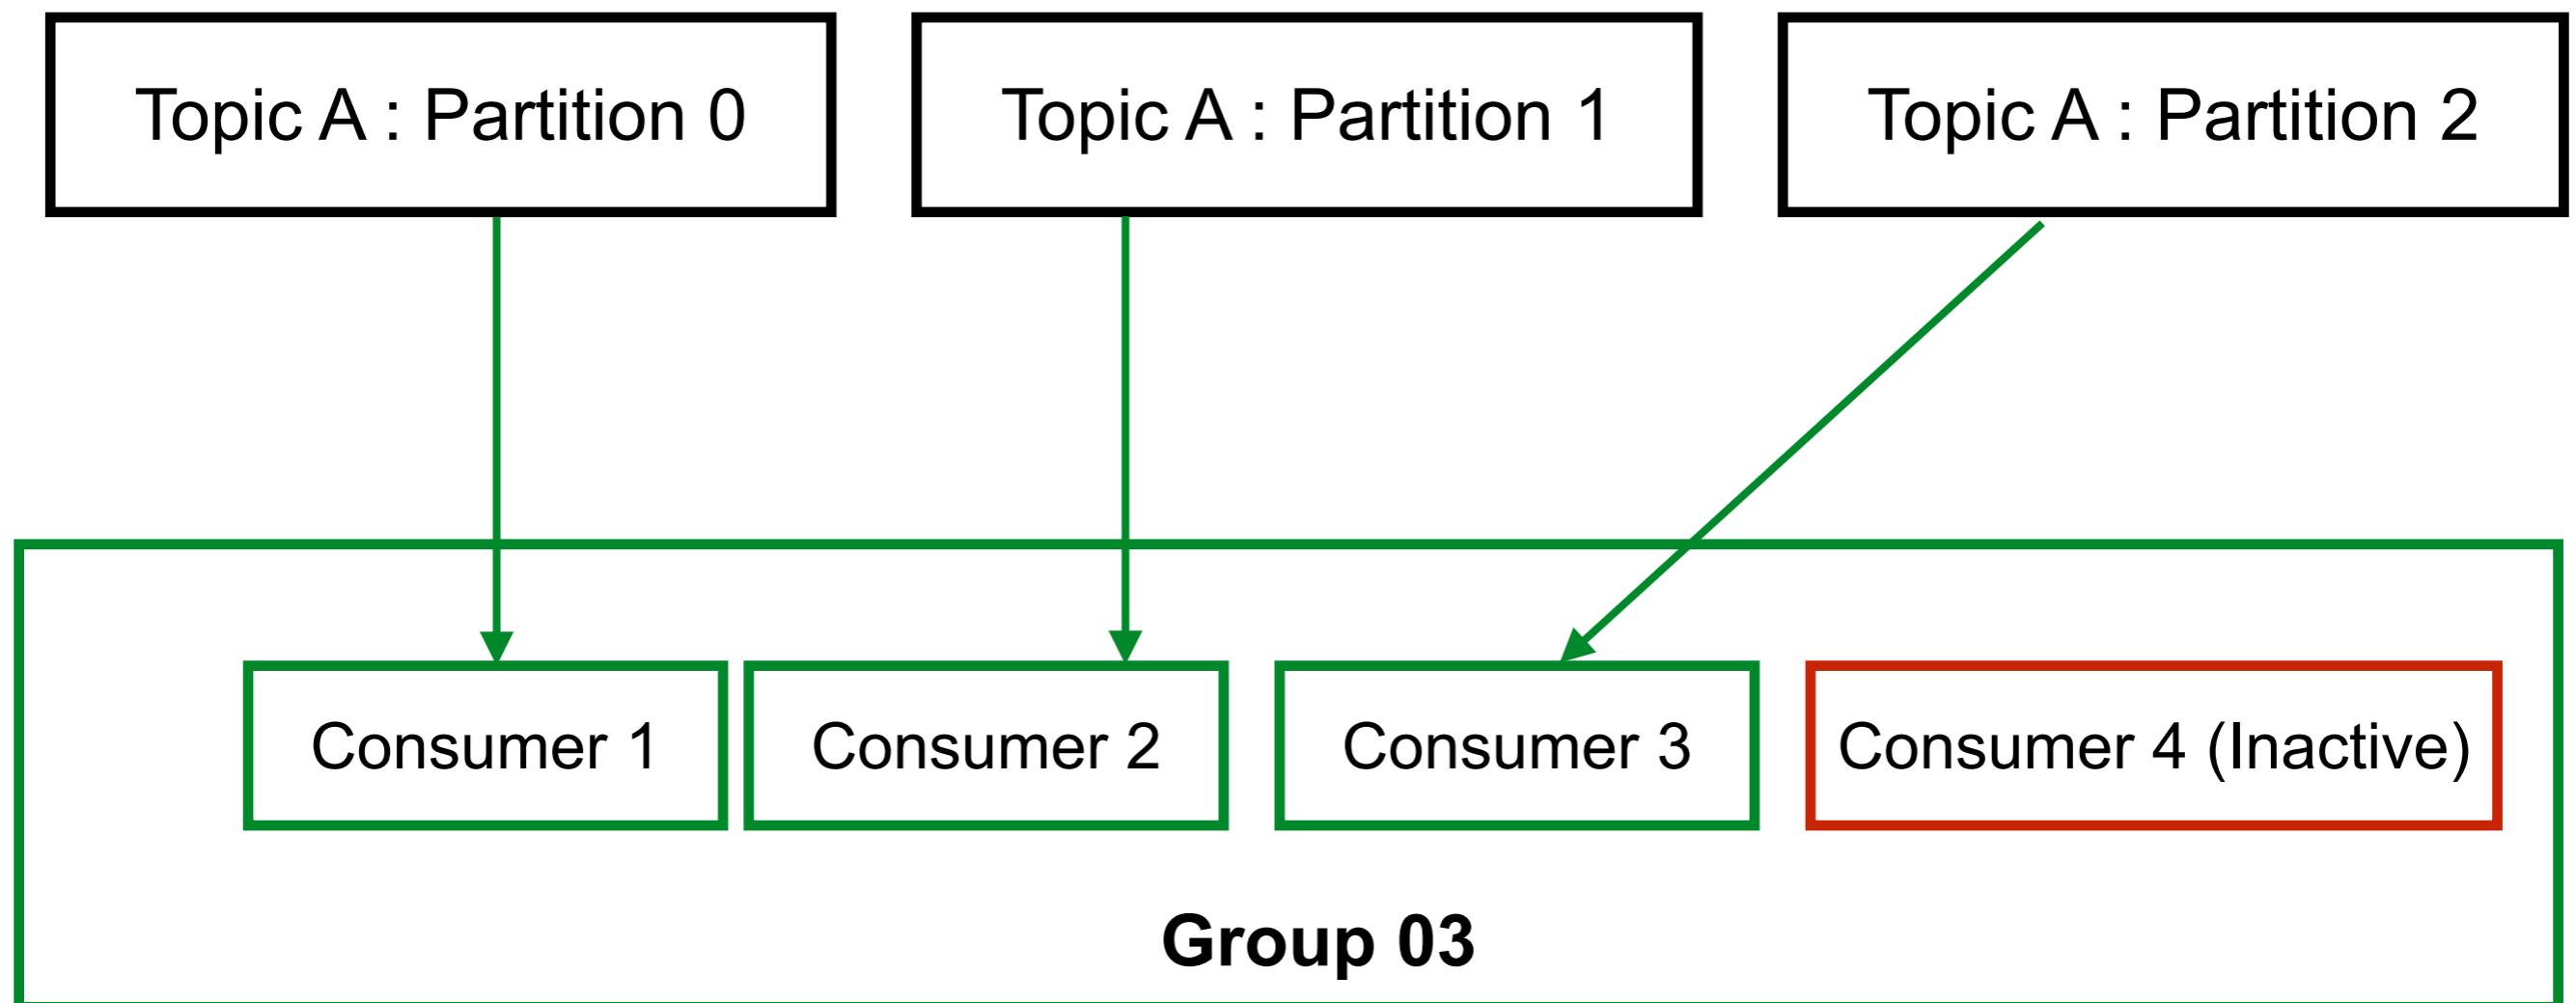


Kafka

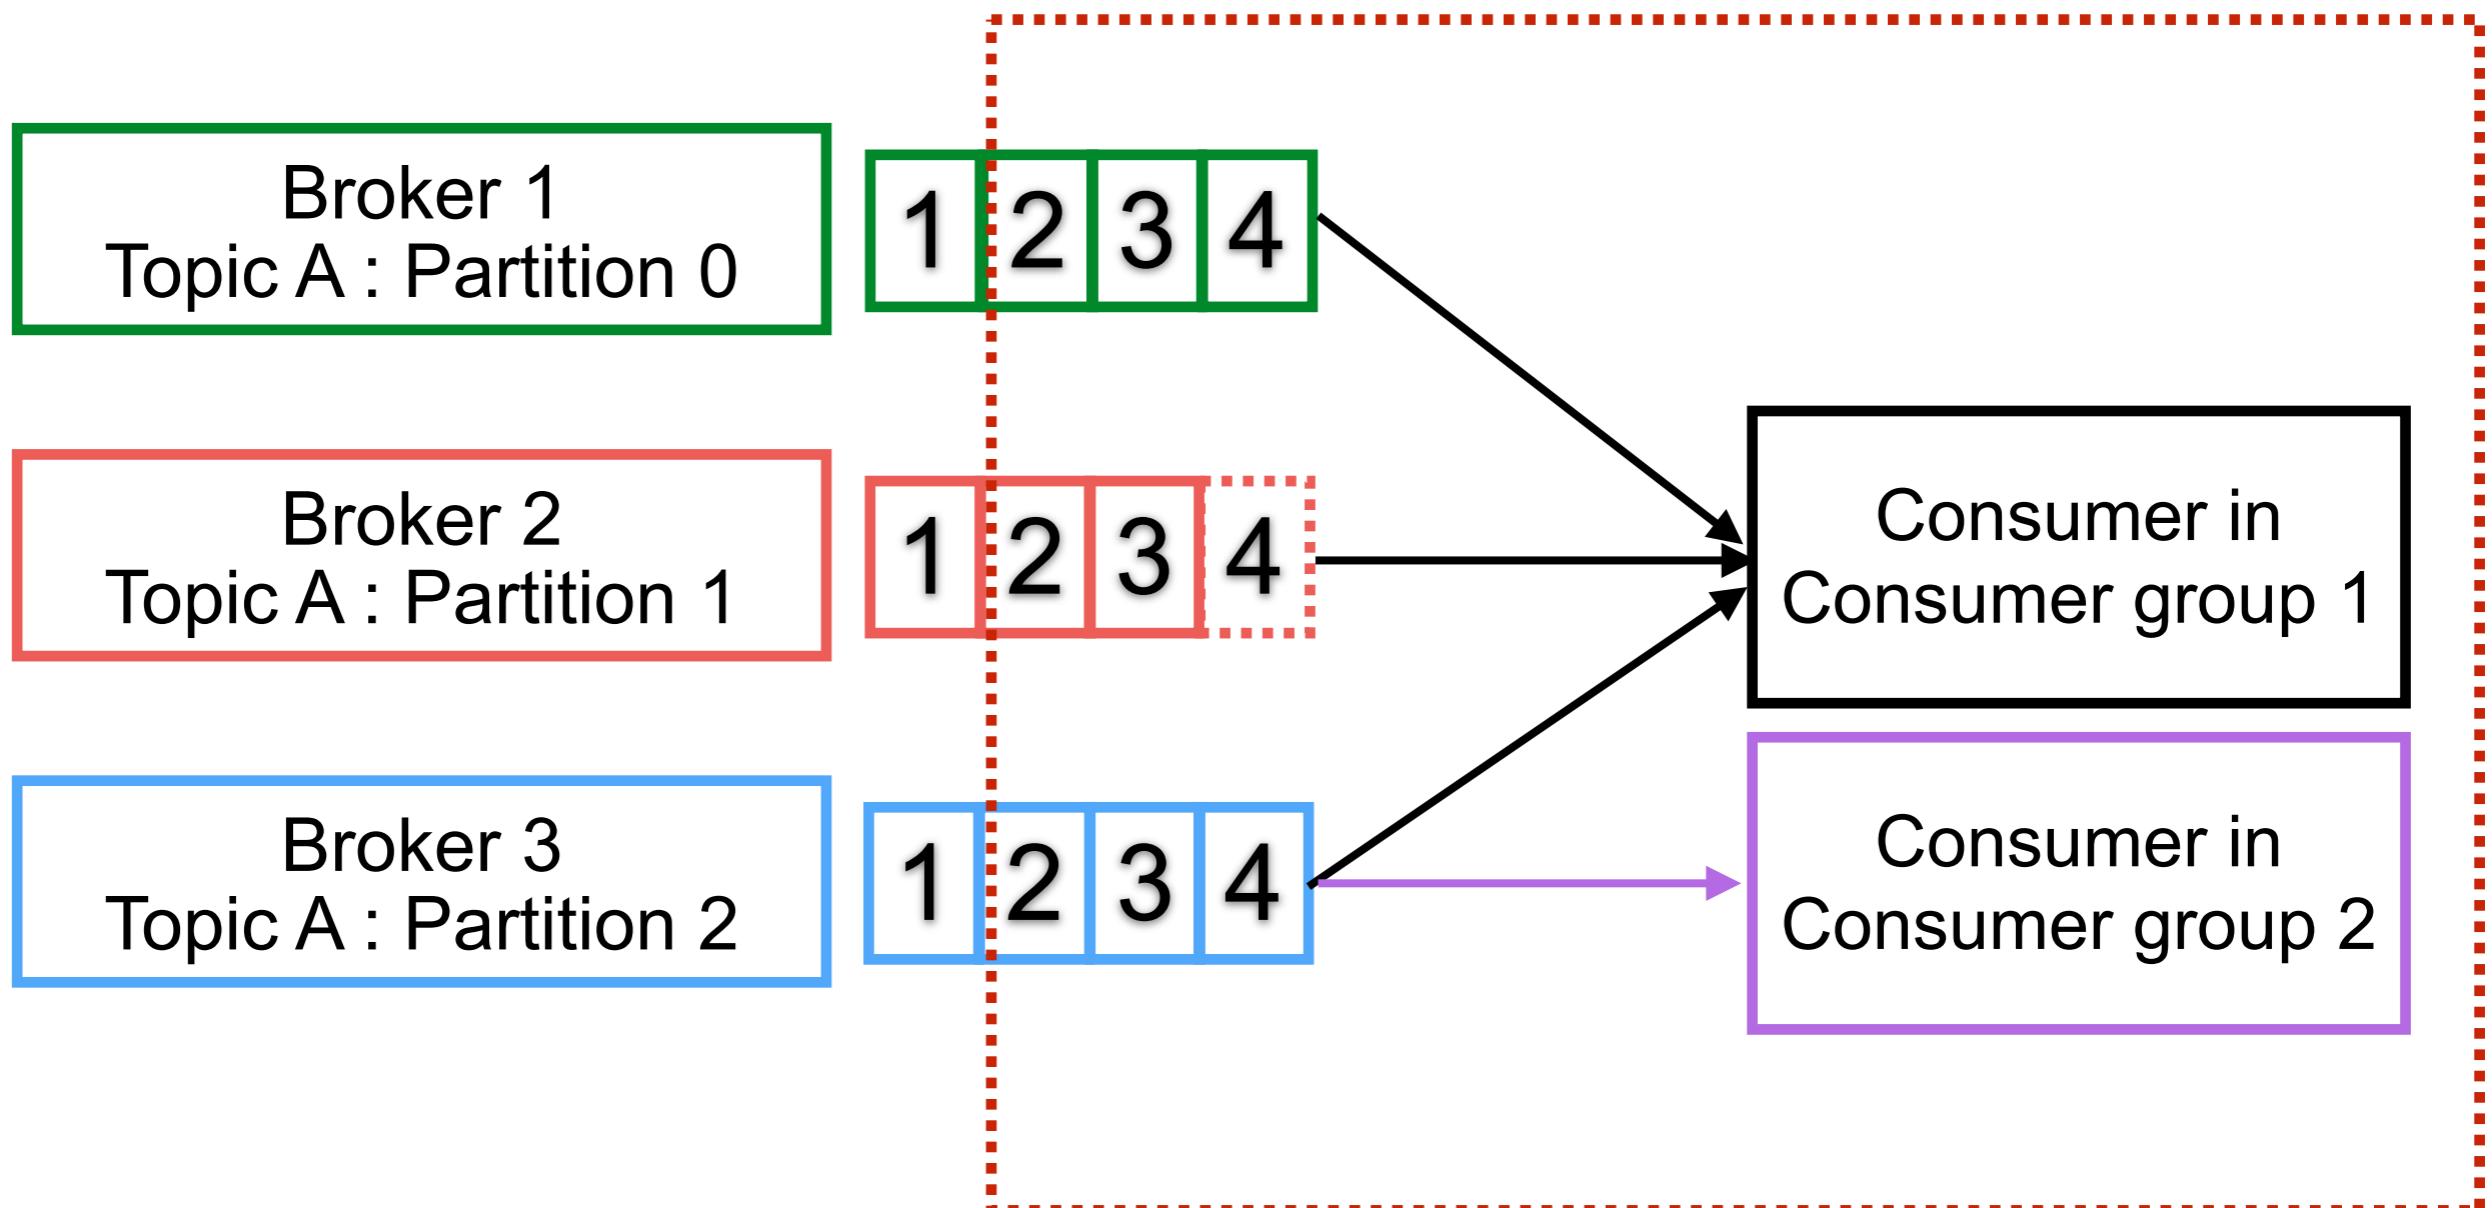
© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# IF Consumers > partitions ?

Some consumers will **inactive**



# Consumers offsets



# Consumers offsets

Kafka store the offset at which a consumer group has been reading

The offsets committed live in topic named  
“`_consumer_offsets`”

When consumer in a group has processed data received from Kafka,  
it should be **committing the offsets**



# When to commit the offset ?



# Delivery semantics for consumer

At most once

At lease once (preferred)

Exactly once

<https://docs.confluent.io/kafka/design/delivery-semantics.html>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# 1. At most once

Offsets are committed as soon as the message is received

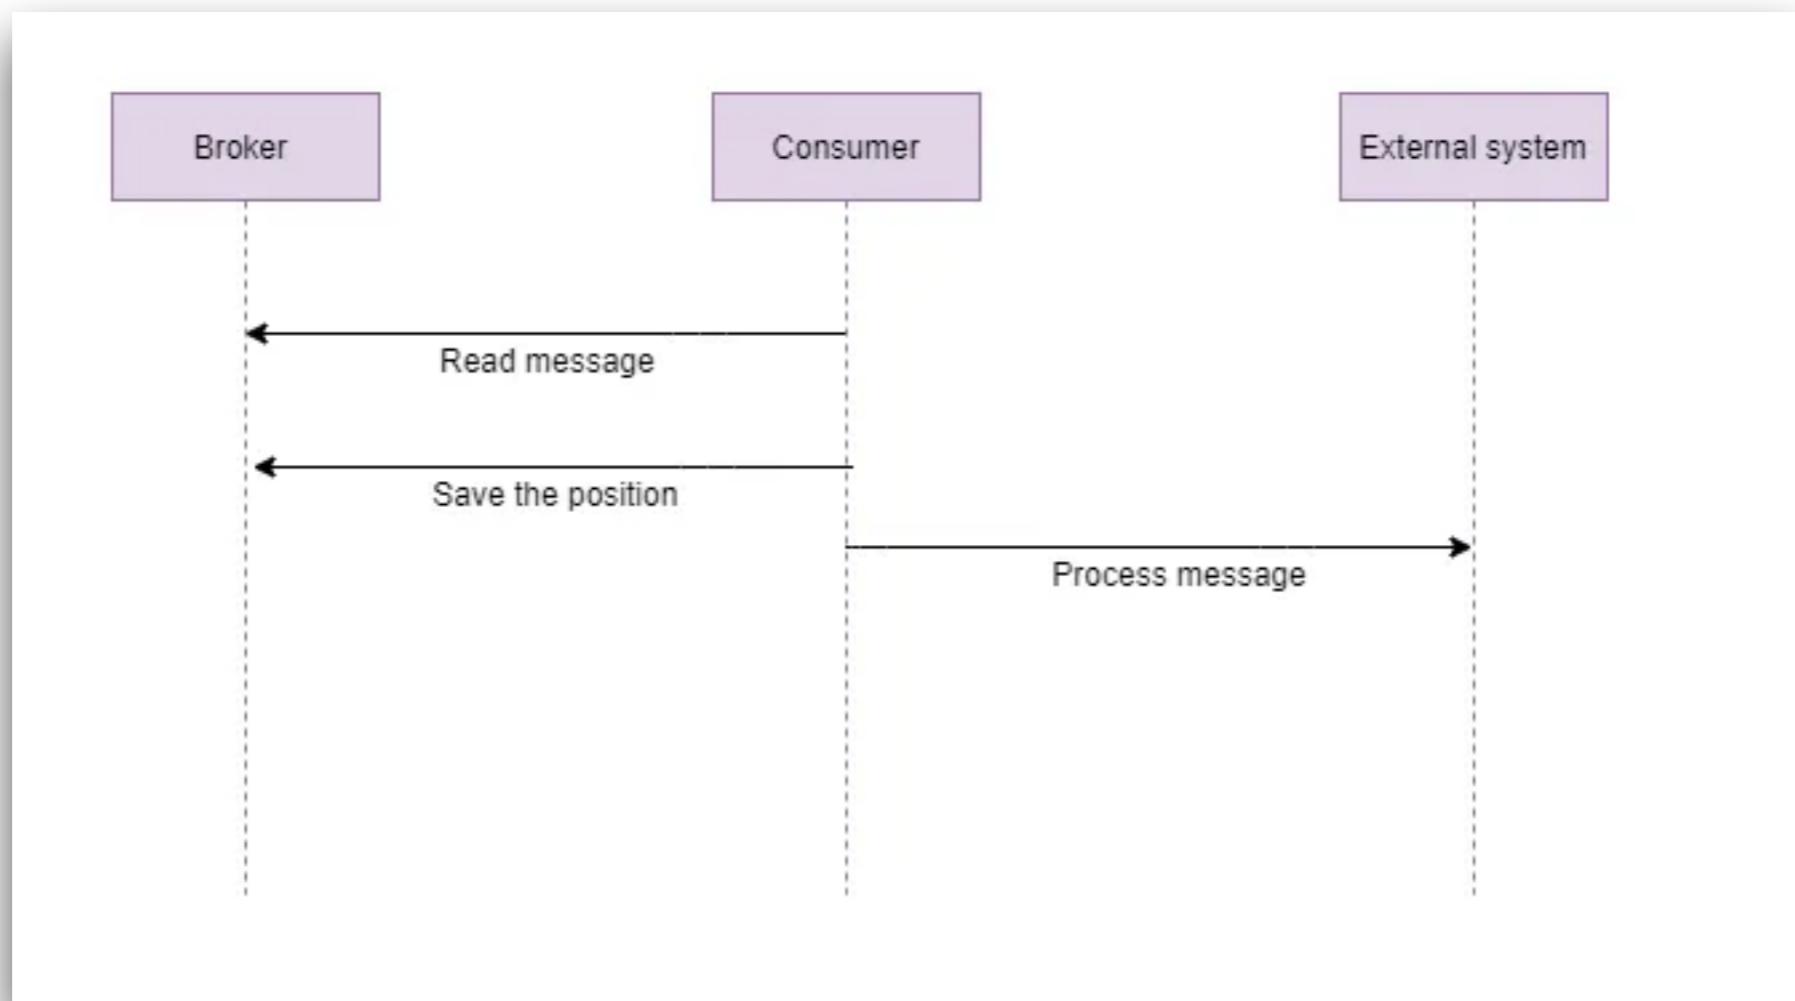
If processing go wrong, the message will be **loss!!**

**Fire and forgot pattern**

Lowest latency



# At most once

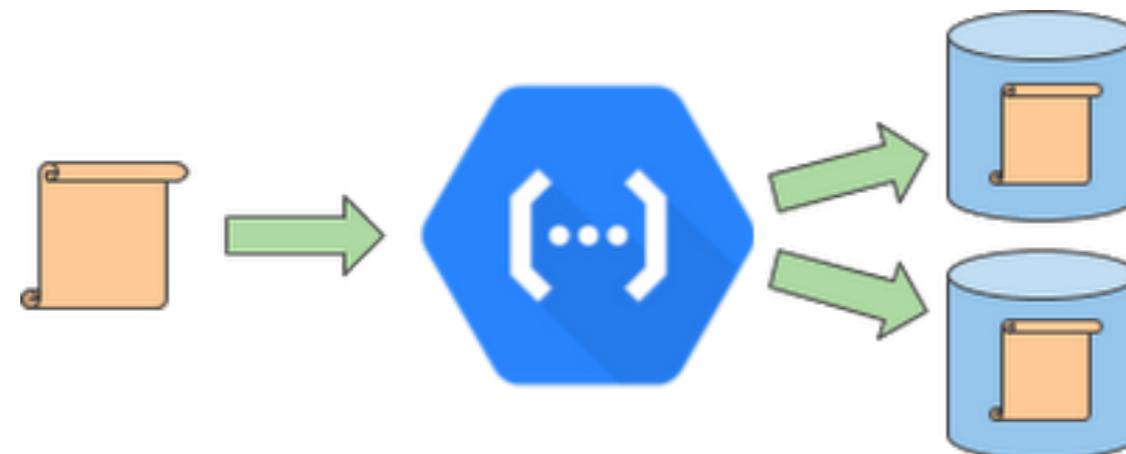


## 2. At lease once (1)

Offsets are committed after the message is processed

Messages are never lost but may be **redelivered**

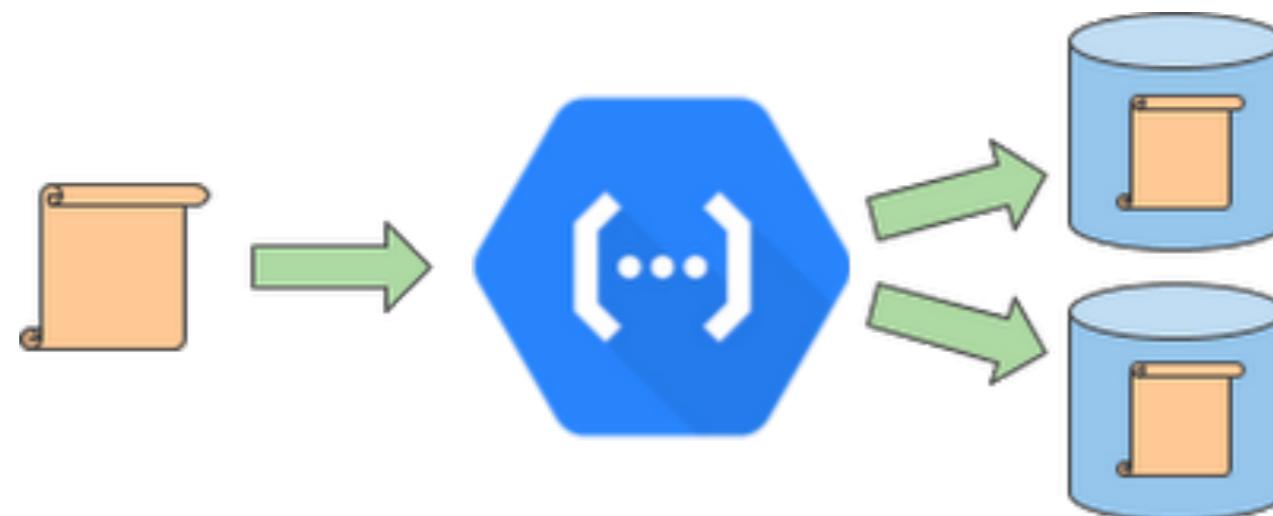
If processing go wrong, the message will be **read again**



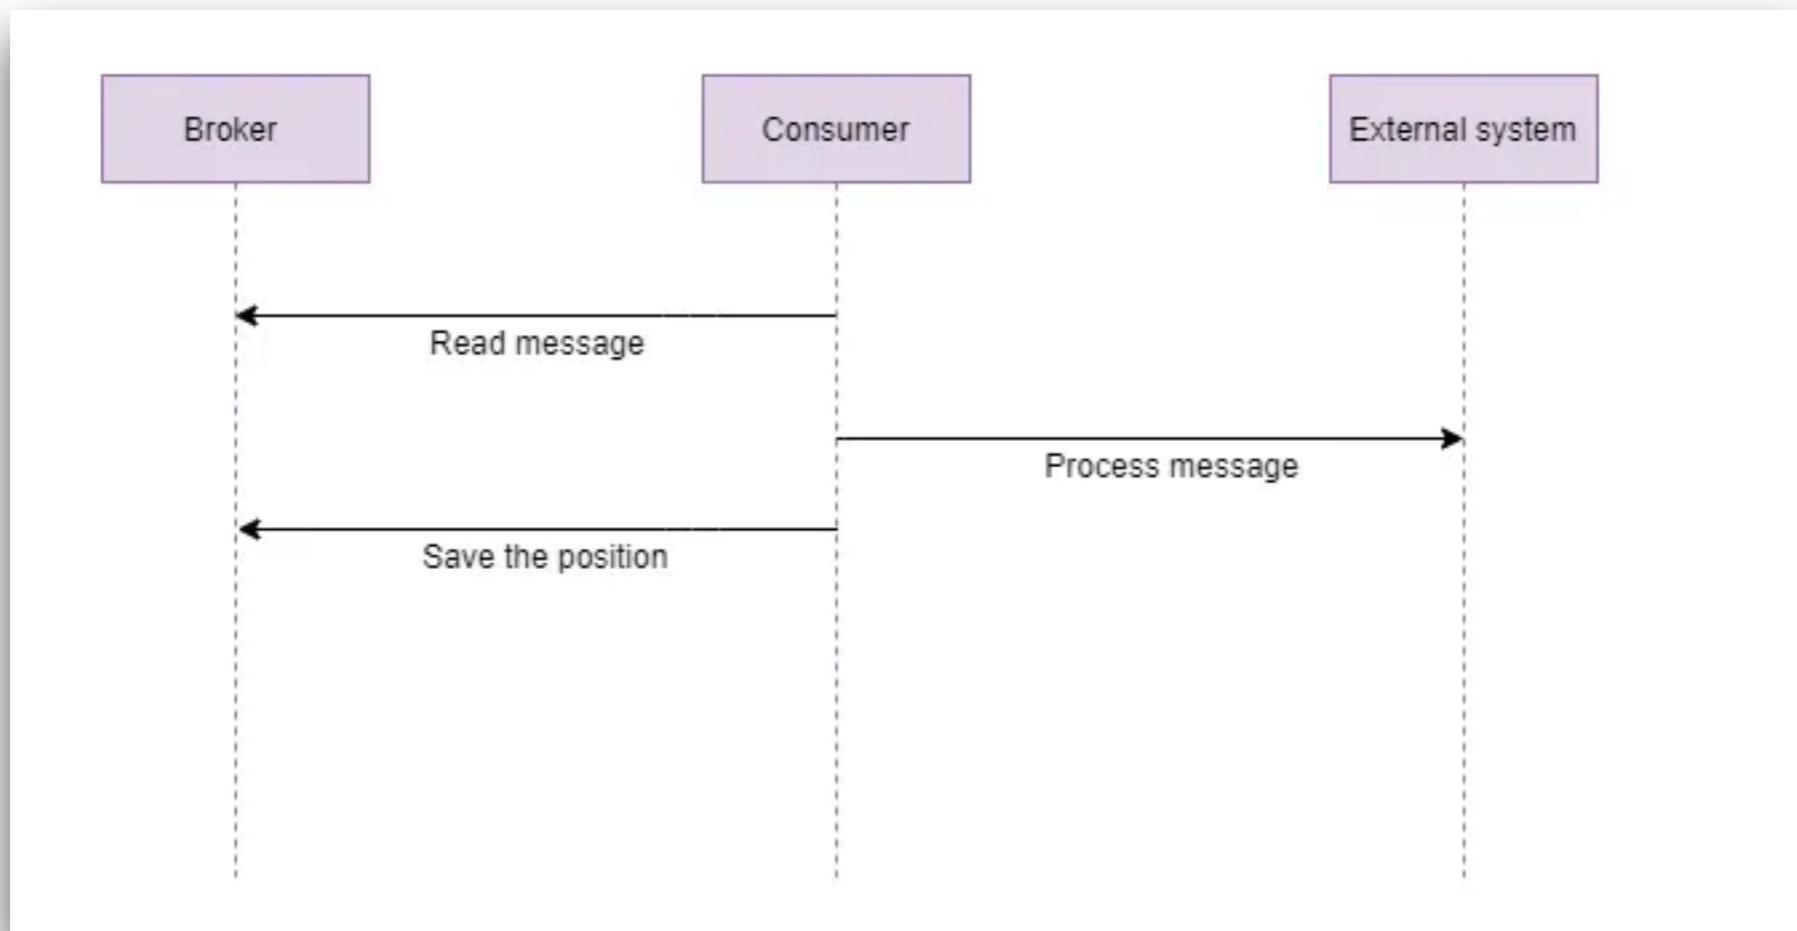
## 2. At lease once (2)

Make sure your processing is **idempotent**

*Processing again the message not impact to your system !!*



# At lease once



# 3. Exactly once

Each message is delivered once and only once  
Can be achieved for Kafka (Workflow, Stream API)

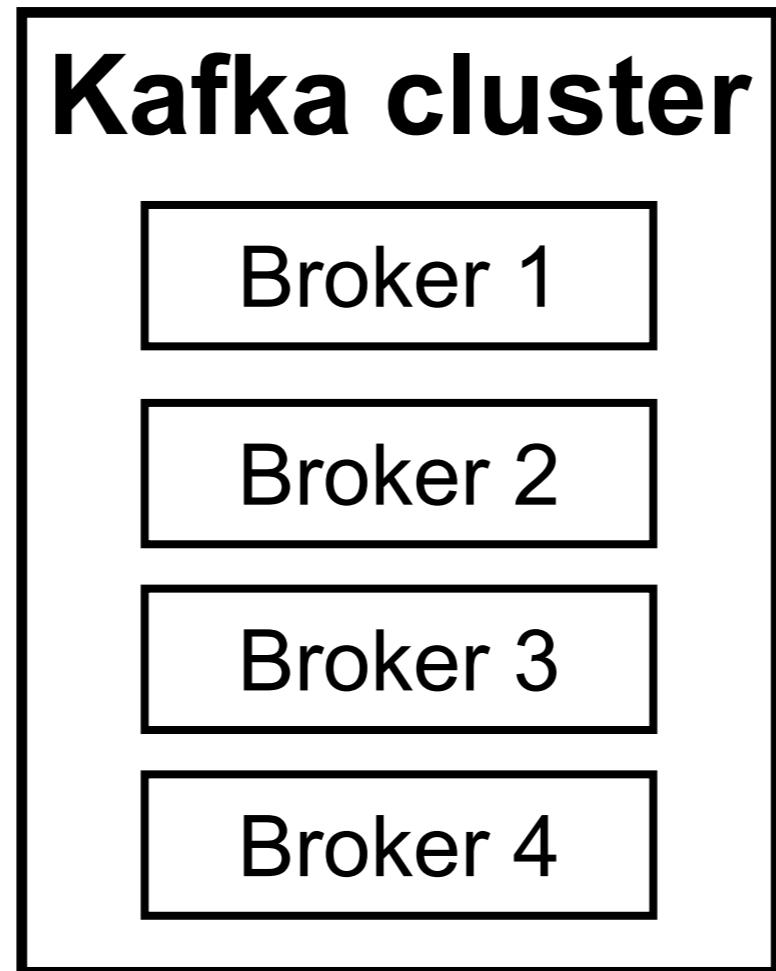
Transactional delivery  
Resent with idempotent



# Summary of Kafka



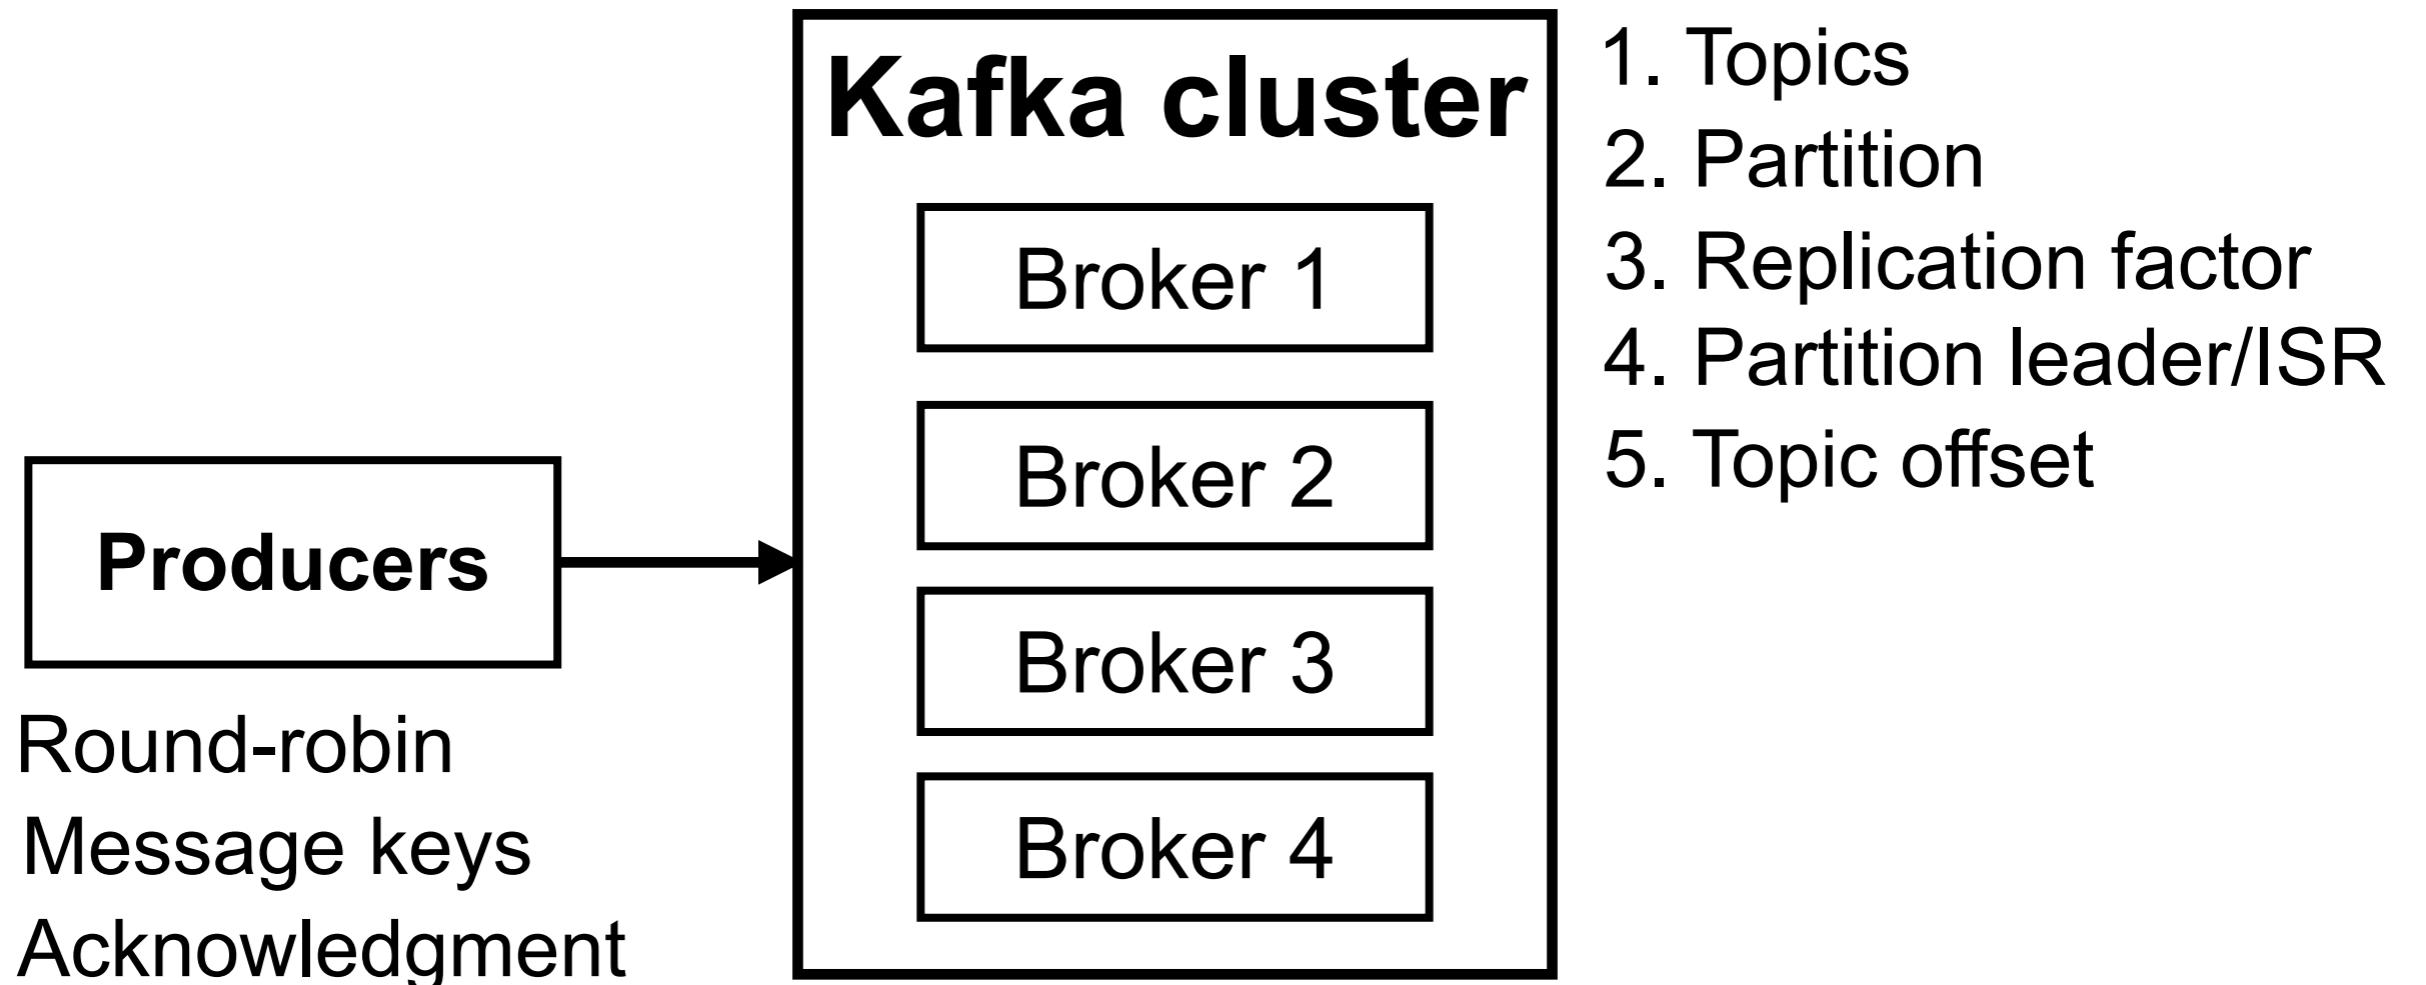
# Kafka concepts



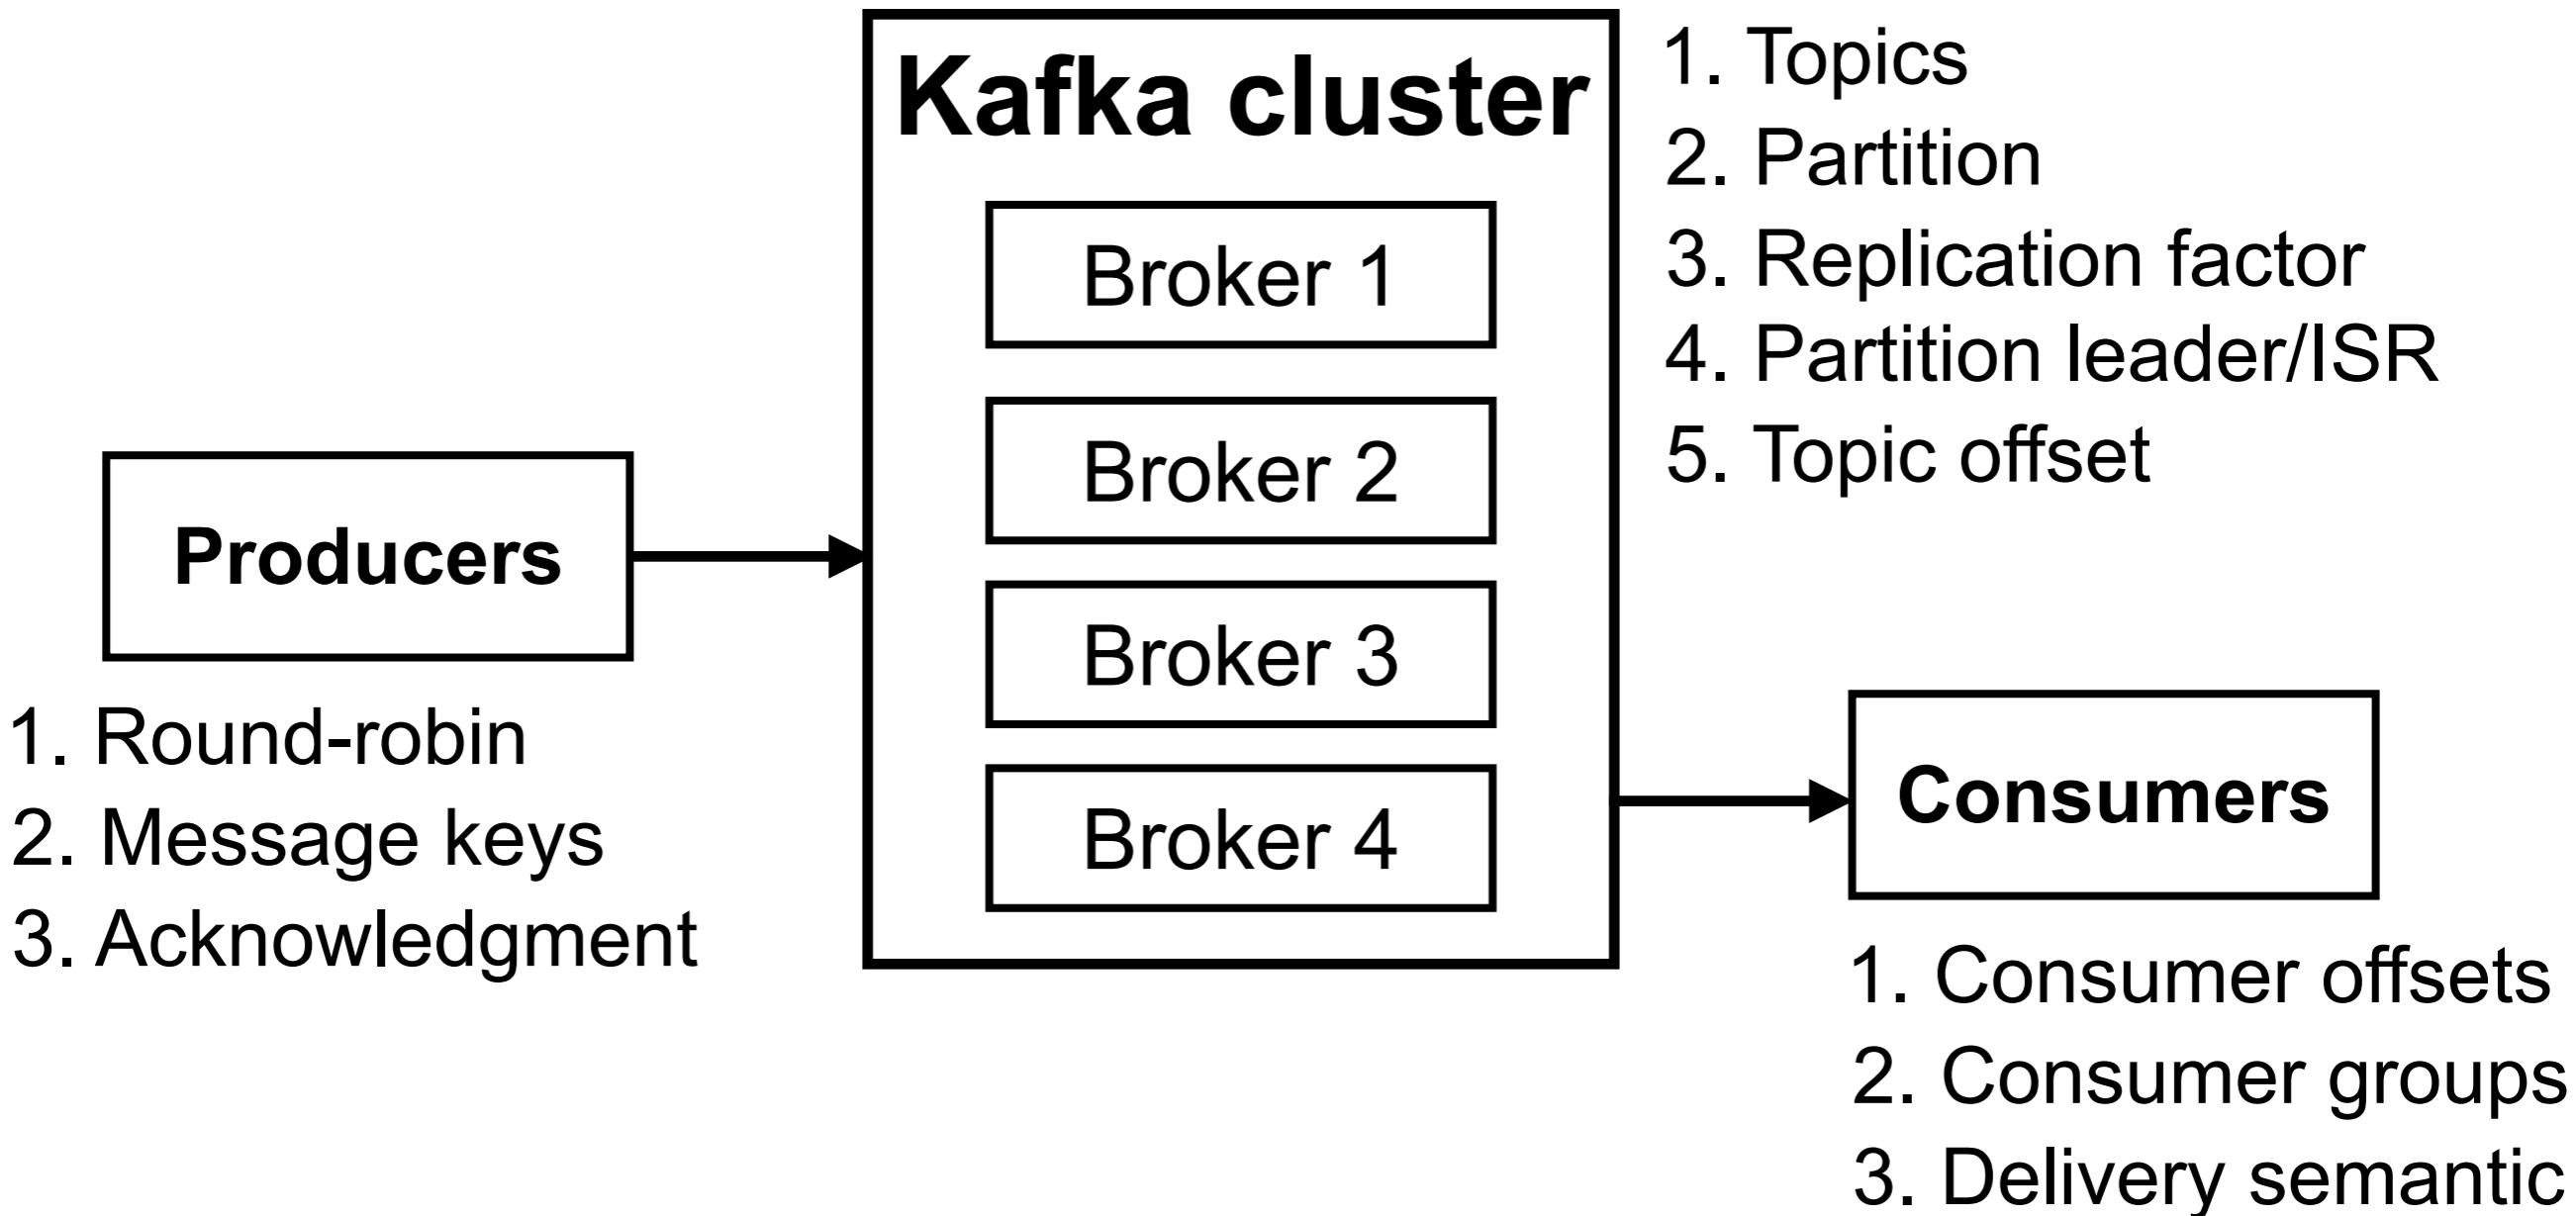
1. Topics
2. Partition
3. Replication factor
4. Partition leader/ISR
5. Topic offset



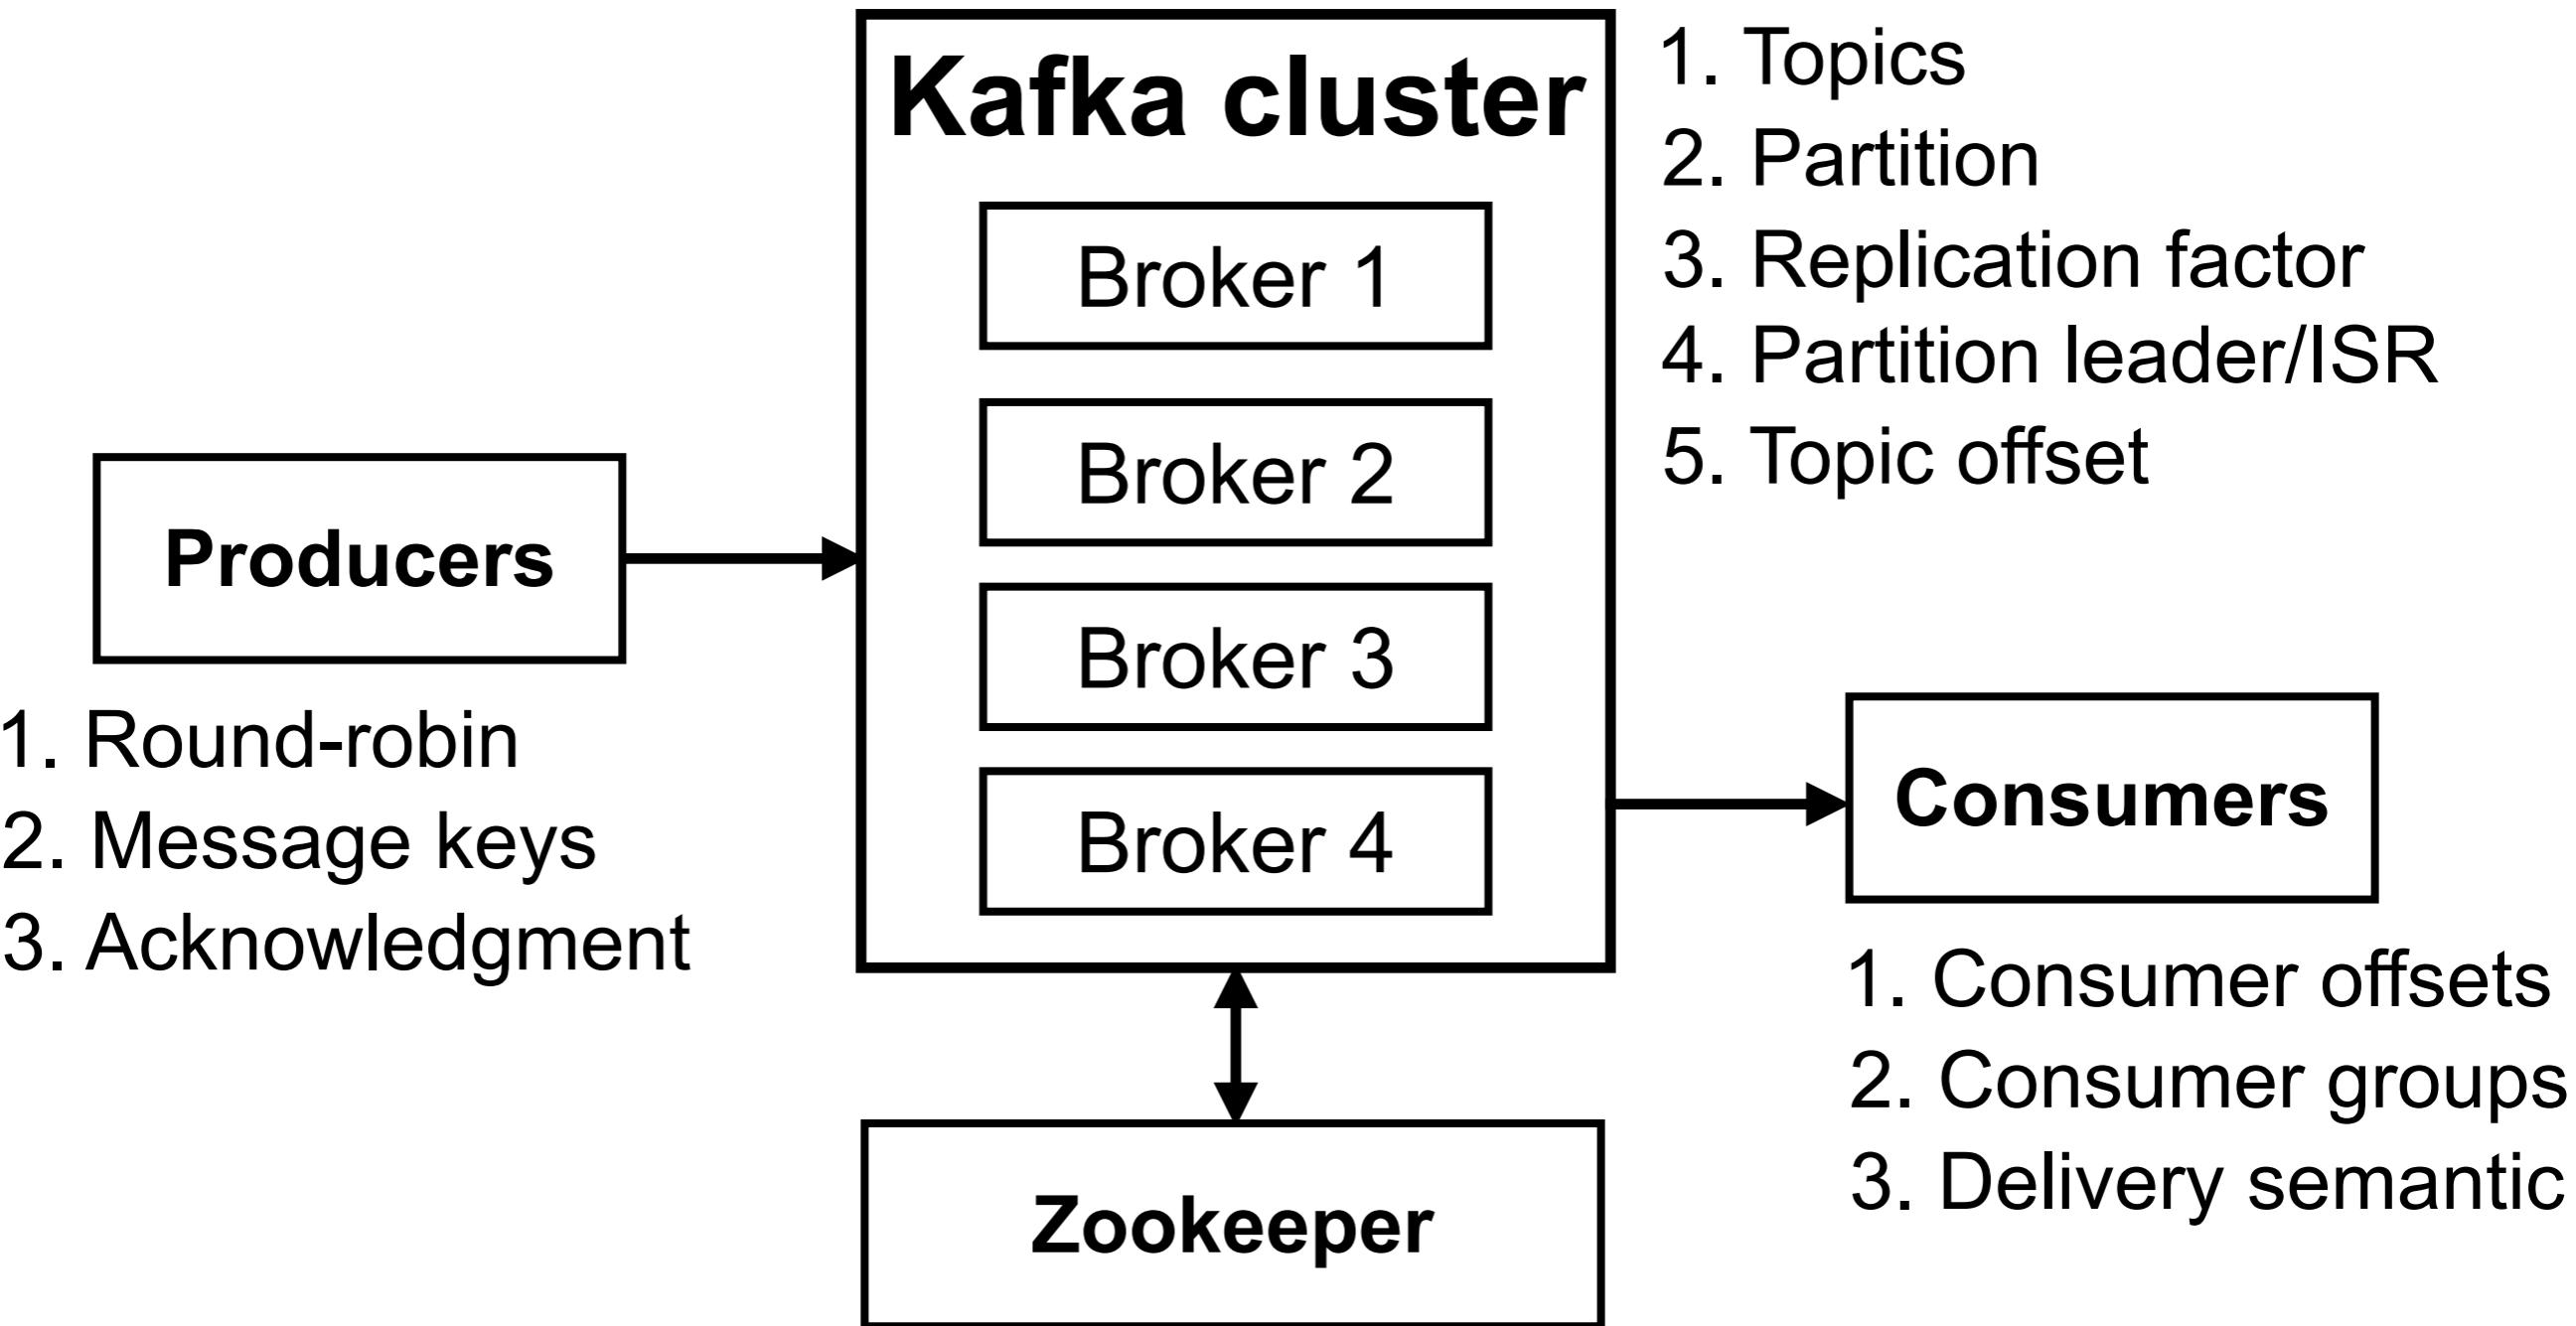
# Kafka concepts



# Kafka concepts



# Kafka concepts

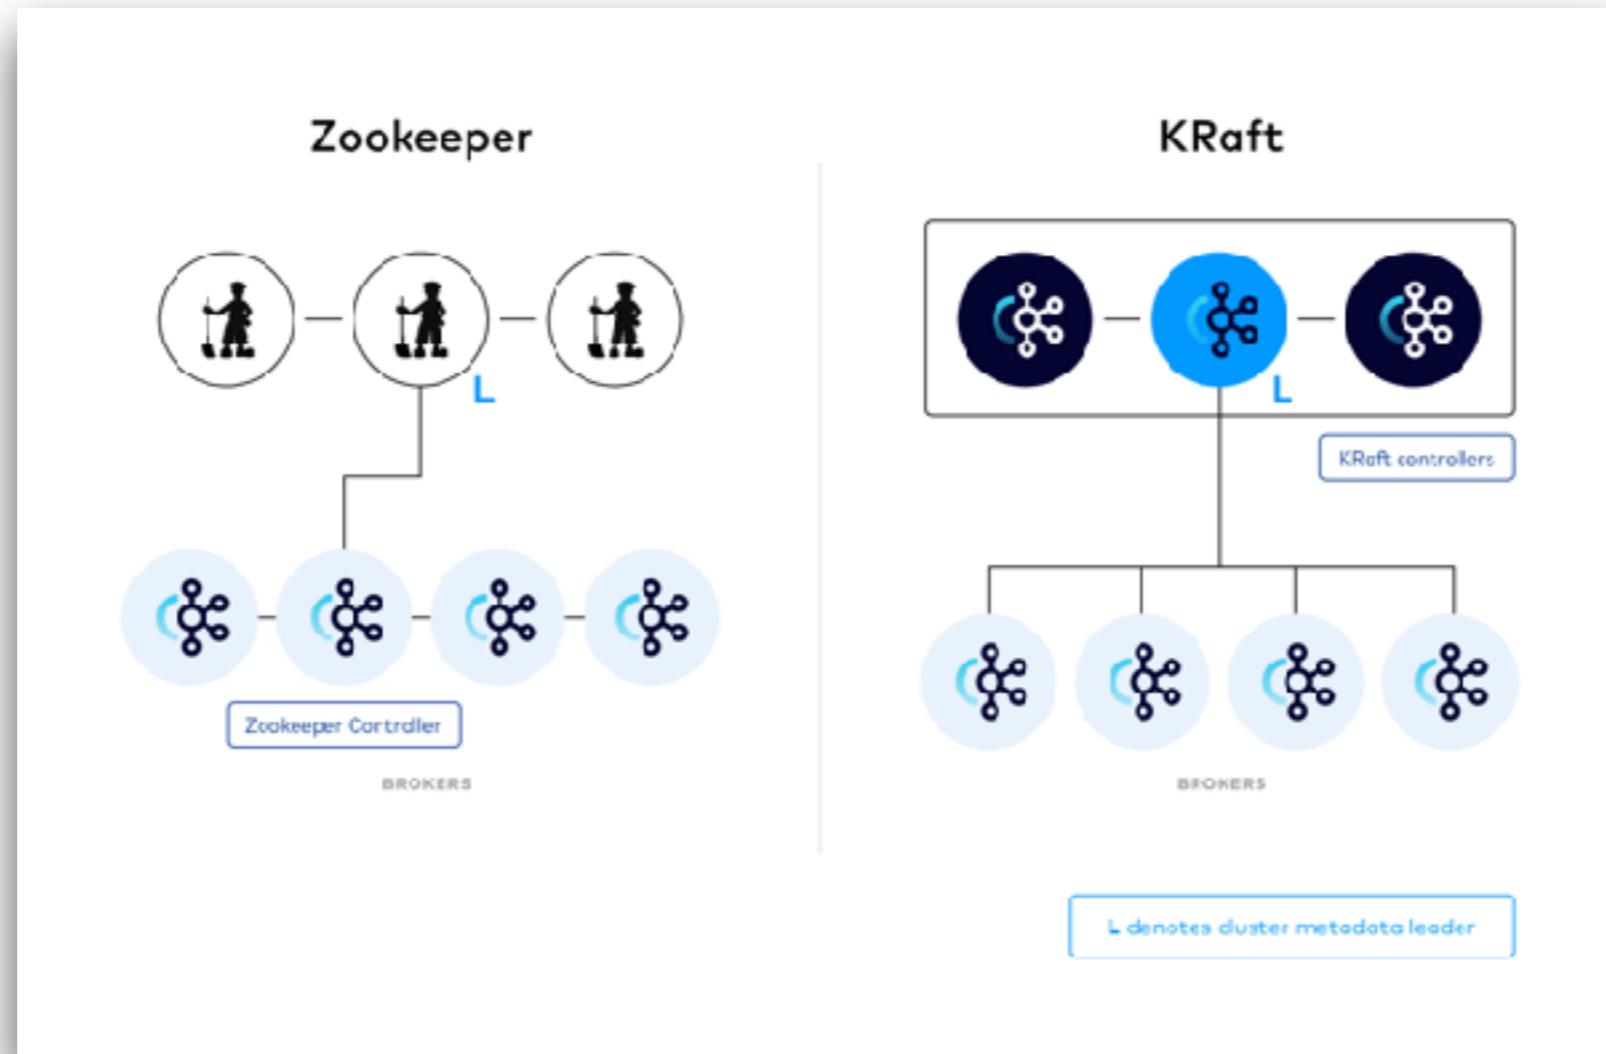


# Kraft



# Kafka without Zookeeper (Kraft)

## Kafka Raft



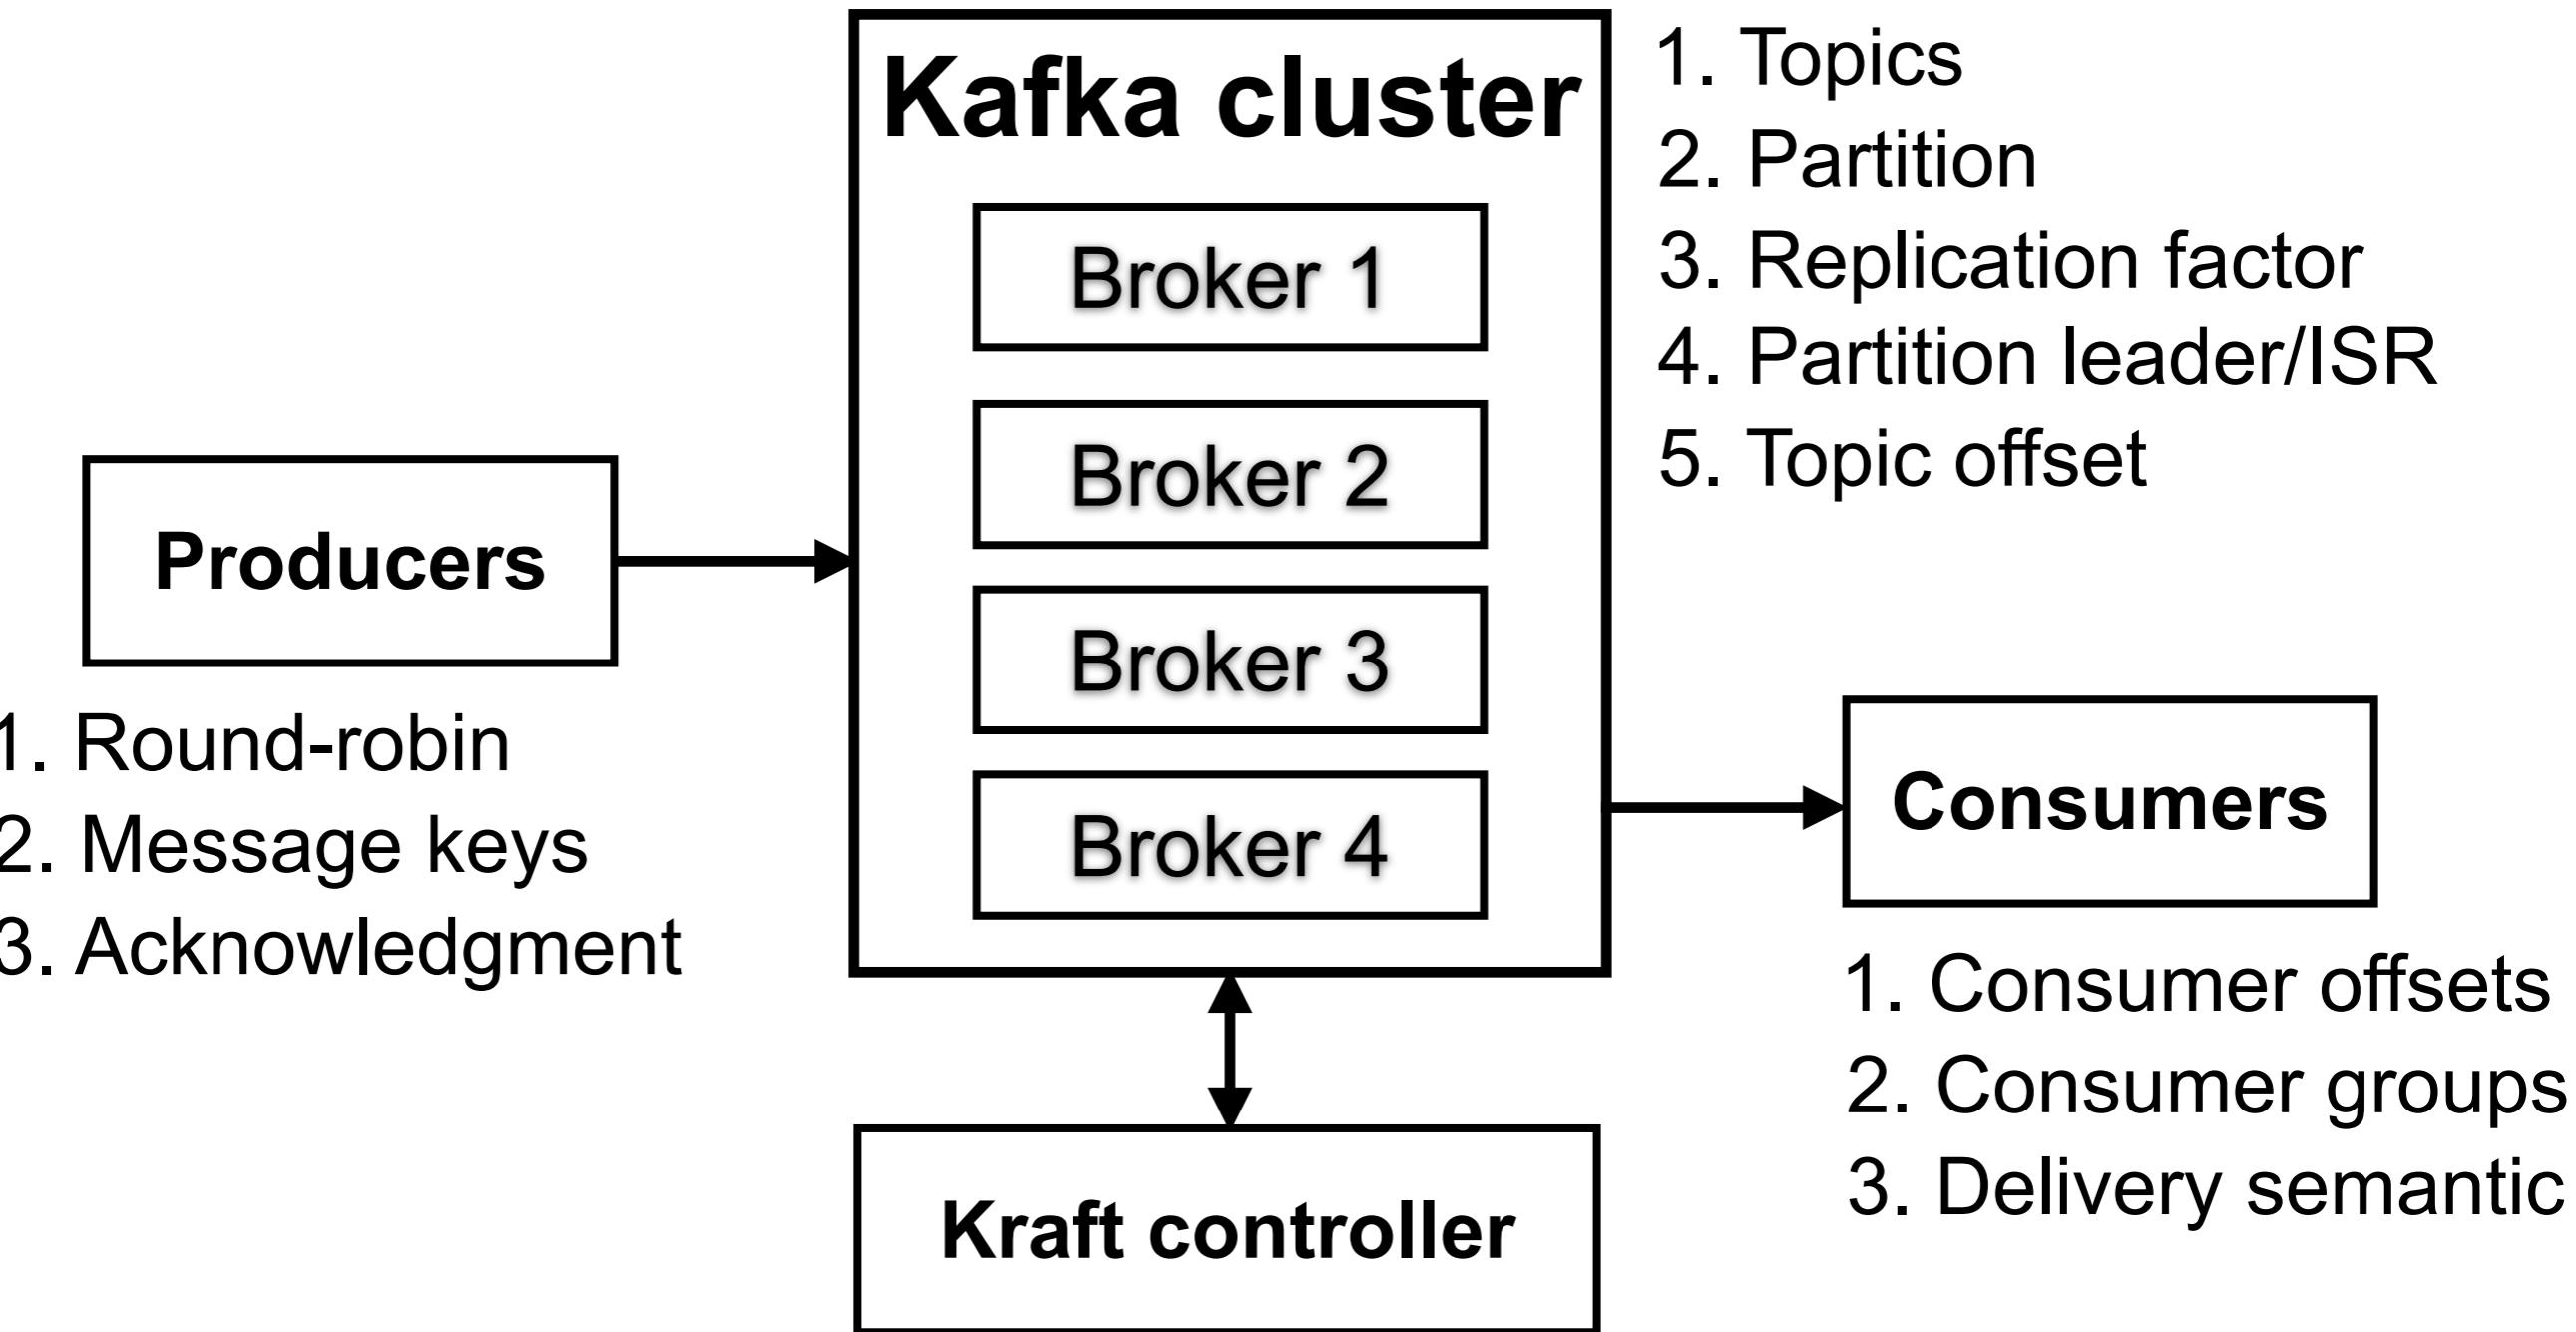
<https://developer.confluent.io/learn/kraft/>



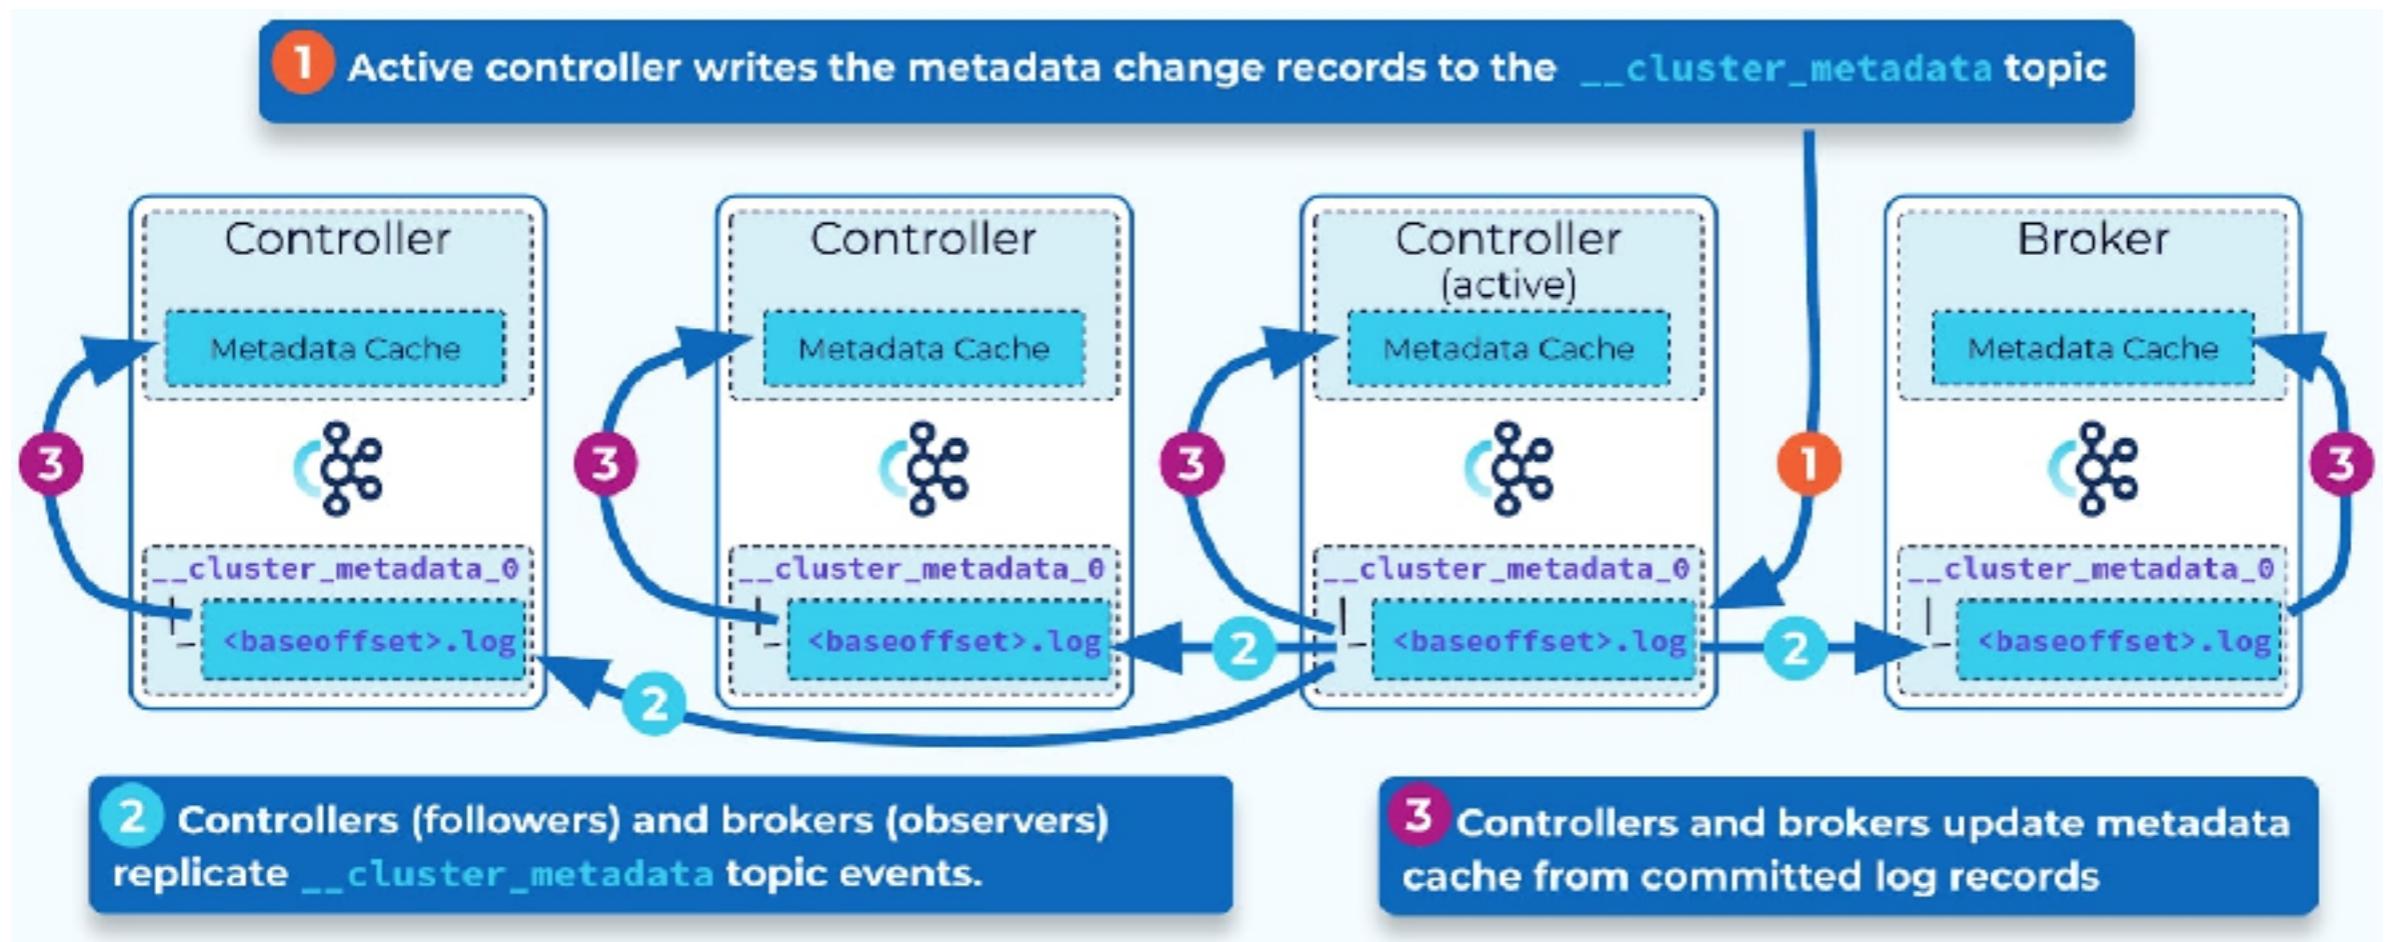
Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Kraft



# Kraft



<https://www.confluent.io/fr-fr/blog/what-is-kraft-and-how-do-you-use-it/>



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# Working with Schema Registry



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

99

# Schema Registry

Server process external to Kafka brokers

Maintain a database of schemas

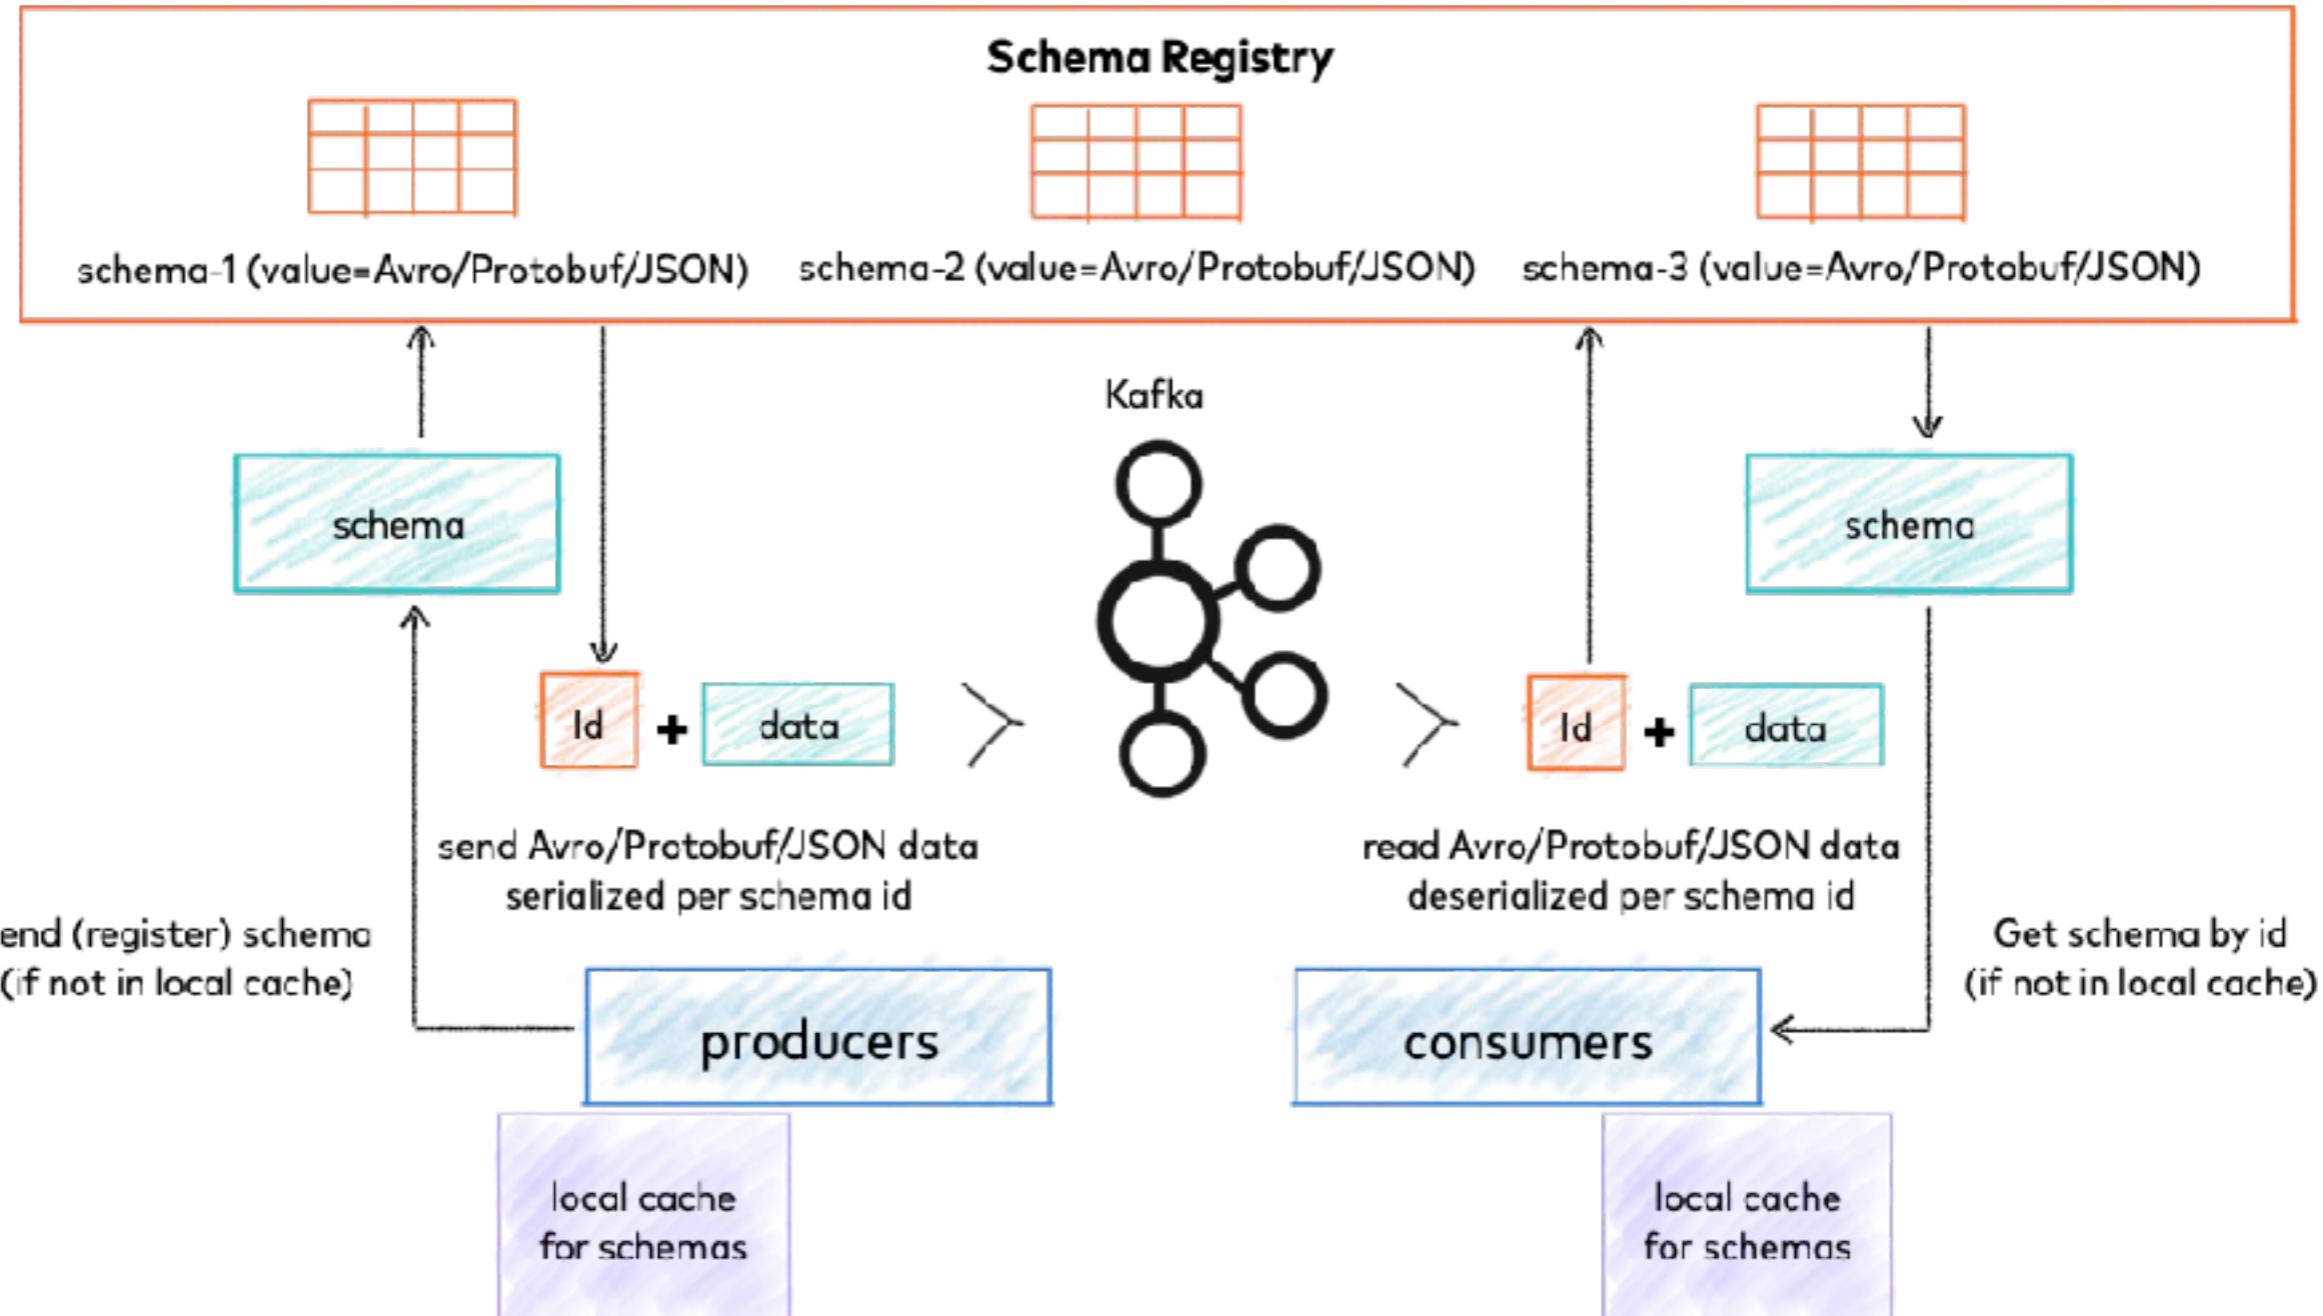
Provider APIs for producer and consumer

Compatible of message in producer and consumer

**Schema change !!  
Validation ...**



# Schema Registry



<https://docs.confluent.io/platform/current/schema-registry/>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Steps in Schema Registry

Define schema compatibility rules per topic

Producer API prevents incompatible messages

Consumer API prevents incompatible messages



# Supported formats

JSON  
Schema

Avro

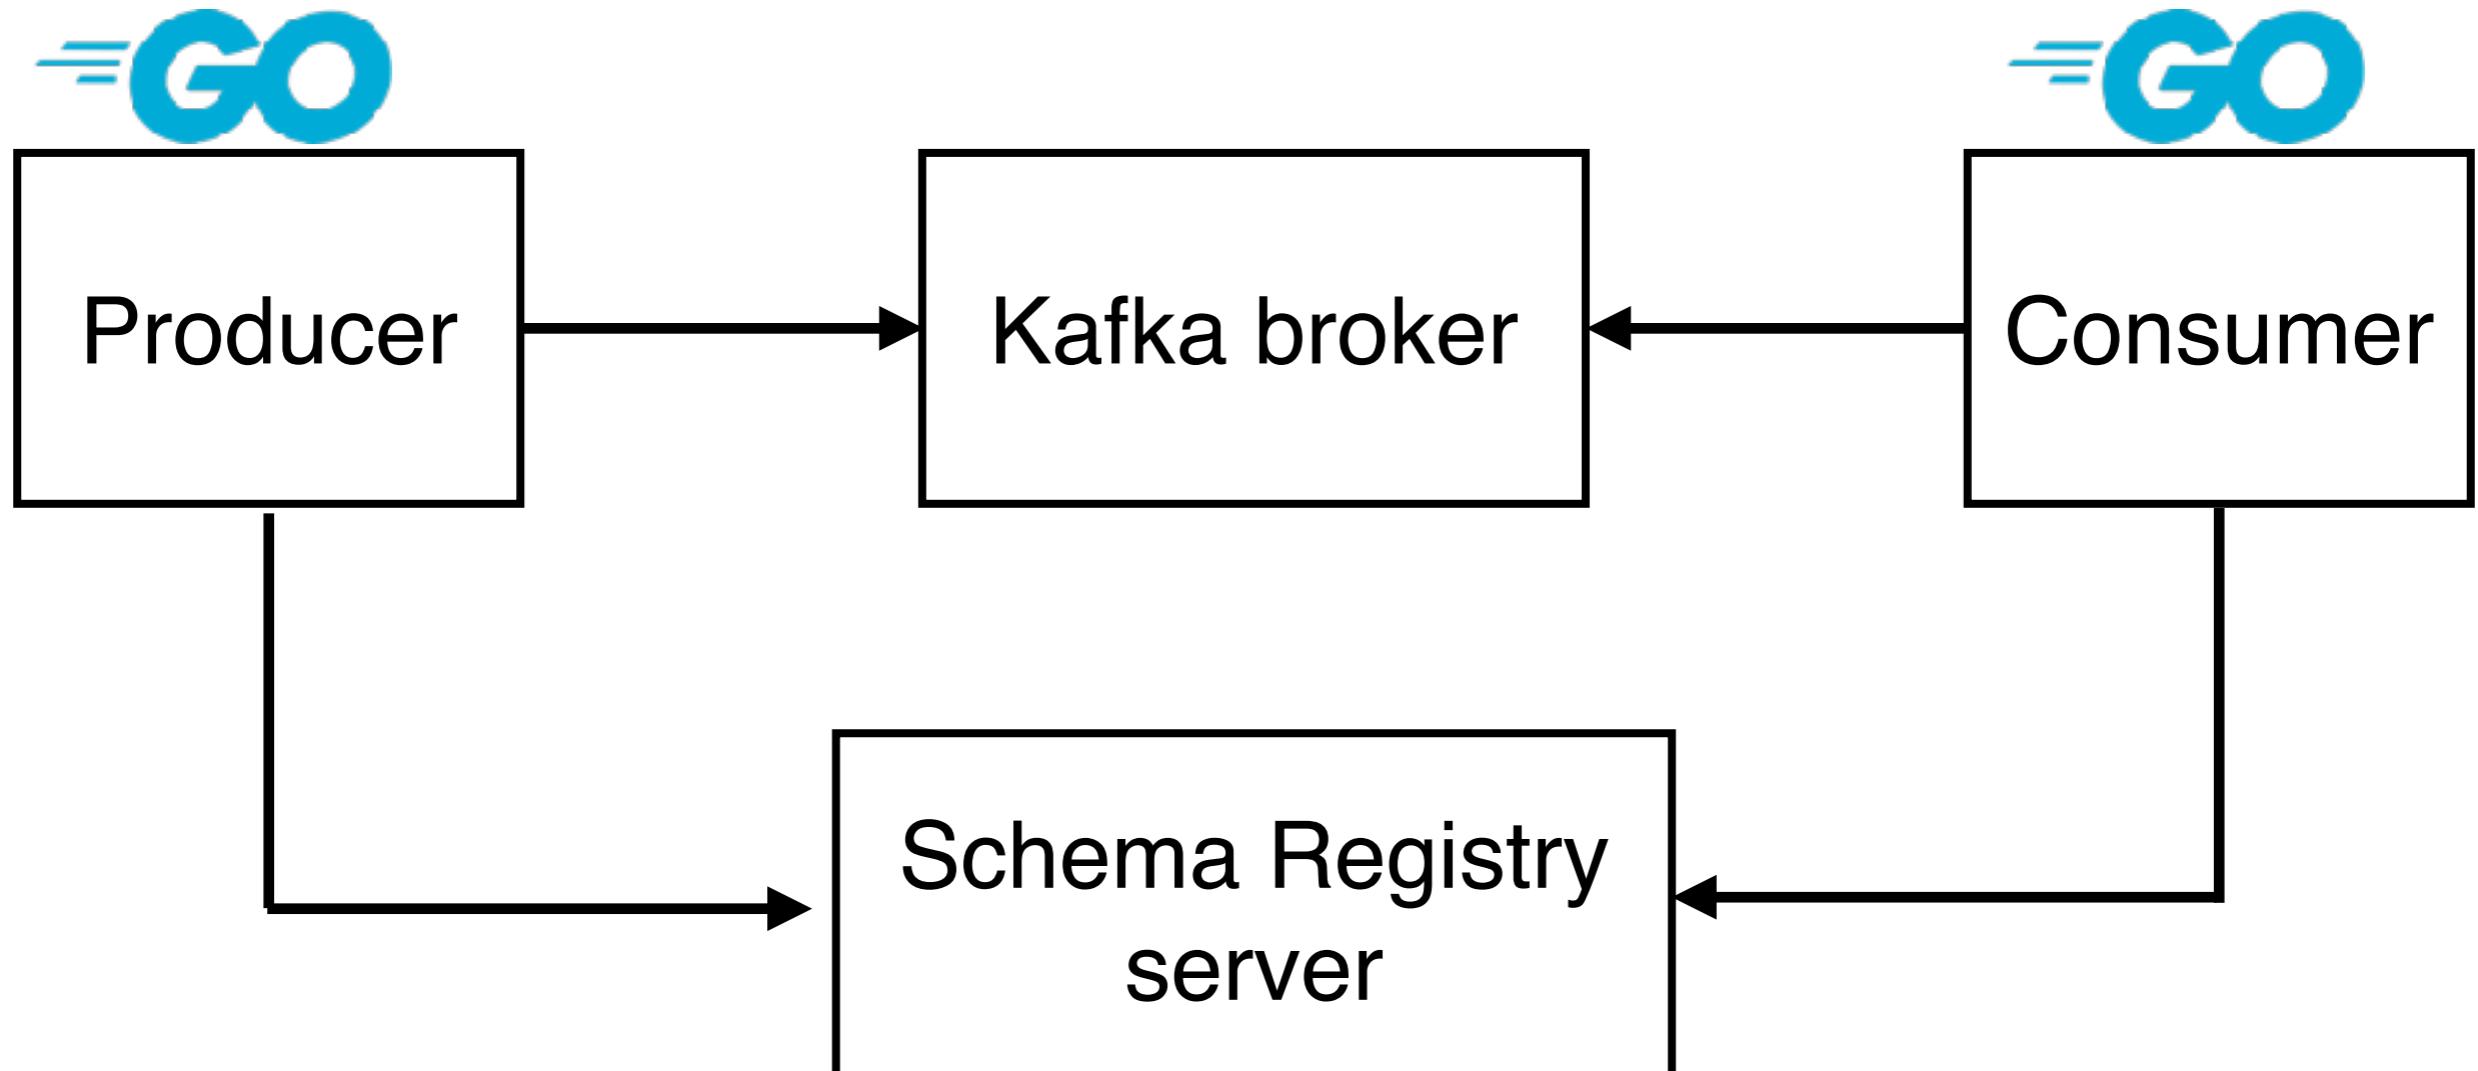
Protocol  
Buffers



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Workshop



<https://github.com/up1/course-kafka-2024/tree/main/workshop/basic>



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# ksqldb



Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

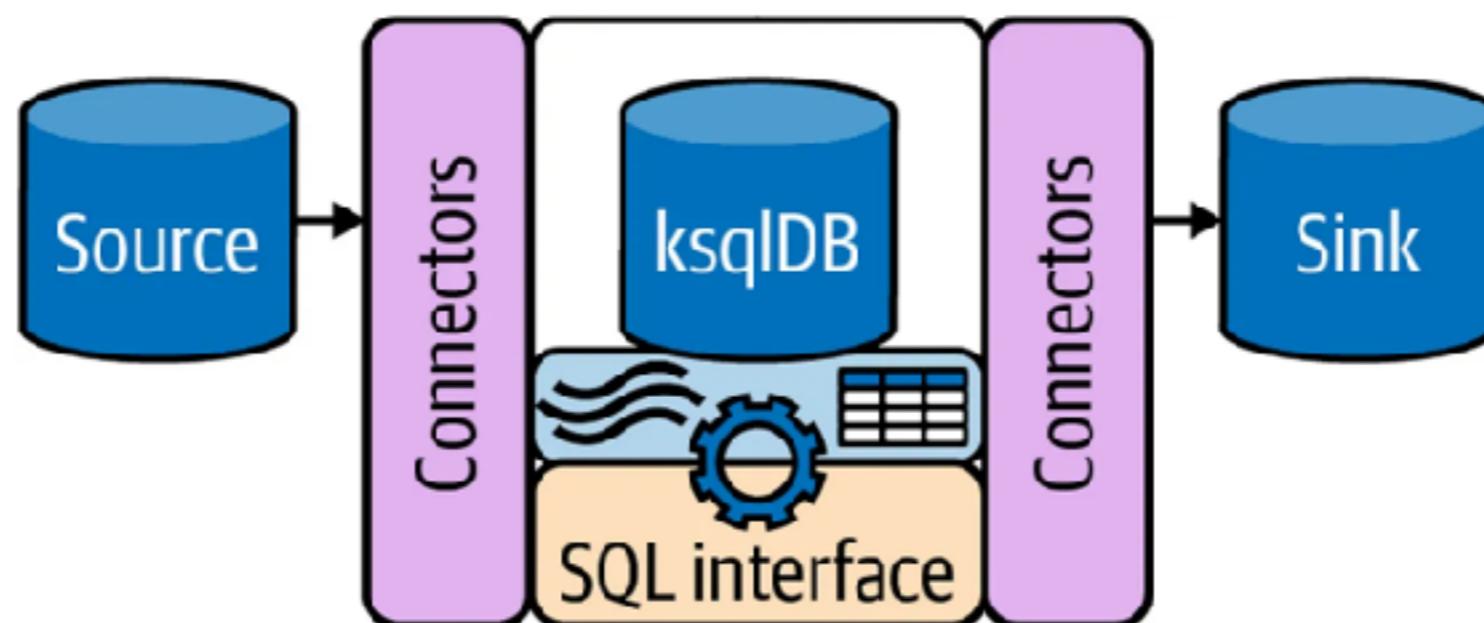
105

# ksqlDB

Kafka native database for **stream processing**

Used like SQL

Working with Kafka connect



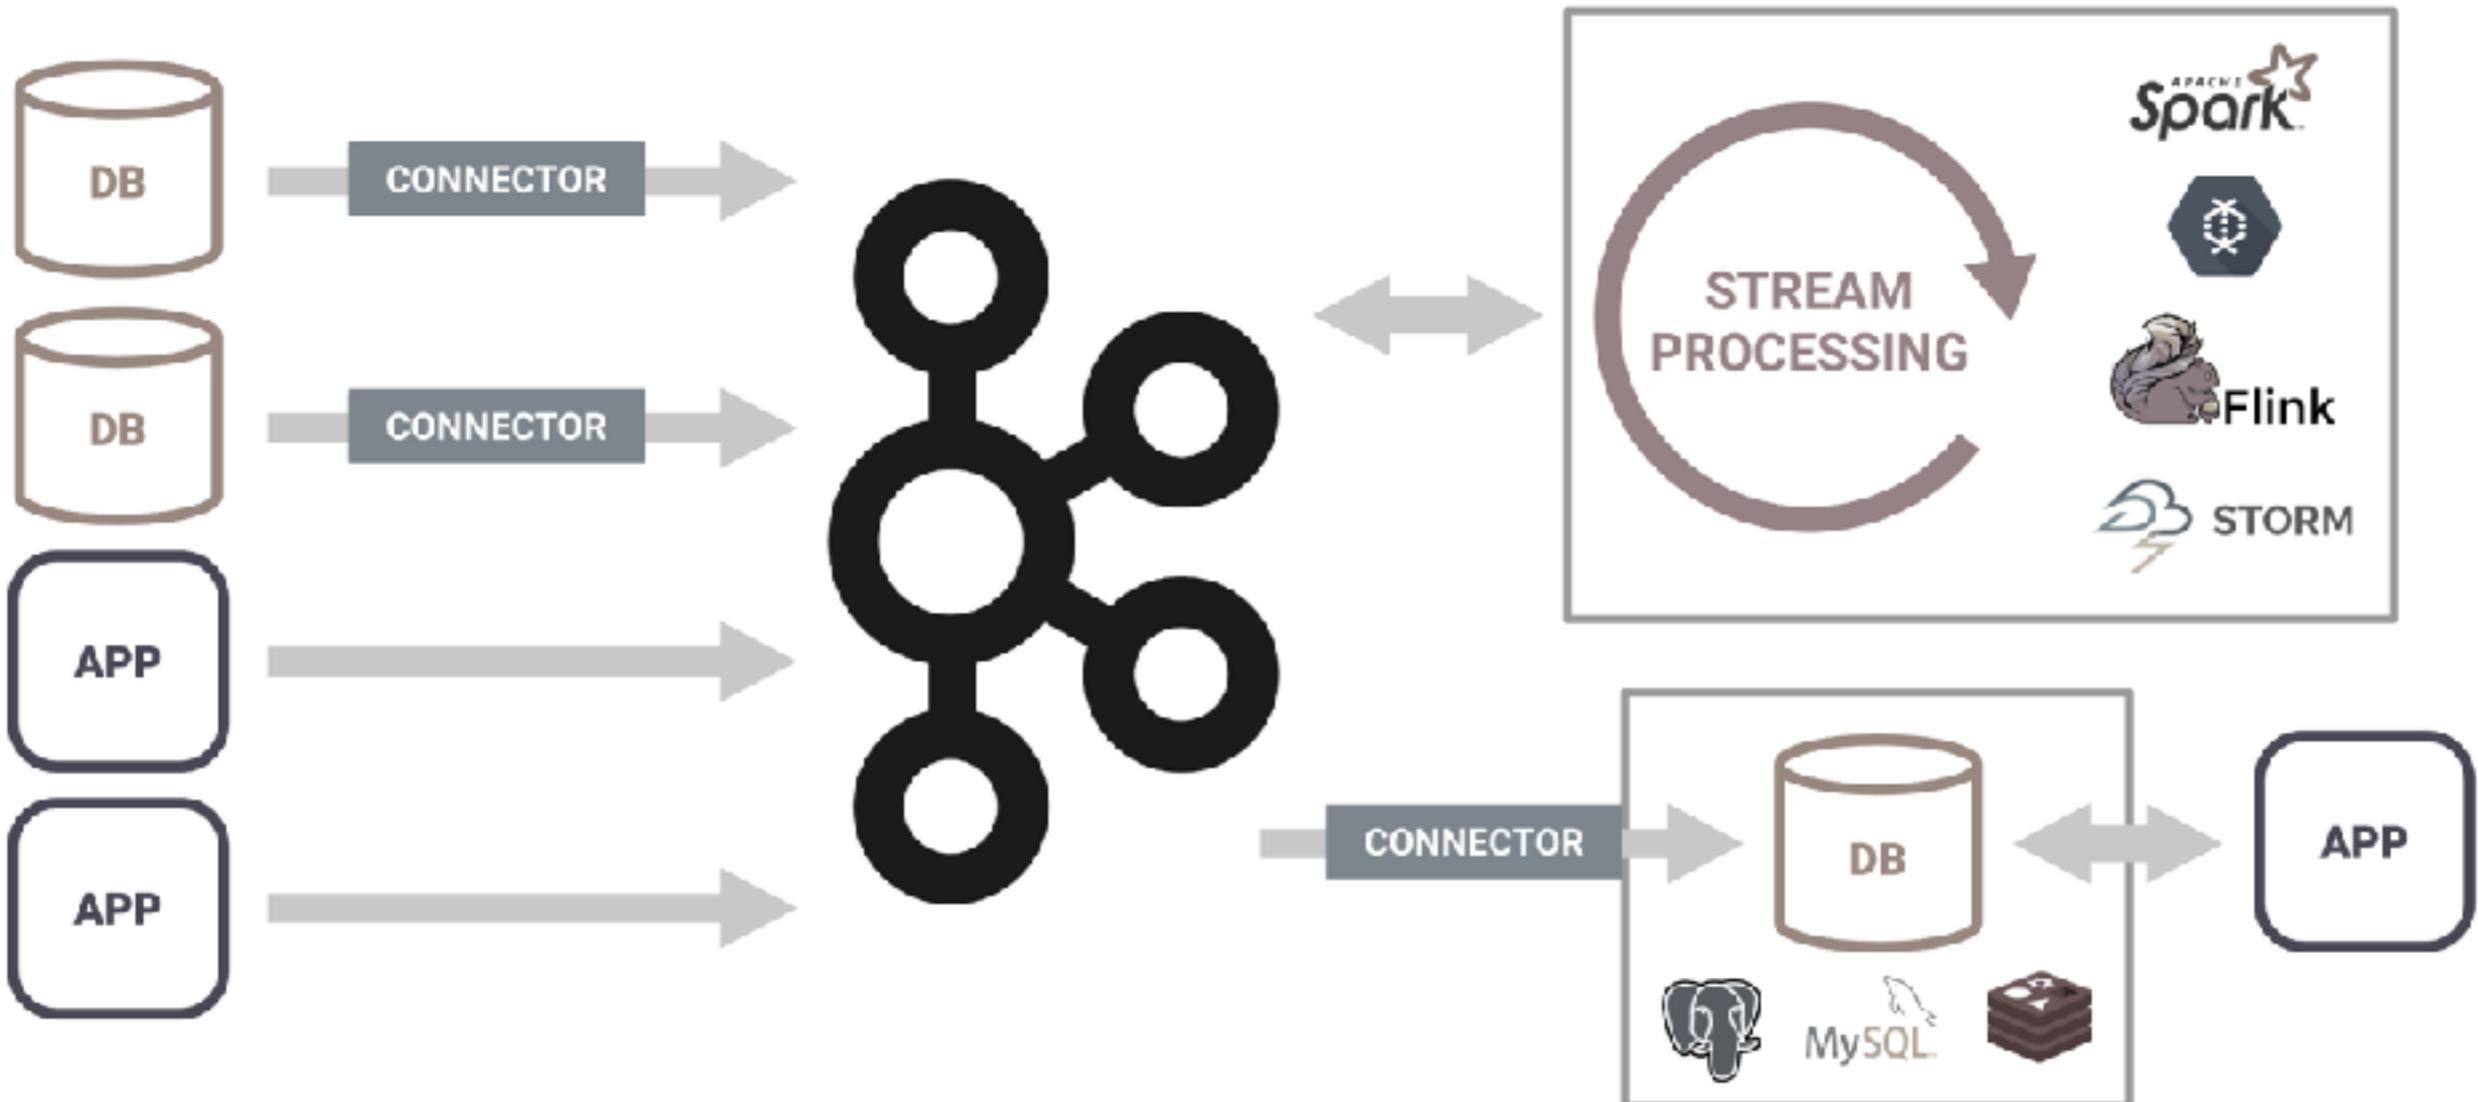
<https://www.confluent.io/product/ksqldb/>



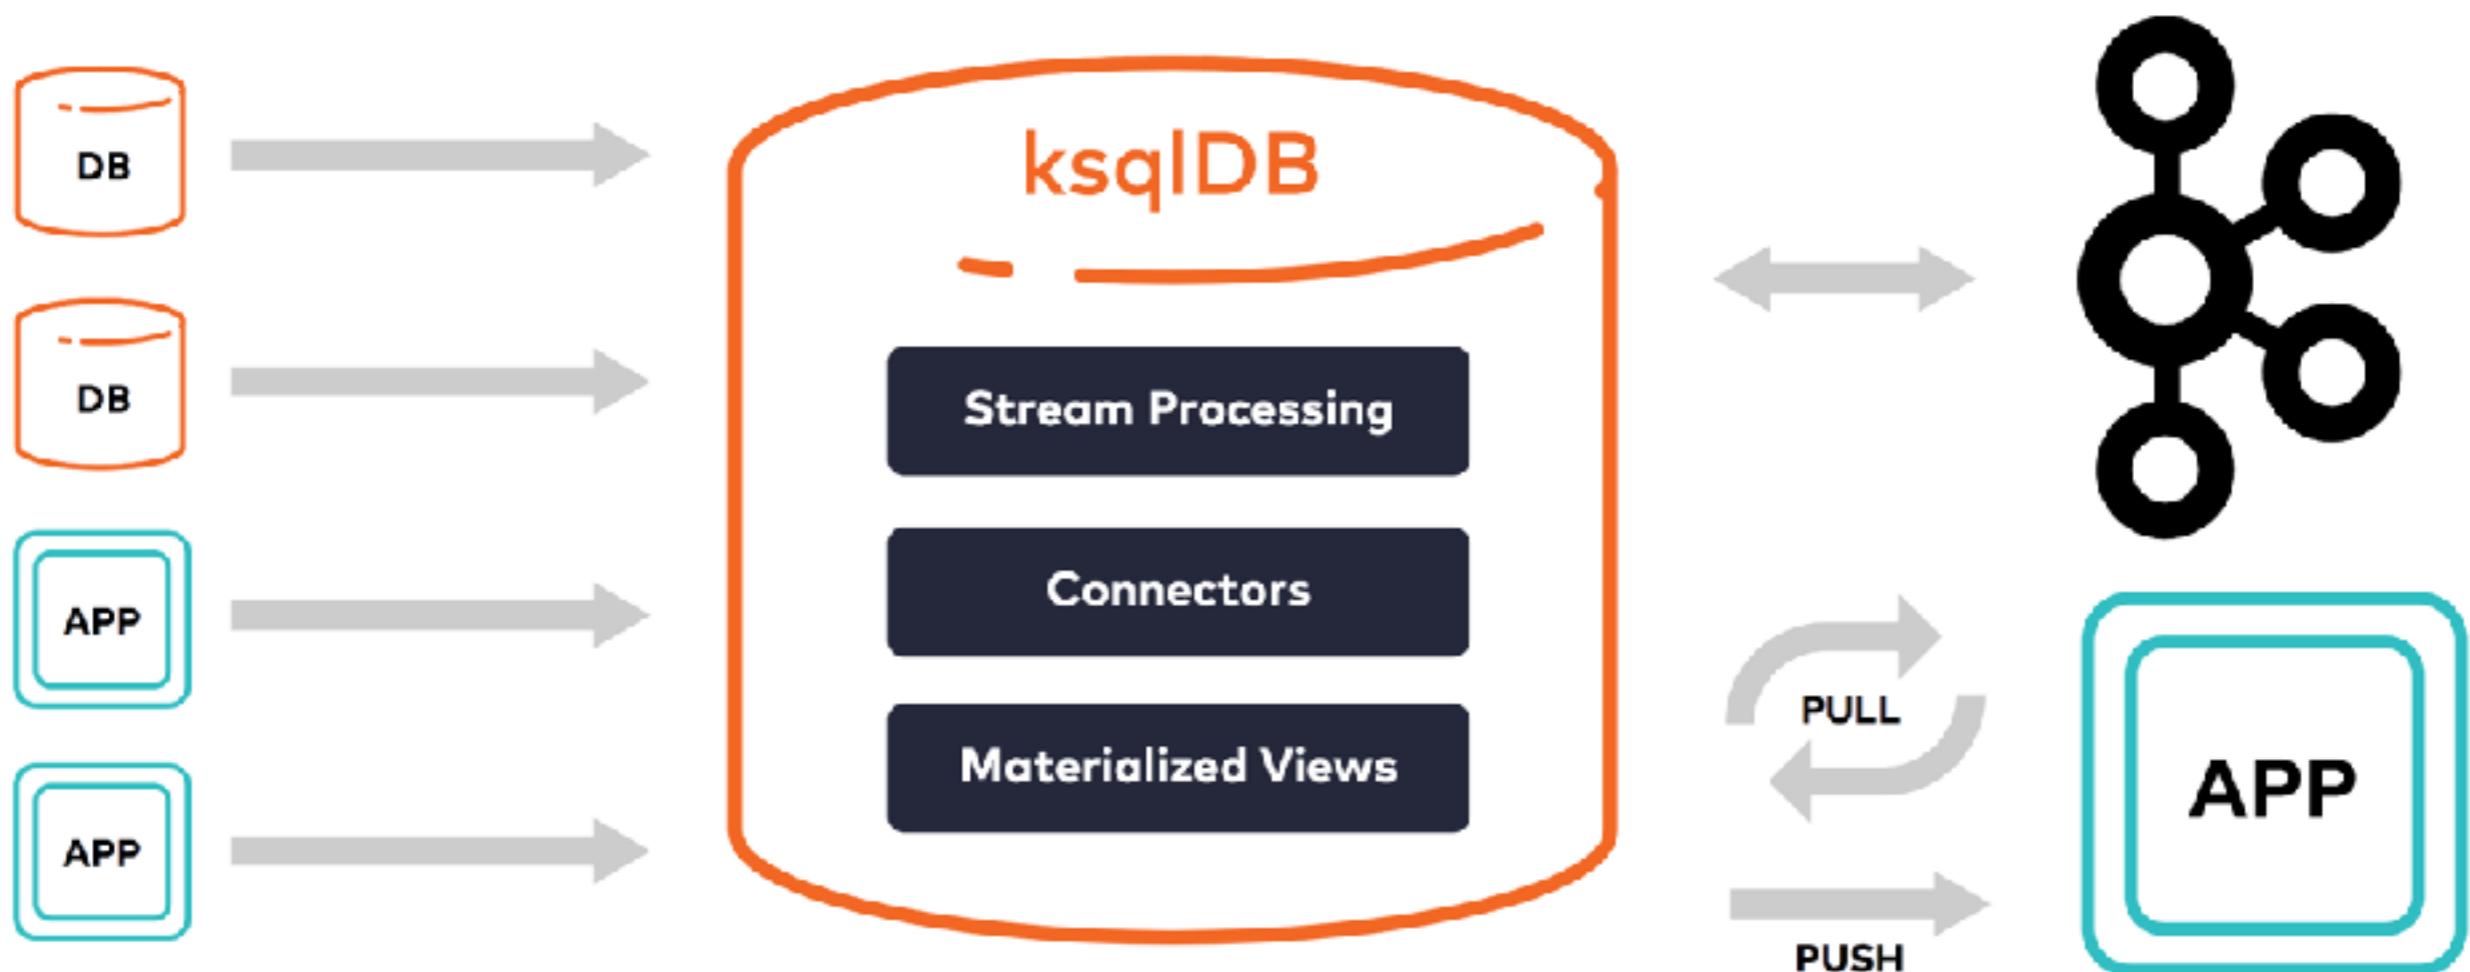
Kafka

© 2017 - 2025 Siam Chamnankit Company Limited. All rights reserved.

# Kafka



# ksqlDB



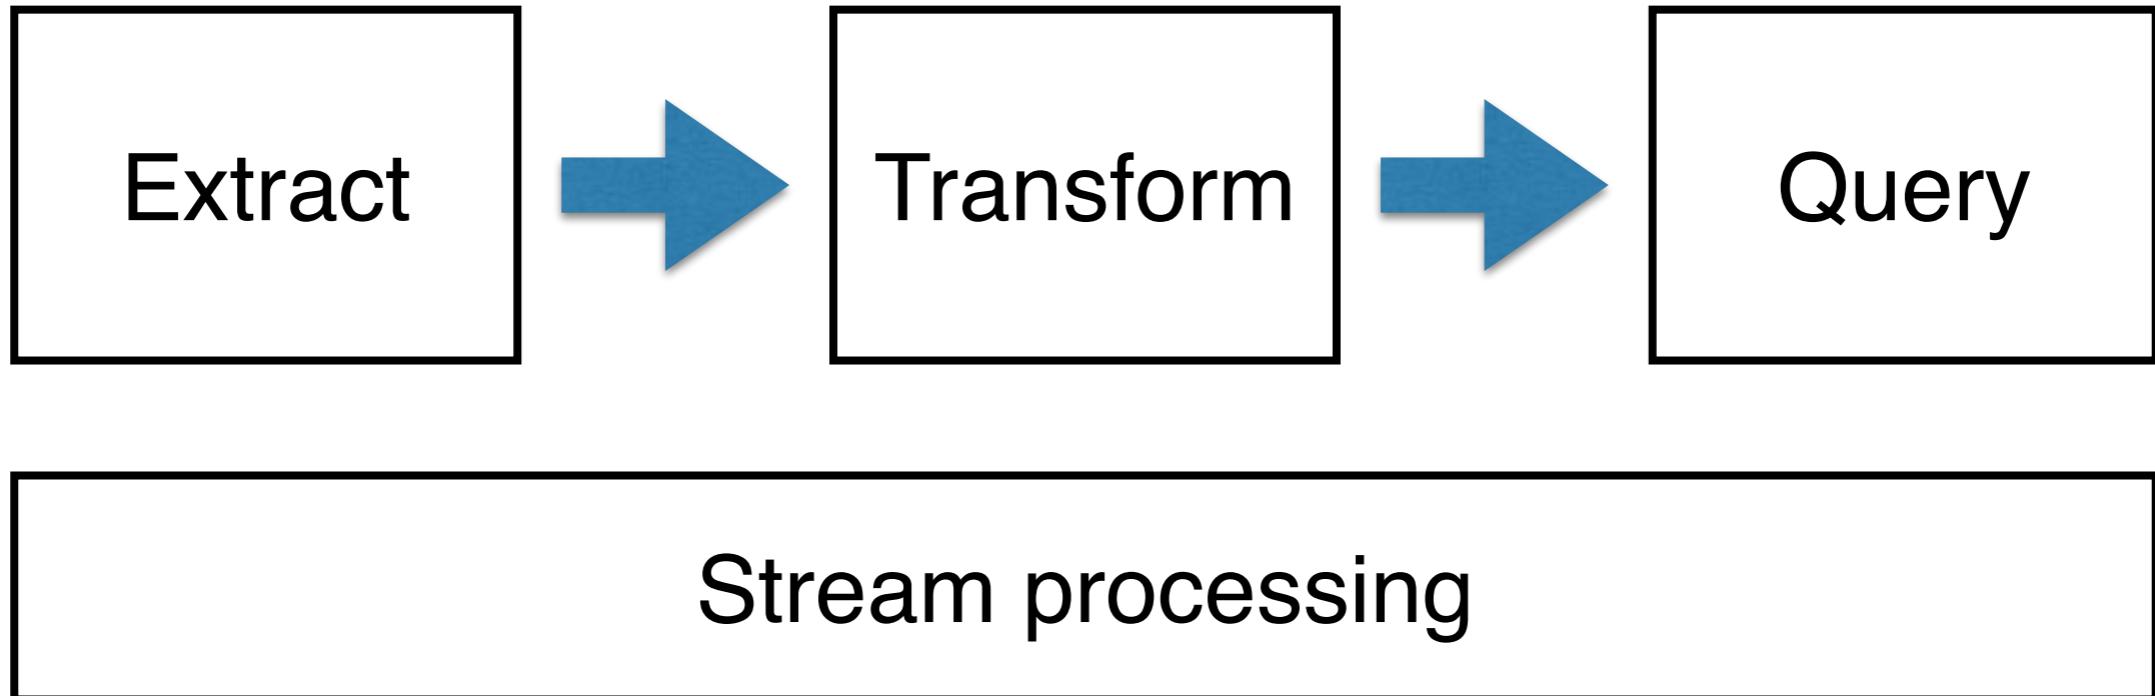
<https://developer.confluent.io/courses/ksqldb/intro/>



Kafka

© 2017 - 2025 Siam Chamnkit Company Limited. All rights reserved.

# ksqlDB

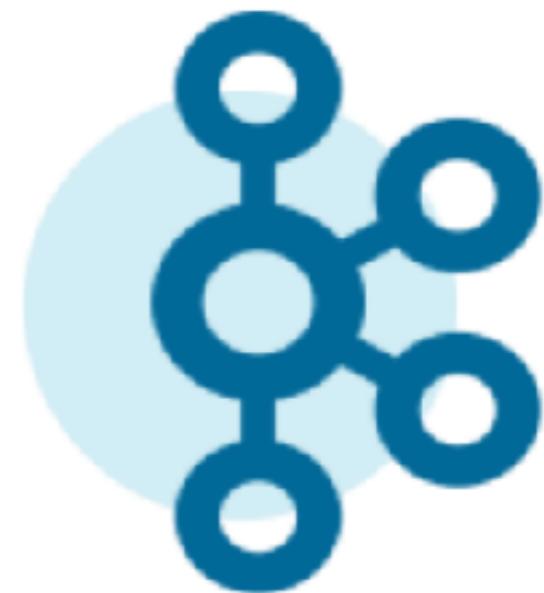


# ksqldb



*Compute*

# Kafka



*Storage*



# Plan for Failures of Kafka



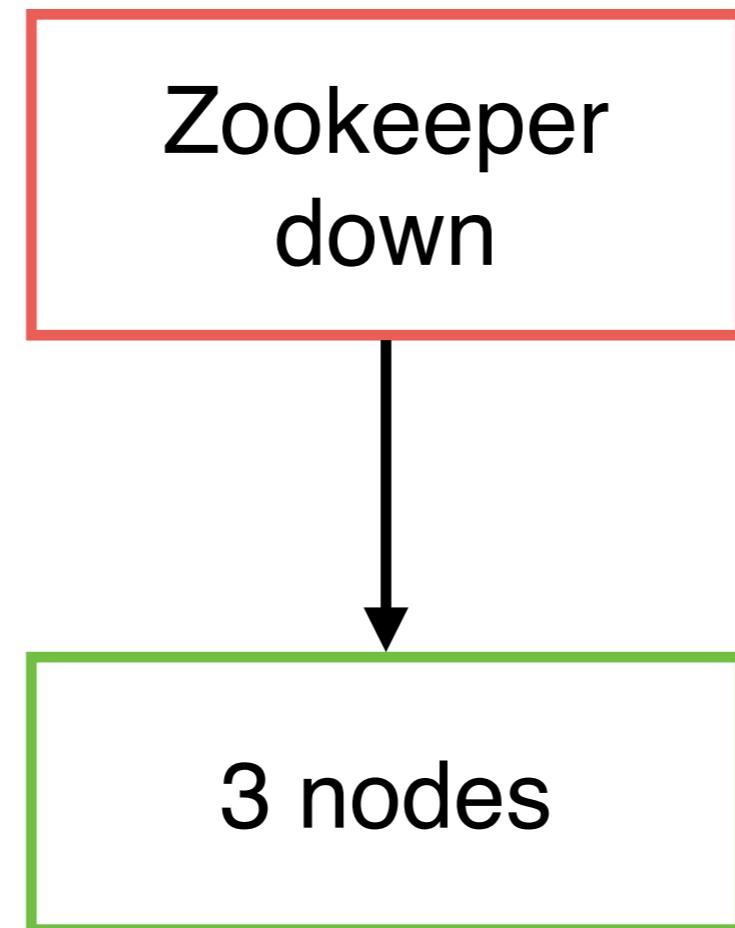
# Plan for Failures of Kafka

Zookeeper

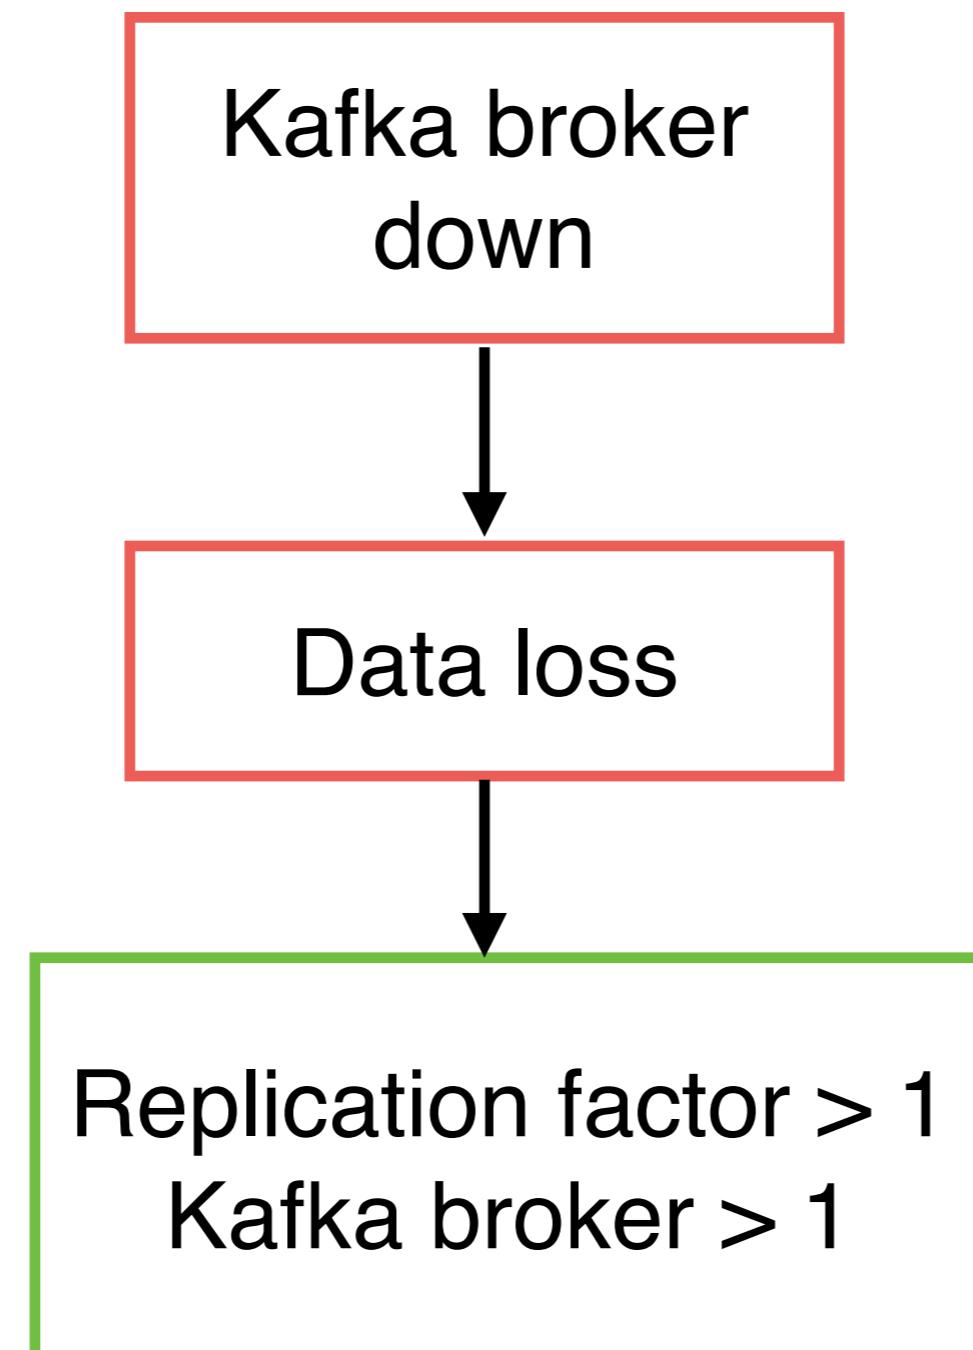
Kafka broker



# Zookeeper



# Kafka broker



# Rule of Thumb

**Replication factor (RF) should typically be less than or equal to the number of brokers**

**Recommended in production = 3**

High Availability

Fault tolerance  
With low overhead



# Ratio guideline

Number of Brokers	Recommended Replication Factor (RF)	Broker-to-RF Ratio
3	3	1:1
4	3	4:3
5	3	5:3
6	3	2:1
10	3-5	2:1 to 10:5
15	3-5	3:1 to 15:5



# Key Considerations (1)

More RF, more storage/network overhead

$RF = 3$

Overhead =  $3 * \# \text{ of partitions}$

Disk

Memory

Network



# Key Considerations (2)

## Broker failure

$$RF = N + 1$$

N = number of broker



# Key Considerations (3)

## Partition distribution

More partition, more RF  
==  
more overhead



# Key Considerations (4)

Write latency of producer

Ack = 0  
Ack = 1  
Ack = all



# Improve write latency of producer

- Increase # of partition
- Reduce batch size and linker
- Increase buffer size (reduce network call)
- Use compression



# Key Considerations (5)

Read latency of consumer

Fetch min

Fetch max

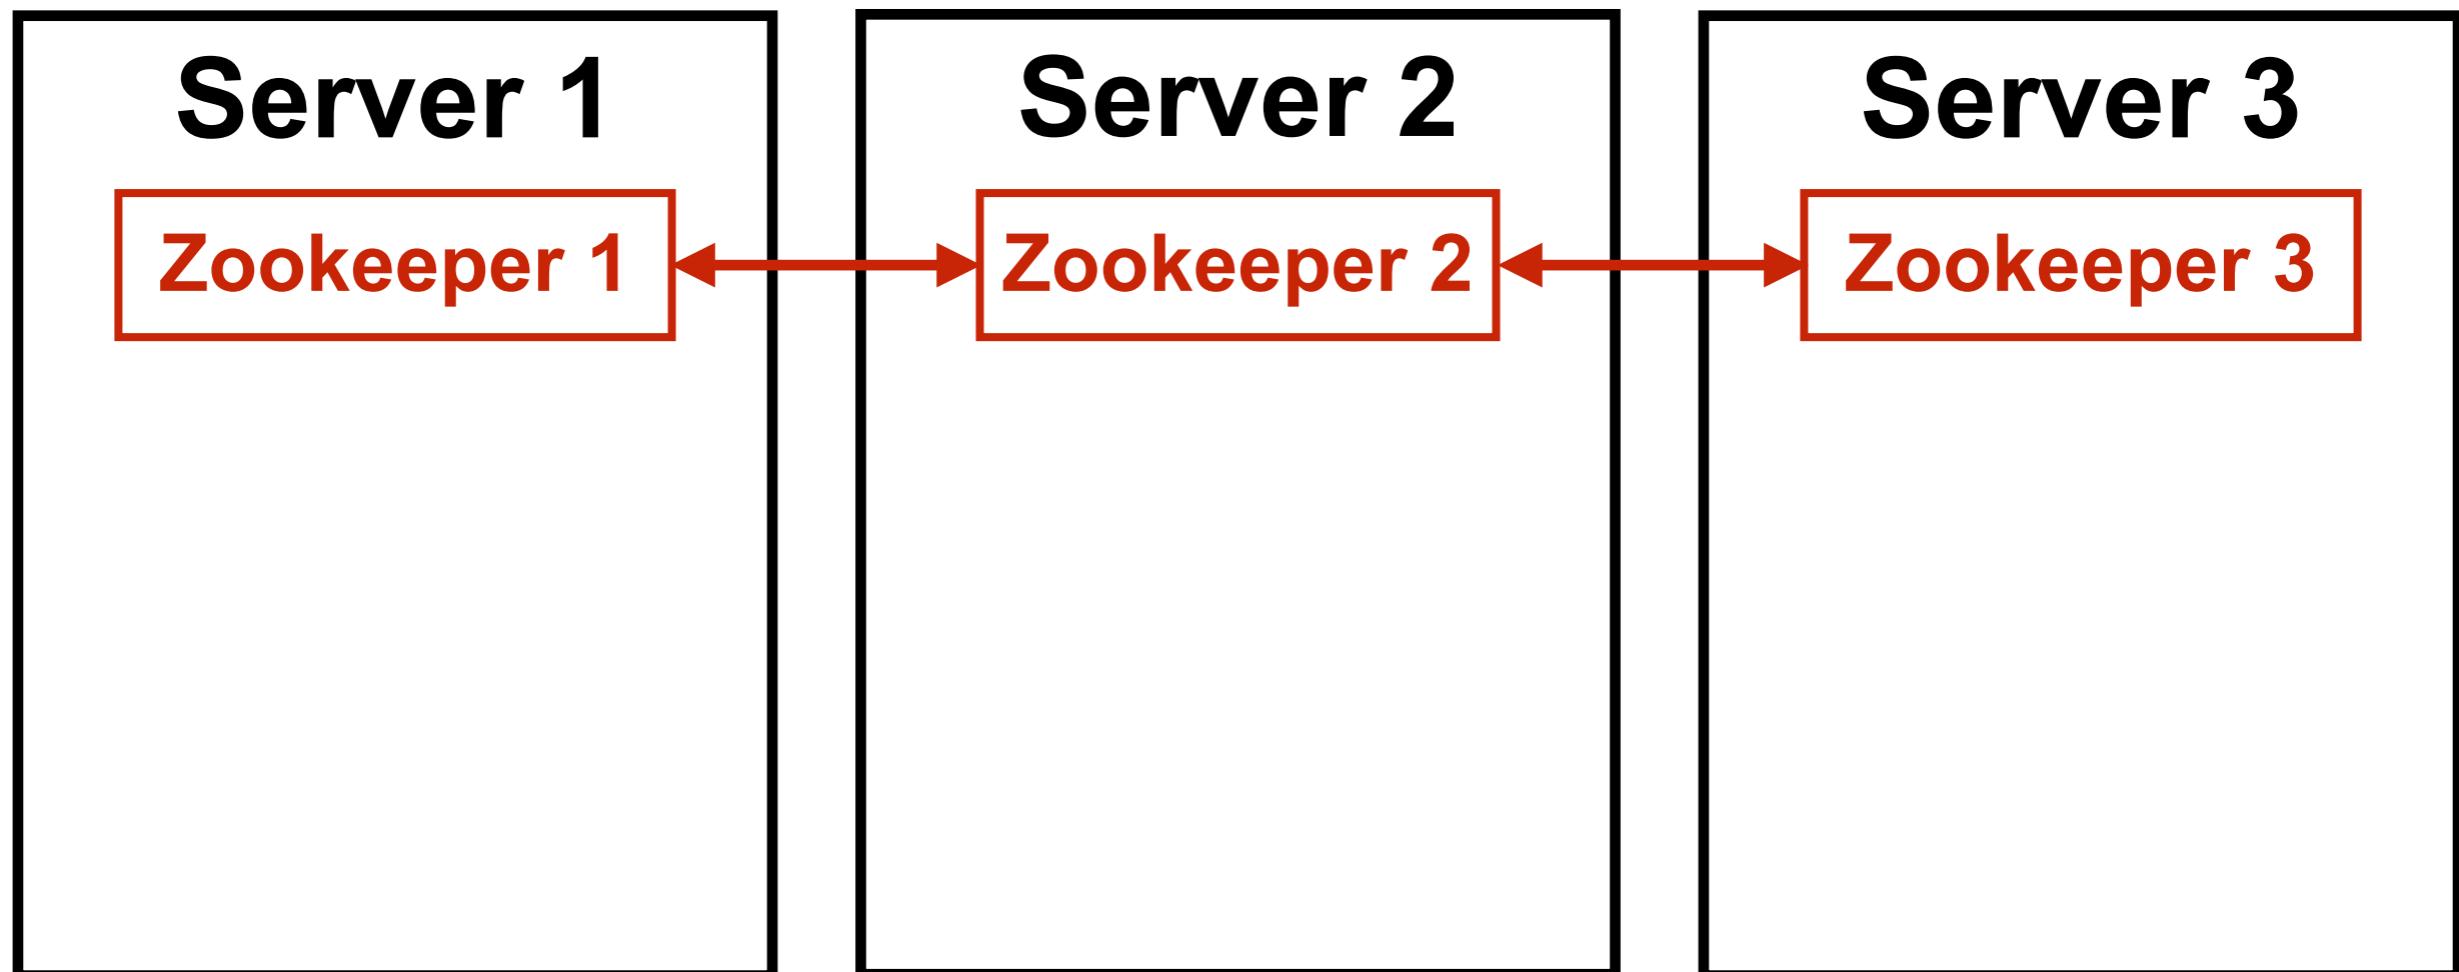
Max poll record



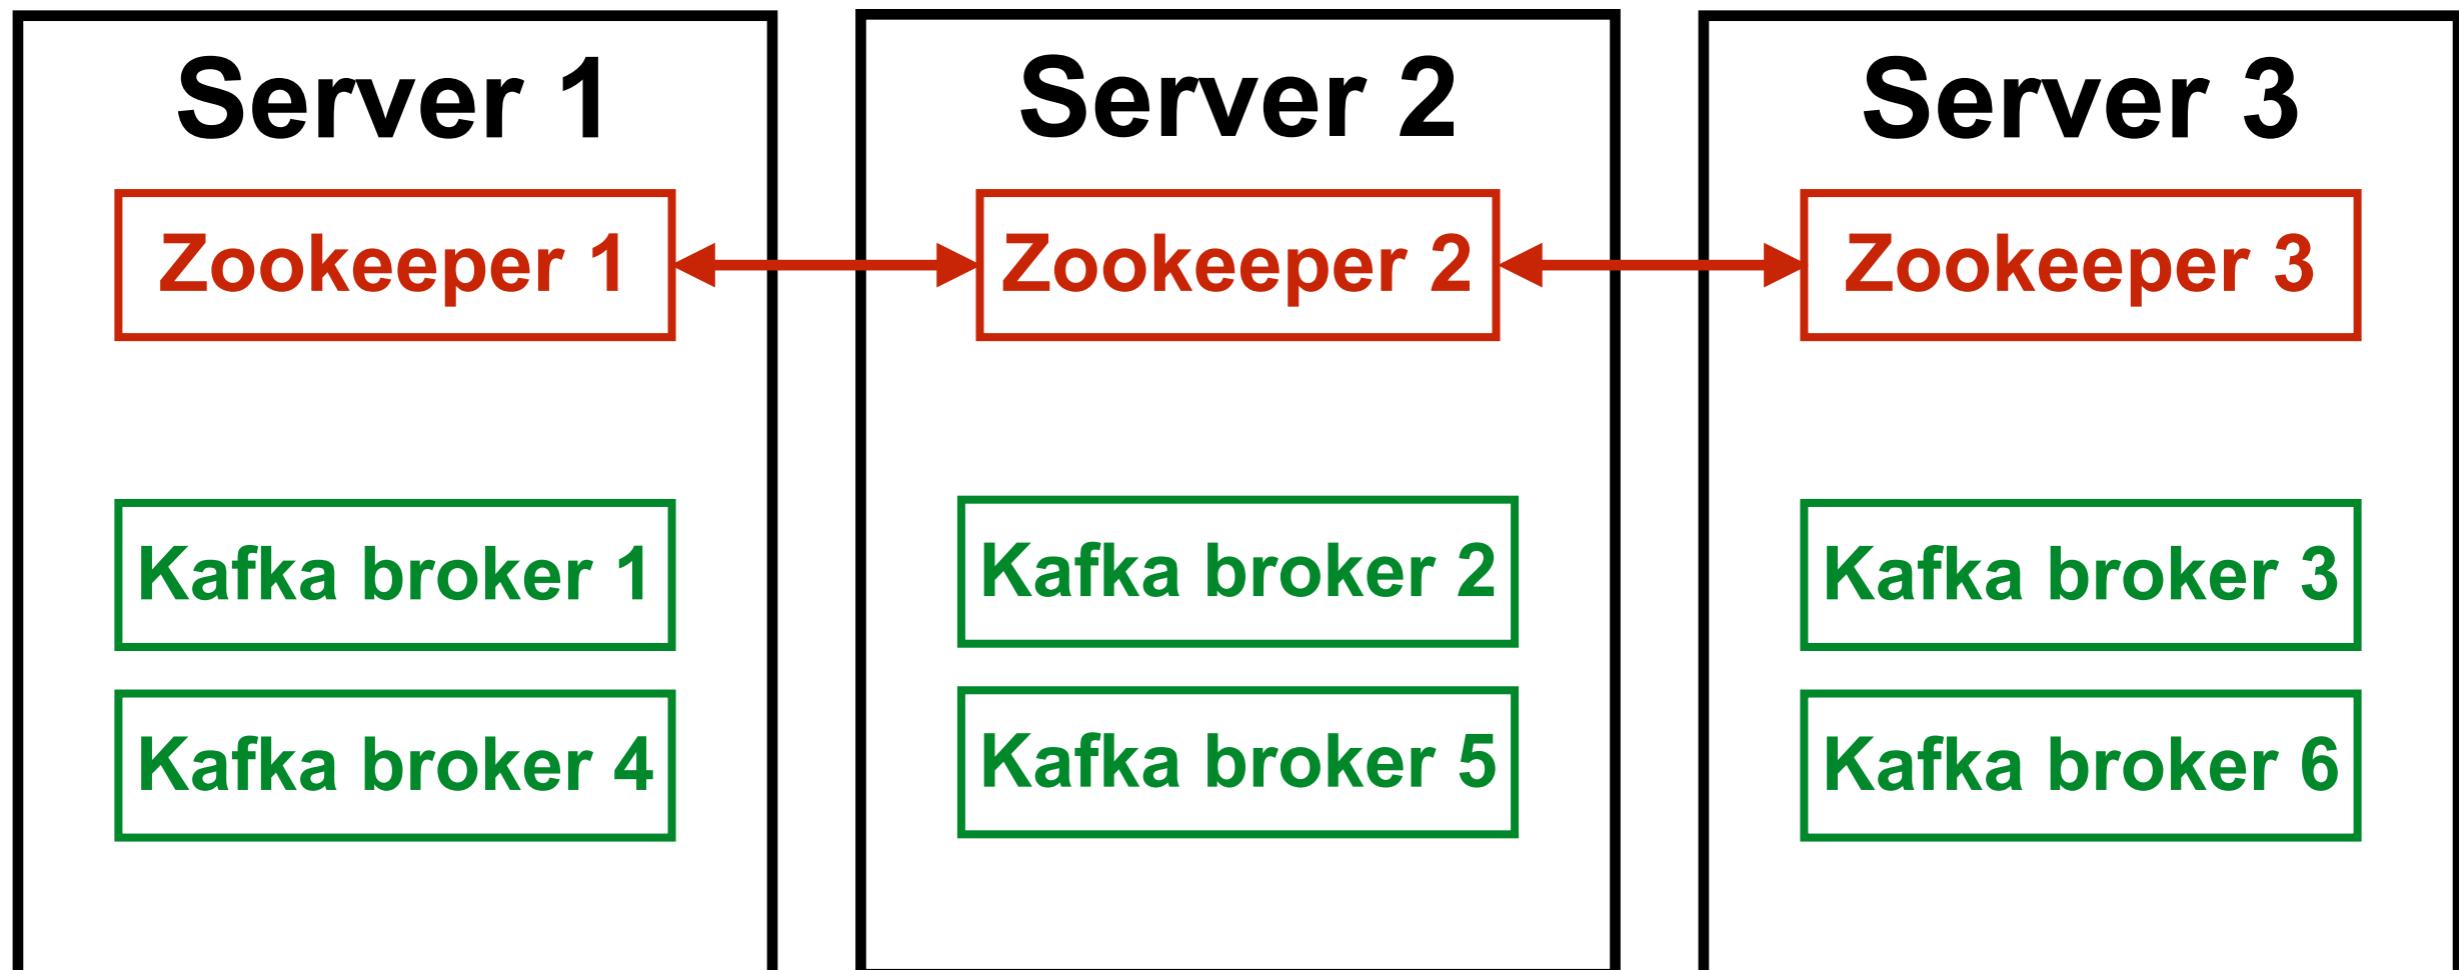
# Monitoring Kafka



# Kafka cluster



# Kafka cluster



# Kafka cluster

It's not easy to setup a cluster

Need to isolate each Zookeeper and broker on separate servers

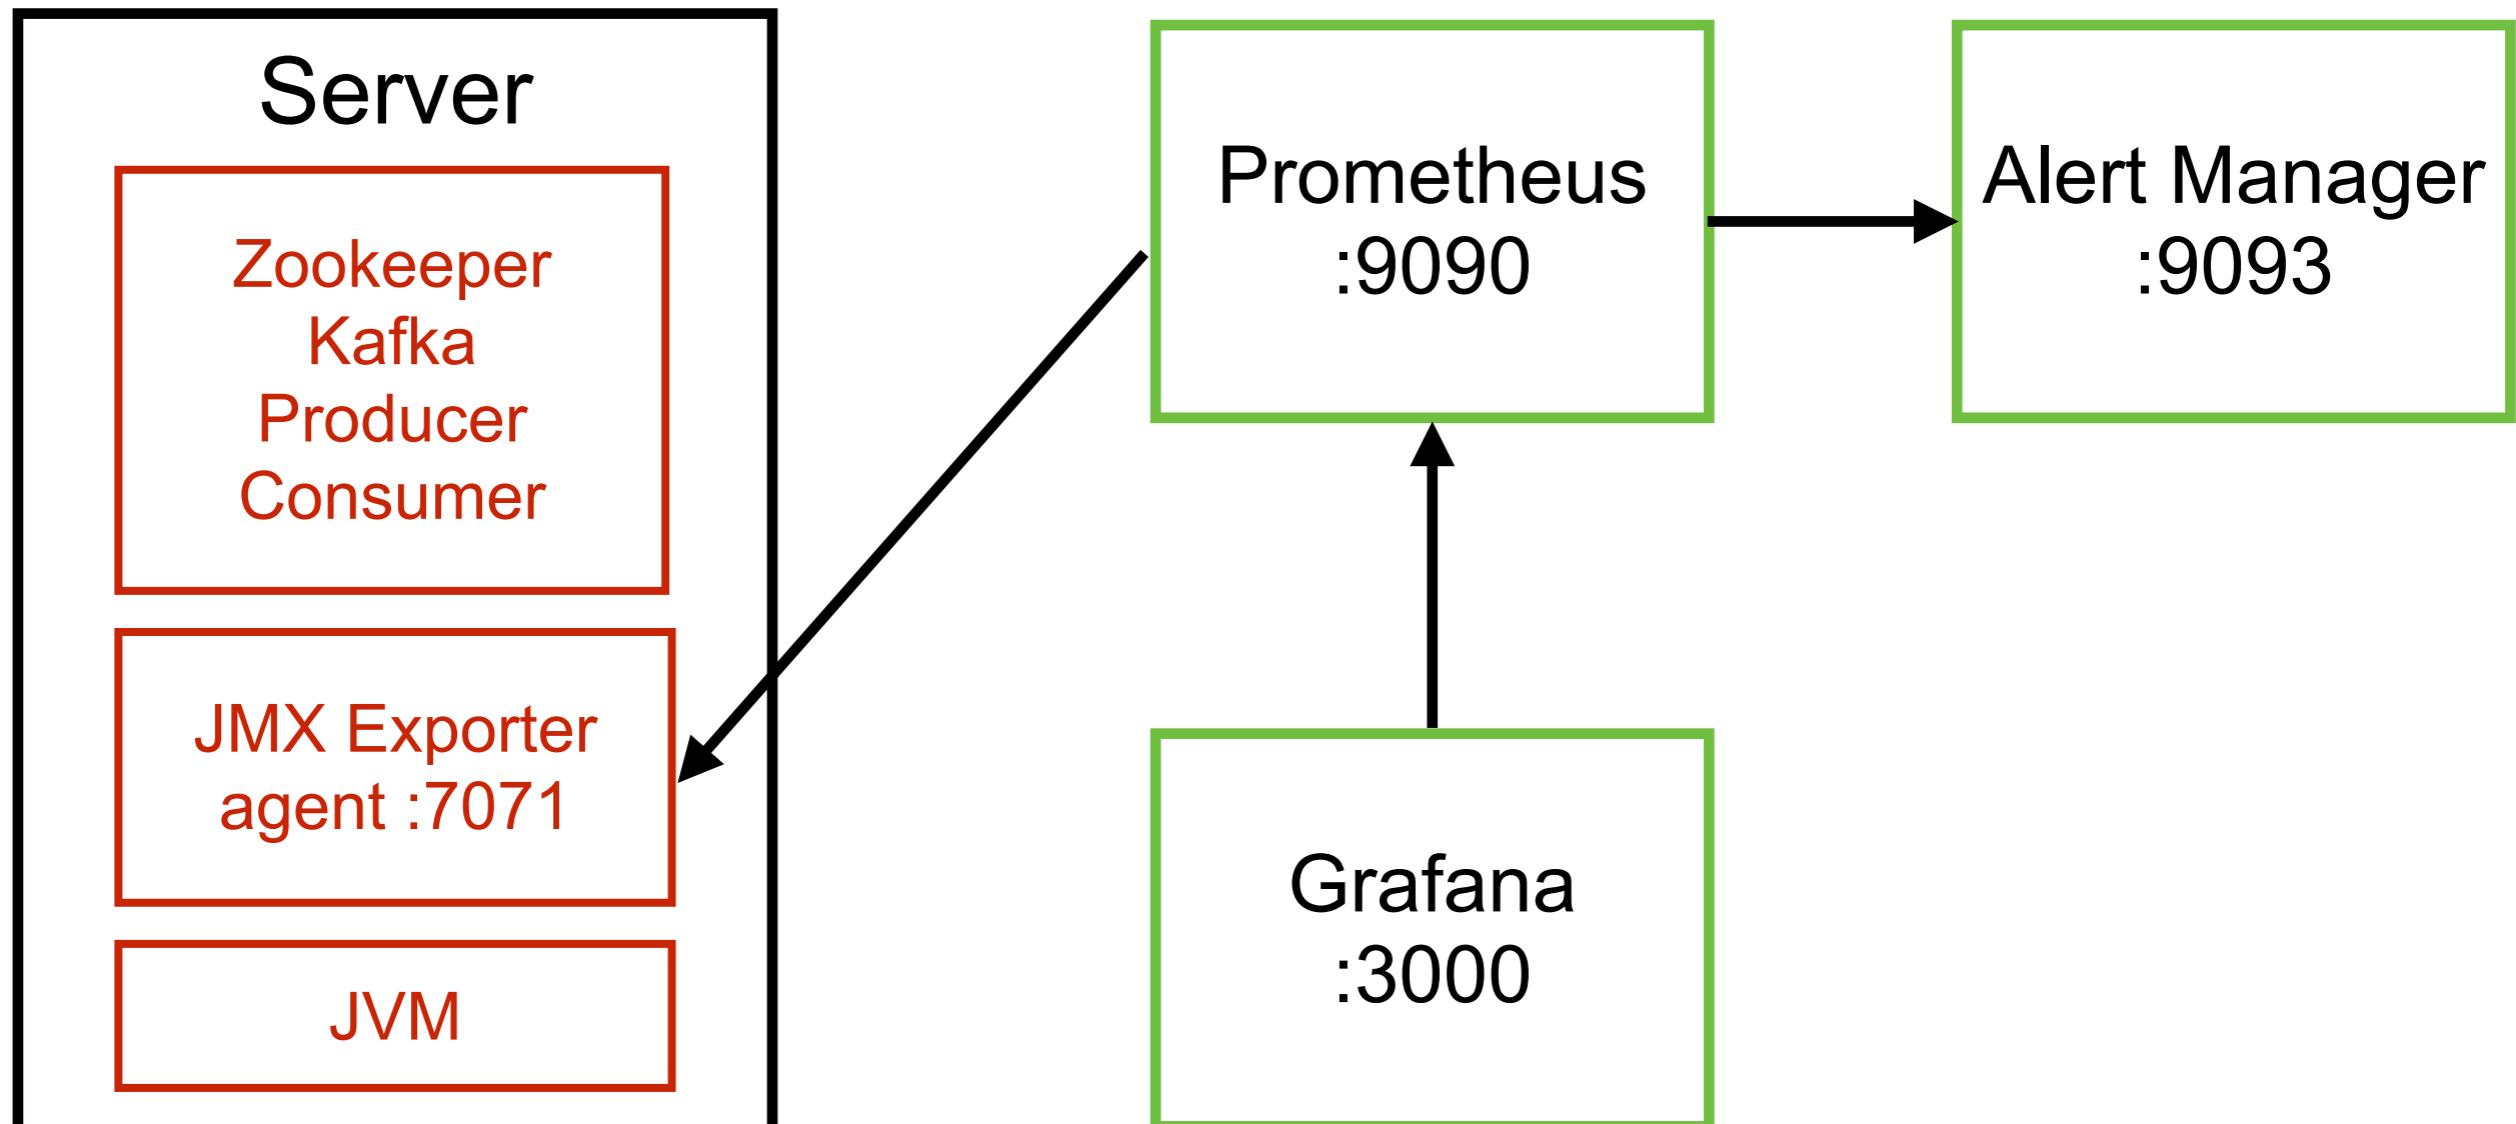
Need to implement a **Monitoring system**

Need skill of **operation teams**

Need a good **Kafka Admin**



# Kafka monitoring



<https://github.com/up1/course-kafka-2024/tree/main/workshop/monitoring>



# Tips #1

Please monitor Your Kafka !!



# Monitoring ?

Broker health  
Message delivery  
Performance  
Capacity



# Kafka monitoring

Kafka exposes **metrics** through JMX

These metrics are very important for monitoring



# Kafka monitoring

JConsole  
JMX/Metrics integration  
Burrow



# Keep Kafka metrics ?

ELK stack  
Prometheus  
Confluent control center  
Datadog  
NewRelic  
etc...



# Important metrics

## Under replicated partitions

# of partitions are have problem with the ISR

May indicate a high load on the system

## Request handlers

Utilization of threads for I/O, network

Utilization of Kafka broker



# Important metrics

## Request timing

How long it takes to reply to requests

Low is better, latency will improved

<https://kafka.apache.org/documentation/#monitoring>

<https://docs.confluent.io/current/kafka/monitoring.html>

<https://www.datadoghq.com/blog/monitoring-kafka-performance-metrics/>

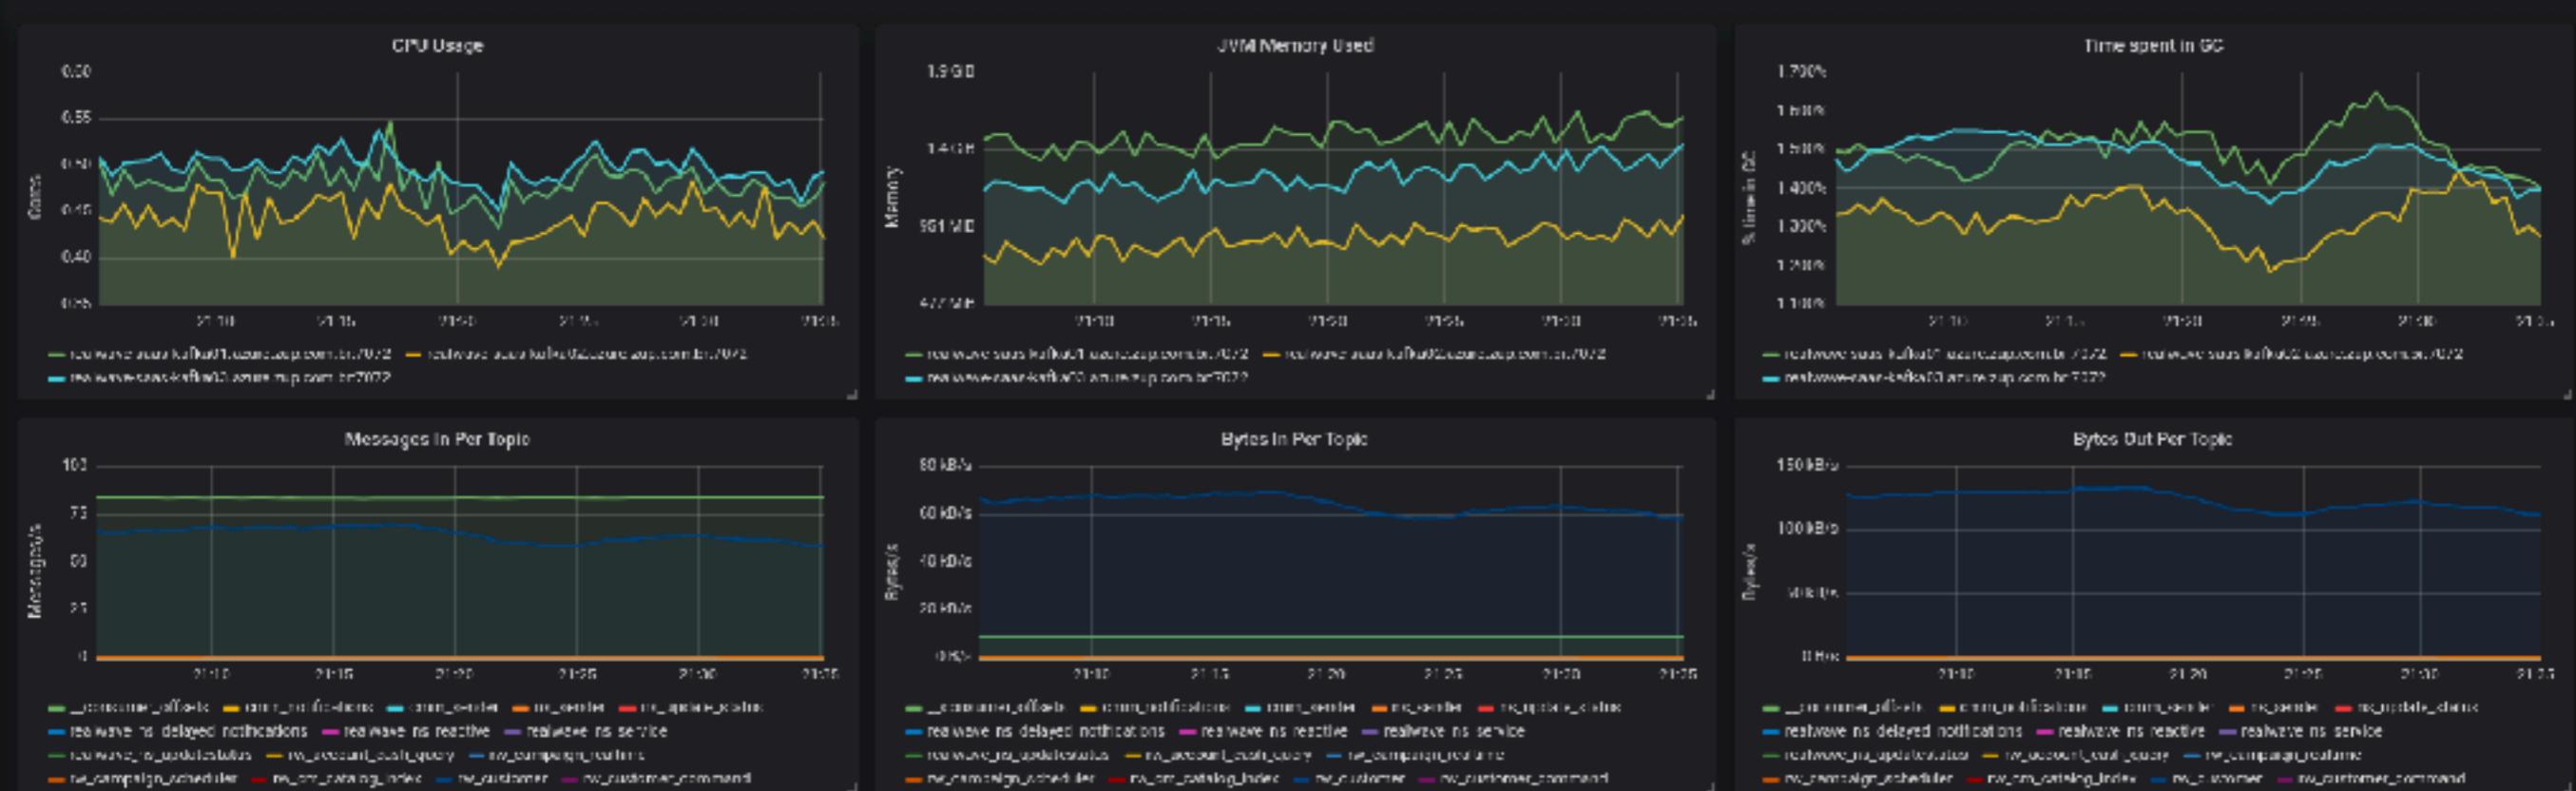


# Tips #2

Have a dashboard that lets you know  
**“Everything is OK ?”**  
(in one look)



# Dashboard



# Dashboard



Kafka

137

# Tips #3

Don't watch the dashboard  
Alert on what is important

*Only restart if you know why this will fix the issue !!*



# Capacity planning

CPU

Network and thread pool usages

Request latency

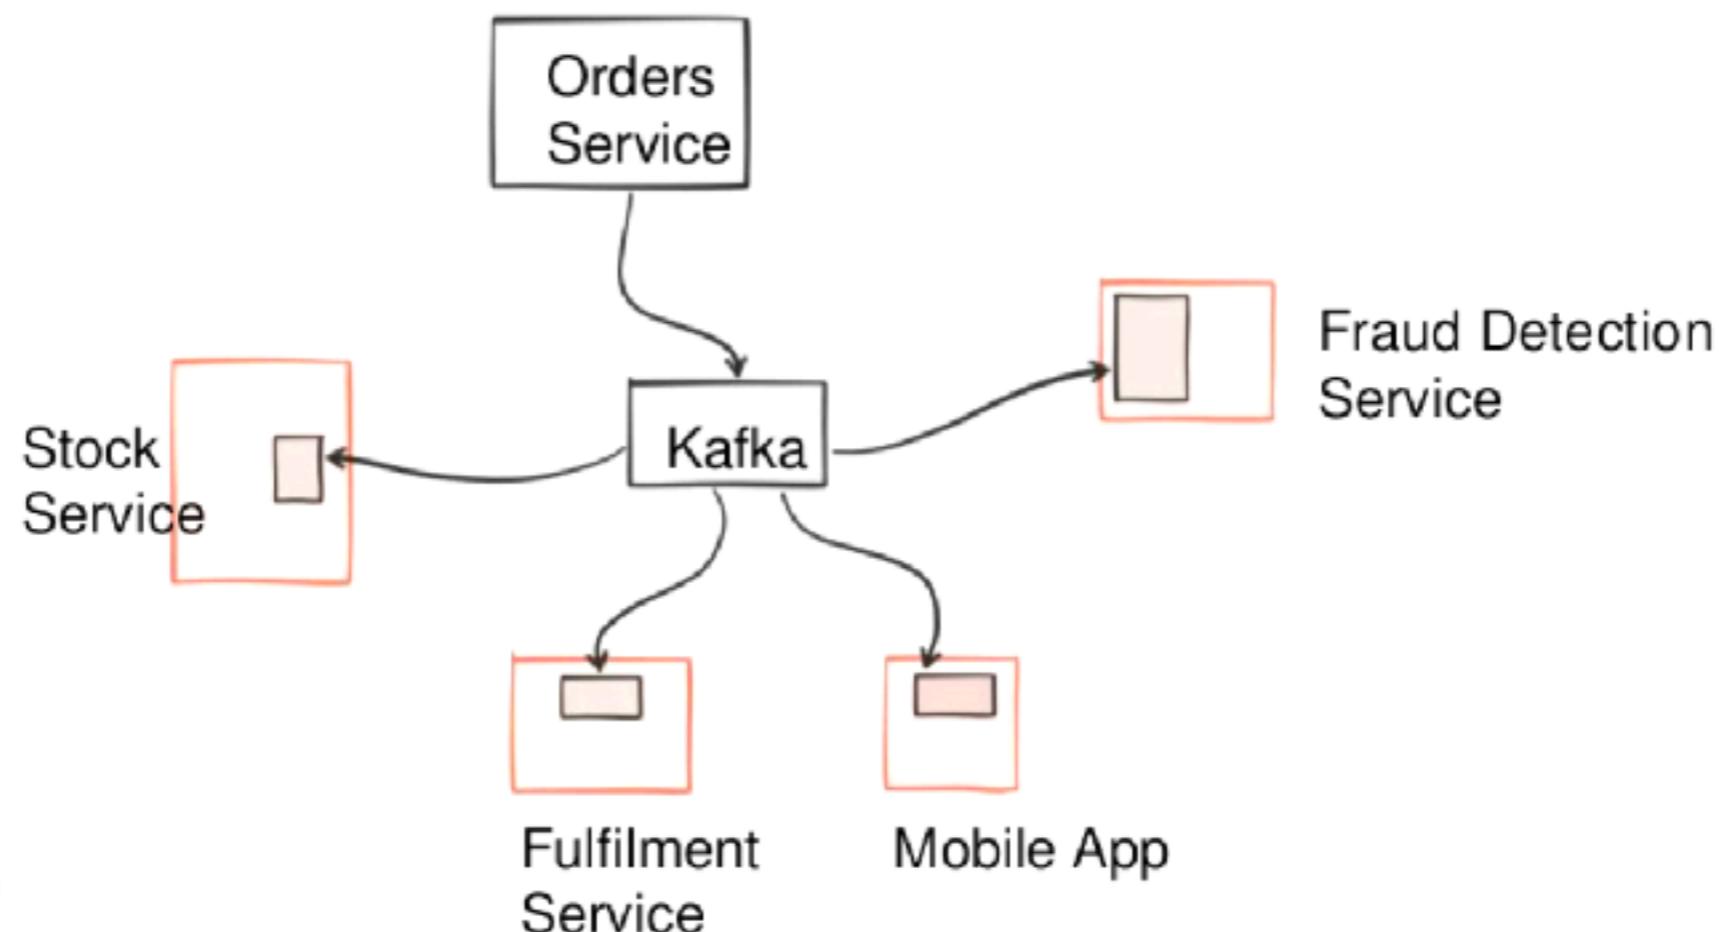
Network utilization

Disk utilization



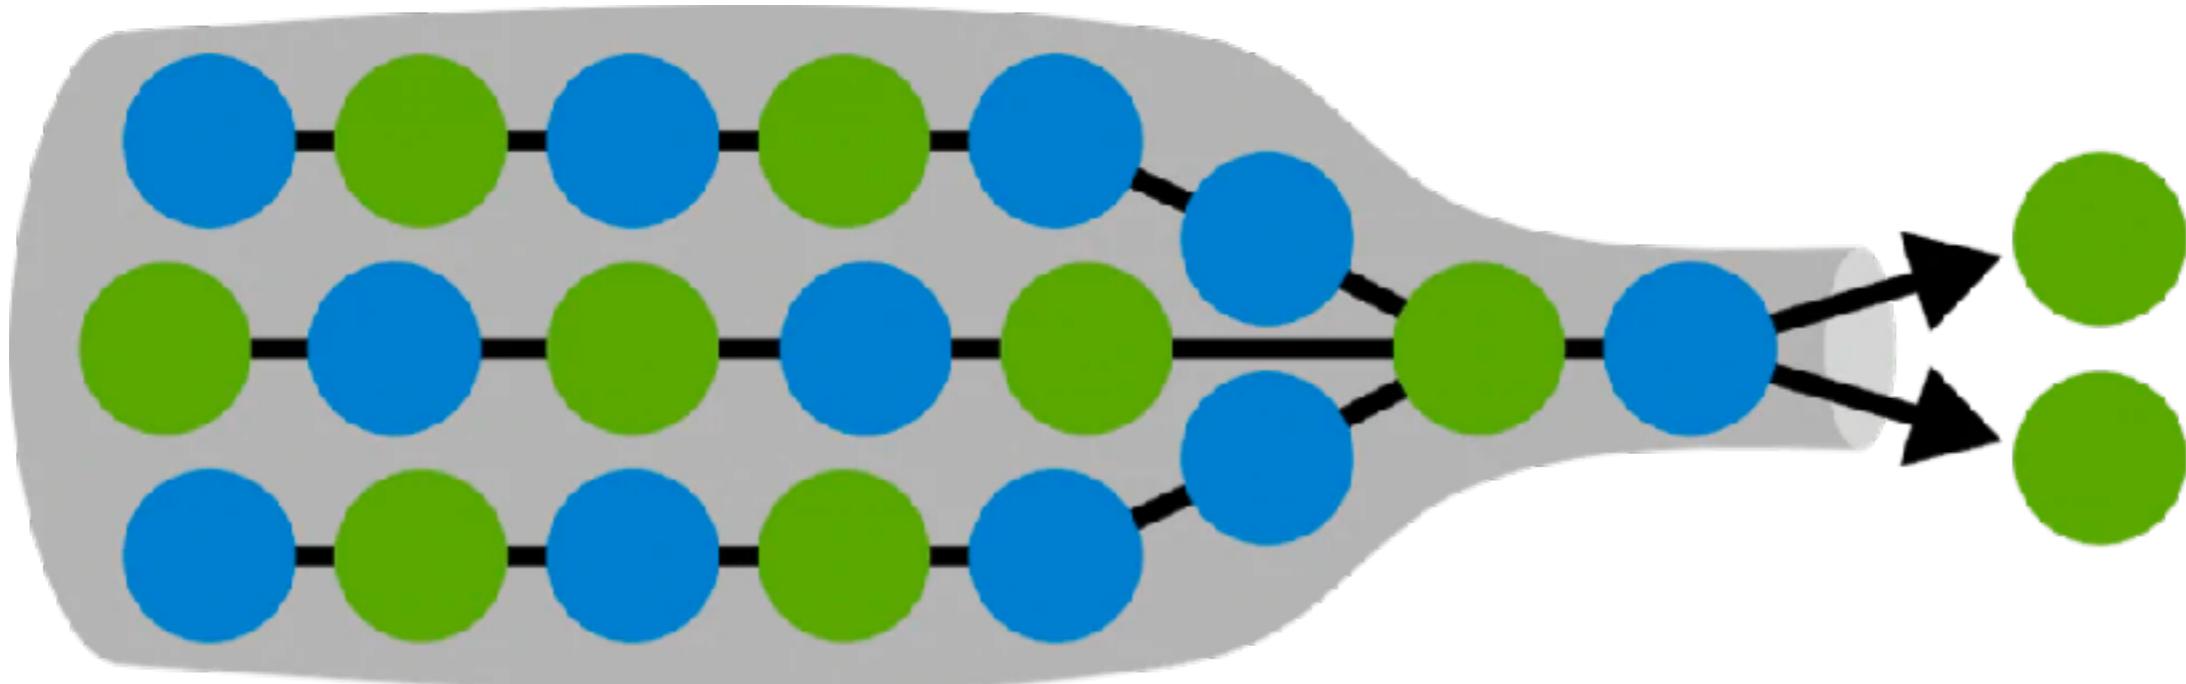
# Tips #4

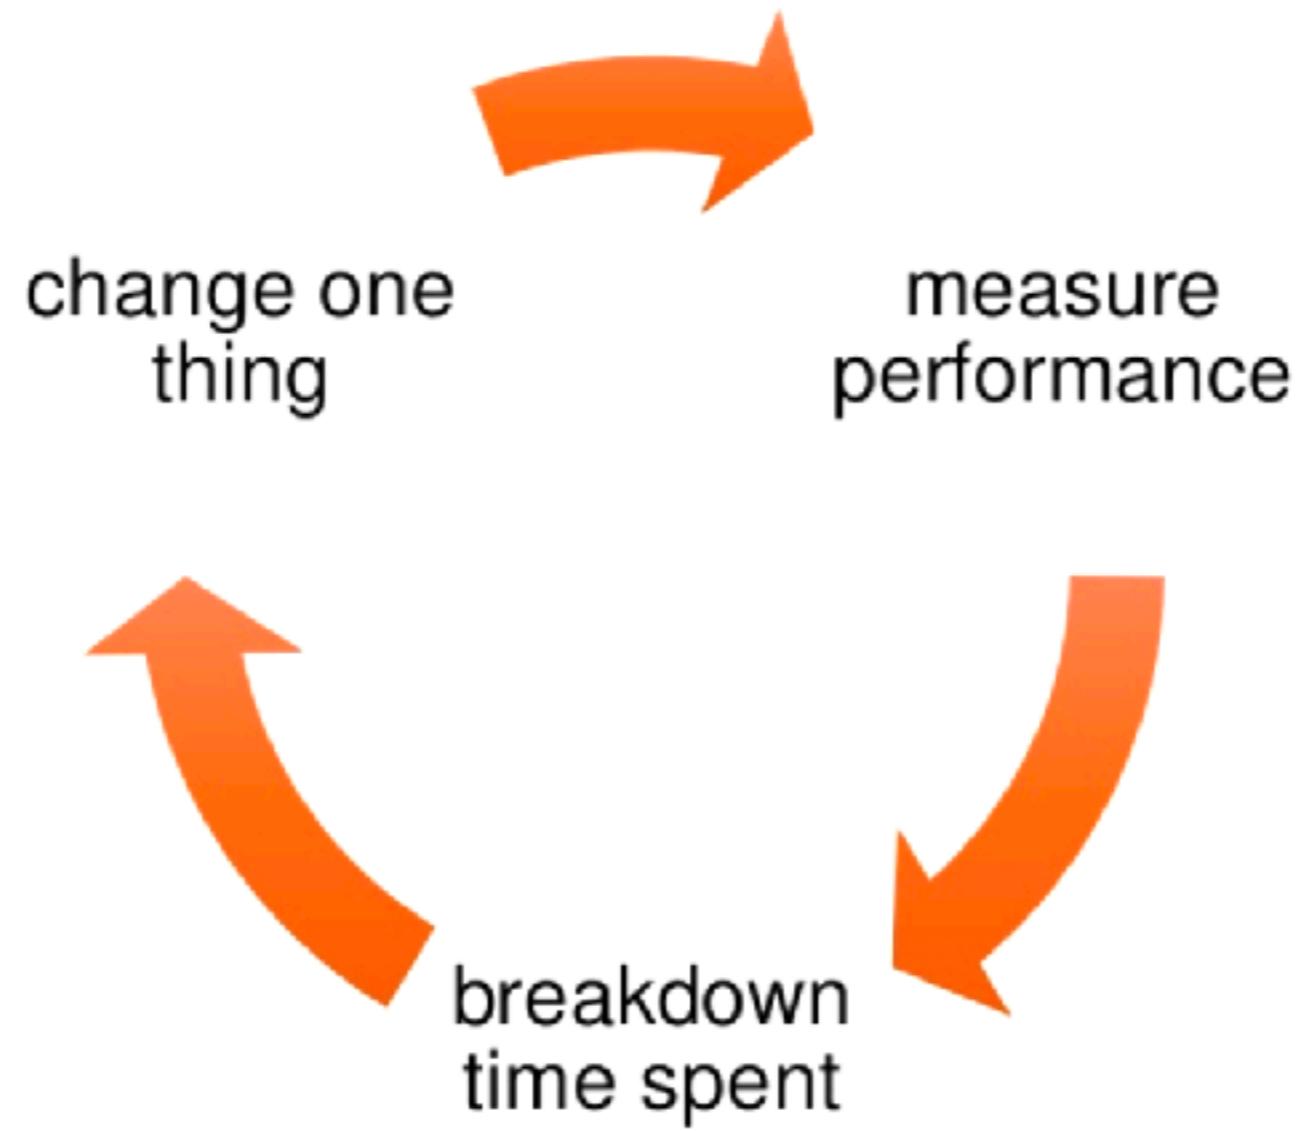
Monitoring brokers isn't enough  
**Need to monitor events**



# Tips #4

## Tuning the bottlenecks





# Life cycle of request

Client send request to broker

Network thread get request and put on queue

I/O thread/handler pick up request and process

(Read and write from/to disk)

Waiting for other brokers to ack messages

Put response on queue

Network thread send response to client



# Kafka for operation



# Kafka operations

Rolling restart of brokers

Update configurations

Rebalancing partitions

Increase replication factor

Add/replace/remove a broker

Upgrade a Kafka cluster with zero downtime



# Security in Kafka



# Security in Kafka

Authentication (SSL/SASL)

Authorization with ACL (Access Control List)

Data encryption



# Plain text vs SSL

